

sbt Reference Manual

Contents

Preface	15
Install	15
Getting Started	15
Features of sbt	16
Also	16
General Information	16
Credits	16
Community Plugins	17
sbt Organization	17
Community Ivy Repository	17
Available Plugins	17
Test plugins	18
Community Repository Policy	24
Bintray For Plugins	24
Create an account on Bintray	25
Create a repository for your sbt plugins	25
Add the bintray-sbt plugin to your build.	25
Make a release	26
Linking your package to the sbt organization	27
Linking your package to the sbt organization (sbt org admins)	27
Summary	27
Setup Notes	27

Do not put <code>sbt-launch.jar</code> on your classpath.	27
Terminal encoding	28
JVM heap, permgen, and stack sizes	28
Boot directory	28
HTTP/HTTPS/FTP Proxy	28
Deploying to Sonatype	29
First - PGP Signatures	29
Second - Maven Publishing Settings	30
Third - POM Metadata	30
Fourth - Adding credentials	31
Finally - Publish	32
Summary	32
Changes	34
0.13.5-RC1 to 0.13.5-RC2	34
0.13.2 to 0.13.5	34
0.13.1 to 0.13.2	35
0.13.0 to 0.13.1	35
0.12.4 to 0.13.0	37
sbt 0.13.0 Changes	37
Overview	37
Details of major changes	39
sbt 0.12.0 Changes	44
Details of major changes from 0.11.2 to 0.12.0	45
scala-library.jar	48
Older Changes	49
0.12.3 to 0.12.4	49
0.12.2 to 0.12.3	49
0.12.1 to 0.12.2	50
0.12.0 to 0.12.1	51
0.11.3 to 0.12.0	52
0.11.2 to 0.11.3	53

0.11.1 to 0.11.2	53
0.11.0 to 0.11.1	54
0.10.1 to 0.11.0	55
0.10.0 to 0.10.1	56
0.7.7 to 0.10.0	56
0.7.5 to 0.7.7	57
0.7.4 to 0.7.5	57
0.7.3 to 0.7.4	58
0.7.2 to 0.7.3	59
0.7.1 to 0.7.2	59
0.7.0 to 0.7.1	60
0.5.6 to 0.7.0	60
0.5.5 to 0.5.6	62
0.5.4 to 0.5.5	62
0.5.2 to 0.5.4	62
0.5.1 to 0.5.2	64
0.4.6 to 0.5/0.5.1	65
0.4.5 to 0.4.6	66
0.4.3 to 0.4.5	67
0.4 to 0.4.3	68
0.3.7 to 0.4	68
0.3.6 to 0.3.7	70
0.3.5 to 0.3.6	71
0.3.2 to 0.3.5	71
0.3.1 to 0.3.2	72
0.3 to 0.3.1	72
0.2.3 to 0.3	72
0.2.2 to 0.2.3	72
0.2.1 to 0.2.2	73
0.2.0 to 0.2.1	73
0.1.9 to 0.2.0	73

0.1.8 to 0.1.9	73
0.1.7 to 0.1.8	74
0.1.6 to 0.1.7	74
0.1.5 to 0.1.6	75
0.1.4 to 0.1.5	75
0.1.3 to 0.1.4	76
0.1.2 to 0.1.3	76
0.1.1 to 0.1.2	76
0.1 to 0.1.1	76
Migrating from 0.7 to 0.10+	76
Why move to 0.13.5?	77
Preserve <code>project/</code> for 0.7.x project	77
Create <code>build.sbt</code> for 0.13.5	77
Run sbt 0.13.5	78
Switching back to sbt 0.7.x	78
Contributing to sbt	79
Documentation	80
Detailed Topics	80
Using sbt	80
Command Line Reference	80
Notes on the command line	81
Project-level tasks	81
Configuration-level tasks	81
General commands	83
Commands for managing the build definition	84
Command Line Options	84
Console Project	86
Description	86
Accessing settings	86
Evaluating tasks	87
State	87

Cross-building	88
Introduction	88
Publishing Conventions	88
Using Cross-Built Libraries	88
Cross-Building a Project	89
Interacting with the Configuration System	90
Selecting commands, tasks, and settings	91
Discovering Settings and Tasks	92
Triggered Execution	96
Compile	96
Testing	97
Running Multiple Commands	97
Scripts, REPL, and Dependencies	97
Setup	97
Usage	98
Understanding Incremental Recompilation	100
sbt heuristics	101
What is included in the interface of a Scala class	102
How to take advantage of sbt heuristics	104
Further references	107
Configuration	107
Classpaths, sources, and resources	107
Basics	107
Compiler Plugin Support	110
Continuations Plugin Example	111
Version-specific Compiler Plugin Example	111
Configuring Scala	111
Automatically managed Scala	111
Using Scala from a local directory	113
sbt's Scala version	114
Forking	114

Enable forking	115
Change working directory	115
Forked JVM options	116
Java Home	116
Configuring output	117
Configuring Input	117
Direct Usage	117
Global Settings	118
Basic global configuration file	118
Global Settings using a Global Plugin	118
Java Sources	119
Usage	119
Mapping Files	120
Relative to a directory	120
Rebase	121
Flatten	122
Alternatives	122
Local Scala	122
Macro Projects	123
Introduction	123
Defining the Project Relationships	123
Common Interface	125
Distribution	125
Paths	126
Constructing a File	126
Path Finders	127
File Filters	129
Parallel Execution	130
Task ordering	130
Practical constraints	130
Configuration	131

Future work	134
External Processes	136
Usage	136
Running Project Code	138
Problems	138
sbt's Solutions	139
Testing	140
Basics	140
Output	141
Options	142
Additional test configurations	143
JUnit	147
Extensions	147
Dependency Management	148
Artifacts	148
Selecting default artifacts	148
Modifying default artifacts	149
Defining custom artifacts	150
Publishing .war files	151
Using dependencies with artifacts	151
Dependency Management Flow	152
Background	152
Caching and Configuration	152
General troubleshooting steps	153
Plugins	153
Notes	154
Library Management	154
Introduction	154
Manual Dependency Management	154
Automatic Dependency Management	155
Proxy Repositories	165

Overview	165
sbt Configuration	167
<code>~/.sbt/repositories</code>	167
Proxying Ivy Repositories	168
Publishing	168
Define the repository	168
Credentials	170
Cross-publishing	170
Published artifacts	171
Modifying the generated POM	171
Publishing Locally	171
Resolvers	172
Maven	172
Predefined	172
Custom	173
Update Report	176
Filtering a Report and Getting Artifacts	176
Tasks and Commands	179
Tasks	179
Introduction	179
Features	180
Defining a Task	180
Getting values from multiple scopes	183
Advanced Task Operations	187
Dynamic Computations with <code>Def.taskDyn</code>	188
Input Tasks	193
Input Keys	193
Basic Input Task Definition	193
Input Task using Parsers	194
The <code>InputTask</code> type	195
Using other input tasks	195

Preapplying input	196
Get a Task from an InputTask	197
Commands	199
What is a “command”?.	199
Introduction	199
Defining a Command	200
Full Example	201
Parsing and tab completion	203
Basic parsers	203
Built-in parsers	204
Combining parsers	204
Transforming results	205
Controlling tab completion	205
State and actions	206
Command-related data	206
Project-related data	207
Project data	208
Classpaths	209
Running tasks	209
Using State in a task	210
Tasks/Settings: Motivation	210
Application	212
Plugins and Best Practices	212
General Best Practices	212
project/ vs. ~/sbt/	212
Local settings	213
.sbtrc	213
Generated files	213
Don’t hard code	213
Don’t “mutate” files	214
Use absolute paths	214

Parser combinators	215
Plugins	215
Using an auto plugin	216
By Description	217
Plugin dependencies	217
Creating an auto plugin	218
Using a library in a build definition example	224
Best Practices	226
Plugins Best Practices	226
Get your plugins known	226
Don't use default package	227
Use settings and tasks. Avoid commands.	227
Use <code>sbt.AutoPlugin</code>	227
Reuse existing keys	227
Avoid namespace clashes	227
Provide core feature in a plain old Scala object	228
Configuration advices	228
Mucking with <code>globalSettings</code>	231
Sbt Launcher	232
Getting Started with the Sbt Launcher	232
Overview	232
Sbt Launcher Architecture	236
Module Resolution	236
Classloader Caching and Isolation	236
Caching	238
Locking	238
Service Discovery and Isolation	238
Sbt Launcher Configuration	239
Example	239
Variable Substitution	242
Syntax	242

Developer's Guide	243
Core Principles	244
Introduction to build state	244
Settings Architecture	245
Task Architecture	247
Settings Core	247
Example	248
sbt Settings Discussion	250
Setting Initialization	251
Controlling Initialization	253
Build Loaders	255
Custom Resolver	255
Custom Builder	257
Custom Transformer	259
The BuildDependencies type	260
Creating Command Line Applications Using sbt	260
Hello World Example	261
Nightly Builds	263
How to...	264
Classpaths	264
Include a new type of managed artifact on the classpath, such as mar	264
Get the classpath used for compilation	264
Get the runtime classpath, including the project's compiled classes	264
Get the test classpath, including the project's compiled test classes	265
Use packaged jars on classpaths instead of class directories	265
Get all managed jars for a configuration	266
Get the files included in a classpath	266
Get the module and artifact that produced a classpath entry . .	266
Customizing paths	267
Change the default Scala source directory	267

Change the default Java source directory	267
Change the default resource directory	267
Change the default (unmanaged) library directory	268
Disable using the project's base directory as a source directory .	268
Add an additional source directory	268
Add an additional resource directory	269
Include/exclude files in the source directory	269
Include/exclude files in the resource directory	269
Include only certain (unmanaged) libraries	270
Generating files	270
Generate sources	270
Generate resources	272
Inspect the build	272
Show or search help for a command, task, or setting	272
List available tasks	273
List available settings	273
Display the description and type of a setting or task	274
Display the delegation chain of a setting or task	275
Show the list of projects and builds	275
Show the current session (temporary) settings	276
Show basic information about sbt and the current build	276
Show the value of a setting	276
Show the result of executing a task	276
Show the classpath used for compilation or testing	277
Show the main classes detected in a project	277
Show the test classes detected in a project	277
Interactive mode	278
Use tab completion	278
Show more tab completion suggestions	278
Modify the default JLine keybindings	279
Configure the prompt string	279

Use history	279
Change the location of the interactive history file	279
Use the same history for all projects	280
Disable interactive history	280
Run commands before entering interactive mode	280
Configure and use logging	281
View the logging output of the previously executed command . .	281
View the previous logging output of a specific task	282
Show warnings from the previous compilation	283
Change the logging level globally	283
Change the logging level for a specific task, configuration, or project	284
Configure printing of stack traces	284
Print the output of tests immediately instead of buffering	285
Add a custom logger	285
Log messages in a task	285
Project metadata	285
Set the project name	285
Set the project version	286
Set the project organization	286
Set the project's homepage and other metadata	286
Configure packaging	286
Use the packaged jar on classpaths instead of class directory . . .	286
Add manifest attributes	287
Change the file name of a package	287
Modify the contents of the package	288
Running commands	288
Pass arguments to a command or task in batch mode	288
Provide multiple commands to run consecutively	288
Read commands from a file	288
Define an alias for a command or task	289
Quickly evaluate a Scala expression	289

Configure and use Scala	289
Set the Scala version used for building the project	289
Disable the automatic dependency on the Scala library	290
Temporarily switch to a different Scala version	290
Use a local Scala installation for building a project	290
Build a project against multiple Scala versions	290
Enter the Scala REPL with a project's dependencies on the class- path, but not the compiled project classes	290
Enter the Scala REPL with a project's dependencies and com- piled code on the classpath	290
Enter the Scala REPL with plugins and the build definition on the classpath	291
Define the initial commands evaluated when entering the Scala REPL	291
Define the commands evaluated when exiting the Scala REPL . .	291
Use the Scala REPL from project code	292
Generate API documentation	292
Select javadoc or scaladoc	292
Set the options used for generating scaladoc independently of compilation	292
Add options for scaladoc to the compilation options	293
Set the options used for generating javadoc independently of com- pilation	293
Add options for javadoc to the compilation options	293
Enable automatic linking to the external Scaladoc of managed dependencies	293
Enable manual linking to the external Scaladoc of managed de- pendencies	293
Define the location of API documentation for a library	294
Triggered execution	294
Run a command when sources change	294
Run multiple commands when sources change	295
Configure the sources that are checked for changes	295
Set the time interval between checks for changes to sources . . .	295

Examples	295
.sbt build examples	296
.scala build example	300
External Builds	303
Advanced configurations example	303
Advanced command example	305
Frequently Asked Questions	306
Project Information	306
Usage	307
Build definitions	307
Extending sbt	309
Errors	313
Dependency Management	314
Miscellaneous	315
0.7 to 0.10+ Migration	315
My tests all run really fast but some are broken that weren't in 0.7!	316
Index	318
Values and Types	318
Methods	319

Preface

sbt is a build tool for Scala, Java, and [more](#). It requires Java 1.6 or later.

Install

See [Installing sbt](#) for the setup instructions.

Getting Started

To get started, *please read* the [Getting Started Guide](#). You will save yourself a *lot* of time if you have the right understanding of the big picture up-front. All documentation may be found via the table of contents included at the end of every page.

Use [Stack Overflow](#) for questions. Use the [sbt-dev mailing list](#) for discussing sbt development. Use [\[@scala_sbt\]\(https://twitter.com/scala_sbt\)](#) for questions and discussions.

Features of sbt

- Little or no configuration required for simple projects
- Scala-based [build definition](#) that can use the full flexibility of Scala code
- Accurate incremental recompilation using information extracted from the compiler
- Continuous compilation and testing with [triggered execution](#)
- Packages and publishes jars
- Generates documentation with scaladoc
- Supports mixed Scala/[Java](#) projects
- Supports [testing](#) with ScalaCheck, specs, and ScalaTest. JUnit is supported by a plugin.
- Starts the Scala REPL with project classes and dependencies on the class-path
- Modularization supported with [sub-projects](#)
- External project support (list a git repository as a dependency!)
- [Parallel task execution](#), including parallel test execution
- [Library management support](#): inline declarations, external Ivy or Maven configuration files, or manual management

Also

This documentation can be forked on [GitHub](#). Feel free to make corrections and add documentation.

Documentation for 0.7.x has been [archived here](#). This documentation applies to sbt 0.13.5.

See also the [API Documentation](#), [SXR Documentation](#), and the [index of names and types](#).

General Information

This part of the documentation has project “meta-information” such as where to get help, find source code and how to contribute.

Credits

See the [sbt contributors on GitHub](#) and [sbt GitHub organization members](#).

Additionally, these people have contributed ideas, documentation, or code to sbt but are not recorded in either of the above:

- Josh Cough
- Nolan Darilek
- Nathan Hamblen
- Ismael Juma
- Viktor Klang
- David R. MacIver
- Ross McDonald
- Andrew O'Malley
- Jorge Ortiz
- Mikko Peltonen
- Ray Racine
- Stuart Roebuck
- Harshad RJ
- Tony Sloane
- Seth Tisue
- Francisco Treacy
- Vesa Vilhonen

Community Plugins

sbt Organization

The [sbt organization](#) is available for use by any sbt plugin. Developers who contribute their plugins into the community organization will still retain control over their repository and its access. The goal of the sbt organization is to organize sbt software into one central location.

A side benefit to using the sbt organization for projects is that you can use gh-pages to host websites under the <http://scala-sbt.org> domain.

Community Ivy Repository

[Typesafe](#) has provided a freely available [Ivy Repository](#) for sbt projects to use. This Ivy repository is mirrored from the freely available [Bintray service](#). If you'd like to submit your plugin, please follow these instructions: [Bintray For Plugins](#).

Available Plugins

Please feel free to [submit a pull request](#) that adds your plugin to the list.

Plugins for IDEs

- IntelliJ IDEA
 - sbt Plugin to generate IDEA project configuration: <https://github.com/mpeltonen/sbt-idea>
 - IDEA Plugin to embed an sbt Console into the IDE: <https://github.com/orfjackal/idea-sbt-plugin>
- Netbeans (no support to create a new sbt project yet)
 - sbt-netbeans-plugin (older): <https://github.com/remeniuk/sbt-netbeans-plugin>
 - sbt plugin to generate NetBeans configuration: <https://github.com/dcaoyuan/nbsbt>
 - sbt plugin to add scala support to NetBeans: <https://github.com/dcaoyuan/nbscala>
- Eclipse: <https://github.com/typesafehub/sbteclipse>
- Sublime Text: <https://github.com/orrsella/sbt-sublime>
- Ensime: <https://github.com/aemoncannon/ensime-sbt-cmd>
- sbt-mode for Emacs: <https://github.com/hvesalai/sbt-mode>
- sbt-ctags (manage library dependency sources for vim, emacs, sublime) <https://github.com/kalmanb/sbt-ctags>

Web Plugins

- xsbt-web-plugin: <https://github.com/JamesEarlDouglas/xsbt-web-plugin>
- xsbt-webstart: <https://github.com/ritschwumm/xsbt-webstart>
- sbt-appengine: <https://github.com/sbt/sbt-appengine>
- sbt-gwt-plugin: <https://github.com/thunderklaus/sbt-gwt-plugin>
- sbt-cloudbees-plugin: <https://github.com/timperrett/sbt-cloudbees-plugin>
- sbt-jelastic-deploy: <https://github.com/casualjim/sbt-jelastic-deploy>
- sbt-elasticbeanstalk (Deploy WAR files to AWS Elastic Beanstalk): <https://github.com/sqs/sbt-elasticbeanstalk>
- sbt-cloudformation (AWS CloudFormation templates and stacks management): <https://github.com/tptodorov/sbt-cloudformation>

Test plugins

- junit_xml_listener: https://github.com/ijuma/junit_xml_listener
- sbt-growl-plugin: <https://github.com/softprops/sbt-growl-plugin>
- sbt-teamcity-test-reporting-plugin: <https://github.com/guardian/sbt-teamcity-test-reporting-plugin>
- xsbt-cucumber-plugin: <https://github.com/skipoleschris/xsbt-cucumber-plugin>

- sbt-multi-jvm: <https://github.com/typesafehub/sbt-multi-jvm>
- sbt-testng-interface: <https://github.com/sbt/sbt-testng-interface>

Static Code Analysis plugins

- cpd4sbt: <https://github.com/sbt/cpd4sbt> (copy/paste detection, works for Scala, too)
- findbugs4sbt: <https://github.com/sbt/findbugs4sbt> (FindBugs only supports Java projects atm)
- scalastyle: <https://github.com/scalastyle/scalastyle-sbt-plugin> (Scalastyle - static code checker for Scala)
- sbt-scapegoat: <https://github.com/sksamuel/sbt-scapegoat> (Scapegoat - Scala static code analysis)
- sbt-stats: <https://github.com/orrsella/sbt-stats> (simple, extensible source code statistics)

One jar plugins

- sbt-assembly: <https://github.com/sbt/sbt-assembly>
- xsbt-proguard-plugin: <https://github.com/adamw/xsbt-proguard-plugin>
- sbt-deploy: <https://github.com/reaktor/sbt-deploy>
- sbt-appbundle (os x standalone): <https://github.com/sbt/sbt-appbundle>
- sbt-onejar (Packages your project using One-JAR™): <https://github.com/sbt/sbt-onejar>

Frontend development plugins

- coffeescripted-sbt: <https://github.com/softprops/coffeescripted-sbt>
- less-sbt (for less-1.3.0): <https://github.com/softprops/less-sbt>
- sbt-less-plugin (it uses less-1.3.0): <https://github.com/btd/sbt-less-plugin>
- sbt-emberjs: <https://github.com/stefri/sbt-emberjs>
- sbt-closure: <https://github.com/eltimn/sbt-closure>
- sbt-imagej: <https://github.com/jpsacha/sbt-imagej>
- sbt-yui-compressor: <https://github.com/indrajitr/sbt-yui-compressor>
- sbt-requirejs: <https://github.com/scalatra/sbt-requirejs>
- sbt-vaadin-plugin: <https://github.com/henrikerola/sbt-vaadin-plugin>
- sbt-purescript: <https://github.com/eamelink/sbt-purescript>
- sbt-jasmine-plugin (Run javascript tests with jasmine within sbt): <https://github.com/joescii/sbt-jasmine-plugin>

Game development plugins

- sbt-lwjgl-plugin (Light Weight Java Game Library): <https://github.com/philcali/sbt-lwjgl-plugin>
- sbt-scage-plugin (Scala Game Engine): <https://github.com/mvallerie/sbt-scage-plugin>

Release plugins

- sbt-aether-plugin (Published artifacts using Sonatype Aether): <https://github.com/arktekk/sbt-aether-deploy>
- posterous-sbt: <https://github.com/n8han/posterous-sbt>
- sbt-signer-plugin: <https://github.com/rossabaker/sbt-signer-plugin>
- sbt-izpack (generates IzPack an installer): <http://software.clapper.org/sbt-izpack/>
- sbt-ghpages (publishes generated site and api): <https://github.com/sbt/sbt-ghpages>
- sbt-pgp (PGP signing plugin, can generate keys too): <https://github.com/sbt/sbt-pgp>
- sbt-release (customizable release process): <https://github.com/sbt/sbt-release>
- sbt-unique-version (emulates unique snapshots): <https://github.com/sbt/sbt-unique-version>
- sbt-install4j: <https://github.com/jpsacha/sbt-install4j>
- sbt-pack (generates packages with dependent jars and launch scripts): <https://github.com/xerial/sbt-pack>
- sbt-start-script: <https://github.com/sbt/sbt-start-script>
- sbt-native-packager: <https://github.com/sbt/sbt-native-packager>
- sbt-sonatype-plugin (releases to Sonatype Nexus repository) <https://github.com/xerial/sbt-sonatype>
- xitrum-package (collects dependency .jar files for standalone Scala programs): <https://github.com/ngocdaothanh/xitrum-package>

System plugins

- sbt-sh (executes shell commands): <https://github.com/steppenwells/sbt-sh>
- cronish-sbt (interval sbt / shell command execution): <https://github.com/philcali/cronish-sbt>
- git (executes git commands): <https://github.com/sbt/sbt-git>
- svn (execute svn commands): <https://github.com/xuwei-k/sbtsvn>
- sbt-groll (sbt plugin to navigate the Git history): <https://github.com/sbt/sbt-groll>
- sbt-twt (twitter processor for sbt): <https://github.com/sbt/sbt-twt>

Code generator plugins

- sbt-haxe (Compiling [Haxe](#) to Java): <https://bitbucket.org/qforce/sbt-haxe>
- sbt-scalabuff (Google Protocol Buffers with native scala support thru ScalaBuff): <https://github.com/sbt/sbt-scalabuff>
- sbt-fmpp (FreeMarker Scala/Java Templating): <https://github.com/sbt/sbt-fmpp>
- sbt-scalaxb (XSD and WSDL binding): <https://github.com/eed3si9n/scalaxb>
- sbt-protobuf (Google Protocol Buffers): <https://github.com/sbt/sbt-protobuf>
- sbt-cppp (Cross-Project Protobuf Plugin for Sbt): <https://github.com/Atty/sbt-cppp>
- sbt-avro (Apache Avro): <https://github.com/cavorite/sbt-avro>
- sbt-xjc (XSD binding, using [JAXB XJC](#)): <https://github.com/sbt/sbt-xjc>
- xsbt-scalate-generate (Generate/Precompile Scalate Templates): <https://github.com/backchatio/xsbt-scalate-generate>
- sbt-antlr (Generate Java source code based on ANTLR3 grammars): <https://github.com/stefri/sbt-antlr>
- sbt-antlr4 (Antlr4 runner for generating Java source code): <https://github.com/ihji/sbt-antlr4>
- xsbt-reflect (Generate Scala source code for project name and version): <https://github.com/ritschwumm/xsbt-reflect>
- sbt-buildinfo (Generate Scala source for any settings): <https://github.com/sbt/sbt-buildinfo>
- lifty (Brings scaffolding to sbt): <https://github.com/lifty/lifty>
- sbt-thrift (Thrift Code Generation): <https://github.com/bigtoast/sbt-thrift>
- xsbt-hginfo (Generate Scala source code for Mercurial repository information): https://bitbucket.org/lukas_pustina/xsbt-hginfo
- sbt-scalashim (Generate Scala shim like `sys.error`): <https://github.com/sbt/sbt-scalashim>
- sbtend (Generate Java source code from [xtend](#)): <https://github.com/xuwei-k/sbtend>
- sbt-boilerplate (generating scala.Tuple/Function related boilerplate code): <https://github.com/sbt/sbt-boilerplate>
- sbt-fxml (Generates controller classes for JavaFX FXML files): <https://bitbucket.org/phdoerfler/sbt-fxml>
- sbt-clojure (Compiling Clojure code): <https://github.com/Geal/sbt-clojure>

Database plugins

- flyway-sbt (Flyway - The agile database migration framework): <http://flywaydb.org/getstarted/firststeps/sbt.html>
- sbt-liquibase (Liquibase RDBMS database migrations): <https://github.com/bigtoast/sbt-liquibase>
- sbt-dbdeploy (dbdeploy, a database change management tool): <https://github.com/mr-ken/sbt-dbdeploy>

Documentation plugins

- sbt-lwm (Convert lightweight markup files, e.g., Markdown and Textile, to HTML): <http://software.clapper.org/sbt-lwm/>
- sbt-site (Site generation for SBT): <https://github.com/sbt/sbt-site>
- Laika (Template-based site generation, Markdown, reStructuredText, no external tools): <http://planet42.github.io/Laika/>
- literator-plugin (Converts sources into markdown documents): <https://github.com/laughedelic/literator>
- sbt-class-diagram (Create a class diagram) <https://github.com/xuwei-k/sbt-class-diagram>

Library dependency plugins

- ls-sbt (An sbt interface for ls.implicit.ly): <https://github.com/softprops/ls>
- sbt-dependency-graph (Creates a graphml file of the dependency tree): <https://github.com/jrudolph/sbt-dependency-graph>
- sbt-dirty-money (Cleans Ivy2 cache): <https://github.com/sbt/sbt-dirty-money>
- sbt-updates (Checks Maven repos for dependency updates): <https://github.com/rtimush/sbt-updates>
- sbt-lock (Locks library versions for reproducible build): <https://github.com/tkawachi/sbt-lock>
- sbt-versions (Checks for updated versions of your dependencies): <https://github.com/sksamuel/sbt-versions>

Utility plugins

- sbt-process-runner (Run your own applications from SBT console) <https://github.com/whysoserious/sbt-process-runner>
- jot (Write down your ideas lest you forget them) <https://github.com/softprops/jot>
- np (Dead simple new project directory generation): <https://github.com/softprops/np>
- npt (Creates new project skeletons based on templates): <https://github.com/reikje/npt>

- sbt-editsource (A poor man's *sed*(1), for sbt): <http://software.clapper.org/sbt-editsource/>
- sbt-conflict-classes (Show conflict classes from classpath): <https://github.com/todesking/sbt-conflict-classes>
- sbt-cross-building (Simplifies building your plugins for multiple versions of sbt): <https://github.com/jrudolph/sbt-cross-building>
- sbt-doge (aggregates tasks across subprojects and their `crossScalaVersions`): <https://github.com/sbt/sbt-doge>
- sbt-revolver (Triggered restart, hot reloading): <https://github.com/spray/sbt-revolver>
- sbt-scalaedit (Open and upgrade ScalaEdit (text editor)): <https://github.com/kjellwinblad/sbt-scalaedit-plugin>
- sbt-man (Looks up scaladoc): <https://github.com/sbt/sbt-man>
- sbt-taglist (Looks for TODO-tags in the sources): <https://github.com/johanandren/sbt-taglist>
- migration-manager: <https://github.com/typesafehub/migration-manager>
- sbt-scalariform (adding support for source code formatting using Scalariform): <https://github.com/sbt/sbt-scalariform>
- sbt-aspectj: <https://github.com/sbt/sbt-aspectj>
- sbt-properties: <https://github.com/sbt/sbt-properties>
- sbt-multi-publish (publish to more than one repository simultaneously): <https://github.com/davidharcombe/sbt-multi-publish>
- sbt-about-plugins (shows some details about plugins loaded): <https://github.com/jozic/sbt-about-plugins>
- sbt-one-log (make Log dependency easy): <https://github.com/zavakid/sbt-one-log>
- sbt-git-stamp (include git metadata in MANIFEST.MF file in artifact): <https://bitbucket.org/pkaeding/sbt-git-stamp>
- fm-sbt-s3-resolver (Resolve and Publish using Amazon S3): <https://github.com/frugalmechanic/fm-sbt-s3-resolver>
- sbt-notebook (Adds scala-notebook capabilities to sbt projects): <https://github.com/alexarchambault/sbt-notebook>

Code coverage plugins

- sbt-scct: <https://github.com/sqality/sbt-scct>
- sbt-scoverage: <https://github.com/scoverage/sbt-scoverage>
- jacoco4sbt: <https://github.com/sbt/jacoco4sbt>
- xsbt-coveralls-plugin: <https://github.com/theon/xsbt-coveralls-plugin>

Android plugin

- android-plugin: <https://github.com/jberkel/android-plugin>

- android-sdk-plugin: <https://github.com/pfn/android-sdk-plugin>

Build interoperability plugins

- ant4sbt: <https://github.com/sbt/ant4sbt>

OSGi plugin

- sbtosgi: <https://github.com/typesafehub/sbtosgi>

Plugin bundles

- tl-os-sbt-plugins (Version, Release, and Package Management, Play 2.0 and Git utilities) : <https://github.com/trafficland/tl-os-sbt-plugins>

Community Repository Policy

The community repository has the following guideline for artifacts published to it:

1. All published artifacts are the authors own work or have an appropriate license which grants distribution rights.
2. All published artifacts come from open source projects, that have an open patch acceptance policy.
3. All published artifacts are placed under an organization in a DNS domain for which you have the permission to use or are an owner (scala-sbt.org is available for sbt plugins).
4. All published artifacts are signed by a committer of the project (coming soon).

Bintray For Plugins

This is currently in Beta mode.

sbt hosts their community plugin repository on [Bintray](#). Bintray is a repository hosting site, similar to github, which allows users to contribute their own plugins, while sbt can aggregate them together in a common repository.

This document walks you through the means to create your own repository for hosting your sbt plugins and then linking them into the sbt shared repository. This will make your plugins available for all sbt users without additional configuration (besides declaring a dependency on your plugin).

To do this, we need to perform the following steps:

Create an account on Bintray

First, go to <http://bintray.com>. Click on the sign in link on the top left, and then the sign up button.

Note: If you had an account on `repo.scala-sbt.org` previous, please use the same email address when you create this account.

Create a repository for your sbt plugins

Now, we'll create a repository to host our personal sbt plugins. In bintray, create a generic repository called `sbt-plugins`.

First, go to your user page and click on the `new repository` link:

You should see the following dialog:

Fill it out similarly to the above image, the settings are:

- Name: `sbt-plugins`
- Type: `Generic`
- Desc: `My sbt plugins`
- Tags: `sbt`

Once this is done, you can begin to configure your `sbt-plugins` to publish to bintray.

Add the bintray-sbt plugin to your build.

First, add the `bintray-sbt` to your plugin build.

First, create a `project/bintray.sbt` file

```
resolvers += Resolver.url(
  "bintray-sbt-plugin-releases",
  url("http://dl.bintray.com/content/sbt/sbt-plugin-releases"))(
  Resolver.ivyStylePatterns)

addSbtPlugin("me.lessis" % "bintray-sbt" % "0.1.1")
```

Next, make sure your `build.sbt` file has the following settings :

```
import bintray.Keys._

sbtPlugin := true
```

```

name := "<YOUR PLUGIN HERE>"

organization := "<INSERT YOUR ORG HERE>"

version := "<YOUR PLUGIN VERSION HERE>"

publishMavenStyle := false

bintrayPublishSettings

repository in bintray := "sbt-plugins"

// This is an example. bintray-sbt requires licenses to be specified
// (using a canonical name).
licenses += ("Apache-2.0", url("http://www.apache.org/licenses/LICENSE-2.0.html"))

bintrayOrganization in bintray := None

```

Make sure your project has a valid license specified, as well as unique name and organization.

Make a release

Note: bintray does not support snapshots. We recommend using [git-revisions](#) supplied by the [sbt-git](#) plugin.

Once your build is configured, open the sbt console in your build and run

```
sbt> publish
```

The plugin will ask you for your credentials. If you don't know where they are, you can find them on [Bintray](#).

1. Login to the website with your credentials.
2. Click on your username
3. Click on edit profile
4. Click on API Key

This will get you your password. The bintray-sbt plugin will save your API key for future use.

NOTE: We have to do this before we can link our package to the sbt org.

Linking your package to the sbt organization

Now that your plugin is packaged on bintray, you can include it in the community sbt repository. To do so, go to the [Community sbt repository](#) screen.

1. Click the green **include my package** button and select your plugin.
2. Search for your plugin by name and click on the link.
3. Your request should be automatically filled out, just click send
4. Shortly, one of the sbt repository admins will approve your link request.

From here on, any releases of your plugin will automatically appear in the community sbt repository. Congratulations and thank you so much for your contributions!

Linking your package to the sbt organization (sbt org admins)

If you're a member of the sbt organization on bintray, you can link your package to the sbt organization, but via a different means. To do so, first navigate to the plugin you wish to include and click on the link button:

After clicking this you should see a link like the following:

Click on the **sbt/sbt-plugin-releases** repository and you're done! Any future releases will be included in the sbt-plugin repository.

Summary

After setting up the repository, all new releases will automatically be included the sbt-plugin-releases repository, available for all users. When you create a new plugin, after the initial release you'll have to link it to the sbt community repository, but the rest of the setup should already be completed. Thanks for you contributions and happy hacking.

Setup Notes

Some notes on how to set up your sbt script.

Do not put sbt-launch.jar on your classpath.

Do *not* put sbt-launch.jar in your \$SCALA_HOME/lib directory, your project's lib directory, or anywhere it will be put on a classpath. It isn't a library.

Terminal encoding

The character encoding used by your terminal may differ from Java's default encoding for your platform. In this case, you will need to add the option `-Dfile.encoding=<encoding>` in your `sbt` script to set the encoding, which might look like:

```
java -Dfile.encoding=UTF8
```

JVM heap, permgen, and stack sizes

If you find yourself running out of permgen space or your workstation is low on memory, adjust the JVM configuration as you would for any application. For example a common set of memory-related options is:

```
java -Xmx1536M -Xss1M -XX:+CMSClassUnloadingEnabled -XX:MaxPermSize=256m
```

Boot directory

`sbt-launch.jar` is just a bootstrap; the actual meat of sbt, and the Scala compiler and standard library, are downloaded to the shared directory `$HOME/.sbt/boot/`.

To change the location of this directory, set the `sbt.boot.directory` system property in your `sbt` script. A relative path will be resolved against the current working directory, which can be useful if you want to avoid sharing the boot directory between projects. For example, the following uses the pre-0.11 style of putting the boot directory in `project/boot/`:

```
java -Dsbt.boot.directory=project/boot/
```

HTTP/HTTPS/FTP Proxy

On Unix, sbt will pick up any HTTP, HTTPS, or FTP proxy settings from the standard `http_proxy`, `https_proxy`, and `ftp_proxy` environment variables. If you are behind a proxy requiring authentication, your `sbt` script must also pass flags to set the `http.proxyUser` and `http.proxyPassword` properties for HTTP, `ftp.proxyUser` and `ftp.proxyPassword` properties for FTP, or `https.proxyUser` and `https.proxyPassword` properties for HTTPS.

For example,

```
java -Dhttp.proxyUser=username -Dhttp.proxyPassword=mypassword
```

On Windows, your script should set properties for proxy host, port, and if applicable, username and password. For example, for HTTP:

```
java -Dhttp.proxyHost=myproxy -Dhttp.proxyPort=8080 -Dhttp.proxyUser=username -Dhttp.proxyPas
```

Replace `http` with `https` or `ftp` in the above command line to configure HTTPS or FTP.

Deploying to Sonatype

Deploying to sonatype is easy! Just follow these simple steps:

First - PGP Signatures

You'll need to PGP sign your artifacts for the Sonatype repository. Don't worry, there's a [plugin for that](#). Follow the instructions for the plugin and you'll have PGP signed artifacts in no time.

Note: The plugin is a jvm-only solution to generate PGP keys and sign artifacts. It can work with the GPG command line tool, but the command line is not needed.

If the command to generate your key fails execute the following commands and remove the displayed files:

```
> show */*:pgpSecretRing
[info] /home/username/.sbt/.gnupg/secring.gpg
> show */*:pgpPublicRing
[info] /home/username/.sbt/.gnupg/pubring.gpg
```

If your PGP key has not yet been distributed to the keyserver pool, i.e., you've just generated it, you'll need to publish it. You can do so using the [sbt-pgp](#) plugin:

```
pgp-cmd send-key keyname hkp://pool.sks-keyservers.net/
```

(where keyname is the name or email address used when creating the key or hexadecimal identifier for the key.)

If you see no output from `sbt-pgp` then the key name specified was not found.

If it fails to run the `SendKey` command you can try another server (for example: `hkp://keyserver.ubuntu.com`). A list of servers can be found at [the status page](#) of `sks-keyservers.net`.

Second - Maven Publishing Settings

To publish to a maven repository, you'll need to configure a few settings so that the correct metadata is generated.

```
publishMavenStyle := true
```

is used to ensure POMs are generated and pushed. Next, you have to set up the repositories you wish to push too. Luckily, Sonatype's OSSRH uses the same URLs for everyone:

```
publishTo := {  
  val nexus = "https://oss.sonatype.org/"  
  if (isSnapshot.value)  
    Some("snapshots" at nexus + "content/repositories/snapshots")  
  else  
    Some("releases" at nexus + "service/local/staging/deploy/maven2")  
}
```

Another good idea is to not publish your test artifacts (this is the default):

```
publishArtifact in Test := false
```

Third - POM Metadata

Now, we want to control what's available in the `pom.xml` file. This file describes our project in the maven repository and is used by indexing services for search and discover. This means it's important that `pom.xml` should have all information we wish to advertise as well as required info!

First, let's make sure no repositories show up in the POM file. To publish on maven-central, all *required* artifacts must also be hosted on maven central. However, sometimes we have optional dependencies for special features. If that's the case, let's remove the repositories for optional dependencies in our artifact:

```
pomIncludeRepository := { _ => false }
```

Next, the POM metadata that isn't generated by sbt must be added. This is done through the `pomExtra` configuration option:

```
pomExtra := (  
  <url>http://jsuereth.com/scala-arm</url>  
  <licenses>
```

```

<license>
  <name>BSD-style</name>
  <url>http://www.opensource.org/licenses/bsd-license.php</url>
  <distribution>repo</distribution>
</license>
</licenses>
<scm>
  <url>git@github.com:jsuereth/scala-arm.git</url>
  <connection>scm:git:git@github.com:jsuereth/scala-arm.git</connection>
</scm>
<developers>
  <developer>
    <id>jsuereth</id>
    <name>Josh Suereth</name>
    <url>http://jsuereth.com</url>
  </developer>
</developers>

```

Specifically, the `url`, `license`, `scm.url`, `scm.connection` and `developer` sections are required. The above is an example from the [scala-arm](#) project.

Note that sbt will automatically inject `licenses` and `url` nodes if they are already present in your build file. Thus an alternative to the above `pomExtra` is to include the following entries:

```

licenses := Seq("BSD-style" -> url("http://www.opensource.org/licenses/bsd-license.php"))

homepage := Some(url("http://jsuereth.com/scala-arm"))

```

This might be advantageous if those keys are used also by other plugins (e.g. `ls`). You **cannot use both** the sbt `licenses` key and the `licenses` section in `pomExtra` at the same time, as this will produce duplicate entries in the final POM file, leading to a rejection in Sonatype's staging process.

The full format of a `pom.xml` file is [outlined here](#).

Fourth - Adding credentials

The credentials for your Sonatype OSSRH account need to be added somewhere. Common convention is a `~/.sbt/0.13/sonatype.sbt` file with the following:

```

credentials += Credentials("Sonatype Nexus Repository Manager",
  "oss.sonatype.org",
  "<your username>",
  "<your password>")

```

Note: The first two strings must be "Sonatype Nexus Repository Manager" and "oss.sonatype.org" for Ivy to use the credentials.

Finally - Publish

In sbt, run `publishSigned` and you should see something like the following:

```
> publishSigned
Please enter your PGP passphrase> *****
[info] Packaging /home/josh/projects/typesafe/scala-arm/target/scala-2.9.1/scala-arm_2.9.1-1
[info] Wrote /home/josh/projects/typesafe/scala-arm/target/scala-2.9.1/scala-arm_2.9.1-1.2.p
[info] Packaging /home/josh/projects/typesafe/scala-arm/target/scala-2.9.1/scala-arm_2.9.1-1
[info] Packaging /home/josh/projects/typesafe/scala-arm/target/scala-2.9.1/scala-arm_2.9.1-1
[info] :: delivering :: com.jsuereth#scala-arm_2.9.1;1.2 :: 1.2 :: release :: Mon Jan 23 13:16:5
[info] Done packaging.
[info] Done packaging.
[info] Done packaging.
[info] delivering ivy file to /home/josh/projects/typesafe/scala-arm/target/scala-2.9.1/ivy-1
[info] published scala-arm_2.9.1 to https://oss.sonatype.org/service/local/staging/deploy/ma
[info] published scala-arm_2.9.1 to https://oss.sonatype.org/service/local/staging/deploy/ma
[info] published scala-arm_2.9.1 to https://oss.sonatype.org/service/local/staging/deploy/ma
[info] published scala-arm_2.9.1 to https://oss.sonatype.org/service/local/staging/deploy/ma
[info] published scala-arm_2.9.1 to https://oss.sonatype.org/service/local/staging/deploy/ma
[info] published scala-arm_2.9.1 to https://oss.sonatype.org/service/local/staging/deploy/ma
[info] published scala-arm_2.9.1 to https://oss.sonatype.org/service/local/staging/deploy/ma
[info] published scala-arm_2.9.1 to https://oss.sonatype.org/service/local/staging/deploy/ma
[info] published scala-arm_2.9.1 to https://oss.sonatype.org/service/local/staging/deploy/ma
[success] Total time: 9 s, completed Jan 23, 2012 1:17:03 PM
```

After publishing you have to follow the [Release workflow of nexus](#). `sbt-sonatype plugin` allows the release workflow procedures to be performed directly from sbt.

Note: Staged releases allow testing across large projects of independent releases before pushing the full project.

Note: An error message of `PGPException: checksum mismatch at 0 of 20` indicates that you got the passphrase wrong. We have found at least on OS X that there may be issues with characters outside the 7-bit ASCII range (e.g. Umlauts). If you are absolutely sure that you typed the right phrase and the error doesn't disappear, try changing the passphrase.

Summary

To get your project hosted on Sonatype (and Maven Central), you will need to:

- Have GPG key pair, with published public key,
- An sbt file with your Sonatype credentials *that is not pushed to the VCS*,
- Add the [sbt-pgp plugin](#) to sign the artefacts,
- Modify `build.sbt` with the required elements in the generated POM.

Starting with a project that is not being published, you'll need to install GPG, generate and publish your key. Switching to sbt, you'll then need to:

`~/.sbt/sonatype.sbt` This file (kept *outside the VCS*) contains the Sonatype credentials settings:

```
credentials += Credentials("Sonatype Nexus Repository Manager",
                           "oss.sonatype.org",
                           "your-sonatype-username",
                           "your-sonatype-password")
```

`~/.sbt/plugins/gpg.sbt` The [sbt-pgp plugin](#) allows you to sign and publish your artefacts by running `publishSigned` in sbt:

```
addSbtPlugin("com.typesafe.sbt" % "sbt-pgp" % "0.8")
```

`build.sbt` Finally, you'll need to tweak the generated POM in your `build.sbt`. The tweaks include specifying the project's authors, URL, SCM and many others:

```
publishTo := {
  val nexus = "https://oss.sonatype.org/"
  if (isSnapshot.value)
    Some("snapshots" at nexus + "content/repositories/snapshots")
  else
    Some("releases" at nexus + "service/local/staging/deploy/maven2")
}

publishMavenStyle := true

publishArtifact in Test := false

pomIncludeRepository := { _ => false }

pomExtra := (
  <url>http://your.project.url</url>
  <licenses>
```

```

    <license>
      <name>BSD-style</name>
      <url>http://www.opensource.org/licenses/bsd-license.php</url>
      <distribution>repo</distribution>
    </license>
  </licenses>
  <scm>
    <url>git@github.com:your-account/your-project.git</url>
    <connection>scm:git:git@github.com:your-account/your-project.git</connection>
  </scm>
  <developers>
    <developer>
      <id>you</id>
      <name>Your Name</name>
      <url>http://your.url</url>
    </developer>
  </developers>
)

```

Changes

0.13.5-RC1 to 0.13.5-RC2

- Fixes auto plugins not detecting object `autoImport`. (gh-1314)

0.13.2 to 0.13.5

- The Scala version for sbt and sbt plugins is now 2.10.4. This is a compatible version bump.
- Added a new setting `testResultLogger` to allow customisation of logging of test results. (gh-1225)
- When `test` is run and there are no tests available, omit logging output. Especially useful for aggregate modules. `test-only` et al unaffected. (gh-1185)
- sbt now uses minor-patch version of ivy 2.4 (`org.scala-sbt.ivy:ivy:2.4.0-sbt-`)
- `sbt.Plugin` deprecated in favor of `sbt.AutoPlugin`
- name-hashing incremental compiler now supports scala macros.
- `testResultLogger` is now configured.
- sbt-server hooks for task cancellation.
- Add `JUnitXmlReportPlugin` which generates junit-xml-reports for all tests.

0.13.1 to 0.13.2

- Adding new name-hashing feature to incremental compiler. Alters how scala dependencies are tracked, reducing number of recompiles necessary.
- Added the ability to launch servers via the sbt-launcher.
- Added `.previous` feature on tasks which can load the pervious value.
- Added an `all` command which can run more than tasks in parallel.
- Exposed the 'overwrite' flags from ivy. Added warning if overwriting a release version.
- Improve the error message when credentials are not found in Ivy.
- Improve task macros to handle more scala constructs.
- Fix `last` and `export` tasks to read from the correct stream.
- Fix issue where ivy's `.+` dependency ranges were not correctly translated to maven.
- Override security manager to ignore file permissions (performance issue)
- 2.11 compatibility fixes
- Launcher can now handle ivy's `.+` revisions.
- SessionSettings now correctly overwrite existing settings.
- Adding a simple Logic system for inclusionary/dependency logic of plugins.
- Improve build hooks for `LoggerReporter` and `TaskProgress`.
- Serialize incremental compiler analysis into text-file format.
- Issue a warning when generating Paths and separate already exists in the path.
- Migrate to Ivy 2.3.0-final.
- Docs: Use bintray as default repository host
- Docs: improved docs on test groups.
- Docs: updated documentation on the Launcher.
- Docs: started architecture document.

0.13.0 to 0.13.1

- The Scala version for sbt and sbt plugins is now 2.10.3. This is a compatible version bump.
- New method `toTask` on `Initialize[InputTask[T]]` to apply the full input and get a plain task out.
- Improved performance of `inspect tree`
- Work around various issues with Maven local repositories, including resolving `-SNAPSHOTs` from them. (gh-321)
- Better representation of no cross-version suffix in suffix conflict error message: now shows `<none>` instead of just `_`
- `TrapExit` support for multiple, concurrent managed applications. Now enabled by default for all `run`-like tasks. (gh-831)

- Add minimal support for class file formats 51.0, 52.0 in incremental compiler. (gh-842)
- Allow main class to be non-public. (gh-883)
- Convert `-classpath` to `CLASSPATH` when forking on Windows and length exceeds a heuristic maximum. (gh-755)
- `scalacOptions` for `.scala` build definitions are now also used for `.sbt` files
- `error`, `warn`, `info`, `debug` commands to set log level and `--error`, ... to set the level before the project is loaded. (gh-806)
- `sLog` settings that provides a `Logger` for use by settings. (gh-806)
- Early commands: any command prefixed with `--` gets moved before other commands on startup and doesn't force sbt into batch mode.
- Deprecate internal `-`, `--`, and `---` commands in favor of `onFailure`, `sbtClearOnFailure`, and `resumeFromFailure`.
- `makePom` no longer generates `<type>` elements for standard classifiers. (gh-728)
- Fix many instances of the Turkish `i` bug.
- Read `https+ftp` proxy environment variables into system properties where Java will use them. (gh-886)
- The `Process` methods that are redirection-like no longer discard the exit code of the input. This addresses an inconsistency with `Fork`, where using the `CustomOutput OutputStrategy` makes the exit code always zero.
- Recover from failed `reload` command in the scripted sbt handler.
- Parse external `pom.xml` with `CustomPomParser` to handle multiple definitions. (gh-758)
- Improve key collision error message (gh-877)
- Display the source position of an undefined setting.
- Respect the `-nowarn` option when compiling Scala sources.
- Improve forked test debugging by listing tests run by sbt in debug output. (gh-868)
- Fix scaladoc cache to track changes to `-doc-root-content` (gh-837)
- Incremental compiler: Internal refactoring in preparation for name-hashing (gh-936)
- Incremental compiler: improved cache loading/saving speed by internal file names (gh-931)
- Docs: many contributed miscellaneous fixes and additions
- Docs: link to page source now at the bottom of the page
- Docs: sitemap now automatically generated
- Docs: custom role enables links from a key name in the docs to the val in [Keys](#)
- Docs: restore `sxr` support and fix links to `sxr`'d sources. (gh-863)

0.12.4 to 0.13.0

The changes for 0.13.0 are listed on a separate page. See [sbt 0.13.0 changes](#).

sbt 0.13.0 Changes

Overview

Features, fixes, changes with compatibility implications

- Moved to Scala 2.10 for sbt and build definitions.
- Support for plugin configuration in `project/plugins/` has been removed. It was deprecated since 0.11.2.
- Dropped support for tab completing the right side of a setting for the `set` command. The new task macros make this tab completion obsolete.
- The convention for keys is now camelCase only. Details below.
- Fixed the default classifier for tests to be `tests` for proper Maven compatibility.
- The global settings and plugins directories are now versioned. Global settings go in `~/.sbt/0.13/` and global plugins in `~/.sbt/0.13/plugins/` by default. Explicit overrides, such as via the `sbt.global.base` system property, are still respected. (gh-735)
- sbt no longer canonicalizes files passed to scalac. (gh-723)
- sbt now enforces that each project must have a unique `target` directory.
- sbt no longer overrides the Scala version in dependencies. This allows independent configurations to depend on different Scala versions and treats Scala dependencies other than `scala-library` as normal dependencies. However, it can result in resolved versions other than `scalaVersion` for those other Scala libraries.
- JLine is now configured differently for Cygwin. See [Installing sbt](#).
- Jline and Ansi codes work better on Windows now. CI servers might have to explicitly disable Ansi codes via `-Dsbt.log.format=false`.
- JLine now tries to respect `~/.inputrc`.
- Forked tests and runs now use the project's base directory as the current working directory.
- `compileInputs` is now defined in `(Compile,compile)` instead of just `Compile`
- The result of running tests is now [Tests.Output](#).

Features

- Use the repositories in `boot.properties` as the default project resolvers. Add `bootOnly` to a repository in `boot.properties` to specify that it should not be used by projects by default. (Josh S., gh-608)

- Support vals and defs in .sbt files. Details below.
- Support defining Projects in .sbt files: vals of type Project are added to the Build. Details below.
- New syntax for settings, tasks, and input tasks. Details below.
- Automatically link to external API scaladocs of dependencies by setting `autoAPIMappings := true`. This requires at least Scala 2.10.1 and for dependencies to define `apiURL` for their scaladoc location. Mappings may be manually added to the `apiMappings` task as well.
- Support setting Scala home directory temporary using the switch command: `++ scala-version=/path/to/scala/home`. The `scala-version` part is optional, but is used as the version for any managed dependencies.
- Add `publishM2` task for publishing to `~/.m2/repository`. (gh-485)
- Use a default root project aggregating all projects if no root is defined. (gh-697)
- New API for getting tasks and settings from multiple projects and configurations. See the new section [getting values from multiple scopes](#).
- Enhanced test interface for better support of test framework features. (Details pending.)
- `export` command
 - For tasks, prints the contents of the ‘export’ stream. By convention, this should be the equivalent command line(s) representation. `compile`, `doc`, and `console` show the approximate command lines for their execution. Classpath tasks print the classpath string suitable for passing as an option.
 - For settings, directly prints the value of a setting instead of going through the logger

Fixes

- sbt no longer tries to warn on dependency conflicts. Configure a [conflict manager](#) instead. (gh-709)
- Run test Cleanup and Setup when forking. The test ClassLoader is not available because it is in another jvm.

Improvements

- Run the API extraction phase after the compiler's `pickler` phase instead of `typer` to allow compiler plugins after `typer`. (Adriaan M., gh-609)
- Record defining source position of settings. `inspect` shows the definition location of all settings contributing to a defined value.
- Allow the root project to be specified explicitly in `Build.rootProject`.
- Tasks that need a directory for storing cache information can now use the `cacheDirectory` method on `streams`. This supersedes the `cacheDirectory` setting.
- The environment variables used when forking `run` and `test` may be set via `envVars`, which is a `Task[Map[String,String]]`. (gh-665)
- Restore class files after an unsuccessful compilation. This is useful when an error occurs in a later incremental step that requires a fix in the originally changed files.
- Better auto-generated IDs for default projects. (gh-554)
- Fork run directly with 'java' to avoid additional class loader from 'scala' command. (gh-702)
- Make `autoCompilerPlugins` support compiler plugins defined in a internal dependency (only if `exportJars := true` due to scalac limitations)
- Track ancestors of non-private templates and use this information to require fewer, smaller intermediate incremental compilation steps.
- `autoCompilerPlugins` now supports compiler plugins defined in a internal dependency. The plugin project must define `exportJars := true`. Depend on the plugin with `...dependsOn(... % Configurations.CompilerPlugin)`.
- Add utilities for debugging API representation extracted by the incremental compiler. (Grzegorz K., gh-677, gh-793)
- `consoleProject` unifies the syntax for getting the value of a setting and executing a task. See [Console Project](#).

Other

- The source layout for the sbt project itself follows the package name to accommodate to Eclipse users. (Grzegorz K., gh-613)

Details of major changes

camelCase Key names The convention for key names is now camelCase only instead of camelCase for Scala identifiers and hyphenated, lower-case on the command line. camelCase is accepted for existing hyphenated key names and the hyphenated form will still be accepted on the command line for those existing tasks and settings declared with hyphenated names. Only camelCase will be shown for tab completion, however.

New key definition methods There are new methods that help avoid duplicating key names by declaring keys as:

```
val myTask = taskKey[Int]("A (required) description of myTask.")
```

The name will be picked up from the val identifier by the implementation of the taskKey macro so there is no reflection needed or runtime overhead. Note that a description is mandatory and the method taskKey begins with a lowercase t. Similar methods exist for keys for settings and input tasks: settingKey and inputKey.

New task/setting syntax First, the old syntax is still supported with the intention of allowing conversion to the new syntax at your leisure. There may be some incompatibilities and some may be unavoidable, but please report any issues you have with an existing build.

The new syntax is implemented by making :=, +=, and ++= macros and making these the only required assignment methods. To refer to the value of other settings or tasks, use the value method on settings and tasks. This method is a stub that is removed at compile time by the macro, which will translate the implementation of the task/setting to the old syntax.

For example, the following declares a dependency on scala-reflect using the value of the scalaVersion setting:

```
libraryDependencies += "org.scala-lang" % "scala-reflect" % scalaVersion.value
```

The value method is only allowed within a call to :=, +=, or ++=. To construct a setting or task outside of these methods, use Def.task or Def.setting. For example,

```
val reflectDep = Def.setting { "org.scala-lang" % "scala-reflect" % scalaVersion.value }  
  
libraryDependencies += reflectDep.value
```

A similar method parsed is defined on Parser[T], Initialize[Parser[T]] (a setting that provides a parser), and Initialize[State => Parser[T]] (a setting that uses the current State to provide a Parser[T]). This method can be used when defining an input task to get the result of user input.

```
myInputTask := {  
  // Define the parser, which is the standard space-delimited arguments parser.  
  val args = Def.spaceDelimited("<args>").parsed  
  // Demonstrates using a setting value and a task result:
```



```

println("Project name: " + name.value)
println("Classpath: " + (fullClasspath in Compile).value.map(_.file))
println("Arguments:")
for(arg <- args) println("  " + arg)
}

```

For details, see [Input Tasks](#).

To expect a task to fail and get the failing exception, use the `failure` method instead of `value`. This provides an `Incomplete` value, which wraps the exception. To get the result of a task whether or not it succeeds, use `result`, which provides a `Result[T]`.

Dynamic settings and tasks (`flatMap`) have been cleaned up. Use the `Def.taskDyn` and `Def.settingDyn` methods to define them (better name suggestions welcome). These methods expect the result to be a task and setting, respectively.

.sbt format enhancements `vals` and `defs` are now allowed in `.sbt` files. They must follow the same rules as settings concerning blank lines, although multiple definitions may be grouped together. For example,

```

val n = "widgets"
val o = "org.example"

name := n

organization := o

```

All definitions are compiled before settings, but it will probably be best practice to put definitions together. Currently, the visibility of definitions is restricted to the `.sbt` file it is defined in. They are not visible in `consoleProject` or the `set` command at this time, either. Use Scala files in `project/` for visibility in all `.sbt` files.

`vals` of type `Project` are added to the `Build` so that multi-project builds can be defined entirely in `.sbt` files now. For example,

```

lazy val a = Project("a", file("a")).dependsOn(b)

lazy val b = Project("b", file("sub")).settings(
  version := "1.0"
)

```

Currently, it only makes sense to define these in the root project's `.sbt` files.

A shorthand for defining Projects is provided by a new macro called `project`. This requires the constructed Project to be directly assigned to a `val`. The name of this `val` is used for the project ID and base directory. The base directory can be changed with the `in` method. The previous example can also be written as:

```
lazy val a = project.dependsOn(b)

lazy val b = project in file("sub") settings(
  version := "1.0"
)
```

This macro is also available for use in Scala files.

Control over automatically added settings sbt loads settings from a few places in addition to the settings explicitly defined by the `Project.settings` field. These include plugins, global settings, and `.sbt` files. The new `Project.autoSettings` method configures these sources: whether to include them for the project and in what order.

`Project.autoSettings` accepts a sequence of values of type `AddSettings`. Instances of `AddSettings` are constructed from methods in the `AddSettings` companion object. The configurable settings are per-user settings (from `~/.sbt`, for example), settings from `.sbt` files, and plugin settings (project-level only). The order in which these instances are provided to `autoSettings` determines the order in which they are appended to the settings explicitly provided in `Project.settings`.

For `.sbt` files, `AddSettings.defaultSbtFiles` adds the settings from all `.sbt` files in the project's base directory as usual. The alternative method `AddSettings.sbtFiles` accepts a sequence of `Files` that will be loaded according to the standard `.sbt` format. Relative files are resolved against the project's base directory.

Plugin settings may be included on a per-Plugin basis by using the `AddSettings.plugins` method and passing a `Plugin => Boolean`. The settings controlled here are only the automatic per-project settings. Per-build and global settings will always be included. Settings that plugins require to be manually added still need to be added manually.

For example,

```
import AddSettings._

lazy val root = Project("root", file(".")) autoSettings(
  userSettings, allPlugins, sbtFiles(file("explicit/a.txt"))
)
```

```

lazy val sub = Project("sub", file("Sub")) autoSettings(
  defaultSbtFiles, plugins(includePlugin)
)

def includePlugin(p: Plugin): Boolean =
  p.getClass.getName.startsWith("org.example.")

```

Resolving Scala dependencies Scala dependencies (like `scala-library` and `scala-compiler`) are now resolved via the normal `update` task. This means:

1. Scala jars won't be copied to the boot directory, except for those needed to run sbt.
2. Scala SNAPSHOTs behave like normal SNAPSHOTs. In particular, running `update` will properly re-resolve the dynamic revision.
3. Scala jars are resolved using the same repositories and configuration as other dependencies.
4. Scala dependencies are not resolved via `update` when `scalaHome` is set, but are instead obtained from the configured directory.
5. The Scala version for sbt will still be resolved via the repositories configured for the launcher.

sbt still needs access to the compiler and its dependencies in order to run `compile`, `console`, and other Scala-based tasks. So, the Scala compiler jar and dependencies (like `scala-reflect.jar` and `scala-library.jar`) are defined and resolved in the `scala-tool` configuration (unless `scalaHome` is defined). By default, this configuration and the dependencies in it are automatically added by sbt. This occurs even when dependencies are configured in a `pom.xml` or `ivy.xml` and so it means that the version of Scala defined for your project must be resolvable by the resolvers configured for your project.

If you need to manually configure where sbt gets the Scala compiler and library used for compilation, the REPL, and other Scala tasks, do one of the following:

1. Set `scalaHome` to use the existing Scala jars in a specific directory. If `autoScalaLibrary` is true, the library jar found here will be added to the (unmanaged) classpath.
2. Set `managedScalaInstance := false` and explicitly define `scalaInstance`, which is of type `ScalaInstance`. This defines the compiler, library, and other jars comprising Scala. If `autoScalaLibrary` is true, the library jar from the defined `ScalaInstance` will be added to the (unmanaged) classpath.

The [Configuring Scala](#) page provides full details.

sbt 0.12.0 Changes

Features, fixes, changes with compatibility implications

- The cross versioning convention has changed for Scala versions 2.10 and later as well as for sbt plugins.
- When invoked directly, ‘update’ will always perform an update (gh-335)
- The sbt plugins repository is added by default for plugins and plugin definitions. gh-380
- Plugin configuration directory precedence has changed (see details section below)
- Source dependencies have been fixed, but the fix required changes (see details section below)
- Aggregation has changed to be more flexible (see details section below)
- Task axis syntax has changed from key(for task) to task::key (see details section below)
- The organization for sbt has to changed to `org.scala-sbt` (was: `org.scala-tools.sbt`). This affects users of the scripted plugin in particular.
- `artifactName` type has changed to `(ScalaVersion, Artifact, ModuleID) => String`
- `javacOptions` is now a task
- `session save` overwrites settings in `build.sbt` (when appropriate). gh-369
- `scala-library.jar` is now required to be on the classpath in order to compile Scala code. See the `scala-library.jar` section at the bottom of the page for details.

Features

- Support for forking tests (gh-415)
- `test-quick` (see details section below)
- Support globally overriding repositories (gh-472)
- Added `print-warnings` task that will print unchecked and deprecation warnings from the previous compilation without needing to recompile (Scala 2.10+ only)
- Support for loading an ivy settings file from a URL.
- `projects add/remove <URI>` for temporarily working with other builds
- Enhanced control over parallel execution (see details section below)
- `inspect tree <key>` for calling `inspect` command recursively (gh-274)

Fixes

- Delete a symlink and not its contents when recursively deleting a directory.

- Fix detection of ancestors for java sources
- Fix the resolvers used for `update-sbt-classifiers` (gh-304)
- Fix auto-imports of plugins (gh-412)
- Argument quoting (see details section below)
- Properly reset JLine after being stopped by Ctrl+z (unix only). gh-394

Improvements

- The launcher can launch all released sbt versions back to 0.7.0.
- A more refined hint to run ‘last’ is given when a stack trace is suppressed.
- Use java 7 Redirect.INHERIT to inherit input stream of subprocess (gh-462,gh-327). This should fix issues when forking interactive programs. (@vigdorchik)
- Mirror ivy ‘force’ attribute (gh-361)
- Various improvements to `help` and `tasks` commands as well as new settings command (gh-315)
- Bump jsch version to 0.1.46. (gh-403)
- Improved help commands: `help`, `tasks`, `settings`.
- Bump to JLine 1.0 (see details section below)
- Global repository setting (see details section below)
- Other fixes/improvements: gh-368, gh-377, gh-378, gh-386, gh-387, gh-388, gh-389

Experimental or In-progress

- API for embedding incremental compilation. This interface is subject to change, but already being used in [a branch of the scala-maven-plugin](#).
- Experimental support for keeping the Scala compiler resident. Enable by passing `-Dsbt.resident.limit=n` to sbt, where n is an integer indicating the maximum number of compilers to keep around.
- The [Howto](#) pages on the [new site](#) are at least readable now. There is more content to write and more formatting improvements are needed, so [pull requests are welcome](#).

Details of major changes from 0.11.2 to 0.12.0

Plugin configuration directory In 0.11.0, plugin configuration moved from `project/plugins/` to just `project/`, with `project/plugins/` being deprecated. Only 0.11.2 had a deprecation message, but in all of 0.11.x, the presence of the old style `project/plugins/` directory took precedence over the new style. In 0.12.0, the new style takes precedence. Support for the old style won’t be removed until 0.13.0.

1. Ideally, a project should ensure there is never a conflict. Both styles are still supported; only the behavior when there is a conflict has changed.
2. In practice, switching from an older branch of a project to a new branch would often leave an empty `project/plugins/` directory that would cause the old style to be used, despite there being no configuration there.
3. Therefore, the intention is that this change is strictly an improvement for projects transitioning to the new style and isn't noticed by other projects.

Parsing task axis There is an important change related to parsing the task axis for settings and tasks that fixes gh-202

1. The syntax before 0.12 has been `{build}project/config:key(for task)`
2. The proposed (and implemented) change for 0.12 is `{build}project/config:task::key`
3. By moving the task axis before the key, it allows for easier discovery (via tab completion) of keys in plugins.
4. It is not planned to support the old syntax.

Aggregation Aggregation has been made more flexible. This is along the direction that has been previously discussed on the mailing list.

1. Before 0.12, a setting was parsed according to the current project and only the exact setting parsed was aggregated.
2. Also, tab completion did not account for aggregation.
3. This meant that if the setting/task didn't exist on the current project, parsing failed even if an aggregated project contained the setting/task.
4. Additionally, if `compile:package` existed for the current project, `*:package` existed for an aggregated project, and the user requested 'package' to run (without specifying the configuration), `*:package` wouldn't be run on the aggregated project (because it isn't the same as the `compile:package` key that existed on the current project).
5. In 0.12, both of these situations result in the aggregated settings being selected. For example,
 1. Consider a project `root` that aggregates a subproject `sub`.
 2. `root` defines `*:package`.
 3. `sub` defines `compile:package` and `compile:compile`.
 4. Running `root/package` will run `root/*:package` and `sub/compile:package`
 5. Running `root/compile` will run `sub/compile:compile`
6. This change was made possible in part by the change to task axis parsing.

Parallel Execution Fine control over parallel execution is supported as described here: [Parallel Execution](#).

1. The default behavior should be the same as before, including the `parallelExecution` settings.
2. The new capabilities of the system should otherwise be considered experimental.
3. Therefore, `parallelExecution` won't be deprecated at this time.

Source dependencies A fix for issue gh-329 is included in 0.12.0. This fix ensures that only one version of a plugin is loaded across all projects. There are two parts to this.

1. The version of a plugin is fixed by the first build to load it. In particular, the plugin version used in the root build (the one in which sbt is started in) always overrides the version used in dependencies.
2. Plugins from all builds are loaded in the same class loader.

Additionally, Sanjin's patches to add support for hg and svn URIs are included.

1. sbt uses subversion to retrieve URIs beginning with `svn` or `svn+ssh`. An optional fragment identifies a specific revision to checkout.
2. Because a URI for mercurial doesn't have a mercurial-specific scheme, sbt requires the URI to be prefixed with `hg:` to identify it as a mercurial repository.
3. Also, URIs that end with `.git` are now handled properly.

Cross building The cross version suffix is shortened to only include the major and minor version for Scala versions starting with the 2.10 series and for sbt versions starting with the 0.12 series. For example, `sbinary_2.10` for a normal library or `sbt-plugin_2.10_0.12` for an sbt plugin. This requires forward and backward binary compatibility across incremental releases for both Scala and sbt.

1. This change has been a long time coming, but it requires everyone publishing an open source project to switch to 0.12 to publish for 2.10 or adjust the cross versioned prefix in their builds appropriately.
2. Obviously, using 0.12 to publish a library for 2.10 requires 0.12.0 to be released before projects publish for 2.10.
3. There is now the concept of a binary version. This is a subset of the full version string that represents binary compatibility. That is, equal binary versions implies binary compatibility. All Scala versions prior to 2.10 use the full version for the binary version to reflect previous sbt behavior. For 2.10 and later, the binary version is `<major>.<minor>`.

4. The cross version behavior for published artifacts is configured by the `crossVersion` setting. It can be configured for dependencies by using the `cross` method on `ModuleID` or by the traditional `%%` dependency construction variant. By default, a dependency has cross versioning disabled when constructed with a single `%` and uses the binary Scala version when constructed with `%%`.
5. The `artifactName` function now accepts a type `ScalaVersion` as its first argument instead of a `String`. The full type is now `(ScalaVersion, ModuleID, Artifact) => String`. `ScalaVersion` contains both the full Scala version (such as 2.10.0) as well as the binary Scala version (such as 2.10).
6. The flexible version mapping added by Indrajit has been merged into the `cross` method and the `%%` variants accepting more than one argument have been deprecated. See [Cross Build](#) for details.

Global repository setting Define the repositories to use by putting a standalone `[repositories]` section (see the [sbt Launcher](#) page) in `~/.sbt/repositories` and pass `-Dsbt.override.build.repos=true` to `sbt`. Only the repositories in that file will be used by the launcher for retrieving `sbt` and `Scala` and by `sbt` when retrieving project dependencies. (@jsuereth)

test-quick `test-quick` (gh-393) runs the tests specified as arguments (or all tests if no arguments are given) that:

1. have not been run yet OR
2. failed the last time they were run OR
3. had any transitive dependencies recompiled since the last successful run

Argument quoting Argument quoting (gh-396) from the interactive mode works like Scala string literals.

1. `> command "arg with spaces,\n escapes interpreted"`
2. `> command ""arg with spaces,\n escapes not interpreted""`
3. For the first variant, note that paths on Windows use backslashes and need to be escaped (`\`). Alternatively, use the second variant, which does not interpret escapes.
4. For using either variant in batch mode, note that a shell will generally require the double quotes themselves to be escaped.

scala-library.jar

`sbt` versions prior to 0.12.0 provided the location of `scala-library.jar` to `scalac` even if `scala-library.jar` wasn't on the classpath. This allowed compiling Scala

code without `scala-library` as a dependency, for example, but this was a misfeature. Instead, the Scala library should be declared as **provided**:

```
// Don't automatically add the scala-library dependency
// in the 'compile' configuration
autoScalaLibrary := false

libraryDependencies += "org.scala-lang" % "scala-library" % "2.9.2" % "provided"
```

Older Changes

0.12.3 to 0.12.4

- Work around URI problems with encoding and resolving. (gh-725)
- Allow `-cp` argument to **apply** command to be quoted. (gh-724)
- Make **sbtBinaryVersion** use the new approach for 0.13 and later to support cross-building plugins.
- Pull **sbtDependency** version from **sbtVersion** to facilitate cross-building plugins.
- Proper support for stashing on-failure handlers. (gh-732)
- Include files with zip extension in unmanaged jars. (gh-750)
- Only add automatically detected plugins to options once. (gh-757)
- Properly handle failure in a multi-command that includes **reload**. (gh-732)
- Fix unsynchronized caching of Scala class loaders that could result in Scala classes being loaded in multiple class loaders.
- Incremental compiler: remove resident compiler code (wasn't used and was a compatibility liability)
- Incremental compiler: properly track **abstract override** modifier. (gh-726)
- Incremental compiler: do not normalize types in the api extraction phase. (gh-736)
- Ivy cache: account for **localOnly** when cache subclass overrides **isChanging**
- Ivy cache: fix corruption when developing sbt or sbt plugins. (gh-768)
- Ivy cache: invalidate when artifact download fails to avoid locking into bad resolver. (gh-760)
- Ivy cache: use publication date from metadata instead of original file's last modified time when deleting out of date artifacts. (gh-764)

0.12.2 to 0.12.3

- Allow **cleanKeepFiles** to contain directories

- Disable Ivy debug-level logging for performance. (gh-635)
- Invalidate artifacts not recorded in the original metadata when a module marked as changing changes. (gh-637, gh-641)
- Ivy Artifact needs wildcard configuration added if no explicit ones are defined. (gh-439)
- Right precedence of sbt.boot.properties lookup, handle qualifier correctly. (gh-651)
- Mark the tests failed exception as having already provided feedback.
- Handle exceptions not caught by the test framework when forking. (gh-653)
- Support `reload plugins` after ignoring a failure to load a project.
- Workaround for os deadlock detection at the process level. (gh-650)
- Fix for dependency on class file corresponding to a package. (Grzegorz K., gh-620)
- Fix incremental compilation problem with package objects inheriting from invalidated sources in a subpackage.
- Use Ivy's default name for the resolution report so that links to other configurations work.
- Include jars from java.ext.dirs in incremental classpath. (gh-678)
- Multi-line prompt text offset issue (Jibbers42, gh-625)
- Added `xml:space="preserve"` attribute to extraDependencyAttributes XML Block for publishing poms for plugins dependent on other plugins (Brendan M., gh-645)
- Tag the actual test task and not a later task. (gh-692)
- Make exclude-classifiers per-user instead of per-build. (gh-634)
- Load global plugins in their own class loader and replace the base loader with that. (gh-272)
- Demote the default conflict warnings to the debug level. These will be removed completely in 0.13. (gh-709)
- Fix Ivy cache issues when multiple resolvers are involved. (gh-704)

0.12.1 to 0.12.2

- Support `-Yrange`. (Lex S., gh-607)
- Only make one call to test frameworks per test name. (gh-520)
- Add `-cp` option to the `apply` method to make adding commands from an external program easier.
- Stable representation of refinement typerefs. This fixes unnecessary re-compilations in some cases. (Adriaan M., gh-610)
- Disable aggregation for `run-main`. (gh-606)
- Concurrent restrictions: Untagged should be set based on the task's tags, not the tags of all tasks.
- When preserving the last modified time of files, convert negative values to

0

- Use `java.lang.Throwable.setStackTrace` when sending exceptions back from forked tests. (Eugene V., gh-543)
- Don't merge dependencies with mismatched transitive/force/changing values. (gh-582)
- Filter out null parent files when deleting empty directories. (Eugene V., gh-589)
- Work around File constructor not accepting URIs for UNC paths. (gh-564)
- Split ForkTests react() out to workaround SI-6526 (avoids a stackoverflow in some forked test situations)
- Maven-style ivy repo support in the launcher config (Eric B., gh-585)
- Compare external binaries with canonical files (nau, gh-584)
- Call System.exit after the main thread is finished. (Eugene V., gh-565)
- Abort running tests on the first failure to communicate results back to the main process. (Eugene V., gh-557)
- Don't let the right side of the alias command fail the parse. (gh-572)
- API extraction: handle any type that is annotated, not just the spec'd simple type. (gh-559)
- Don't try to look up the class file for a package. (gh-620)

0.12.0 to 0.12.1

Dependency management fixes:

- Merge multiple dependency definitions for the same ID. Workaround for gh-468, gh-285, gh-419, gh-480.
- Don't write section of pom if scope is 'compile'.
- Ability to properly match on artifact type. Fixes gh-507 (Thomas).
- Force **update** to run on changes to last modified time of artifacts or cached descriptor (part of fix for gh-532). It may also fix issues when working with multiple local projects via 'publish-local' and binary dependencies.
- Per-project resolution cache that deletes cached files before update. Notes:
 - The resolution cache differs from the repository cache and does not contain dependency metadata or artifacts.
 - The resolution cache contains the generated ivy files, properties, and resolve reports for the project.
 - There will no longer be individual files directly in `~/.ivy2/cache/`
 - Resolve reports are now in `target/resolution-cache/reports/`, viewable with a browser.
 - Cache location includes extra attributes so that cross builds of a plugin do not overwrite each other. Fixes gh-532.

Three stage incremental compilation:

- As before, the first step recompiles sources that were edited (or otherwise directly invalidated).
- The second step recompiles sources from the first step whose API has changed, their direct dependencies, and sources forming a cycle with these sources.
- The third step recompiles transitive dependencies of sources from the second step whose API changed.
- Code relying mainly on composition should see decreased compilation times with this approach.
- Code with deep inheritance hierarchies and large cycles between sources may take longer to compile.
- **last compile** will show cycles that were processed in step 2. Reducing large cycles of sources shown here may decrease compile times.

Miscellaneous fixes and improvements:

- Various test forking fixes. Fixes gh-512, gh-515.
- Proper isolation of build definition classes. Fixes gh-536, gh-511.
- **orbit** packaging should be handled like a standard jar. Fixes gh-499.
- In **I0.copyFile**, limit maximum size transferred via NIO. Fixes gh-491.
- Add OSX JNI library extension in **includeFilter** by default. Fixes gh-500. (Indrajit)
- Translate **show x y** into **;show x ;show y** . Fixes gh-495.
- Clean up temporary directory on exit. Fixes gh-502.
- **set** prints the scopes+keys it defines and affects.
- Tab completion for **set** (experimental).
- Report file name when an error occurs while opening a corrupt zip file in incremental compilation code. (James)
- Defer opening logging output files until an actual write. Helps reduce number of open file descriptors.
- Back all console loggers by a common console interface that merges (overwrites) consecutive Resolving xxxx ... lines when ansi codes are enabled (as first done by Play).

Forward-compatible-only change (not present in 0.12.0):

- **sourcesInBase** setting controls whether sources in base directory are included. Fixes gh-494.

0.11.3 to 0.12.0

The changes for 0.12.0 are listed on a separate page. See [sbt 0.12.0 changes](#).

0.11.2 to 0.11.3

Dropping scala-tools.org:

- The sbt group ID is changed to `org.scala-sbt` (from `org.scala-tools.sbt`). This means you must use a 0.11.3 launcher to launch 0.11.3.
- The convenience objects `ScalaToolsReleases` and `ScalaToolsSnapshots` now point to <https://oss.sonatype.org/content/repositories/releases> and [.../snapshots](https://oss.sonatype.org/content/repositories/snapshots)
- The launcher no longer includes `scala-tools.org` repositories by default and instead uses the Sonatype OSS snapshots repository for Scala snapshots.
- The `scala-tools.org` releases repository is no longer included as an application repository by default. The Sonatype OSS repository is *not* included by default in its place.

Other fixes:

- Compiler interface works with 2.10
- `maxErrors` setting is no longer ignored
- Correct test count. gh-372 (Eugene)
- Fix file descriptor leak in process library (Daniel)
- Buffer url input stream returned by Using. gh-437
- Jsch version bumped to 0.1.46. gh-403
- JUnit test detection handles ancestors properly (Indrajit)
- Avoid unnecessarily re-resolving plugins. gh-368
- Substitute variables in explicit version strings and custom repository definitions in launcher configuration
- Support setting `sbt.version` from system property, which overrides setting in a properties file. gh-354
- Minor improvements to command/key suggestions

0.11.1 to 0.11.2

Notable behavior change:

- The local Maven repository has been removed from the launcher's list of default repositories, which is used for obtaining sbt and Scala dependencies. This is motivated by the high probability that including this repository was causing the various problems some users have with the launcher not finding some dependencies (gh-217).

Fixes:

- gh-257 Fix invalid classifiers in pom generation (Indrajit)
- gh-255 Fix scripted plugin descriptor (Artyom)
- Fix forking git on windows (Stefan, Josh)
- gh-261 Fix whitespace handling for semicolon-separated commands
- gh-263 Fix handling of dependencies with an explicit URL
- gh-272 Show deprecation message for `project/plugins/`

0.11.0 to 0.11.1

Breaking change:

- The scripted plugin is now in the `sbt` package so that it can be used from a named package

Notable behavior change:

- By default, there is more logging during update: one line per dependency resolved and two lines per dependency downloaded. This is to address the appearance that sbt hangs on larger 'update's.

Fixes and improvements:

- Show help for a key with `help <key>`
- gh-21 Reduced memory and time overhead of incremental recompilation with signature hash based approach.
- Rotate global log so that only output since last prompt is displayed for last
- gh-169 Add support for exclusions with `excludeAll` and `exclude` methods on `ModuleID`. (Indrajit)
- gh-235 Checksums configurable for launcher
- gh-246 Invalidate `update` when `update` is invalidated for an internal project dependency
- gh-138 Include plugin sources and docs in `update-sbt-classifiers`
- gh-219 Add `cleanupCommands` setting to specify commands to run before interpreter exits
- gh-46 Fix regression in caching missing classifiers for `update-classifiers` and `update-sbt-classifiers`.
- gh-228 Set `connectInput` to true to connect standard input to forked run
- gh-229 Limited task execution interruption using `ctrl+c`
- gh-220 Properly record source dependencies from separate compilation runs in the same step.
- gh-214 Better default behavior for `classpathConfiguration` for external Ivy files

- gh-212 Fix transitive plugin dependencies.
- gh-222 Generate section in make-pom. (Jan)
- Build resolvers, loaders, and transformers.
- Allow project dependencies to be modified by a setting (`buildDependencies`) but with the restriction that new builds cannot be introduced.
- gh-174, gh-196, gh-201, gh-204, gh-207, gh-208, gh-226, gh-224, gh-253

0.10.1 to 0.11.0

Major Improvements:

- Move to 2.9.1 for project definitions and plugins
- Drop support for 2.7
- Settings overhaul, mainly to make API documentation more usable
- Support using native libraries in `run` and `test` (but not console, for example)
- Automatic plugin cross-versioning. Use

```
addSbtPlugin("group" % "name" % "version")
```

in `project/plugins.sbt` instead of `libraryDependencies += ...` See [Plugins](#) for details

Fixes and Improvements:

- Display all undefined settings at once, instead of only the first one
- Deprecate separate `classpathFilter`, `defaultExcludes`, and `sourceFilter` keys in favor of `includeFilter` and `excludeFilter` explicitly scoped by `unmanagedSources`, `unmanagedResources`, or `unmanagedJars` as appropriate (Indrajit)
- Default to using shared boot directory in `~/.sbt/boot/`
- Can put contents of `project/plugins/` directly in `project/` instead. Will likely deprecate `plugins/` directory
- Key display is context sensitive. For example, in a single project, the build and project axes will not be displayed
- gh-114, gh-118, gh-121, gh-132, gh-135, gh-157: Various settings and error message improvements
- gh-115: Support configuring checksums separately for `publish` and `update`
- gh-118: Add `about` command

- gh-118, gh-131: Improve `last` command. Aggregate `last <task>` and display all recent output for `last`
- gh-120: Support read-only external file projects (Fred)
- gh-128: Add `skip` setting to override recompilation change detection
- gh-139: Improvements to pom generation (Indrajit)
- gh-140, gh-145: Add standard manifest attributes to binary and source jars (Indrajit)
- Allow sources used for `doc` generation to be different from sources for `compile`
- gh-156: Made `package` an alias for `package-bin`
- gh-162: handling of optional dependencies in pom generation

0.10.0 to 0.10.1

Some of the more visible changes:

- Support “provided” as a valid configuration for inter-project dependencies gh-53
- Try out some better error messages for `build.sbt` in a few common situations gh-58
- Drop “Incomplete tasks ...” line from error messages. gh-32
- Better handling of `javac` logging. gc-74
- Warn when reload discards session settings
- Cache failing classifiers, making ‘update-classifiers’ a practical replacement for `withSources()`
- Global settings may be provided in `~/.sbt/build.sbt` gh-52
- No need to define `"sbtPlugin := true"` in `project/plugins/` or `~/.sbt/plugins/`
- Provide statistics and list of evicted modules in `UpdateReport`
- Scope use of ‘transitive-classifiers’ by ‘update-sbt-classifiers’ and ‘update-classifiers’ for separate configuration.
- Default project ID includes a hash of base directory to avoid collisions in simple cases.
- ‘extra-loggers’ setting to make it easier to add loggers
- Associate `ModuleID`, `Artifact` and `Configuration` with a classpath entry (`moduleID`, `artifact`, and `configuration` keys). gh-41
- Put `httpClient` on Ivy’s classpath, which seems to speed up ‘update’.

0.7.7 to 0.10.0

Major redesign, only prominent changes listed.

- Project definitions in Scala 2.8.1

- New configuration system: See [.sbt build example](#), [.scala build definition](#) and [.sbt build definition](#).
- New task engine: [Tasks](#)
- New multiple project support: [.scala build definition](#)
- More aggressive incremental recompilation for both Java and Scala sources
- Merged plugins and processors into improved plugins system: [Plugins](#)
- [Web application](#) and webstart support moved to plugins instead of core features
- Fixed all of the issues in (Google Code) issue #44
- Managed dependencies automatically updated when configuration changes
- `update-sbt-classifiers` and `update-classifiers` tasks for retrieving sources and/or javadocs for dependencies, transitively
- Improved [artifact handling and configuration][Artifacts]
- Tab completion parser combinators for commands and input tasks: [Commands](#)
- No project creation prompts anymore
- Moved to GitHub: <http://github.com/harrah/xsbt>

0.7.5 to 0.7.7

- Workaround for Scala issue [#4426](#)
- Fix issue 156

0.7.4 to 0.7.5

- Joonas's update to work with Jetty 7.1 logging API changes.
- Updated to work with Jetty 7.2 WebAppClassLoader binary incompatibility (issue 129).
- Provide application and boot classpaths to tests and 'run'ning code according to <http://gist.github.com/404272>
- Fix `provided` configuration. It is no longer included on the classpath of dependent projects.
- Scala 2.8.1 is the default version used when starting a new project.
- Updated to [Ivy 2.2.0](#).
- Trond's patches that allow configuring [jetty-env.xml](#) and [webdefault.xml](#)
- Doug's [patch](#) to make 'projects' command show an asterisk next to current project
- Fixed issue 122
- Implemented issue 118
- Patch from Viktor and Ross for issue 123
- (RC1) Patch from Jorge for issue 100
- (RC1) Fix `<packaging>` type

0.7.3 to 0.7.4

- prefix continuous compilation with run number for better feedback when logging level is 'warn'
- Added `pomIncludeRepository(repo: MavenRepository): Boolean` that can be overridden to exclude local repositories by default
- Added `pomPostProcess(pom: Node): Node` to make advanced manipulation of the default pom easier (pomExtra already covers basic cases)
- Added `reset` command to reset JLine terminal. This needs to be run after suspending and then resuming sbt.
- Installer plugin is now a proper subproject of sbt.
- Plugins can now only be Scala sources. BND should be usable in a plugin now.
- More accurate detection of invalid test names. Invalid test names now generate an error and prevent the test action from running instead of just logging a warning.
- Fix issue with using 2.8.0.RC1 compiler in tests.
- Precompile compiler interface against 2.8.0.RC2
- Add `consoleOptions` for specifying options to the console. It defaults to `compileOptions`.
- Properly support sftp/ssh repositories using key-based authentication. See the updated section of the [Resolvers](#) page.
- `def ivyUpdateLogging = UpdateLogging.DownloadOnly | Full | Quiet`. Default is `DownloadOnly`. Full will log metadata resolution and provide a final summary.
- `offline` property for disabling checking for newer dynamic revisions (like -SNAPSHOT). This allows working offline with remote snapshots. Not honored for plugins yet.
- History commands: `!!`, `!?string`, `!-n`, `!n`, `!string`, `!:n`, `!:` Run `!` to see help.
- New section in launcher configuration `[ivy]` with a single label `cache-directory`. Specify this to change the cache location used by the launcher.
- New label `classifiers` under `[app]` to specify classifiers of additional artifacts to retrieve for the application.
- Honor `-Xfatal-warnings` option added to compiler in 2.8.0.RC2.
- Make `scaladocTask` a `fileTask` so that it runs only when `index.html` is older than some input source.
- Made it easier to create default `test-*` tasks with different options
- Sort input source files for consistency, addressing scalac's issues with source file ordering.
- Derive Java source file from name of class file when no `SourceFile` attribute is present in the class file. Improves tracking when `-g:none` option is used.
- Fix `FileUtilities.unzip` to be tail-recursive again.

0.7.2 to 0.7.3

- Fixed issue with `scala.library.jar` not being on `javac`'s classpath
- Fixed buffered logging for parallel execution
- Fixed `test-*` tab completion being permanently set on first completion
- Works with Scala 2.8 trunk again.
- Launcher: Maven local repository excluded when the Scala version is a snapshot. This should fix issues with out of date Scala snapshots.
- The compiler interface is precompiled against common Scala versions (for this release, 2.7.7 and 2.8.0.Beta1).
- Added `PathFinder.distinct`
- Running multiple commands at once at the interactive prompt is now supported. Prefix each command with `;`.
- Run and return the output of a process as a String with `!!` or as a (blocking) `Stream[String]` with `lines`.
- Java tests + Annotation detection
- Test frameworks can now specify annotation fingerprints. Specify the names of annotations and sbt discovers classes with the annotations on it or one of its methods. Use version 0.5 of the test-interface.
- Detect subclasses and annotations in Java sources (really, their class files)
- Discovered is new root of hierarchy representing discovered subclasses + annotations. `TestDefinition` no longer fulfills this role.
- `TestDefinition` is modified to be `name+Fingerprint` and represents a runnable test. It need not be `Discovered`, but could be file-based in the future, for example.
- Replaced `testDefinitionClassNames` method with `fingerprints` in `CompileConfiguration`.
- Added `foundAnnotation` to `AnalysisCallback`
- Added `Runner2`, `Fingerprint`, `AnnotationFingerprint`, and `SubclassFingerprint` to the test-interface. Existing test frameworks should still work. Implement `Runner2` to use fingerprints other than `SubclassFingerprint`.

0.7.1 to 0.7.2

- `Process.apply` no longer uses `CommandParser`. This should fix issues with the android-plugin.
- Added `sbt.impl.Arguments` for parsing a command like a normal action (for Processors)
- Arguments are passed to `javac` using an argument file (`@`)
- Added `webappUnmanaged: PathFinder` method to `DefaultWebProject`. Paths selected by this `PathFinder` will not be pruned by `prepare-webapp`

and will not be packaged by package. For example, to exclude the GAE datastore directory:

```
override def webappUnmanaged = (temporaryWarPath / "WEB-INF" / "appengine-generated"
```

- Added some String generation methods to `PathFinder`: `toString` for debugging and `absString` and `relativeString` for joining the absolute (relative) paths by the platform separator.
- Made tab completors lazier to reduce startup time.
- Fixed `console-project` for custom subprojects
- `Processor` split into `Processor/BasicProcessor`. `Processor` provides high level of integration with command processing. `BasicProcessor` operates on a `Project` but does not affect command processing.
- Can now use `Launcher` externally, including launching `sbt` outside of the official jar. This means a `Project` can now be created from tests.
- Works with Scala 2.8 trunk
- Fixed logging level behavior on subprojects.
- All sbt code is now at <http://github.com/harrah/xsbt> in one project.

0.7.0 to 0.7.1

- Fixed Jetty 7 support to work with JRebel
- Fixed make-pom to generate valid dependencies section

0.5.6 to 0.7.0

- Unified batch and interactive commands. All commands that can be executed at interactive prompt can be run from the command line. To run commands and then enter interactive prompt, make the last command 'shell'.
- Properly track certain types of synthetic classes, such as for comprehension with >30 clauses, during compilation.
- Jetty 7 support
- Allow launcher in the project root directory or the `lib` directory. The jar name must have the form 'sbt-launch.jar' in order to be excluded from the classpath.
- Stack trace detail can be controlled with 'on', 'off', 'nosbt', or an integer level. 'nosbt' means to show stack frames up to the first sbt method. An integer level denotes the number of frames to show for each cause. This feature is courtesy of Tony Sloane.

- New action ‘test-run’ method that is analogous to ‘run’, but for test classes.
- New action ‘clean-plugins’ task that clears built plugins (useful for plugin development).
- Can provide commands from a file with new command: `<filename`
- Can provide commands over loopback interface with new command: `<port`
- Scala version handling has been completely redone.
- The version of Scala used to run sbt (currently 2.7.7) is decoupled from the version used to build the project.
- Changing between Scala versions on the fly is done with the command: `++<version>`
- Cross-building is quicker. The project definition does not need to be re-compiled against each version in the cross-build anymore.
- Scala versions are specified in a space-delimited list in the `build.scala.versions` property.
- Dependency management:
- `make-pom` task now uses custom pom generation code instead of Ivy’s pom writer.
- Basic support for writing out Maven-style repositories to the pom
- Override the ‘pomExtra’ method to provide XML (`scala.xml.NodeSeq`) to insert directly into the generated pom.
- Complete control over repositories is now possible by overriding `ivyRepositories`.
- The interface to Ivy can be used directly.
- Test framework support is now done through a uniform test interface. Implications:
- New versions of `specs`, `ScalaCheck`, and `ScalaTest` are supported as soon as they are released.
- Support is better, since the test framework authors provide the implementation.
- Arguments can be passed to the test framework. For example: `{{{ > test-only your.test -a -b -c }}}}`
- Can provide custom task start and end delimiters by defining the system properties `sbt.start.delimiter` and `sbt.end.delimiter`.
- Revamped launcher that can launch Scala applications, not just sbt
- Provide a configuration file to the launcher and it can download the application and its dependencies from a repository and run it.
- sbt’s configuration can be customized. For example,
- The `sbt` version to use in projects can be fixed, instead of read from `project/build.properties`.
- The default values used to create a new project can be changed.
- The repositories used to fetch `sbt` and its dependencies, including Scala, can be configured.

- The location `sbt` is retrieved to is configurable. For example, `/home/user/.ivy2/sbt/` could be used instead of `project/boot/`.

0.5.5 to 0.5.6

- Support specs specifications defined as classes
- Fix specs support for 1.6
- Support ScalaTest 1.0
- Support ScalaCheck 1.6
- Remove remaining uses of structural types

0.5.4 to 0.5.5

- Fixed problem with classifier support and the corresponding test
- No longer need `"->default"` in configurations (automatically mapped).
- Can specify a specific nightly of Scala 2.8 to use (for example: `2.8.0-20090910.003346--+`)
- Experimental support for searching for project (`-Dsbt.boot.search=none | only | root-first | nearest`)
- Fix issue where last path component of local repository was dropped if it did not exist.
- Added support for configuring repositories on a per-module basis.
- Unified batch-style and interactive-style commands. All commands that were previously interactive-only should be available batch-style. ‘reboot’ does not pick up changes to ‘scala.version’ properly, however.

0.5.2 to 0.5.4

- Many logging related changes and fixes. Added `FilterLogger` and cleaned up interaction between `Logger`, scripted testing, and the builder projects. This included removing the `recordingDepth` hack from `Logger`. `Logger` buffering is now enabled/disabled per thread.
- Fix `compileOptions` being fixed after the first compile
- Minor fixes to output directory checking
- Added `defaultLoggingLevel` method for setting the initial level of a project’s `Logger`
- Cleaned up internal approach to adding extra default configurations like plugin
- Added `syncPathsTask` for synchronizing paths to a target directory

- Allow multiple instances of Jetty (new `jettyRunTasks` can be defined with different ports)
- `jettyRunTask` accepts configuration in a single configuration wrapper object instead of many parameters
- Fix web application class loading (issue #35) by using `jettyClasspath=testClasspath—jettyRunClasspath` for loading Jetty. A better way would be to have a `jetty` configuration and have `jettyClasspath=managedClasspath('jetty')`, but this maintains compatibility.
- Copy resources to `target/resources` and `target/test-resources` using `copyResources` and `copyTestResources` tasks. Properly include all resources in web applications and classpaths (issue #36). `mainResources` and `testResources` are now the definitive methods for getting resources.
- Updated for 2.8 (`sbt` now compiles against September 11, 2009 nightly build of Scala)
- Fixed issue with position of `^` in compile errors
- Changed order of repositories (local, shared, Maven Central, user, Scala Tools)
- Added Maven Central to resolvers used to find Scala library/compiler in launcher
- Fixed problem that prevented detecting user-specified subclasses
- Fixed exit code returned when exception thrown in main thread for `TrapExit`
- Added `javap` task to `DefaultProject`. It has tab completion on compiled project classes and the run classpath is passed to `javap` so that library classes are available. Examples: :

```
> javap your.Clazz
> javap -c scala.List
```

- Added `exec` task. Mixin `Exec` to project definition to use. This forks the command following `exec`. Examples: :

```
> exec echo Hi
> exec find src/main/scala -iname *.scala -exec wc -l {} ;
```

- Added `sh` task for users with a unix-style shell available (runs `/bin/sh -c <arguments>`). Mixin `Exec` to project definition to use. Example: :

```
> sh find src/main/scala -iname *.scala | xargs cat | wc -l
```

- Proper dependency graph actions (previously was an unsupported prototype): `graph-src` and `graph-pkg` for source dependency graph and quasi-package dependency graph (based on source directories and source dependencies)
- Improved Ivy-related code to not load unnecessary default settings
- Fixed issue #39 (sources were not relative in src package)
- Implemented issue #38 (`InstallProject` with ‘install’ task)
- Vesa’s patch for configuring the output of forked Scala/Java and processes
- Don’t buffer logging of forked `run` by default
- Check `Project.terminateWatch` to determine if triggered execution should stop for a given keypress.
- Terminate triggered execution only on ‘enter’ by default (previously, any keypress stopped it)
- Fixed issue #41 (parent project should not declare jar artifact)
- Fixed issue #42 (search parent directories for `ivysettings.xml`)
- Added support for extra attributes with Ivy. Use `extra(key -> value)` on `ModuleIDs` and `Artifacts`. To define for a project’s ID: :

```
override def projectID = super.projectID extra(key -> value)
```

To specify in a dependency: :

```
val dep = normalID extra(key -> value)
```

0.5.1 to 0.5.2

- Fixed problem where dependencies of `sbt` plugins were not on the compile classpath
- Added `execTask` that runs an `sbt.ProcessBuilder` when invoked
- Added implicit conversion from `scala.xml.Elem` to `sbt.ProcessBuilder` that takes the element’s text content, trims it, and splits it around whitespace to obtain the command.
- Processes can now redirect standard input (see `run` with Boolean argument or `!<` operator on `ProcessBuilder`), off by default
- Made scripted framework a plugin and scripted tests now go in `src/sbt-test` by default

- Can define and use an sbt test framework extension in a project
- Fixed `run` action swallowing exceptions
- Fixed tab completion for method tasks for multi-project builds
- Check that tasks in `compoundTask` do not reference static tasks
- Make `toString` of `Paths` in subprojects relative to root project directory
- `crossScalaVersions` is now inherited from parent if not specified
- Added `scala-library.jar` to the `javac` classpath
- Project dependencies are added to published `ivy.xml`
- Added dependency tracking for Java sources using classfile parsing (with the usual limitations)
- Added `Process.cat` that will send contents of `URLs` and `Files` to standard output. Alternatively, `cat` can be used on a single `URL` or `File`. Example:
:

```
import java.net.URL
import java.io.File
val spde = new URL("http://technically.us/spde/About")
val dispatch = new URL("http://databinder.net/dispatch/About")
val build = new File("project/build.properties")
cat(spde, dispatch, build) #| "grep -i scala" !
```

0.4.6 to 0.5/0.5.1

- Fixed `ScalaTest` framework dropping stack traces
- Publish only public configurations by default
- Loader now adds `.m2/repository` for downloading Scala jars
- Can now fork the compiler and runner and the runner can use a different working directory.
- Maximum compiler errors shown is now configurable
- Fixed rebuilding and republishing released versions of `sbt` against new Scala versions (attempt #2)
- Fixed snapshot reversion handling (Ivy needs changing pattern set on cache, apparently)
- Fixed handling of default configuration when `useMavenConfiguration` is `true`
- Cleanup on `Environment`, `Analysis`, `Conditional`, `MapUtilities`, and more...
- Tests for `Environment`, source dependencies, library dependency management, and more...

- Dependency management and multiple Scala versions
- Experimental plugin for producing project bootstrapper in a self-extracting jar
- Added ability to directly specify `URL` to use for dependency with the `from(url: URL)` method defined on `ModuleID`
- Fixed issue #30
- Support cross-building with `+` when running batch actions
- Additional flattening for project definitions: sources can go either in `project/build/src` (recursively) or `project/build` (flat)
- Fixed manual `reboot` not changing the version of Scala when it is manually set
- Fixed tab completion for cross-building
- Fixed a class loading issue with web applications

0.4.5 to 0.4.6

- Publishing to `ssh/sftp/filesystem` repository supported
- Exception traces are printed by default
- Fixed warning message about no `Class-Path` attribute from showing up for run
- Fixed `package-project` operation
- Fixed `Path.fromFile`
- Fixed issue with external process output being lost when sent to a `Buffered-Logger` with `parallelExecution` enabled.
- Preserve history across `clean`
- Fixed issue with making relative path in jar with wrong separator
- Added cross-build functionality (prefix action with `+`).
- Added methods `scalaLibraryJar` and `scalaCompilerJar` to `FileUtilities`
- Include project dependencies for `deliver/publish`
- Add Scala dependencies for `make-pom/deliver/publish`, which requires these to depend on package
- Properly add compiler jar to run/test classpaths when main sources depend on it
- `TestFramework` root `ClassLoader` filters compiler classes used by sbt, which is required for projects using the compiler.
- Better access to dependencies:
- `mainDependencies` and `testDependencies` provide an analysis of the dependencies of your code as determined during compilation
- `scalaJars` is deprecated, use `mainDependencies.scalaJars` instead (provides a `PathFinder`, which is generally more useful)
- Added `jettyPort` method to `DefaultWebProject`.
- Fixed `package-project` to exclude `project/boot` and `project/build/target`
- Support specs 1.5.0 for Scala 2.7.4 version.

- Parallelization at the subtask level
- Parallel test execution at the suite/specification level.

0.4.3 to 0.4.5

- Sorted out repository situation in loader
- Added support for `http_proxy` environment variable
- Added `download` method from Nathan to `FileUtilities` to retrieve the contents of a URL.
- Added special support for compiler plugins, see [compiler plugins support](#) page.
- `reload` command in scripted tests will now properly handle success/failure
- Very basic support for Java sources: Java sources under `src/main/java` and `src/test/java` will be compiled.
- `parallelExecution` defaults to value in parent project if there is one.
- Added ‘console-project’ that enters the Scala interpreter with the current Project bound to the variable project.
- The default Ivy cache manager is now configured with `useOrigin=true` so that it doesn’t cache artifacts from the local filesystem.
- For users building from trunk, if a project specifies a version of sbt that ends in `-SNAPSHOT`, the loader will update sbt every time it starts up. The trunk version of sbt will always end in `-SNAPSHOT` now.
- Added automatic detection of classes with main methods for use when `mainClass` is not explicitly specified in the project definition. If exactly one main class is detected, it is used for run and package. If multiple main classes are detected, the user is prompted for which one to use for run. For package, no `Main-Class` attribute is automatically added and a warning is printed.
- Updated build to cross-compile against Scala 2.7.4.
- Fixed `proguard` task in sbt’s project definition
- Added `manifestClassPath` method that accepts the value for the `Class-Path` attribute
- Added `PackageOption` called `ManifestAttributes` that accepts `(java.util.jar.Attributes.Name, String)` or `(String, String)` pairs and adds them to the main manifest attributes
- Fixed some situations where characters would not be echoed at prompts other than main prompt.
- Fixed issue #20 (use `http_proxy` environment variable)
- Implemented issue #21 (native process wrapper)
- Fixed issue #22 (rebuilding and republishing released versions of sbt against new Scala versions, specifically Scala 2.7.4)
- Implemented issue #23 (inherit inline repositories declared in parent project)

0.4 to 0.4.3

- Direct dependencies on Scala libraries are checked for version equality with `scala.version`
- Transitive dependencies on `scala-library` and `scala-compiler` are filtered
- They are fixed by `scala.version` and provided on the classpath by sbt
- To access them, use the `scalaJars` method, `classOf[ScalaObject].getProtectionDomain.getCodeSource`, or `mainCompileConditional.analysis.allExternals`
- The configurations checked/filtered as described above are configurable. Nonstandard configurations are not checked by default.
- Version of `sbt` and Scala printed on startup
- Launcher asks if you want to try a different version if `sbt` or Scala could not be retrieved.
- After changing `scala.version` or `sbt.version` with `set`, note is printed that reboot is required.
- Moved managed dependency actions to `BasicManagedProject` (update is now available on `ParentProject`)
- Cleaned up `sbt`'s build so that you just need to do `update` and full-build to build from source. The trunk version of sbt will be available for use from the loader.
- The loader is now a subproject.
- For development, you'll still want the usual actions (such as `package`) for the main builder and proguard to build the loader.
- Fixed analysis plugin improperly including traits/abstract classes in subclass search
- `ScalaProjects` already had everything required to be parent projects: flipped the switch to enable it
- Proper method task support in scripted tests (`package` group tests rightly pass again)
- Improved tests in loader that check that all necessary libraries were downloaded properly

0.3.7 to 0.4

- Fixed issue with `build.properties` being unnecessarily updated in subprojects when loading.
- Added method to compute the SHA-1 hash of a `String`
- Added `pack200` methods
- Added initial process interface
- Added initial webstart support
- Added `gzip` methods
- Added `sleep` and `newer` commands to scripted testing.

- Scripted tests now test the version of `sbt` being built instead of the version doing the building.
- `testResources` is put on the test classpath instead of `testResourcesPath`
- Added `jetty-restart`, which does `jetty-stop` and then `jetty-run`
- Added automatic reloading of default web application
- Changed packaging behaviors (still likely to change)
- Inline configurations now allowed (can be used with configurations in inline XML)
- Split out some code related to managed dependencies from `BasicScalaProject` to new class `BasicManagedProject`
- Can specify that maven-like configurations should be automatically declared
- Fixed problem with nested modules being detected as tests
- `testResources`, `integrationTestResources`, and `mainResources` should now be added to appropriate classpaths
- Added project organization as a property that defaults to inheriting from the parent project.
- Project creation now prompts for the organization.
- Added method tasks, which are top-level actions with parameters.
- Made `help`, `actions`, and `methods` commands available to batch-style invocation.
- Applied Mikko's two fixes for webstart and fixed problem with `pack200+sign`. Also, fixed nonstandard behavior when `gzip` enabled.
- Added `control` method to `Logger` for action lifecycle logging
- Made standard logging level convenience methods final
- Made `BufferedLogger` have a per-actor buffer instead of a global buffer
- Added a `SynchronizedLogger` and a `MultiLogger` (intended to be used with the yet unwritten `FileLogger`)
- Changed method of atomic logging to be a method `logAll` accepting `List[LogEvent]` instead of `doSynchronized`
- Improved action lifecycle logging
- Parallel logging now provides immediate feedback about starting an action
- General cleanup, including removing unused classes and methods and reducing dependencies between classes
- `run` is now a method task that accepts options to pass to the main method (`runOptions` has been removed, `runTask` is no longer interactive, and `run` no longer starts a console if `mainClass` is undefined)
- Major task execution changes:
- Tasks automatically have implicit dependencies on tasks with the same name in dependent projects
- Implicit dependencies on interactive tasks are ignored, explicit dependencies produce an error
- Interactive tasks must be executed directly on the project on which they

are defined

- Method tasks accept input arguments (`Array[String]`) and dynamically create the task to run
- Tasks can depend on tasks in other projects
- Tasks are run in parallel breadth-first style
- Added **test-only** method task, which restricts the tests to run to only those passed as arguments.
- Added **test-failed** method task, which restricts the tests to run. First, only tests passed as arguments are run. If no tests are passed, no filtering is done. Then, only tests that failed the previous run are run.
- Added **test-quick** method task, which restricts the tests to run. First, only tests passed as arguments are run. If no tests are passed, no filtering is done. Then, only tests that failed the previous run or had a dependency change are run.
- Added launcher that allows declaring version of sbt/scala to build project with.
- Added tab completion with ~
- Added basic tab completion for method tasks, including **test-***
- Changed default pack options to be the default options of `Pack200.Packer`
- Fixed ~ behavior when action doesn't exist

0.3.6 to 0.3.7

- Improved classpath methods
- Refactored various features into separate project traits
- `ParentProject` can now specify dependencies
- Support for `optional` scope
- More API documentation
- Test resource paths provided on classpath for testing
- Added some missing read methods in `FileUtilities`
- Added scripted test framework
- Change detection using hashes of files
- Fixed problem with manifests not being generated (bug #14)
- Fixed issue with scala-tools repository not being included by default (again)
- Added option to set ivy cache location (mainly for testing)
- `trace` is no longer a logging level but a flag enabling/disabling stack traces
- `Project.loadProject` and related methods now accept a `Logger` to use
- Made hidden files and files that start with `'.'` excluded by default (`'.*'` is required because subversion seems to not mark `.svn` directories hidden on Windows)
- Implemented exit codes
- Added continuous compilation command `cc`

0.3.5 to 0.3.6

- Fixed bug #12.
- Compiled with 2.7.2.

0.3.2 to 0.3.5

- Fixed bug #11.
- Fixed problem with dependencies where source jars would be used instead of binary jars.
- Fixed scala-tools not being used by default for inline configurations.
- Small dependency management error message correction
- Slight refactoring for specifying whether scala-tools releases gets added to configured resolvers
- Separated repository/dependency overriding so that repositories can be specified inline for use with ivy.xml or pom.xml files
- Added ability to specify Ivy XML configuration in Scala.
- Added `clean-cache` action for deleting Ivy's cache
- Some initial work towards accessing a resource directory from tests
- Initial tests for `Path`
- Some additional `FileUtilities` methods, some `FileUtilities` method adjustments and some initial tests for `FileUtilities`
- A basic framework for testing `ReflectUtilities`, not run by default because of run time
- Minor cleanup to `Path` and added non-empty check to path components
- Catch additional exceptions in `TestFramework`
- Added `copyTask` task creation method.
- Added `jetty-run` action and added ability to package war files.
- Added `jetty-stop` action.
- Added `console-quick` action that is the same as `console` but doesn't compile sources first.
- Moved some custom `ClassLoaders` to `ClasspathUtilities` and improved a check.
- Added ability to specify hooks to call before `sbt` shuts down.
- Added `zip`, `unzip` methods to `FileUtilities`
- Added `append` equivalents to `write*` methods in `FileUtilities`
- Added first draft of integration testing
- Added batch command `compile-stats`
- Added methods to create tasks that have basic conditional execution based on declared sources/products of the task
- Added `newerThan` and `olderThan` methods to `Path`
- Added `reload` action to reread the project definition without losing the performance benefits of an already running jvm

- Added **help** action to tab completion
- Added handling of (effectively empty) scala source files that create no class files: they are always interpreted as modified.
- Added prompt to retry project loading if compilation fails
- **package** action now uses **fileTask** so that it only executes if files are out of date
- fixed **ScalaTest** framework wrapper so that it fails the **test** action if tests fail
- Inline dependencies can now specify configurations

0.3.1 to 0.3.2

- Compiled jar with Java 1.5.

0.3 to 0.3.1

- Fixed bugs #8, #9, and #10.

0.2.3 to 0.3

- Version change only for first release.

0.2.2 to 0.2.3

- Added tests for **Dag**, **NameFilter**, **Version**
- Fixed handling of trailing *s in **GlobFilter** and added some error-checking for control characters, which **Pattern** doesn't seem to like
- Fixed **Analysis.allProducts** implementation
- It previously returned the sources instead of the generated classes
- Will only affect the count of classes (it should be correct now) and the debugging of missed classes (erroneously listed classes as missed)
- Made some implied preconditions on **BasicVersion** and **OpaqueVersion** explicit
- Made increment version behavior in **ScalaProject** easier to overload
- Added **Seq[..Option]** alternative to **...Option*** for tasks
- Documentation generation fixed to use latest value of version
- Fixed **BasicVersion.incrementMicro**
- Fixed test class loading so that **sbt** can test the version of sbt being developed (previously, the classes from the executing version of sbt were tested)

0.2.1 to 0.2.2

- Package name is now a call-by-name parameter for the package action
- Fixed release action calling compile multiple times

0.2.0 to 0.2.1

- Added some action descriptions
- jar name now comes from normalized name (lowercased and spaces to dashes)
- Some cleanups related to creating filters
- Path should only 'get' itself if the underlying file exists to be consistent with other PathFinders
- Added `---` operator for `PathFinder` that excludes paths from the `PathFinder` argument
- Removed `***` operator on `PathFinder`
- `**` operator on `PathFinder` matches all descendents or self that match the `NameFilter` argument
- The above should fix bug #6
- Added version increment and release actions.
- Can now build sbt with sbt. Build scripts `build` and `clean` will still exist.

0.1.9 to 0.2.0

- Implemented typed properties and access to system properties
- Renamed `metadata` directory to `project`
- Information previously in `info` file now obtained by properties:
- `info.name --> name`
- `info.currentVersion --> version`
- Concrete `Project` subclasses should have a constructor that accepts a single argument of type `ProjectInfo` (argument dependencies: `Iterable[Project]` has been merged into `ProjectInfo`)

0.1.8 to 0.1.9

- Better default implementation of `allSources`.
- Generate warning if two jars on classpath have the same name.
- Upgraded to specs 1.4.0
- Upgraded to `ScalaCheck` 1.5
- Changed some update options to be final vals instead of objects.
- Added some more API documentation.
- Removed release action.

- Split compilation into separate main and test compilations.
- A failure in a `ScalaTest` run now fails the test action.
- Implemented reporters for `compile/scaladoc`, `ScalaTest`, `ScalaCheck`, and specs that delegate to the appropriate `sbt.Logger`.

0.1.7 to 0.1.8

- Improved configuring of tests to exclude.
- Simplified version handling.
- Task `&&` operator properly handles dependencies of tasks it combines.
- Changed method of inline library dependency declarations to be simpler.
- Better handling of errors in parallel execution.

0.1.6 to 0.1.7

- Added graph action to generate dot files (for graphviz) from dependency information (work in progress).
- Options are now passed to tasks as varargs.
- Redesigned `Path` properly, including `PathFinder` returning a `Set[Path]` now instead of `Iterable[Path]`.
- Moved paths out of `ScalaProject` and into `BasicProjectPaths` to keep path definitions separate from task definitions.
- Added initial support for managing third-party libraries through the update task, which must be explicitly called (it is not a dependency of compile or any other task). This is experimental, undocumented, and known to be incomplete.
- Parallel execution implementation at the project level, disabled by default. To enable, add: `scala override def parallelExecution = true` to your project definition. In order for logging to make sense, all project logging is buffered until the project is finished executing. Still to be done is some sort of notification of project execution (which ones are currently executing, how many remain)
- `run` and `console` are now specified as “interactive” actions, which means they are only executed on the project in which they are defined when called directly, and not on all dependencies. Their dependencies are still run on dependent projects.
- Generalized conditional tasks a bit. Of note is that analysis is no longer required to be in `metadata/analysis`, but is now in `target/analysis` by default.
- Message now displayed when project definition is recompiled on startup
- Project no longer inherits from `Logger`, but now has a `log` member.
- Dependencies passed to `project` are checked for null (may help with errors related to initialization/circular dependencies)

- Task dependencies are checked for null
- Projects in a multi-project configuration are checked to ensure that output paths are different (check can be disabled)
- Made **update** task globally synchronized because Ivy is not thread-safe.
- Generalized test framework, directly invoking frameworks now (used reflection before).
- Moved license files to licenses/
- Added support for **specs** and some support for **ScalaTest** (the test action doesn't fail if ScalaTest tests fail).
- Added **specs**, **ScalaCheck**, **ScalaTest** jars to lib/
- These are now required for compilation, but are optional at runtime.
- Added the appropriate licenses and notices.
- Options for **update** action are now taken from updateOptions member.
- Fixed **SbtManager** inline dependency manager to work properly.
- Improved Ivy configuration handling (not compiled with test dependencies yet though).
- Added case class implementation of **SbtManager** called SimpleManager.
- Project definitions not specifying dependencies can now use just a single argument constructor.

0.1.5 to 0.1.6

- **run** and **console** handle **System.exit** and multiple threads in user code under certain circumstances (see [running project code](#)).

0.1.4 to 0.1.5

- Generalized interface with plugin (see Analysis Callback)
- Split out task implementations and paths from **Project** to **ScalaProject**
- Subproject support (changed required project constructor signature: see `sbt/DefaultProject.scala`)
- Can specify dependencies between projects
- Execute tasks across multiple projects
- Classpath of all dependencies included when compiling
- Proper inter-project source dependency handling
- Can change to a project in an interactive session to work only on that project (and its dependencies)
- External dependency handling
- Tracks non-source dependencies (compiled classes and jars)
- Requires each class to be provided by exactly one classpath element (This means you cannot have two versions of the same class on the classpath, e.g. from two versions of a library)

- Changes in a project propagate the right source recompilations in dependent projects
- Consequences:
- Recompilation when changing java/scala version
- Recompilation when upgrading libraries (again, as indicated in the second point, situations where you have library-1.0.jar and library-2.0.jar on the classpath at the same time are not handled predictably. Replacing library-1.0.jar with library-2.0.jar should work as expected.)
- Changing sbt version will recompile project definitions

0.1.3 to 0.1.4

- Autodetection of Project definitions.
- Simple tab completion/history in an interactive session with JLine
- Added descriptions for most actions

0.1.2 to 0.1.3

- Dependency management between tasks and auto-discovery tasks.
- Should work on Windows.

0.1.1 to 0.1.2

- Should compile/build on Java 1.5
- Fixed run action implementation to include scala library on classpath
- Made project configuration easier

0.1 to 0.1.1

- Fixed handling of source files without a package
- Added easy project setup

Migrating from 0.7 to 0.10+

The assumption here is that you are familiar with sbt 0.7 but new to sbt 0.13.5. sbt 0.13.5's many new capabilities can be a bit overwhelming, but this page should help you migrate to 0.13.5 with a minimum of fuss.

Why move to 0.13.5?

1. Faster builds (because it is smarter at re-compiling only what it must)
2. Easier configuration. For simple projects a single `build.sbt` file in your root directory is easier to create than `project/build/MyProject.scala` was.
3. No more `lib_managed` directory, reducing disk usage and avoiding backup and version control hassles.
4. `update` is now much faster and it's invoked automatically by sbt.
5. Terser output. (Yet you can ask for more details if something goes wrong.)

Step 1: Read the Getting Started Guide for sbt 0.13.5 Reading the [Getting Started Guide](#) will probably save you a lot of confusion.

Step 2: Install sbt 0.13.5 Download sbt 0.13.5 as described on [the setup page](#).

You can run 0.13.5 the same way that you run 0.7.x, either simply:

```
$ java -jar sbt-launch.jar
```

Or (as most users do) with a shell script, as described on [the setup page](#).

For more details see [the setup page](#).

Step 3: A technique for switching an existing project Here is a technique for switching an existing project to 0.13.5 while retaining the ability to switch back again at will. Some builds, such as those with subprojects, are not suited for this technique, but if you learn how to transition a simple project it will help you do a more complex one next.

Preserve project/ for 0.7.x project

Rename your `project/` directory to something like `project-old`. This will hide it from sbt 0.13.5 but keep it in case you want to switch back to 0.7.x.

Create build.sbt for 0.13.5

Create a `build.sbt` file in the root directory of your project. See [.sbt build definition](#) in the Getting Started Guide, and for [simple examples](#). If you have a simple project then converting your existing project file to this format is largely a matter of re-writing your dependencies and maven archive declarations in a modified yet familiar syntax.

This `build.sbt` file combines aspects of the old `project/build/ProjectName.scala` and `build.properties` files. It looks like a property file, yet contains Scala code in a special format.

A `build.properties` file like:

```
#Project properties
#Fri Jan 07 15:34:00 GMT 2011
project.organization=org.myproject
project.name=My Project
sbt.version=0.7.7
project.version=1.0
def.scala.version=2.7.7
build.scala.versions=2.8.1
project.initialize=false
```

Now becomes part of your `build.sbt` file with lines like:

```
name := "My Project"

version := "1.0"

organization := "org.myproject"

scalaVersion := "2.9.2"
```

Currently, a `project/build.properties` is still needed to explicitly select the sbt version. For example:

Run sbt 0.13.5

Now launch sbt. If you're lucky it works and you're done. For help debugging, see below.

Switching back to sbt 0.7.x

If you get stuck and want to switch back, you can leave your `build.sbt` file alone. sbt 0.7.x will not understand or notice it. Just rename your 0.13.5 project directory to something like `project10` and rename the backup of your old project from `project-old` to `project` again.

FAQs There's a section in the [FAQ](#) about migration from 0.7 that covers several other important points.

Contributing to sbt

Below is a running list of potential areas of contribution. This list may become out of date quickly, so you may want to check on the [sbt-dev mailing list](#) if you are interested in a specific topic.

1. There are plenty of possible visualization and analysis opportunities.
 - ‘compile’ produces an Analysis of the source code containing
 - Source dependencies
 - Inter-project source dependencies
 - Binary dependencies (jars + class files)
 - data structure representing the [API](#) of the source code There is some code already for generating dot files that isn’t hooked up, but graphing dependencies and inheritance relationships is a general area of work.
 - ‘update’ produces an [Update Report](#) mapping Configuration/ModuleID/Artifact to the retrieved File
 - Ivy produces more detailed XML reports on dependencies. These come with an XSL stylesheet to view them, but this does not scale to large numbers of dependencies. Working on this is pretty straightforward: the XML files are created in `~/.ivy2` and the `.xml` and `.css` are there as well, so you don’t even need to work with sbt. Other approaches described in [the email thread](#)
 - Tasks are a combination of static and dynamic graphs and it would be useful to view the graph of a run
 - Settings are a static graph and there is code to generate the dot files, but isn’t hooked up anywhere.
2. There is support for dependencies on external projects, like on GitHub. To be more useful, this should support being able to update the dependencies. It is also easy to extend this to other ways of retrieving projects. Support for svn and hg was a recent contribution, for example.
3. Dependency management: see [adept](#)
4. If you like parsers, sbt commands and input tasks are written using custom parser combinators that provide tab completion and error handling. Among other things, the efficiency could be improved.
5. The javap task hasn’t been reintegrated
6. Implement enhanced 0.11-style warn/debug/info/error/trace commands. Currently, you set it like any other setting:

```
set logLevel := Level.Warn
```

```
or set logLevel in Test := Level.Warn
```

You could make commands that wrap this, like:

```
warn test:run
```

Also, trace is currently an integer, but should really be an abstract data type.

7. Each sbt version has more aggressive incremental compilation and reproducing bugs can be difficult. It would be helpful to have a mode that generates a diff between successive compilations and records the options passed to scalac. This could be replayed or inspected to try to find the cause.

Documentation

1. There's a lot to do with this documentation. If you check it out from git, there's a directory called Dormant with some content that needs going through.
2. the main page mentions external project references (e.g. to a git repo) but doesn't have anything to link to that explains how to use those.
3. API docs are much needed.
4. Find useful answers or types/methods/values in the other docs, and pull references to them up into /faq or /Name-Index so people can find the docs. In general the /faq should feel a bit more like a bunch of pointers into the regular docs, rather than an alternative to the docs.
5. A lot of the pages could probably have better names, and/or little 2-4 word blurbs to the right of them in the sidebar.

Detailed Topics

This part of the documentation has pages documenting particular sbt topics in detail. Before reading anything in here, you will need the information in the [Getting Started Guide](#) as a foundation.

Other resources include the [How to](#) and [Developer's Guide](#) sections in this reference, and the [API Documentation](#)

Using sbt

This part of the documentation has pages documenting particular sbt topics in detail. Before reading anything in here, you will need the information in the [Getting Started Guide](#) as a foundation.

Command Line Reference

This page is a relatively complete list of command line options, commands, and tasks you can use from the sbt interactive prompt or in batch mode. See

[Running](#) in the Getting Started Guide for an intro to the basics, while this page has a lot more detail.

Notes on the command line

- There is a technical distinction in sbt between *tasks*, which are “inside” the build definition, and *commands*, which manipulate the build definition itself. If you’re interested in creating a command, see [Commands](#). This specific sbt meaning of “command” means there’s no good general term for “thing you can type at the sbt prompt”, which may be a setting, task, or command.
- Some tasks produce useful values. The `toString` representation of these values can be shown using `show <task>` to run the task instead of just `<task>`.
- In a multi-project build, execution dependencies and the aggregate setting control which tasks from which projects are executed. See [multi-project builds](#).

Project-level tasks

- `clean` Deletes all generated files (the `target` directory).
- `publishLocal` Publishes artifacts (such as jars) to the local Ivy repository as described in Publishing.
- `publish` Publishes artifacts (such as jars) to the repository defined by the `publishTo` setting, described in Publishing.
- `update` Resolves and retrieves external dependencies as described in [library dependencies](#).

Configuration-level tasks

Configuration-level tasks are tasks associated with a configuration. For example, `compile`, which is equivalent to `compile:compile`, compiles the main source code (the `compile` configuration). `test:compile` compiles the test source code (test `test` configuration). Most tasks for the `compile` configuration have an equivalent in the `test` configuration that can be run using a `test:` prefix.

- `compile` Compiles the main sources (in the `src/main/scala` directory). `test:compile` compiles test sources (in the `src/test/scala/` directory).
- `console` Starts the Scala interpreter with a classpath including the compiled sources, all jars in the `lib` directory, and managed libraries. To return to sbt, type `:quit`, `Ctrl+D` (Unix), or `Ctrl+Z` (Windows). Similarly, `test:console` starts the interpreter with the test classes and classpath.

- **consoleQuick** Starts the Scala interpreter with the project's compile-time dependencies on the classpath. `test:consoleQuick` uses the test dependencies. This task differs from `console` in that it does not force compilation of the current project's sources.
- **consoleProject** Enters an interactive session with sbt and the build definition on the classpath. The build definition and related values are bound to variables and common packages and values are imported. See the [consoleProject documentation](#) for more information.
- **doc** Generates API documentation for Scala source files in `src/main/scala` using scaladoc. `test:doc` generates API documentation for source files in `src/test/scala`.
- **package** Creates a jar file containing the files in `src/main/resources` and the classes compiled from `src/main/scala`. `test:package` creates a jar containing the files in `src/test/resources` and the class compiled from `src/test/scala`.
- **packageDoc** Creates a jar file containing API documentation generated from Scala source files in `src/main/scala`. `test:packageDoc` creates a jar containing API documentation for test sources files in `src/test/scala`.
- **packageSrc**: Creates a jar file containing all main source files and resources. The packaged paths are relative to `src/main/scala` and `src/main/resources`. Similarly, `test:packageSrc` operates on test source files and resources.
- **run <argument>*** Runs the main class for the project in the same virtual machine as sbt. The main class is passed the arguments provided. Please see [Running Project Code](#) for details on the use of `System.exit` and multithreading (including GUIs) in code run by this action. `test:run` runs a main class in the test code.
- **runMain <main-class> <argument>*** Runs the specified main class for the project in the same virtual machine as sbt. The main class is passed the arguments provided. Please see [Running Project Code](#) for details on the use of `System.exit` and multithreading (including GUIs) in code run by this action. `test:runMain` runs the specified main class in the test code.
- **test** Runs all tests detected during test compilation. See [Testing](#) for details.
- **testOnly <test>*** Runs the tests provided as arguments. `*` (will be) interpreted as a wildcard in the test name. See [Testing](#) for details.
- **testQuick <test>*** Runs the tests specified as arguments (or all tests if no arguments are given) that:
 1. have not been run yet OR
 2. failed the last time they were run OR
 3. had any transitive dependencies recompiled since the last successful run `*` (will be) interpreted as a wildcard in the test name. See [Testing](#) for details.

General commands

- **exit** or **quit** End the current interactive session or build. Additionally, Ctrl+D (Unix) or Ctrl+Z (Windows) will exit the interactive prompt.
- **help** `<command>` Displays detailed help for the specified command. If the command does not exist, help lists detailed help for commands whose name or description match the argument, which is interpreted as a regular expression. If no command is provided, displays brief descriptions of the main commands. Related commands are tasks and settings.
- **projects** [`add|remove` `<URI>`] List all available projects if no arguments provided or adds/removes the build at the provided URI. (See [multi-project builds](#) for details on multi-project builds.)
- **project** `<project-id>` Change the current project to the project with ID `<project-id>`. Further operations will be done in the context of the given project. (See [multi-project builds](#) for details on multiple project builds.)
- **~** `<command>` Executes the project specified action or method whenever source files change. See [Triggered Execution](#) for details.
- **< filename** Executes the commands in the given file. Each command should be on its own line. Empty lines and lines beginning with '#' are ignored
- **+** `<command>` Executes the project specified action or method for all versions of Scala defined in the `crossScalaVersions` setting.
- **++** `<version|home-directory>` `<command>` Temporarily changes the version of Scala building the project and executes the provided command. `<command>` is optional. The specified version of Scala is used until the project is reloaded, settings are modified (such as by the `set` or `session` commands), or `++` is run again. `<version>` does not need to be listed in the build definition, but it must be available in a repository. Alternatively, specify the path to a Scala installation.
- **;** `A` **;** `B` Execute A and if it succeeds, run B. Note that the leading semicolon is required.
- **eval** `<Scala-expression>` Evaluates the given Scala expression and returns the result and inferred type. This can be used to set system properties, as a calculator, to fork processes, etc ... For example:

```
> eval System.setProperty("demo", "true")
> eval 1+1
> eval "ls -l" !
```

Commands for managing the build definition

- **reload** [**plugins**|**return**] If no argument is specified, reloads the build, recompiling any build or plugin definitions as necessary. **reload plugins** changes the current project to the build definition project (in `project/`). This can be useful to directly manipulate the build definition. For example, running **clean** on the build definition project will force snapshots to be updated and the build definition to be recompiled. **reload return** changes back to the main project.
- **set** <setting-expression> Evaluates and applies the given setting definition. The setting applies until sbt is restarted, the build is reloaded, or the setting is overridden by another **set** command or removed by the **session** command. See [.sbt build definition](#) and [inspecting settings](#) for details.
- **session** <command> Manages session settings defined by the **set** command. It can persist settings configured at the prompt. See [Inspecting-Settings](#) for details.
- **inspect** <setting-key> Displays information about settings, such as the value, description, defining scope, dependencies, delegation chain, and related settings. See [Inspecting Settings](#) for details.

Command Line Options

System properties can be provided either as JVM options, or as SBT arguments, in both cases as **-Dprop=value**. The following properties influence SBT execution. Also see [sbt launcher](#).

Property

Values

Default

Meaning

`sbt.log.noformat`

Boolean

false

If true, disable ANSI color codes. Useful on build servers or terminals that do not support color.

`sbt.global.base`

Directory

`~/.sbt/0.13`

The directory containing global settings and plugins

sbt.ivy.home

Directory

~/ivy2

The directory containing the local Ivy repository and artifact cache

sbt.boot.directory

Directory

~/sbt/boot

Path to shared boot directory

sbt.main.class

String

xsbt.inc.debug

Boolean

false

sbt.extraClasspath

Classpath Entries

(jar files or directories) that are added to sbt's classpath. Note that the entries are delimited by comma, e.g.: entry1, entry2,... See also resource in the sbt launcher documentation.

sbt.version

Version

0.13.5

sbt version to use, usually taken from project/build.properties.

sbt.boot.properties

File

The path to find the sbt boot properties file. This can be a relative path, relative to the sbt base directory, the users home directory or the location of the sbt jar file, or it can be an absolute path or an absolute file URI.

sbt.override.build.repos

Boolean

false

If true, repositories configured in a build definition are ignored and the repositories configured for the launcher are used instead. See sbt.repository.config and the sbt launcher documentation.

sbt.repository.config

File

~/sbt/repositories

A file containing the repositories to use for the launcher. The format is the same as a [repositories] section for a sbt launcher configuration file. This setting is typically used in conjunction with setting `sbt.override.build.repos` to true (see previous row and the sbt launcher documentation).

Console Project

Description

The `consoleProject` task starts the Scala interpreter with access to your project definition and to `sbt`. Specifically, the interpreter is started up with these commands already executed:

```
import sbt._
import Process._
import Keys._
import <your-project-definition>._
import currentState._
import extracted._
import cpHelpers._
```

For example, running external processes with sbt's process library (to be included in the standard library in Scala 2.9):

```
> "tar -zcvf project-src.tar.gz src" !
> "find project -name *.jar" !
> "cat build.sbt" #| "grep version" #> new File("sbt-version") !
> "grep -r null src" #|| "echo null-free" !
> uri("http://databinder.net/dispatch/About").toURL #> file("About.html") !
```

`consoleProject` can be useful for creating and modifying your build in the same way that the Scala interpreter is normally used to explore writing code. Note that this gives you raw access to your build. Think about what you pass to `IO.delete`, for example.

Accessing settings

To get a particular setting, use the form:

```
> val value = (<key> in <scope>).eval
```

Examples

```
> IO.delete( (classesDirectory in Compile).eval )
```

Show current compile options:

```
> (scalacOptions in Compile).eval foreach println
```

Show additionally configured repositories.

```
> resolvers.eval foreach println
```

Evaluating tasks

To evaluate a task (and its dependencies), use the same form:

```
> val value = (<key> in <scope>).eval
```

Examples Show all repositories, including defaults.

```
> fullResolvers.eval foreach println
```

Show the classpaths used for compilation and testing:

```
> (fullClasspath in Compile).eval.files foreach println  
> (fullClasspath in Test).eval.files foreach println
```

State

The current `build State` is available as `currentState`. The contents of `currentState` are imported by default and can be used without qualification.

Examples Show the remaining commands to be executed in the build (more interesting if you invoke `consoleProject` like `; consoleProject ; clean ; compile`):

```
> remainingCommands
```

Show the number of currently registered commands:

```
> definedCommands.size
```

Cross-building

Introduction

Different versions of Scala can be binary incompatible, despite maintaining source compatibility. This page describes how to use `sbt` to build and publish your project against multiple versions of Scala and how to use libraries that have done the same.

Publishing Conventions

The underlying mechanism used to indicate which version of Scala a library was compiled against is to append `_<scala-version>` to the library's name. For Scala 2.10.0 and later, the binary version is used. For example, `dispatch` becomes `dispatch_2.8.1` for the variant compiled against Scala 2.8.1 and `dispatch_2.10` when compiled against 2.10.0, 2.10.0-M1 or any 2.10.x version. This fairly simple approach allows interoperability with users of Maven, Ant and other build tools.

The rest of this page describes how `sbt` handles this for you as part of cross-building.

Using Cross-Built Libraries

To use a library built against multiple versions of Scala, double the first `%` in an inline dependency to be `%%`. This tells `sbt` that it should append the current version of Scala being used to build the library to the dependency's name. For example:

```
libraryDependencies += "net.databinder" %% "dispatch" % "0.8.0"
```

A nearly equivalent, manual alternative for a fixed version of Scala is:

```
libraryDependencies += "net.databinder" % "dispatch_2.10" % "0.8.0"
```

or for Scala versions before 2.10:

```
libraryDependencies += "net.databinder" % "dispatch_2.8.1" % "0.8.0"
```


Cross-Building a Project

Define the versions of Scala to build against in the `crossScalaVersions` setting. Versions of Scala 2.8.0 or later are allowed. For example, in a `.sbt` build definition:

```
crossScalaVersions := Seq("2.8.2", "2.9.2", "2.10.0")
```

To build against all versions listed in `build.scala.versions`, prefix the action to run with `+`. For example:

```
> + package
```

A typical way to use this feature is to do development on a single Scala version (no `+` prefix) and then cross-build (using `+`) occasionally and when releasing. The ultimate purpose of `+` is to cross-publish your project. That is, by doing:

```
> + publish
```

you make your project available to users for different versions of Scala. See [Publishing](#) for more details on publishing your project.

In order to make this process as quick as possible, different output and managed dependency directories are used for different versions of Scala. For example, when building against Scala 2.10.0,

- `./target/` becomes `./target/scala_2.1.0/`
- `./lib_managed/` becomes `./lib_managed/scala_2.10/`

Packaged jars, wars, and other artifacts have `_<scala-version>` appended to the normal artifact ID as mentioned in the Publishing Conventions section above.

This means that the outputs of each build against each version of Scala are independent of the others. `sbt` will resolve your dependencies for each version separately. This way, for example, you get the version of Dispatch compiled against 2.8.1 for your 2.8.1 build, the version compiled against 2.10 for your 2.10.x builds, and so on. You can have fine-grained control over the behavior for different Scala versions by using the `cross` method on `ModuleID`. These are equivalent:

```
"a" % "b" % "1.0"  
"a" % "b" % "1.0" cross CrossVersion.Disabled
```

These are equivalent:

```
"a" %% "b" % "1.0"
"a" % "b" % "1.0" cross CrossVersion.binary
```

This overrides the defaults to always use the full Scala version instead of the binary Scala version:

```
"a" % "b" % "1.0" cross CrossVersion.full
```

This uses a custom function to determine the Scala version to use based on the binary Scala version:

```
"a" % "b" % "1.0" cross CrossVersion.binaryMapped {
  case "2.9.1" => "2.9.0" // remember that pre-2.10, binary=full
  case "2.10"  => "2.10.0" // useful if a%b was released with the old style
  case x      => x
}
```

This uses a custom function to determine the Scala version to use based on the full Scala version:

```
"a" % "b" % "1.0" cross CrossVersion.fullMapped {
  case "2.9.1" => "2.9.0"
  case x      => x
}
```

A custom function is mainly used when cross-building and a dependency isn't available for all Scala versions or it uses a different convention than the default.

As a final note, you can use ++ <version> to temporarily switch the Scala version currently being used to build. <version> should be either a version for Scala published to a repository, as in ++ 2.10.0 or the path to a Scala home directory, as in ++ /path/to/scala/home. See [Command Line Reference](#) for details.

Interacting with the Configuration System

Central to sbt is the new configuration system, which is designed to enable extensive customization. The goal of this page is to explain the general model behind the configuration system and how to work with it. The Getting Started Guide (see [.sbt files](#)) describes how to define settings; this page describes interacting with them and exploring them at the command line.

Selecting commands, tasks, and settings

A fully-qualified reference to a setting or task looks like:

```
{<build-uri>}<project-id>/config:inkey::key
```

This “scoped key” reference is used by commands like **last** and **inspect** and when selecting a task to run. Only **key** is usually required by the parser; the remaining optional pieces select the scope. These optional pieces are individually referred to as scope axes. In the above description, `{<build-uri>}` and `<project-id>/` specify the project axis, `config:` is the configuration axis, and `inkey` is the task-specific axis. Unspecified components are taken to be the current project (project axis) or auto-detected (configuration and task axes). An asterisk (*) is used to explicitly refer to the **Global** context, as in `*/*:key`.

Selecting the configuration In the case of an unspecified configuration (that is, when the `config:` part is omitted), if the key is defined in **Global**, that is selected. Otherwise, the first configuration defining the key is selected, where order is determined by the project definition’s `configurations` member. By default, this ordering is `compile`, `test`, ...

For example, the following are equivalent when run in a project root in the build in `/home/user/sample/`:

```
> compile
> compile:compile
> root/compile
> root/compile:compile
> {file:/home/user/sample/}root/compile:compile
```

As another example, `run` by itself refers to `compile:run` because there is no global `run` task and the first configuration searched, `compile`, defines a `run`. Therefore, to reference the `run` task for the **Test** configuration, the configuration axis must be specified like `test:run`. Some other examples that require the explicit `test:` axis:

```
> test:consoleQuick
> test:console
> test:doc
> test:package
```

Task-specific Settings Some settings are defined per-task. This is used when there are several related tasks, such as `package`, `packageSrc`, and `packageDoc`, in the same configuration (such as `compile` or `test`). For package tasks, their settings are the files to package, the options to use, and the output file to produce. Each package task should be able to have different values for these settings.

This is done with the task axis, which selects the task to apply a setting to. For example, the following prints the output jar for the different package tasks.

```
> package::artifactPath
[info] /home/user/sample/target/scala-2.8.1.final/demo_2.8.1-0.1.jar

> packageSrc::artifactPath
[info] /home/user/sample/target/scala-2.8.1.final/demo_2.8.1-0.1-src.jar

> packageDoc::artifactPath
[info] /home/user/sample/target/scala-2.8.1.final/demo_2.8.1-0.1-doc.jar

> test:package::artifactPath
[info] /home/user/sample/target/scala-2.8.1.final/root_2.8.1-0.1-test.jar
```

Note that a single colon `:` follows a configuration axis and a double colon `::` follows a task axis.

Discovering Settings and Tasks

This section discusses the `inspect` command, which is useful for exploring relationships between settings. It can be used to determine which setting should be modified in order to affect another setting, for example.

Value and Provided By The first piece of information provided by `inspect` is the type of a task or the value and type of a setting. The following section of output is labeled “Provided by”. This shows the actual scope where the setting is defined. For example,

```
> inspect libraryDependencies
[info] Setting: scala.collection.Seq[sbt.ModuleID] = List(org.scalaz:scalaz-core:6.0-SNAPSHOT
[info] Provided by:
[info] {file:/home/user/sample/}root/*:libraryDependencies
...
```

This shows that `libraryDependencies` has been defined on the current project (`{file:/home/user/sample/}root`) in the global configuration (`*`). For a task like `update`, the output looks like:

```
> inspect update
[info] Task: sbt.UpdateReport
[info] Provided by:
[info] {file:/home/user/sample/}root/*:update
...
```

Related Settings The “Related” section of `inspect` output lists all of the definitions of a key. For example,

```
> inspect compile
...
[info] Related:
[info] test:compile
```

This shows that in addition to the requested `compile:compile` task, there is also a `test:compile` task.

Dependencies Forward dependencies show the other settings (or tasks) used to define a setting (or task). Reverse dependencies go the other direction, showing what uses a given setting. `inspect` provides this information based on either the requested dependencies or the actual dependencies. Requested dependencies are those that a setting directly specifies. Actual settings are what those dependencies get resolved to. This distinction is explained in more detail in the following sections.

Requested Dependencies As an example, we’ll look at `console`:

```
> inspect console
...
[info] Dependencies:
[info] compile:console::fullClasspath
[info] compile:console::scalacOptions
[info] compile:console::initialCommands
[info] compile:console::cleanupCommands
[info] compile:console::compilers
[info] compile:console::taskTemporary-directory
[info] compile:console::scalaInstance
[info] compile:console::streams
...
```

This shows the inputs to the `console` task. We can see that it gets its classpath and options from `fullClasspath` and `scalacOptions` (for `console`). The information provided by the `inspect` command can thus assist in finding the right

setting to change. The convention for keys, like `console` and `fullClasspath`, is that the Scala identifier is camel case, while the String representation is lowercase and separated by dashes. The Scala identifier for a configuration is uppercase to distinguish it from tasks like `compile` and `test`. For example, we can infer from the previous example how to add code to be run when the Scala interpreter starts up:

```
> set initialCommands in Compile in console := "import mypackage._"
> console
...
import mypackage._
...
```

`inspect` showed that `console` used the setting `compile:console::initialCommands`. Translating the `initialCommands` string to the Scala identifier gives us `initialCommands`. `compile` indicates that this is for the main sources. `console::` indicates that the setting is specific to `console`. Because of this, we can set the initial commands on the `console` task without affecting the `consoleQuick` task, for example.

Actual Dependencies `inspect actual <scoped-key>` shows the actual dependency used. This is useful because delegation means that the dependency can come from a scope other than the requested one. Using `inspect actual`, we see exactly which scope is providing a value for a setting. Combining `inspect actual` with plain `inspect`, we can see the range of scopes that will affect a setting. Returning to the example in Requested Dependencies,

```
> inspect actual console
...
[info] Dependencies:
[info]   compile:scalacOptions
[info]   compile:fullClasspath
[info]   *:scalaInstance
[info]   */*:initialCommands
[info]   */*:cleanupCommands
[info]   */*:taskTemporaryDirectory
[info]   *:console::compilers
[info]   compile:console::streams
...
```

For `initialCommands`, we see that it comes from the global scope (`*/*`). Combining this with the relevant output from `inspect console`:

```
compile:console::initialCommands
```

we know that we can set `initialCommands` as generally as the global scope, as specific as the current project's `console` task scope, or anything in between. This means that we can, for example, set `initialCommands` for the whole project and will affect `console`:

```
> set initialCommands := "import mypackage._"
...
```

The reason we might want to set it here this is that other console tasks will use this value now. We can see which ones use our new setting by looking at the reverse dependencies output of `inspect actual`:

```
> inspect actual initialCommands
...
[info] Reverse dependencies:
[info]   test:console
[info]   compile:consoleQuick
[info]   compile:console
[info]   test:consoleQuick
[info]   *:consoleProject
...
```

We now know that by setting `initialCommands` on the whole project, we affect all console tasks in all configurations in that project. If we didn't want the initial commands to apply for `consoleProject`, which doesn't have our project's classpath available, we could use the more specific task axis:

```
> set initialCommands in console := "import mypackage._"
> set initialCommands in consoleQuick := "import mypackage._"
```

or configuration axis:

```
> set initialCommands in Compile := "import mypackage._"
> set initialCommands in Test := "import mypackage._"
```

The next part describes the `Delegates` section, which shows the chain of delegation for scopes.

Delegates A setting has a key and a scope. A request for a key in a scope A may be delegated to another scope if A doesn't define a value for the key. The delegation chain is well-defined and is displayed in the `Delegates` section of the `inspect` command. The `Delegates` section shows the order in which scopes are searched when a value is not defined for the requested key.

As an example, consider the initial commands for `console` again:

```

> inspect console::initialCommands
...
[info] Delegates:
[info] *:console::initialCommands
[info] *:initialCommands
[info] {./}/*:console::initialCommands
[info] {./}/*:initialCommands
[info] /*:console::initialCommands
[info] /*:initialCommands
...

```

This means that if there is no value specifically for `*:console::initialCommands`, the scopes listed under Delegates will be searched in order until a defined value is found.

Triggered Execution

You can make a command run when certain files change by prefixing the command with `~`. Monitoring is terminated when **enter** is pressed. This triggered execution is configured by the `watch` setting, but typically the basic settings `watchSources` and `pollInterval` are modified.

- `watchSources` defines the files for a single project that are monitored for changes. By default, a project watches resources and Scala and Java sources.
- `watchTransitiveSources` then combines the `watchSources` for the current project and all execution and classpath dependencies (see [.scala build definition](#) for details on interProject dependencies).
- `pollInterval` selects the interval between polling for changes in milliseconds. The default value is 500 ms.

Some example usages are described below.

Compile

The original use-case was continuous compilation:

```

> ~ test:compile

> ~ compile

```


Testing

You can use the triggered execution feature to run any command or task. One use is for test driven development, as suggested by Erick on the mailing list.

The following will poll for changes to your source code (main or test) and run `testOnly` for the specified test.

```
> ~ testOnly example.TestA
```

Running Multiple Commands

Occasionally, you may need to trigger the execution of multiple commands. You can use semicolons to separate the commands to be triggered.

The following will poll for source changes and run `clean` and `test`.

```
> ~ ;clean ;test
```

Scripts, REPL, and Dependencies

sbt has two alternative entry points that may be used to:

- Compile and execute a Scala script containing dependency declarations or other sbt settings
- Start up the Scala REPL, defining the dependencies that should be on the classpath

These entry points should be considered experimental. A notable disadvantage of these approaches is the startup time involved.

Setup

To set up these entry points, you can either use [conscript](#) or manually construct the startup scripts. In addition, there is a [setup script](#) for the script mode that only requires a JRE installed.

Setup with Conscript Install [conscript](#).

```
$ cs sbt/sbt --branch 0.12.0
```

This will create two scripts: `screpl` and `scalas`.

Manual Setup Duplicate your standard `sbt` script, which was set up according to [Setup](#), as `scalas` and `screpl` (or whatever names you like).

`scalas` is the script runner and should use `sbt.ScriptMain` as the main class, by adding the `-Dsbt.main.class=sbt.ScriptMain` parameter to the `java` command. Its command line should look like:

```
$ java -Dsbt.main.class=sbt.ScriptMain -Dsbt.boot.directory=/home/user/.sbt/boot -jar sbt-launch.jar
```

For the REPL runner `screpl`, use `sbt.ConsoleMain` as the main class:

```
$ java -Dsbt.main.class=sbt.ConsoleMain -Dsbt.boot.directory=/home/user/.sbt/boot -jar sbt-launch.jar
```

In each case, `/home/user/.sbt/boot` should be replaced with wherever you want sbt's boot directory to be; you might also need to give more memory to the JVM via `-Xms512M -Xmx1536M` or similar options, just like shown in [Setup](#).

Usage

sbt Script runner The script runner can run a standard Scala script, but with the additional ability to configure sbt. sbt settings may be embedded in the script in a comment block that opens with `/**`.

Example Copy the following script and make it executable. You may need to adjust the first line depending on your script name and operating system. When run, the example should retrieve Scala, the required dependencies, compile the script, and run it directly. For example, if you name it `dispatch_example.scala`, you would do on Unix:

```
chmod u+x dispatch_example.scala
./dispatch_example.scala

#!/usr/bin/env scalas
!#

/**
  scalaVersion := "2.9.0-1"

  libraryDependencies += Seq(
    "net.databinder" %% "dispatch-twitter" % "0.8.3",
    "net.databinder" %% "dispatch-http" % "0.8.3"
  )
*/
```

```

import dispatch.{ json, Http, Request }
import dispatch.twitter.Search
import json.{ Js, JsObject }

def process(param: JsObject) = {
  val Search.text(txt)      = param
  val Search.from_user(usr) = param
  val Search.created_at(time) = param

  "(" + time + ")" + usr + ": " + txt
}

Http.x((Search("#scala") lang "en") ~> (_ map process foreach println))

```

sbt REPL with dependencies The arguments to the REPL mode configure the dependencies to use when starting up the REPL. An argument may be either a jar to include on the classpath, a dependency definition to retrieve and put on the classpath, or a resolver to use when retrieving dependencies.

A dependency definition looks like:

```
organization%module%revision
```

Or, for a cross-built dependency:

```
organization%%module%revision
```

A repository argument looks like:

```
"id at url"
```

Example: To add the Sonatype snapshots repository and add Scalaz 7.0-SNAPSHOT to REPL classpath:

```
$ screpl "sonatype-releases at https://oss.sonatype.org/content/repositories/snapshots/" "org
```

This syntax was a quick hack. Feel free to improve it. The relevant class is [IvyConsole](#).

Understanding Incremental Recompilation

Compiling Scala code is slow, and sbt makes it often faster. By understanding how, you can even understand how to make compilation even faster. Modifying source files with many dependencies might require recompiling only those source files—which might take, say, 5 seconds—instead of all the dependencies—which might take, say, 2 minutes. Often you can control which will be your case and make development much faster by some simple coding practices.

In fact, improving Scala compilation times is one major goal of sbt, and conversely the speedups it gives are one of the major motivations to use it. A significant portion of sbt sources and development efforts deals with strategies for speeding up compilation.

To reduce compile times, sbt uses two strategies:

1. reduce the overhead for restarting Scalac;
2. implement smart and transparent strategies for incremental recompilation, so that only modified files and the needed dependencies are recompiled.
3. sbt runs Scalac always in the same virtual machine. If one compiles source code using sbt, keeps sbt alive, modifies source code and triggers a new compilation, this compilation will be faster because (part of) Scalac will have already been JIT-compiled. In the future, sbt will reintroduce support for reusing the same compiler instance, similarly to fsc.
4. When a source file `A.scala` is modified, sbt goes to great effort to recompile other source files depending on `A.scala` only if required - that is, only if the interface of `A.scala` was modified. With other build management tools (especially for Java, like ant), when a developer changes a source file in a non-binary-compatible way, he needs to manually ensure that dependencies are also recompiled - often by manually running the clean command to remove existing compilation output; otherwise compilation might succeed even when dependent class files might need to be recompiled. What is worse, the change to one source might make dependencies incorrect, but this is not discovered automatically: One might get a compilation success with incorrect source code. Since Scala compile times are so high, running clean is particularly undesirable.

By organizing your source code appropriately, you can minimize the amount of code affected by a change. sbt cannot determine precisely which dependencies have to be recompiled; the goal is to compute a conservative approximation, so that whenever a file must be recompiled, it will, even though we might recompile extra files.

sbt heuristics

sbt tracks source dependencies at the granularity of source files. For each source file, sbt tracks files which depend on it directly; if the **interface** of classes, objects or traits in a file changes, all files dependent on that source must be recompiled. At the moment sbt uses the following algorithm to calculate source files dependent on a given source file:

- dependencies introduced through inheritance are included *transitively*; a dependency is introduced through inheritance if a class/trait in one file inherits from a trait/class in another file
- all other direct dependencies are included; other dependencies are also called “member reference” dependencies because they are introduced by referring to a member (class, method, type, etc.) defined in some other source file

Here’s an example illustrating the definition above:

```
//A.scala
class A {
  def foo: Int = 123
}

//B.scala
class B extends A

//C.scala
class C extends B

//D.scala
class D(a: A)

//E.scala
class E(d: D)
```

There are the following dependencies through inheritance:

```
B.scala -> A.scala
C.scala -> B.scala
```

There are also the following member reference dependencies:

```
D.scala -> A.scala
E.scala -> D.scala
```

Now if the interface of `A.scala` is changed the following files will get invalidated: `B.scala`, `C.scala`, `D.scala`. Both `B.scala` and `C.scala` were included through transitive closure of inheritance dependencies. The `E.scala` was not included because `E.scala` doesn't depend directly on `A.scala`.

The distinction between dependencies by inheritance or member reference is a new feature in sbt 0.13 and is responsible for improved recompilation times in many cases where deep inheritance chains are not used extensively.

sbt does not instead track dependencies to source code at the granularity of individual output `.class` files, as one might hope. Doing so would be incorrect, because of some problems with sealed classes (see below for discussion).

Dependencies on binary files are different - they are tracked both on the `.class` level and on the source file level. Adding a new implementation of a sealed trait to source file `A` affects all clients of that sealed trait, and such dependencies are tracked at the source file level.

Different sources are moreover recompiled together; hence a compile error in one source implies that no bytecode is generated for any of those. When a lot of files need to be recompiled and the compile fix is not clear, it might be best to comment out the offending location (if possible) to allow other sources to be compiled, and then try to figure out how to fix the offending location—this way, trying out a possible solution to the compile error will take less time, say 5 seconds instead of 2 minutes.

What is included in the interface of a Scala class

It is surprisingly tricky to understand which changes to a class require recompiling its clients. The rules valid for Java are much simpler (even if they include some subtle points as well); trying to apply them to Scala will prove frustrating. Here is a list of a few surprising points, just to illustrate the ideas; this list is not intended to be complete.

1. Since Scala supports named arguments in method invocations, the name of method arguments are part of its interface.
2. Adding a method to a trait requires recompiling all implementing classes. The same is true for most changes to a method signature in a trait.
3. Calls to `super.methodName` in traits are resolved to calls to an abstract method called `fullyQualifiedTraitName$$$super$methodName`; such methods only exist if they are used. Hence, adding the first call to `super.methodName` for a specific `methodName` changes the interface. At present, this is not yet handled—see gh-466.
4. `sealed` hierarchies of case classes allow to check exhaustiveness of pattern matching. Hence pattern matches using case classes must depend on the complete hierarchy - this is one reason why dependencies cannot be easily tracked at the class level (see Scala issue [SI-2559](#) for an example.)

Debugging an interface representation If you see spurious incremental recompilations or you want understand what changes to an extracted interface cause incremental recompilation then sbt 0.13 has the right tools for that.

In order to debug the interface representation and its changes as you modify and recompile source code you need to do two things:

1. Enable incremental compiler's `apiDebug` option.
2. Add [diff-utils library](#) to sbt's classpath. Check documentation of `sbt.extraClasspath` system property in the [Command-Line-Reference](#).

warning

Enabling the `apiDebug` option increases significantly

memory consumption and degrades performance of the incremental compiler. The underlying reason is that in order to produce meaningful debugging information about interface differences incremental compiler has to retain the full representation of the interface instead of just hash sum as it does by default.

Keep this option enabled when you are debugging incremental compiler problem only.

Below is complete transcript which shows how to enable interface debugging in your project. First, we download the `diffutils` jar and pass it to sbt:

```
curl -O https://java-diff-utils.googlecode.com/files/diffutils-1.2.1.jar
sbt -Dsbt.extraClasspath=diffutils-1.2.1.jar
[info] Loading project definition from /Users/grek/tmp/sbt-013/project
[info] Set current project to sbt-013 (in build file:/Users/grek/tmp/sbt-013/)
> set incOptions := incOptions.value.copy(apiDebug = true)
[info] Defining *:incOptions
[info] The new value will be used by compile:incCompileSetup, test:incCompileSetup
[info] Reapplying settings...
[info] Set current project to sbt-013 (in build file:/Users/grek/tmp/sbt-013/)
```

Let's suppose you have the following source code in `Test.scala`:

```
class A {
  def b: Int = 123
}
```

compile it and then change the `Test.scala` file so it looks like:

```
class A {
  def b: String = "abc"
}
```

and run `compile` task again. Now if you run `last compile` you should see the following lines in the debugging log

```
> last compile
[...]
[debug] Detected a change in a public API:
[debug] --- /Users/grek/tmp/sbt-013/Test.scala
[debug] +++ /Users/grek/tmp/sbt-013/Test.scala
[debug] @@ -23,7 +23,7 @@
[debug] ~inherited~ final def ##(): scala.this#Int
[debug] ~inherited~ final def synchronized[ java.lang.Object.T0 >: scala.this#Nothing <: scala
[debug] ~inherited~ final def $isInstanceOf[ java.lang.Object.T0 >: scala.this#Nothing <: scal
[debug] ~inherited~ final def $asInstanceOf[ java.lang.Object.T0 >: scala.this#Nothing <: scal
[debug] def <init>(): this#A
[debug] -def b: scala.this#Int
[debug] +def b: java.lang.this#String
[debug] }
```

You can see an unified diff of two interface textual representations. As you can see, the incremental compiler detected a change to the return type of `b` method.

How to take advantage of sbt heuristics

The heuristics used by sbt imply the following user-visible consequences, which determine whether a change to a class affects other classes.

XXX Please note that this part of the documentation is a first draft; part of the strategy might be unsound, part of it might be not yet implemented.

1. Adding, removing, modifying **private** methods does not require recompilation of client classes. Therefore, suppose you add a method to a class with a lot of dependencies, and that this method is only used in the declaring class; marking it private will prevent recompilation of clients. However, this only applies to methods which are not accessible to other classes, hence methods marked with `private[this]`; methods which are private to a package, marked with `private[name]`, are part of the API.
2. Modifying the interface of a non-private method requires recompiling all clients, even if the method is not used.
3. Modifying one class does require recompiling dependencies of other classes defined in the same file (unlike said in a previous version of this guide). Hence separating different classes in different source files might reduce recompilations.

4. Adding a method which did not exist requires recompiling all clients, counterintuitively, due to complex scenarios with implicit conversions. Hence in some cases you might want to start implementing a new method in a separate, new class, complete the implementation, and then cut-n-paste the complete implementation back into the original source.
5. Changing the implementation of a method should *not* affect its clients, unless the return type is inferred, and the new implementation leads to a slightly different type being inferred. Hence, annotating the return type of a non-private method explicitly, if it is more general than the type actually returned, can reduce the code to be recompiled when the implementation of such a method changes. (Explicitly annotating return types of a public API is a good practice in general.)

All the above discussion about methods also applies to fields and members in general; similarly, references to classes also extend to objects and traits.

Why changing the implementation of a method might affect clients, and why type annotations help This section explains why relying on type inference for return types of public methods is not always appropriate. However this is an important design issue, so we cannot give fixed rules. Moreover, this change is often invasive, and reducing compilation times is not often a good enough motivation. That is why we discuss also some of the implications from the point of view of binary compatibility and software engineering.

Consider the following source file `A.scala`:

```
import java.io._
object A {
  def openFiles(list: List[File]) =
    list.map(name => new FileWriter(name))
}
```

Let us now consider the public interface of trait `A`. Note that the return type of method `openFiles` is not specified explicitly, but computed by type inference to be `List[FileWriter]`. Suppose that after writing this source code, we introduce client code and then modify `A.scala` as follows:

```
import java.io._
object A {
  def openFiles(list: List[File]) =
    Vector(list.map(name => new BufferedWriter(new FileWriter(name)))): _*
}
```

Type inference will now compute as result type `Vector[BufferedWriter]`; in other words, changing the implementation lead to a change of the public interface, with two undesirable consequences:

1. Concerning our topic, client code needs to be recompiled, since changing the return type of a method, in the JVM, is a binary-incompatible interface change.
2. If our component is a released library, using our new version requires recompiling all client code, changing the version number, and so on. Often not good, if you distribute a library where binary compatibility becomes an issue.
3. More in general, client code might now even be invalid. The following code will for instance become invalid after the change:

```
val res: List[FileWriter] = A.openFiles(List(new File("foo.input")))
```

Also the following code will break:

```
val a: Seq[Writer] = new BufferedWriter(new FileWriter("bar.input"))
A.openFiles(List(new File("foo.input")))
```

How can we avoid these problems?

Of course, we cannot solve them in general: if we want to alter the interface of a module, breakage might result. However, often we can remove *implementation details* from the interface of a module. In the example above, for instance, it might well be that the intended return type is more general - namely `Seq[Writer]`. It might also not be the case - this is a design choice to be decided on a case-by-case basis. In this example I will assume however that the designer chooses `Seq[Writer]`, since it is a reasonable choice both in the above simplified example and in a real-world extension of the above code.

The client snippets above will now become

```
val res: Seq[Writer] =
  A.openFiles(List(new File("foo.input")))

val a: Seq[Writer] =
  new BufferedWriter(new FileWriter("bar.input")) +:
  A.openFiles(List(new File("foo.input")))
```

XXX the rest of the section must be reintegrated or dropped: In general, changing the return type of a method might be source-compatible, for instance if the new type is more specific, or if it is less specific, but still more specific than the type required by clients (note however that making the type more specific might still invalidate clients in non-trivial scenarios involving for instance type inference or implicit conversions—for a more specific type, too many implicit conversions might be available, leading to ambiguity); however, the bytecode

for a method call includes the return type of the invoked method, hence the client code needs to be recompiled.

Hence, adding explicit return types on classes with many dependencies might reduce the occasions where client code needs to be recompiled. Moreover, this is in general a good development practice when interface between different modules become important—specifying such interface documents the intended behavior and helps ensuring binary compatibility, which is especially important when the exposed interface is used by other software component.

Why adding a member requires recompiling existing clients In Java adding a member does not require recompiling existing valid source code. The same should seemingly hold also in Scala, but this is not the case: implicit conversions might enrich class `Foo` with method `bar` without modifying class `Foo` itself (see discussion in issue gh-288 - XXX integrate more). However, if another method `bar` is introduced in class `Foo`, this method should be used in preference to the one added through implicit conversions. Therefore any class depending on `Foo` should be recompiled. One can imagine more fine-grained tracking of dependencies, but this is currently not implemented.

Further references

The incremental compilation logic is implemented in <https://github.com/sbt/sbt/blob/0.13/compile/inc/src/main/scala/inc/Incremental.scala>. Some related documentation for sbt 0.7 is available at: <https://code.google.com/p/simple-build-tool/wiki/ChangeDetectionAndTesting>. Some discussion on the incremental recompilation policies is available in issue gh-322 and gh-288.

Configuration

This part of the documentation has pages documenting particular sbt topics in detail. Before reading anything in here, you will need the information in the [Getting Started Guide](#) as a foundation.

Classpaths, sources, and resources

This page discusses how sbt builds up classpaths for different actions, like `compile`, `run`, and `test` and how to override or augment these classpaths.

Basics

In sbt 0.10 and later, classpaths now include the Scala library and (when declared as a dependency) the Scala compiler. Classpath-related settings

and tasks typically provide a value of type `Classpath`. This is an alias for `Seq[Attributed[File]]`. `Attributed` is a type that associates a heterogeneous map with each classpath entry. Currently, this allows sbt to associate the `Analysis` resulting from compilation with the corresponding classpath entry and for managed entries, the `ModuleID` and `Artifact` that defined the dependency.

To explicitly extract the raw `Seq[File]`, use the `files` method implicitly added to `Classpath`:

```
val cp: Classpath = ...
val raw: Seq[File] = cp.files
```

To create a `Classpath` from a `Seq[File]`, use `classpath` and to create an `Attributed[File]` from a `File`, use `Attributed.blank`:

```
val raw: Seq[File] = ...
val cp: Classpath = raw.classpath

val rawFile: File = ..
val af: Attributed[File] = Attributed.blank(rawFile)
```

Unmanaged vs managed Classpaths, sources, and resources are separated into two main categories: unmanaged and managed. Unmanaged files are manually created files that are outside of the control of the build. They are the inputs to the build. Managed files are under the control of the build. These include generated sources and resources as well as resolved and retrieved dependencies and compiled classes.

Tasks that produce managed files should be inserted as follows:

```
sourceGenerators in Compile +=
  generate( (sourceManaged in Compile).value / "some_directory")
```

In this example, `generate` is some function of type `File => Seq[File]` that actually does the work. So, we are appending a new task to the list of main source generators (`sourceGenerators in Compile`).

To insert a named task, which is the better approach for plugins:

```
val mySourceGenerator = taskKey[Seq[File]](...)

mySourceGenerator in Compile :=
  generate( (sourceManaged in Compile).value / "some_directory")

sourceGenerators in Compile += (mySourceGenerator in Compile).task
```

The `task` method is used to refer to the actual task instead of the result of the task.

For resources, there are similar keys `resourceGenerators` and `resourceManaged`.

Excluding source files by name The project base directory is by default a source directory in addition to `src/main/scala`. You can exclude source files by name (`butler.scala` in the example below) like:

```
excludeFilter in unmanagedSources := "butler.scala"
```

Read more on [How to exclude .scala source file in project folder - Google Groups](#)

External vs internal Classpaths are also divided into internal and external dependencies. The internal dependencies are inter-project dependencies. These effectively put the outputs of one project on the classpath of another project.

External classpaths are the union of the unmanaged and managed classpaths.

Keys For classpaths, the relevant keys are:

- `unmanagedClasspath`
- `managedClasspath`
- `externalDependencyClasspath`
- `internalDependencyClasspath`

For sources:

- `unmanagedSources` These are by default built up from `unmanagedSourceDirectories`, which consists of `scalaSource` and `javaSource`.
- `managedSources` These are generated sources.
- `sources` Combines `managedSources` and `unmanagedSources`.
- `sourceGenerators` These are tasks that generate source files. Typically, these tasks will put sources in the directory provided by `sourceManaged`.

For resources

- `unmanagedResources` These are by default built up from `unmanagedResourceDirectories`, which by default is `resourceDirectory`, excluding files matched by `defaultExcludes`.
- `managedResources` By default, this is empty for standard projects. `sbt` plugins will have a generated descriptor file here.

- **resourceGenerators** These are tasks that generate resource files. Typically, these tasks will put resources in the directory provided by `resourceManaged`.

Use the [inspect command](#) for more details.

See also a related [StackOverflow answer](#).

Example You have a standalone project which uses a library that loads `xxx.properties` from classpath at run time. You put `xxx.properties` inside directory “`config`”. When you run “`sbt run`”, you want the directory to be in classpath.

```
unmanagedClasspath in Runtime += baseDirectory.value / "config"
```

Compiler Plugin Support

There is some special support for using compiler plugins. You can set `autoCompilerPlugins` to `true` to enable this functionality.

```
autoCompilerPlugins := true
```

To use a compiler plugin, you either put it in your unmanaged library directory (`lib/` by default) or add it as managed dependency in the `plugin` configuration. `addCompilerPlugin` is a convenience method for specifying `plugin` as the configuration for a dependency:

```
addCompilerPlugin("org.scala-tools.sxr" %% "sxr" % "0.3.0")
```

The `compile` and `testCompile` actions will use any compiler plugins found in the `lib` directory or in the `plugin` configuration. You are responsible for configuring the plugins as necessary. For example, Scala X-Ray requires the extra option:

```
// declare the main Scala source directory as the base directory
scalacOptions :=
  scalacOptions.value :+ ("-Psxr:base-directory:" + (scalaSource in Compile).value.getAbsolute
```

You can still specify compiler plugins manually. For example:

```
scalacOptions += "-Xplugin:<path-to-sxr>/sxr-0.3.0.jar"
```

Continuations Plugin Example

Support for continuations in Scala 2.8 is implemented as a compiler plugin. You can use the compiler plugin support for this, as shown here.

```
autoCompilerPlugins := true

addCompilerPlugin("org.scala-lang.plugins" % "continuations" % "2.8.1")

scalacOptions += "-P:continuations:enable"
```

Version-specific Compiler Plugin Example

Adding a version-specific compiler plugin can be done as follows:

```
autoCompilerPlugins := true

libraryDependencies +=
  compilerPlugin("org.scala-lang.plugins" % "continuations" % scalaVersion.value)

scalacOptions += "-P:continuations:enable"
```

Configuring Scala

sbt needs to obtain Scala for a project and it can do this automatically or you can configure it explicitly. The Scala version that is configured for a project will compile, run, document, and provide a REPL for the project code. When compiling a project, sbt needs to run the Scala compiler as well as provide the compiler with a classpath, which may include several Scala jars, like the reflection jar.

Automatically managed Scala

The most common case is when you want to use a version of Scala that is available in a repository. The only required configuration is the Scala version you want to use. For example,

```
scalaVersion := "2.10.0"
```

This will retrieve Scala from the repositories configured via the **resolvers** setting. It will use this version for building your project: compiling, running, scaladoc, and the REPL.

Configuring the scala-library dependency By default, the standard Scala library is automatically added as a dependency. If you want to configure it differently than the default or you have a project with only Java sources, set:

```
autoScalaLibrary := false
```

In order to compile Scala sources, the Scala library needs to be on the classpath. When `autoScalaLibrary` is true, the Scala library will be on all classpaths: test, runtime, and compile. Otherwise, you need to add it like any other dependency. For example, the following dependency definition uses Scala only for tests:

```
autoScalaLibrary := false
```

```
libraryDependencies += "org.scala-lang" % "scala-library" % scalaVersion.value % "test"
```

Configuring additional Scala dependencies When using a Scala dependency other than the standard library, add it as a normal managed dependency. For example, to depend on the Scala compiler,

```
libraryDependencies += "org.scala-lang" % "scala-compiler" % scalaVersion.value
```

Note that this is necessary regardless of the value of the `autoScalaLibrary` setting described in the previous section.

Configuring Scala tool dependencies In order to compile Scala code, run scaladoc, and provide a Scala REPL, sbt needs the `scala-compiler` jar. This should not be a normal dependency of the project, so sbt adds a dependency on `scala-compiler` in the special, private `scala-tool` configuration. It may be desirable to have more control over this in some situations. Disable this automatic behavior with the `managedScalaInstance` key:

```
managedScalaInstance := false
```

This will also disable the automatic dependency on `scala-library`. If you do not need the Scala compiler for anything (compiling, the REPL, scaladoc, etc...), you can stop here. sbt does not need an instance of Scala for your project in that case. Otherwise, sbt will still need access to the jars for the Scala compiler for compilation and other tasks. You can provide them by either declaring a dependency in the `scala-tool` configuration or by explicitly defining `scalaInstance`.

In the first case, add the `scala-tool` configuration and add a dependency on `scala-compiler` in this configuration. The organization is not important, but sbt needs the module name to be `scala-compiler` and `scala-library` in order to handle those jars appropriately. For example,


```

managedScalaInstance := false

// Add the configuration for the dependencies on Scala tool jars
// You can also use a manually constructed configuration like:
//   config("scala-tool").hide
ivyConfigurations += Configurations.ScalaTool

// Add the usual dependency on the library as well on the compiler in the
// 'scala-tool' configuration
libraryDependencies += Seq(
  "org.scala-lang" % "scala-library" % scalaVersion.value,
  "org.scala-lang" % "scala-compiler" % scalaVersion.value % "scala-tool"
)

```

In the second case, directly construct a value of type `ScalaInstance`, typically using a method in the `companion object`, and assign it to `scalaInstance`. You will also need to add the `scala-library` jar to the classpath to compile and run Scala sources. For example,

```

managedScalaInstance := false

scalaInstance := ...

unmanagedJars in Compile += scalaInstance.value.libraryJar

```

Switching to a local Scala version To use a locally built Scala version, configure Scala home as described in the following section. Scala will still be resolved as before, but the jars will come from the configured Scala home directory.

Using Scala from a local directory

The result of building Scala from source is a Scala home directory `<base>/build/pack/` that contains a subdirectory `lib/` containing the Scala library, compiler, and other jars. The same directory layout is obtained by downloading and extracting a Scala distribution. Such a Scala home directory may be used as the source for jars by setting `scalaHome`. For example,

```
scalaHome := Some(file("/home/user/scala-2.10/"))
```

By default, `lib/scala-library.jar` will be added to the unmanaged classpath and `lib/scala-compiler.jar` will be used to compile Scala sources and provide

a Scala REPL. No managed dependency is recorded on `scala-library`. This means that Scala will only be resolved from a repository if you explicitly define a dependency on Scala or if Scala is depended on indirectly via a dependency. In these cases, the artifacts for the resolved dependencies will be substituted with jars in the Scala home `lib/` directory.

Mixing with managed dependencies As an example, consider adding a dependency on `scala-reflect` when `scalaHome` is configured:

```
scalaHome := Some(file("/home/user/scala-2.10/"))

libraryDependencies += "org.scala-lang" % "scala-reflect" % scalaVersion.value
```

This will be resolved as normal, except that sbt will see if `/home/user/scala-2.10/lib/scala-reflect.jar` exists. If it does, that file will be used in place of the artifact from the managed dependency.

Using unmanaged dependencies only Instead of adding managed dependencies on Scala jars, you can directly add them. The `scalaInstance` task provides structured access to the Scala distribution. For example, to add all jars in the Scala home `lib/` directory,

```
scalaHome := Some(file("/home/user/scala-2.10/"))

unmanagedJars in Compile += scalaInstance.value.jars
```

To add only some jars, filter the jars from `scalaInstance` before adding them.

sbt's Scala version

sbt needs Scala jars to run itself since it is written in Scala. sbt uses that same version of Scala to compile the build definitions that you write for your project because they use sbt APIs. This version of Scala is fixed for a specific sbt release and cannot be changed. For sbt 0.13.5, this version is Scala 2.10.3. Because this Scala version is needed before sbt runs, the repositories used to retrieve this version are configured in the sbt [launcher](#).

Forking

By default, the `run` task runs in the same JVM as sbt. Forking is required under [certain circumstances](#), however. Or, you might want to fork Java processes when implementing new tasks.

By default, a forked process uses the same Java and Scala versions being used for the build and the working directory and JVM options of the current process. This page discusses how to enable and configure forking for both **run** and **test** tasks. Each kind of task may be configured separately by scoping the relevant keys as explained below.

Enable forking

The **fork** setting controls whether forking is enabled (true) or not (false). It can be set in the **run** scope to only fork **run** commands or in the **test** scope to only fork **test** commands.

To fork all test tasks (**test**, **testOnly**, and **testQuick**) and run tasks (**run**, **runMain**, **test:run**, and **test:runMain**),

```
fork := true
```

To enable forking run tasks only, set **fork** to **true** in the **run** scope.

```
fork in run := true
```

To only fork **test:run** and **test:runMain**:

```
fork in (Test, run) := true
```

Similarly, set **fork in (Compile,run) := true** to only fork the main **run** tasks. **run** and **runMain** share the same configuration and cannot be configured separately.

To enable forking all **test** tasks only, set **fork** to **true** in the **test** scope:

```
fork in test := true
```

See [Testing](#) for more control over how tests are assigned to JVMs and what options to pass to each group.

Change working directory

To change the working directory when forked, set **baseDirectory in run** or **baseDirectory in test**:

```

// sets the working directory for all `run`-like tasks
baseDirectory in run := file("/path/to/working/directory/")

// sets the working directory for `run` and `runMain` only
baseDirectory in (Compile,run) := file("/path/to/working/directory/")

// sets the working directory for `test:run` and `test:runMain` only
baseDirectory in (Test,run) := file("/path/to/working/directory/")

// sets the working directory for `test`, `testQuick`, and `testOnly`
baseDirectory in test := file("/path/to/working/directory/")

```

Forked JVM options

To specify options to be provided to the forked JVM, set `javaOptions`:

```
javaOptions in run += "-Xmx8G"
```

or specify the configuration to affect only the main or test `run` tasks:

```
javaOptions in (Test,run) += "-Xmx8G"
```

or only affect the `test` tasks:

```
javaOptions in test += "-Xmx8G"
```

Java Home

Select the Java installation to use by setting the `javaHome` directory:

```
javaHome := Some(file("/path/to/jre/"))
```

Note that if this is set globally, it also sets the Java installation used to compile Java sources. You can restrict it to running only by setting it in the `run` scope:

```
javaHome in run := Some(file("/path/to/jre/"))
```

As with the other settings, you can specify the configuration to affect only the main or test `run` tasks or just the `test` tasks.

Configuring output

By default, forked output is sent to the Logger, with standard output logged at the **Info** level and standard error at the **Error** level. This can be configured with the `outputStrategy` setting, which is of type [OutputStrategy](#).

```
// send output to the build's standard output and error
outputStrategy := Some(StdoutOutput)

// send output to the provided OutputStream `someStream`
outputStrategy := Some(CustomOutput(someStream: OutputStream))

// send output to the provided Logger `log` (unbuffered)
outputStrategy := Some(LoggedOutput(log: Logger))

// send output to the provided Logger `log` after the process terminates
outputStrategy := Some(BufferedOutput(log: Logger))
```

As with other settings, this can be configured individually for main or test run tasks or for test tasks.

Configuring Input

By default, the standard input of the sbt process is not forwarded to the forked process. To enable this, configure the `connectInput` setting:

```
connectInput in run := true
```

Direct Usage

To fork a new Java process, use the [Fork API](#). The values of interest are `Fork.java`, `Fork.javac`, `Fork.scala`, and `Fork.scalac`. These are of type [Fork](#) and provide `apply` and `fork` methods. For example, to fork a new Java process, :

```
val options = ForkOptions(...)
val arguments: Seq[String] = ...
val mainClass: String = ...
val exitCode: Int = Fork.java(options, mainClass +: arguments)
```

[ForkOptions](#) defines the Java installation to use, the working directory, environment variables, and more. For example, :

```

val cwd: File = ...
val javaDir: File = ...
val options = ForkOptions(
  envVars = Map("KEY" -> "value"),
  workingDirectory = Some(cwd),
  javaHome = Some(javaDir)
)

```

Global Settings

Basic global configuration file

Settings that should be applied to all projects can go in `~/.sbt/0.13/global.sbt` (or any file in `~/.sbt/0.13` with a `.sbt` extension). Plugins that are defined globally in `~/.sbt/0.13/plugins/` are available to these settings. For example, to change the default `shellPrompt` for your projects:

```
~/.sbt/0.13/global.sbt
```

```

shellPrompt := { state =>
  "sbt (%s)> ".format(Project.extract(state).currentProject.id)
}

```

Global Settings using a Global Plugin

The `~/.sbt/0.13/plugins/` directory is a global plugin project. This can be used to provide global commands, plugins, or other code.

To add a plugin globally, create `~/.sbt/0.13/plugins/build.sbt` containing the dependency definitions. For example:

```
addSbtPlugin("org.example" % "plugin" % "1.0")
```

To change the default `shellPrompt` for every project using this approach, create a local plugin `~/.sbt/0.13/plugins/ShellPrompt.scala`:

```

import sbt._
import Keys._

object ShellPrompt extends Plugin {
  override def settings = Seq(
    shellPrompt := { state =>
      "sbt (%s)> ".format(Project.extract(state).currentProject.id) }
  )
}

```

The `~/.sbt/0.13/plugins/` directory is a full project that is included as an external dependency of every plugin project. In practice, settings and code defined here effectively work as if they were defined in a project's `project/` directory. This means that `~/.sbt/0.13/plugins/` can be used to try out ideas for plugins such as shown in the `shellPrompt` example.

Java Sources

sbt has support for compiling Java sources with the limitation that dependency tracking is limited to the dependencies present in compiled class files.

Usage

- `compile` will compile the sources under `src/main/java` by default.
- `testCompile` will compile the sources under `src/test/java` by default.

Pass options to the Java compiler by setting `javacOptions`:

```
javacOptions += "-g:none"
```

As with options for the Scala compiler, the arguments are not parsed by sbt. Multi-element options, such as `-source 1.5`, are specified like:

```
javacOptions ++= Seq("-source", "1.5")
```

You can specify the order in which Scala and Java sources are built with the `compileOrder` setting. Possible values are from the `CompileOrder` enumeration: `Mixed`, `JavaThenScala`, and `ScalaThenJava`. If you have circular dependencies between Scala and Java sources, you need the default, `Mixed`, which passes both Java and Scala sources to `scalac` and then compiles the Java sources with `javac`. If you do not have circular dependencies, you can use one of the other two options to speed up your build by not passing the Java sources to `scalac`. For example, if your Scala sources depend on your Java sources, but your Java sources do not depend on your Scala sources, you can do:

```
compileOrder := CompileOrder.JavaThenScala
```

To specify different orders for main and test sources, scope the setting by configuration:

```
// Java then Scala for main sources
compileOrder in Compile := CompileOrder.JavaThenScala

// allow circular dependencies for test sources
compileOrder in Test := CompileOrder.Mixed
```

Note that in an incremental compilation setting, it is not practical to ensure complete isolation between Java sources and Scala sources because they share the same output directory. So, previously compiled classes not involved in the current recompilation may be picked up. A clean compile will always provide full checking, however.

By default, sbt includes `src/main/scala` and `src/main/java` in its list of unmanaged source directories. For Java-only projects, the unnecessary Scala directories can be ignored by modifying `unmanagedSourceDirectories`:

```
// Include only src/main/java in the compile configuration
unmanagedSourceDirectories in Compile := (javaSource in Compile).value :: Nil

// Include only src/test/java in the test configuration
unmanagedSourceDirectories in Test := (javaSource in Test).value :: Nil
```

However, there should not be any harm in leaving the Scala directories if they are empty.

Mapping Files

Tasks like `package`, `packageSrc`, and `packageDoc` accept mappings of type `Seq[(File, String)]` from an input file to the path to use in the resulting artifact (jar). Similarly, tasks that copy files accept mappings of type `Seq[(File, File)]` from an input file to the destination file. There are some methods on [PathFinder](#) and [Path](#) that can be useful for constructing the `Seq[(File, String)]` or `Seq[(File, File)]` sequences.

A common way of making this sequence is to start with a `PathFinder` or `Seq[File]` (which is implicitly convertible to `PathFinder`) and then call the `pair` method. See the [PathFinder](#) API for details, but essentially this method accepts a function `File => Option[String]` or `File => Option[File]` that is used to generate mappings.

Relative to a directory

The `Path.relativeTo` method is used to map a `File` to its path `String` relative to a base directory or directories. The `relativeTo` method accepts a base

directory or sequence of base directories to relativize an input file against. The first directory that is an ancestor of the file is used in the case of a sequence of base directories.

For example:

```
import Path.relativeTo
val files: Seq[File] = file("/a/b/C.scala") :: Nil
val baseDirectories: Seq[File] = file("/a") :: Nil
val mappings: Seq[(File,String)] = files pair relativeTo(baseDirectories)

val expected = (file("/a/b/C.scala") -> "b/C.scala") :: Nil
assert( mappings == expected )
```

Rebase

The `Path.rebase` method relativizes an input file against one or more base directories (the first argument) and then prepends a base String or File (the second argument) to the result. As with `relativeTo`, the first base directory that is an ancestor of the input file is used in the case of multiple base directories.

For example, the following demonstrates building a `Seq[(File, String)]` using `rebase`:

```
import Path.rebase
val files: Seq[File] = file("/a/b/C.scala") :: Nil
val baseDirectories: Seq[File] = file("/a") :: Nil
val mappings: Seq[(File,String)] = files pair rebase(baseDirectories, "pre/")

val expected = (file("/a/b/C.scala") -> "pre/b/C.scala" ) :: Nil
assert( mappings == expected )
```

Or, to build a `Seq[(File, File)]`:

```
import Path.rebase
val files: Seq[File] = file("/a/b/C.scala") :: Nil
val baseDirectories: Seq[File] = file("/a") :: Nil
val newBase: File = file("/new/base")
val mappings: Seq[(File,File)] = files pair rebase(baseDirectories, newBase)

val expected = (file("/a/b/C.scala") -> file("/new/base/b/C.scala") ) :: Nil
assert( mappings == expected )
```

Flatten

The `Path.flat` method provides a function that maps a file to the last component of the path (its name). For a File to File mapping, the input file is mapped to a file with the same name in a given target directory. For example:

```
import Path.flat
val files: Seq[File] = file("/a/b/C.scala") :: Nil
val mappings: Seq[(File,String)] = files pair flat

val expected = (file("/a/b/C.scala") -> "C.scala") :: Nil
assert( mappings == expected )
```

To build a `Seq[(File, File)]` using `flat`:

```
import Path.flat
val files: Seq[File] = file("/a/b/C.scala") :: Nil
val newBase: File = file("/new/base")
val mappings: Seq[(File,File)] = files pair flat(newBase)

val expected = (file("/a/b/C.scala") -> file("/new/base/C.scala")) :: Nil
assert( mappings == expected )
```

Alternatives

To try to apply several alternative mappings for a file, use `|`, which is implicitly added to a function of type `A => Option[B]`. For example, to try to relativize a file against some base directories but fall back to flattening:

```
import Path.relativeTo
val files: Seq[File] = file("/a/b/C.scala") :: file("/zzz/D.scala") :: Nil
val baseDirectories: Seq[File] = file("/a") :: Nil
val mappings: Seq[(File,String)] = files pair ( relativeTo(baseDirectories) | flat )

val expected = (file("/a/b/C.scala") -> "b/C.scala") :: (file("/zzz/D.scala") -> "D.scala")
assert( mappings == expected )
```

Local Scala

To use a locally built Scala version, define the `scalaHome` setting, which is of type `Option[File]`. This Scala version will only be used for the build and not for sbt, which will still use the version it was compiled against.

Example:

```
scalaHome := Some(file("/path/to/scala"))
```

Using a local Scala version will override the `scalaVersion` setting and will not work with [cross building](#).

sbt reuses the class loader for the local Scala version. If you recompile your local Scala version and you are using sbt interactively, run

```
> reload
```

to use the new compilation results.

Macro Projects

Introduction

Some common problems arise when working with macros.

1. The current macro implementation in the compiler requires that macro implementations be compiled before they are used. The solution is typically to put the macros in a subproject or in their own configuration.
2. Sometimes the macro implementation should be distributed with the main code that uses them and sometimes the implementation should not be distributed at all.

The rest of the page shows example solutions to these problems.

Defining the Project Relationships

The macro implementation will go in a subproject in the `macro/` directory. The main project in the project's base directory will depend on this subproject and use the macro. This configuration is shown in the following build definition. `project/Build.scala`:

```
import sbt._
import Keys._

object MacroBuild extends Build {
  lazy val main = Project("main", file(".")) dependsOn(macroSub)
  lazy val macroSub = Project("macro", file("macro")) settings(
    libraryDependencies += "org.scala-lang" % "scala-reflect" % scalaVersion.value
  )
}
```

This specifies that the macro implementation goes in `macro/src/main/scala/` and tests go in `macro/src/test/scala/`. It also shows that we need a dependency on the compiler for the macro implementation. As an example macro, we'll use `desugar` from [macrocosm](#). `macro/src/main/scala/demo/Demo.scala`:

```
package demo

import language.experimental.macros
import scala.reflect.macros.Context

object Demo {

  // Returns the tree of `a` after the typer, printed as source code.
  def desugar(a: Any): String = macro desugarImpl

  def desugarImpl(c: Context)(a: c.Expr[Any]) = {
    import c.universe._

    val s = show(a.tree)
    c.Expr(
      Literal(Constant(s))
    )
  }
}
```

`macro/src/test/scala/demo/Usage.scala`:

```
package demo

object Usage {
  def main(args: Array[String]) {
    val s = Demo.desugar(List(1, 2, 3).reverse)
    println(s)
  }
}
```

This can be then be run at the console:

Actual tests can be defined and run as usual with `macro/test`.

The main project can use the macro in the same way that the tests do. For example,

`src/main/scala/MainUsage.scala`:

```
package demo
```

```

object Usage {
  def main(args: Array[String]) {
    val s = Demo.desugar(List(6, 4, 5).sorted)
    println(s)
  }
}

```

Common Interface

Sometimes, the macro implementation and the macro usage should share some common code. In this case, declare another subproject for the common code and have the main project and the macro subproject depend on the new subproject. For example, the project definitions from above would look like:

```

lazy val main = Project("main", file(".")) dependsOn(macroSub, commonSub)
lazy val macroSub = Project("macro", file("macro")) dependsOn(commonSub) settings(
  libraryDependencies += "org.scala-lang" % "scala-reflect" % scalaVersion.value
)
lazy val commonSub = Project("common", file("common"))

```

Code in `common/src/main/scala/` is available for both the macro and main projects to use.

Distribution

To include the macro code with the main code, add the binary and source mappings from the macro subproject to the main project. For example, the main Project definition above would now look like:

```

lazy val main = Project("main", file(".")) dependsOn(macroSub) settings(
  // include the macro classes and resources in the main jar
  mappings in (Compile, packageBin) ++= mappings.in(macroSub, Compile, packageBin).value,
  // include the macro sources in the main source jar
  mappings in (Compile, packageSrc) ++= mappings.in(macroSub, Compile, packageSrc).value
)

```

You may wish to disable publishing the macro implementation. This is done by overriding `publish` and `publishLocal` to do nothing:

```

lazy val macroSub = Project("macro", file("macro")) settings(
  publish := {},
  publishLocal := {}
)

```

The techniques described here may also be used for the common interface described in the previous section.

Paths

This page describes files, sequences of files, and file filters. The base type used is `java.io.File`, but several methods are augmented through implicits:

- `RichFile` adds methods to `File`
- `PathFinder` adds methods to `File` and `Seq[File]`
- `Path` and `IO` provide general methods related to files and I/O.

Constructing a File

sbt 0.10+ uses `java.io.File` to represent a file instead of the custom `sbt.Path` class that was in sbt 0.7 and earlier. sbt defines the alias `File` for `java.io.File` so that an extra import is not necessary. The `file` method is an alias for the single-argument `File` constructor to simplify constructing a new file from a `String`:

```
val source: File = file("/home/user/code/A.scala")
```

Additionally, sbt augments `File` with a `/` method, which is an alias for the two-argument `File` constructor for building up a path:

```
def readme(base: File): File = base / "README"
```

Relative files should only be used when defining the base directory of a `Project`, where they will be resolved properly.

```
val root = Project("root", file("."))
```

Elsewhere, files should be absolute or be built up from an absolute base `File`. The `baseDirectory` setting defines the base directory of the build or project depending on the scope.

For example, the following setting sets the unmanaged library directory to be the “`custom_lib`” directory in a project’s base directory:

```
unmanagedBase := baseDirectory.value / "custom_lib"
```

Or, more concisely:

```
unmanagedBase := baseDirectory.value / "custom_lib"
```

This setting sets the location of the shell history to be in the base directory of the build, irrespective of the project the setting is defined in:

```
historyPath := Some( (baseDirectory in ThisBuild).value / ".history"),
```

Path Finders

A `PathFinder` computes a `Seq[File]` on demand. It is a way to build a sequence of files. There are several methods that augment `File` and `Seq[File]` to construct a `PathFinder`. Ultimately, call `get` on the resulting `PathFinder` to evaluate it and get back a `Seq[File]`.

Selecting descendants The `**` method accepts a `java.io.FileFilter` and selects all files matching that filter.

```
def scalaSources(base: File): PathFinder = (base / "src") ** "*.scala"
```

`get` This selects all files that end in `.scala` that are in `src` or a descendent directory. The list of files is not actually evaluated until `get` is called:

```
def scalaSources(base: File): Seq[File] = {  
  val finder: PathFinder = (base / "src") ** "*.scala"  
  finder.get  
}
```

If the filesystem changes, a second call to `get` on the same `PathFinder` object will reflect the changes. That is, the `get` method reconstructs the list of files each time. Also, `get` only returns `Files` that existed at the time it was called.

Selecting children Selecting files that are immediate children of a subdirectory is done with a single `*`:

```
def scalaSources(base: File): PathFinder = (base / "src") * "*.scala"
```

This selects all files that end in `.scala` that are in the `src` directory.

Existing files only If a selector, such as `/`, `**`, or `*`, is used on a path that does not represent a directory, the path list will be empty:

```
def emptyFinder(base: File) = (base / "lib" / "ivy.jar") * "not_possible"
```

Name Filter The argument to the child and descendent selectors `*` and `**` is actually a `NameFilter`. An implicit is used to convert a `String` to a `NameFilter` that interprets `*` to represent zero or more characters of any value. See the `Name Filters` section below for more information.

Combining PathFinders Another operation is concatenation of PathFinders:

```
def multiPath(base: File): PathFinder =  
  (base / "src" / "main") +++  
  (base / "lib") +++  
  (base / "target" / "classes")
```

When evaluated using `get`, this will return `src/main/`, `lib/`, and `target/classes/`. The concatenated finder supports all standard methods. For example,

```
def jars(base: File): PathFinder =  
  (base / "lib" +++ base / "target") * "*.jar"
```

selects all jars directly in the “lib” and “target” directories.

A common problem is excluding version control directories. This can be accomplished as follows:

```
def sources(base: File) =  
  ( (base / "src") ** "*.scala" ) --- ( (base / "src") ** ".svn" ** "*.scala" )
```

The first selector selects all Scala sources and the second selects all sources that are a descendent of a `.svn` directory. The `---` method removes all files returned by the second selector from the sequence of files returned by the first selector.

Filtering There is a `filter` method that accepts a predicate of type `File => Boolean` and is non-strict:

```
// selects all directories under "src"  
def srcDirs(base: File) = ( (base / "src") ** "*" ) filter { _.isDirectory }  
  
// selects archives (.zip or .jar) that are selected by 'somePathFinder'  
def archivesOnly(base: PathFinder) = base filter ClasspathUtilities.isArchive
```

Empty PathFinder `PathFinder.empty` is a `PathFinder` that returns the empty sequence when `get` is called:

```
assert( PathFinder.empty.get == Seq[File]() )
```


PathFinder to String conversions Convert a `PathFinder` to a `String` using one of the following methods:

- `toString` is for debugging. It puts the absolute path of each component on its own line.
- `absString` gets the absolute paths of each component and separates them by the platform's path separator.
- `getPaths` produces a `Seq[String]` containing the absolute paths of each component

Mappings The packaging and file copying methods in sbt expect values of type `Seq[(File,String)]` and `Seq[(File,File)]`, respectively. These are mappings from the input file to its (`String`) path in the jar or its (`File`) destination. This approach replaces the relative path approach (using the `##` method) from earlier versions of sbt.

Mappings are discussed in detail on the [Mapping-Files](#) page.

File Filters

The argument to `*` and `**` is of type `java.io.FileFilter`. sbt provides combinators for constructing `FileFilters`.

First, a `String` may be implicitly converted to a `FileFilter`. The resulting filter selects files with a name matching the string, with a `*` in the string interpreted as a wildcard. For example, the following selects all Scala sources with the word “Test” in them:

```
def testSrcs(base: File): PathFinder = (base / "src") * "*Test*.scala"
```

There are some useful combinators added to `FileFilter`. The `||` method declares alternative `FileFilters`. The following example selects all Java or Scala source files under “src”:

```
def sources(base: File): PathFinder = (base / "src") ** ("*.scala" || "*.java")
```

The `--` method excludes a files matching a second filter from the files matched by the first:

```
def imageResources(base: File): PathFinder =  
  (base/"src"/"main"/"resources") * ("*.png" -- "logo.png")
```

This will get `right.png` and `left.png`, but not `logo.png`, for example.

Parallel Execution

Task ordering

Task ordering is specified by declaring a task's inputs. Correctness of execution requires correct input declarations. For example, the following two tasks do not have an ordering specified:

```
write := IO.write(file("/tmp/sample.txt"), "Some content.")

read := IO.read(file("/tmp/sample.txt"))
```

sbt is free to execute `write` first and then `read`, `read` first and then `write`, or `read` and `write` simultaneously. Execution of these tasks is non-deterministic because they share a file. A correct declaration of the tasks would be:

```
write := {
  val f = file("/tmp/sample.txt")
  IO.write(f, "Some content.")
  f
}

read := IO.read(write.value)
```

This establishes an ordering: `read` must run after `write`. We've also guaranteed that `read` will read from the same file that `write` created.

Practical constraints

Note: The feature described in this section is experimental. The default configuration of the feature is subject to change in particular.

Background Declaring inputs and dependencies of a task ensures the task is properly ordered and that code executes correctly. In practice, tasks share finite hardware and software resources and can require control over utilization of these resources. By default, sbt executes tasks in parallel (subject to the ordering constraints already described) in an effort to utilize all available processors. Also by default, each test class is mapped to its own task to enable executing tests in parallel.

Prior to sbt 0.12, user control over this process was restricted to:

1. Enabling or disabling all parallel execution (`parallelExecution := false`, for example).

2. Enabling or disabling mapping tests to their own tasks (`parallelExecution` in `Test := false`, for example).

(Although never exposed as a setting, the maximum number of tasks running at a given time was internally configurable as well.)

The second configuration mechanism described above only selected between running all of a project's tests in the same task or in separate tasks. Each project still had a separate task for running its tests and so test tasks in separate projects could still run in parallel if overall execution was parallel. There was no way to restriction execution such that only a single test out of all projects executed.

Configuration

sbt 0.12.0 introduces a general infrastructure for restricting task concurrency beyond the usual ordering declarations. There are two parts to these restrictions.

1. A task is tagged in order to classify its purpose and resource utilization. For example, the compile task may be tagged as `Tags.Compile` and `Tags.CPU`.
2. A list of rules restrict the tasks that may execute concurrently. For example, `Tags.limit(Tags.CPU, 4)` would allow up to four computation-heavy tasks to run at a time.

The system is thus dependent on proper tagging of tasks and then on a good set of rules.

Tagging Tasks In general, a tag is associated with a weight that represents the task's relative utilization of the resource represented by the tag. Currently, this weight is an integer, but it may be a floating point in the future. `Initialize[Task[T]]` defines two methods for tagging the constructed `Task`: `tag` and `tagw`. The first method, `tag`, fixes the weight to be 1 for the tags provided to it as arguments. The second method, `tagw`, accepts pairs of tags and weights. For example, the following associates the `CPU` and `Compile` tags with the `compile` task (with a weight of 1).

```
def myCompileTask = Def.task { ... } tag(Tags.CPU, Tags.Compile)

compile := myCompileTask.value
```

Different weights may be specified by passing tag/weight pairs to `tagw`:

```
def downloadImpl = Def.task { ... } tagw(Tags.Network -> 3)

download := downloadImpl.value
```

Defining Restrictions Once tasks are tagged, the `concurrentRestrictions` setting sets restrictions on the tasks that may be concurrently executed based on the weighted tags of those tasks. This is necessarily a global set of rules, so it must be scoped in `Global`. For example,

```
concurrentRestrictions in Global := Seq(
  Tags.limit(Tags.CPU, 2),
  Tags.limit(Tags.Network, 10),
  Tags.limit(Tags.Test, 1),
  Tags.limitAll( 15 )
)
```

The example limits:

- the number of CPU-using tasks to be no more than 2
- the number of tasks using the network to be no more than 10
- test execution to only one test at a time across all projects
- the total number of tasks to be less than or equal to 15

Note that these restrictions rely on proper tagging of tasks. Also, the value provided as the limit must be at least 1 to ensure every task is able to be executed. sbt will generate an error if this condition is not met.

Most tasks won't be tagged because they are very short-lived. These tasks are automatically assigned the label `Untagged`. You may want to include these tasks in the CPU rule by using the `limitSum` method. For example:

```
...
Tags.limitSum(2, Tags.CPU, Tags.Untagged)
...
```

Note that the limit is the first argument so that tags can be provided as varargs.

Another useful convenience function is `Tags.exclusive`. This specifies that a task with the given tag should execute in isolation. It starts executing only when no other tasks are running (even if they have the exclusive tag) and no other tasks may start execution until it completes. For example, a task could be tagged with a custom tag `Benchmark` and a rule configured to ensure such a task is executed by itself:

```
...
Tags.exclusive(Benchmark)
...
```

Finally, for the most flexibility, you can specify a custom function of type `Map[Tag,Int] => Boolean`. The `Map[Tag,Int]` represents the weighted tags of a set of tasks. If the function returns `true`, it indicates that the set of tasks is allowed to execute concurrently. If the return value is `false`, the set of tasks will not be allowed to execute concurrently. For example, `Tags.exclusive(Benchmark)` is equivalent to the following:

```
...
Tags.customLimit { (tags: Map[Tag,Int]) =>
  val exclusive = tags.getOrElse(Benchmark, 0)
  // the total number of tasks in the group
  val all = tags.getOrElse(Tags.All, 0)
  // if there are no exclusive tasks in this group, this rule adds no restrictions
  exclusive == 0 ||
  // If there is only one task, allow it to execute.
  all == 1
}
...
```

There are some basic rules that custom functions must follow, but the main one to be aware of in practice is that if there is only one task, it must be allowed to execute. `sbt` will generate a warning if the user defines restrictions that prevent a task from executing at all and will then execute the task anyway.

Built-in Tags and Rules Built-in tags are defined in the `Tags` object. All tags listed below must be qualified by this object. For example, `CPU` refers to the `Tags.CPU` value.

The built-in semantic tags are:

- `Compile` - describes a task that compiles sources.
- `Test` - describes a task that performs a test.
- `Publish`
- `Update`
- `Untagged` - automatically added when a task doesn't explicitly define any tags.
- `All` - automatically added to every task.

The built-in resource tags are:

- `Network` - describes a task's network utilization.
- `Disk` - describes a task's filesystem utilization.
- `CPU` - describes a task's computational utilization.

The tasks that are currently tagged by default are:

- `compile` : `Compile`, `CPU`
- `test` : `Test`
- `update` : `Update`, `Network`
- `publish`, `publishLocal` : `Publish`, `Network`

Of additional note is that the default `test` task will propagate its tags to each child task created for each test class.

The default rules provide the same behavior as previous versions of sbt:

```
concurrentRestrictions in Global := {  
  val max = Runtime.getRuntime.availableProcessors  
  Tags.limitAll(if(parallelExecution.value) max else 1) :: Nil  
}
```

As before, `parallelExecution` in `Test` controls whether tests are mapped to separate tasks. To restrict the number of concurrently executing tests in all projects, use:

```
concurrentRestrictions in Global += Tags.limit(Tags.Test, 1)
```

Custom Tags To define a new tag, pass a String to the `Tags.Tag` method. For example:

```
val Custom = Tags.Tag("custom")
```

Then, use this tag as any other tag. For example:

```
def aImpl = Def.task { ... } tag(Custom)
```

```
aCustomTask := aImpl.value
```

```
concurrentRestrictions in Global +=  
  Tags.limit(Custom, 1)
```

Future work

This is an experimental feature and there are several aspects that may change or require further work.

Tagging Tasks Currently, a tag applies only to the immediate computation it is defined on. For example, in the following, the second compile definition has no tags applied to it. Only the first computation is labeled.

```
def myCompileTask = Def.task { ... } tag(Tags.CPU, Tags.Compile)

compile := myCompileTask.value

compile := {
  val result = compile.value
  ... do some post processing ...
}
```

Is this desirable? expected? If not, what is a better, alternative behavior?

Fractional weighting Weights are currently ints, but could be changed to be doubles if fractional weights would be useful. It is important to preserve a consistent notion of what a weight of 1 means so that built-in and custom tasks share this definition and useful rules can be written.

Default Behavior User feedback on what custom rules work for what workloads will help determine a good set of default tags and rules.

Adjustments to Defaults Rules should be easier to remove or redefine, perhaps by giving them names. As it is, rules must be appended or all rules must be completely redefined. Also, tags can only be defined for tasks at the original definition site when using the `:=` syntax.

For removing tags, an implementation of `removeTag` should follow from the implementation of `tag` in a straightforward manner.

Other characteristics The system of a tag with a weight was selected as being reasonably powerful and flexible without being too complicated. This selection is not fundamental and could be enhance, simplified, or replaced if necessary. The fundamental interface that describes the constraints the system must work within is `sbt.ConcurrentRestrictions`. This interface is used to provide an intermediate scheduling queue between task execution (`sbt.Execute`) and the underlying thread-based parallel execution service (`java.util.concurrent.CompletionService`). This intermediate queue restricts new tasks from being forwarded to the `j.u.c.CompletionService` according to the `sbt.ConcurrentRestrictions` implementation. See the [sbt.ConcurrentRestrictions](#) API documentation for details.

External Processes

Usage

`sbt` includes a process library to simplify working with external processes. The library is available without import in build definitions and at the interpreter started by the `consoleProject` task.

To run an external command, follow it with an exclamation mark `!`:

```
"find project -name *.jar" !
```

An implicit converts the `String` to `sbt.ProcessBuilder`, which defines the `!` method. This method runs the constructed command, waits until the command completes, and returns the exit code. Alternatively, the `run` method defined on `ProcessBuilder` runs the command and returns an instance of `sbt.Process`, which can be used to `destroy` the process before it completes. With no arguments, the `!` method sends output to standard output and standard error. You can pass a `Logger` to the `!` method to send output to the `Logger`:

```
"find project -name *.jar" ! log
```

Two alternative implicit conversions are from `scala.xml.Elem` or `List[String]` to `sbt.ProcessBuilder`. These are useful for constructing commands. An example of the first variant from the android plugin:

```
<x> {dxPath.absolutePath} --dex --output={classesDexPath.absolutePath} {classesMinJarPath.al
```

If you need to set the working directory or modify the environment, call `sbt.Process` explicitly, passing the command sequence (command and argument list) or command string first and the working directory second. Any environment variables can be passed as a vararg list of key/value `String` pairs.

```
Process("ls" :: "-l" :: Nil, Path.userHome, "key1" -> value1, "key2" -> value2) ! log
```

Operators are defined to combine commands. These operators start with `#` in order to keep the precedence the same and to separate them from the operators defined elsewhere in `sbt` for filters. In the following operator definitions, `a` and `b` are subcommands.

- `a ##&& b` Execute `a`. If the exit code is nonzero, return that exit code and do not execute `b`. If the exit code is zero, execute `b` and return its exit code.

- `a ## b` Execute `a`. If the exit code is zero, return zero for the exit code and do not execute `b`. If the exit code is nonzero, execute `b` and return its exit code.
- `a #| b` Execute `a` and `b`, piping the output of `a` to the input of `b`.

There are also operators defined for redirecting output to `Files` and input from `Files` and `URLs`. In the following definitions, `url` is an instance of `URL` and `file` is an instance of `File`.

- `a #< url or url #> a` Use `url` as the input to `a`. `a` may be a `File` or a command.
- `a #< file or file #> a` Use `file` as the input to `a`. `a` may be a `File` or a command.
- `a #> file or file #< a` Write the output of `a` to `file`. `a` may be a `File`, `URL`, or a command.
- `a #>> file or file #<< a` Append the output of `a` to `file`. `a` may be a `File`, `URL`, or a command.

There are some additional methods to get the output from a forked process into a `String` or the output lines as a `Stream[String]`. Here are some examples, but see the [ProcessBuilder API](#) for details.

```
val listed: String = "ls" !!
val lines2: Stream[String] = "ls" lines_!
```

Finally, there is a `cat` method to send the contents of `Files` and `URLs` to standard output.

Examples Download a `URL` to a `File`:

```
url("http://databinder.net/dispatch/About") #> file("About.html") !
// or
file("About.html") #< url("http://databinder.net/dispatch/About") !
```

Copy a `File`:

```
file("About.html") #> file("About_copy.html") !
// or
file("About_copy.html") #< file("About.html") !
```

Append the contents of a `URL` to a `File` after filtering through `grep`:

```
url("http://databinder.net/dispatch/About") #> "grep JSON" #>> file("About_JSON") !
// or
file("About_JSON") #<< ( "grep JSON" #< url("http://databinder.net/dispatch/About") ) !
```

Search for uses of `null` in the source directory:

```
"find src -name *.scala -exec grep null {} ;" #| "xargs test -z" #&& "echo null-free" #
```

Use `cat`:

```
val spde = url("http://technically.us/spde/About")
val dispatch = url("http://databinder.net/dispatch/About")
val build = file("project/build.properties")
cat(spde, dispatch, build) #| "grep -i scala" !
```

Running Project Code

The `run` and `console` actions provide a means for running user code in the same virtual machine as `sbt`. This page describes the problems with doing so, how `sbt` handles these problems, what types of code can use this feature, and what types of code must use a [forked jvm](#). Skip to User Code if you just want to see when you should use a [forked jvm](#).

Problems

System.exit User code can call `System.exit`, which normally shuts down the JVM. Because the `run` and `console` actions run inside the same JVM as `sbt`, this also ends the build and requires restarting `sbt`.

Threads User code can also start other threads. Threads can be left running after the main method returns. In particular, creating a GUI creates several threads, some of which may not terminate until the JVM terminates. The program is not completed until either `System.exit` is called or all non-daemon threads terminate.

Deserialization and class loading During deserialization, the wrong class loader might be used for various complex reasons. This can happen in many scenarios, and running under SBT is just one of them. This is discussed for instance in issues [#163](#) and [#136](#). The reason is explained [here](#).

sbt's Solutions

System.exit User code is run with a custom `SecurityManager` that throws a custom `SecurityException` when `System.exit` is called. This exception is caught by sbt. sbt then disposes of all top-level windows, interrupts (not stops) all user-created threads, and handles the exit code. If the exit code is nonzero, `run` and `console` complete unsuccessfully. If the exit code is zero, they complete normally.

Threads sbt makes a list of all threads running before executing user code. After the user code returns, sbt can then determine the threads created by the user code. For each user-created thread, sbt replaces the uncaught exception handler with a custom one that handles the custom `SecurityException` thrown by calls to `System.exit` and delegates to the original handler for everything else. sbt then waits for each created thread to exit or for `System.exit` to be called. sbt handles a call to `System.exit` as described above.

A user-created thread is one that is not in the `system` thread group and is not an AWT implementation thread (e.g. `AWT-XAWT`, `AWT-Window`s). User-created threads include the `AWT-EventQueue-*` thread(s).

User Code Given the above, when can user code be run with the `run` and `console` actions?

The user code cannot rely on shutdown hooks and at least one of the following situations must apply for user code to run in the same JVM:

1. User code creates no threads.
2. User code creates a GUI and no other threads.
3. The program ends when user-created threads terminate on their own.
4. `System.exit` is used to end the program and user-created threads terminate when interrupted.
5. No deserialization is done, or the deserialization code avoids ensures that the right class loader is used, as in <https://github.com/NetLogo/NetLogo/blob/master/src/main/org/nlogo/util/ClassLoaderObjectInputStream.scala> or <https://github.com/scala/scala/blob/master/src/actors/scala/actors/remote/JavaSerializer.scala#L20>.

The requirements on threading and shutdown hooks are required because the JVM does not actually shut down. So, shutdown hooks cannot be run and threads are not terminated unless they stop when interrupted. If these requirements are not met, code must run in a `forked jvm`.

The feature of allowing `System.exit` and multiple threads to be used cannot completely emulate the situation of running in a separate JVM and is intended

for development. Program execution should be checked in a [forked jvm](#) when using multiple threads or `System.exit`.

As of sbt 0.13.1, multiple `run` instances can be managed. There can only be one application that uses AWT at a time, however.

Testing

Basics

The standard source locations for testing are:

- Scala sources in `src/test/scala/`
- Java sources in `src/test/java/`
- Resources for the test classpath in `src/test/resources/`

The resources may be accessed from tests by using the `getResource` methods of `java.lang.Class` or `java.lang.ClassLoader`.

The main Scala testing frameworks ([specs2](#), [ScalaCheck](#), and [ScalaTest](#)) provide an implementation of the common test interface and only need to be added to the classpath to work with sbt. For example, `ScalaCheck` may be used by declaring it as a [managed dependency](#):

```
libraryDependencies += "org.scalacheck" %% "scalacheck" % "1.11.4" % "test"
```

The fourth component `"test"` is the [configuration](#) and means that `ScalaCheck` will only be on the test classpath and it isn't needed by the main sources. This is generally good practice for libraries because your users don't typically need your test dependencies to use your library.

With the library dependency defined, you can then add test sources in the locations listed above and compile and run tests. The tasks for running tests are `test` and `testOnly`. The `test` task accepts no command line arguments and runs all tests:

```
> test
```

testOnly The `testOnly` task accepts a whitespace separated list of test names to run. For example:

```
> testOnly org.example.MyTest1 org.example.MyTest2
```

It supports wildcards as well:

```
> testOnly org.example.*Slow org.example.MyTest1
```

testQuick The `testQuick` task, like `testOnly`, allows to filter the tests to run to specific tests or wildcards using the same syntax to indicate the filters. In addition to the explicit filter, only the tests that satisfy one of the following conditions are run:

- The tests that failed in the previous run
- The tests that were not run before
- The tests that have one or more transitive dependencies, maybe in a different project, recompiled.

Tab completion Tab completion is provided for test names based on the results of the last `test:compile`. This means that a new sources aren't available for tab completion until they are compiled and deleted sources won't be removed from tab completion until a recompile. A new test source can still be manually written out and run using `testOnly`.

Other tasks Tasks that are available for main sources are generally available for test sources, but are prefixed with `test:` on the command line and are referenced in Scala code with `in Test`. These tasks include:

- `test:compile`
- `test:console`
- `test:consoleQuick`
- `test:run`
- `test:runMain`

See [Running](#) for details on these tasks.

Output

By default, logging is buffered for each test source file until all tests for that file complete. This can be disabled by setting `logBuffered`:

```
logBuffered in Test := false
```

Test Reports By default, sbt will generate JUnit XML test reports for all tests in the build, located in the `target/test-reports` directory for a project. This can be disabled by disabling the `JUnitXmlReportPlugin`

```
val myProject = project in file(".") disablePlugins (plugins.JUnitXmlReportPlugin)
```

Options

Test Framework Arguments Arguments to the test framework may be provided on the command line to the `testOnly` tasks following a `--` separator. For example:

```
> testOnly org.example.MyTest -- -verbosity 1
```

To specify test framework arguments as part of the build, add options constructed by `Tests.Argument`:

```
testOptions in Test += Tests.Argument("-verbosity", "1")
```

To specify them for a specific test framework only:

```
testOptions in Test += Tests.Argument(TestFrameworks.ScalaCheck, "-verbosity", "1")
```

Setup and Cleanup Specify setup and cleanup actions using `Tests.Setup` and `Tests.Cleanup`. These accept either a function of type `() => Unit` or a function of type `ClassLoader => Unit`. The variant that accepts a `ClassLoader` is passed the class loader that is (or was) used for running the tests. It provides access to the test classes as well as the test framework classes.

Note: When forking, the `ClassLoader` containing the test classes cannot be provided because it is in another JVM. Only use the `() => Unit` variants in this case.

Examples:

```
testOptions in Test += Tests.Setup( () => println("Setup") )
```

```
testOptions in Test += Tests.Cleanup( () => println("Cleanup") )
```

```
testOptions in Test += Tests.Setup( loader => ... )
```

```
testOptions in Test += Tests.Cleanup( loader => ... )
```

Disable Parallel Execution of Tests By default, sbt runs all tasks in parallel and within the same JVM as sbt itself. Because each test is mapped to a task, tests are also run in parallel by default. To make tests within a given project execute serially: :

```
parallelExecution in Test := false
```

`Test` can be replaced with `IntegrationTest` to only execute integration tests serially. Note that tests from different projects may still execute concurrently.

Filter classes If you want to only run test classes whose name ends with “Test”, use `Tests.Filter`:

```
testOptions in Test := Seq(Tests.Filter(s => s.endsWith("Test")))
```

Forking tests The setting:

```
fork in Test := true
```

specifies that all tests will be executed in a single external JVM. See [Forking](#) for configuring standard options for forking. By default, tests executed in a forked JVM are executed *sequentially*. More control over how tests are assigned to JVMs and what options to pass to those is available with `testGrouping` key. For example in `build.sbt`:

```
import Tests._

{
  def groupByFirst(tests: Seq[TestDefinition]) =
    tests groupBy (_.name(0)) map {
      case (letter, tests) => new Group(letter.toString, tests, SubProcess(Seq("-Dfirst.lett
    }) toSeq

  testGrouping in Test <:= groupByFirst( (definedTests in Test).value )
}
```

The tests in a single group are run sequentially. Control the number of forked JVMs allowed to run at the same time by setting the limit on `Tags.ForkedTestGroup` tag, which is 1 by default. `Setup` and `Cleanup` actions cannot be provided with the actual test class loader when a group is forked.

In addition, forked tests can optionally be run in parallel. This feature is still considered experimental, and may be enabled with the following setting :

```
testForkedParallel in Test := true
```

Additional test configurations

You can add an additional test configuration to have a separate set of test sources and associated compilation, packaging, and testing tasks and settings. The steps are:

- Define the configuration

- Add the tasks and settings
- Declare library dependencies
- Create sources
- Run tasks

The following two examples demonstrate this. The first example shows how to enable integration tests. The second shows how to define a customized test configuration. This allows you to define multiple types of tests per project.

Integration Tests The following full build configuration demonstrates integration tests.

```
import sbt._
import Keys._

object B extends Build {
  lazy val root =
    Project("root", file(".")).
      configs( IntegrationTest ).
      settings( Defaults.itSettings : _* ).
      settings( libraryDependencies += specs )

  lazy val specs = "org.specs2" %% "specs2" % "2.0" % "it,test"
}
```

- `configs(IntegrationTest)` adds the predefined integration test configuration. This configuration is referred to by the name `it`.
- `settings(Defaults.itSettings : _*)` adds compilation, packaging, and testing actions and settings in the `IntegrationTest` configuration.
- `settings(libraryDependencies += specs)` adds `specs` to both the standard test configuration and the integration test configuration `it`. To define a dependency only for integration tests, use “`it`” as the configuration instead of “`it,test`”.

The standard source hierarchy is used:

- `src/it/scala` for Scala sources
- `src/it/java` for Java sources
- `src/it/resources` for resources that should go on the integration test classpath

The standard testing tasks are available, but must be prefixed with `it::`. For example,


```
> it:testOnly org.example.AnIntegrationTest
```

Similarly the standard settings may be configured for the `IntegrationTest` configuration. If not specified directly, most `IntegrationTest` settings delegate to `Test` settings by default. For example, if test options are specified as:

```
testOptions in Test += ...
```

then these will be picked up by the `Test` configuration and in turn by the `IntegrationTest` configuration. Options can be added specifically for integration tests by putting them in the `IntegrationTest` configuration:

```
testOptions in IntegrationTest += ...
```

Or, use `:=` to overwrite any existing options, declaring these to be the definitive integration test options:

```
testOptions in IntegrationTest := Seq(...)
```

Custom test configuration The previous example may be generalized to a custom test configuration.

```
import sbt._
import Keys._

object B extends Build {
  lazy val root =
    Project("root", file(".")).
      configs( FunTest ).
      settings( inConfig(FunTest)(Defaults.testSettings) : _* ).
      settings( libraryDependencies += specs )

  lazy val FunTest = config("fun") extend(Test)
  lazy val specs = "org.specs2" %% "specs2" % "2.0" % "fun"
}
```

Instead of using the built-in configuration, we defined a new one:

```
lazy val FunTest = config("fun") extend(Test)
```

The `extend(Test)` part means to delegate to `Test` for undefined `CustomTest` settings. The line that adds the tasks and settings for the new test configuration is:

```
settings( inConfig(FunTest)(Defaults.testSettings) : _*)
```

This says to add test and settings tasks in the `FunTest` configuration. We could have done it this way for integration tests as well. In fact, `Defaults.itSettings` is a convenience definition: `val itSettings = inConfig(IntegrationTest)(Defaults.testSettings)`.

The comments in the integration test section hold, except with `IntegrationTest` replaced with `FunTest` and `"it"` replaced with `"fun"`. For example, test options can be configured specifically for `FunTest`:

```
testOptions in FunTest += ...
```

Test tasks are run by prefixing them with `fun`:

```
> fun:test
```

Additional test configurations with shared sources An alternative to adding separate sets of test sources (and compilations) is to share sources. In this approach, the sources are compiled together using the same classpath and are packaged together. However, different tests are run depending on the configuration.

```
import sbt._
import Keys._

object B extends Build {
  lazy val root =
    Project("root", file("."))
      .configs( FunTest )
      .settings( inConfig(FunTest)(Defaults.testTasks) : _*)
      .settings(
        libraryDependencies += specs,
        testOptions in Test := Seq(Tests.Filter(unitFilter)),
        testOptions in FunTest := Seq(Tests.Filter(itFilter))
      )

  def itFilter(name: String): Boolean = name endsWith "ITest"
  def unitFilter(name: String): Boolean = (name endsWith "Test") && !itFilter(name)

  lazy val FunTest = config("fun") extend(Test)
  lazy val specs = "org.specs2" %% "specs2" % "2.0" % "test"
}
```

The key differences are:

- We are now only adding the test tasks (`inConfig(FunTest)(Defaults.testTasks)`) and not compilation and packaging tasks and settings.
- We filter the tests to be run for each configuration.

To run standard unit tests, run `test` (or equivalently, `test:test`):

```
> test
```

To run tests for the added configuration (here, `"fun"`), prefix it with the configuration name as before:

```
> fun:test
> fun:testOnly org.example.AFunTest
```

Application to parallel execution One use for this shared-source approach is to separate tests that can run in parallel from those that must execute serially. Apply the procedure described in this section for an additional configuration. Let's call the configuration `serial`:

```
lazy val Serial = config("serial") extend(Test)
```

Then, we can disable parallel execution in just that configuration using:

```
parallelExecution in Serial := false
```

The tests to run in parallel would be run with `test` and the ones to run in serial would be run with `serial:test`.

JUnit

Support for JUnit is provided by [junit-interface](#). To add JUnit support into your project, add the `junit-interface` dependency in your project's main `build.sbt` file.

```
libraryDependencies += "com.novocode" % "junit-interface" % "0.9" % "test"
```

Extensions

This page describes adding support for additional testing libraries and defining additional test reporters. You do this by implementing `sbt` interfaces (described below). If you are the author of the testing framework, you can depend on the test interface as a provided dependency. Alternatively, anyone can provide support for a test framework by implementing the interfaces in a separate project and packaging the project as an `sbt` [Plugin](#).

Custom Test Framework The main Scala testing libraries have built-in support for sbt. To add support for a different framework, implement the [uniform test interface](#).

Custom Test Reporters Test frameworks report status and results to test reporters. You can create a new test reporter by implementing either [TestReportListener](#) or [TestsListener](#).

Using Extensions To use your extensions in a project definition:

Modify the `testFrameworks` setting to reference your test framework:

```
testFrameworks += new TestFramework("custom.framework.ClassName")
```

Specify the test reporters you want to use by overriding the `testListeners` setting in your project definition.

```
testListeners += customTestListener
```

where `customTestListener` is of type `sbt.TestReportListener`.

Dependency Management

This part of the documentation has pages documenting particular sbt topics in detail. Before reading anything in here, you will need the information in the [Getting Started Guide](#) as a foundation.

Artifacts

Selecting default artifacts

By default, the published artifacts are the main binary jar, a jar containing the main sources and resources, and a jar containing the API documentation. You can add artifacts for the test classes, sources, or API or you can disable some of the main artifacts.

To add all test artifacts:

```
publishArtifact in Test := true
```

To add them individually:

```

// enable publishing the jar produced by `test:package`
publishArtifact in (Test, packageBin) := true

// enable publishing the test API jar
publishArtifact in (Test, packageDoc) := true

// enable publishing the test sources jar
publishArtifact in (Test, packageSrc) := true

```

To disable main artifacts individually:

```

// disable publishing the main jar produced by `package`
publishArtifact in (Compile, packageBin) := false

// disable publishing the main API jar
publishArtifact in (Compile, packageDoc) := false

// disable publishing the main sources jar
publishArtifact in (Compile, packageSrc) := false

```

Modifying default artifacts

Each built-in artifact has several configurable settings in addition to `publishArtifact`. The basic ones are `artifact` (of type `SettingKey[Artifact]`), `mappings` (of type `TaskKey[(File,String)]`), and `artifactPath` (of type `SettingKey[File]`). They are scoped by (`<config>`, `<task>`) as indicated in the previous section.

To modify the type of the main artifact, for example:

```

artifact in (Compile, packageBin) := {
  val previous: Artifact = (artifact in (Compile, packageBin)).value
  previous.copy(`type` = "bundle")
}

```

The generated artifact name is determined by the `artifactName` setting. This setting is of type `(ScalaVersion, ModuleID, Artifact) => String`. The `ScalaVersion` argument provides the full Scala version String and the binary compatible part of the version String. The String result is the name of the file to produce. The default implementation is `Artifact.artifactName _`. The function may be modified to produce different local names for artifacts without affecting the published name, which is determined by the `artifact` definition combined with the repository pattern.

For example, to produce a minimal name without a classifier or cross path:

```
artifactName := { (sv: ScalaVersion, module: ModuleID, artifact: Artifact) =>
  artifact.name + "-" + module.revision + "." + artifact.extension
}
```

(Note that in practice you rarely want to drop the classifier.)

Finally, you can get the (`Artifact`, `File`) pair for the artifact by mapping the `packagedArtifact` task. Note that if you don't need the `Artifact`, you can get just the `File` from the package task (`package`, `packageDoc`, or `packageSrc`). In both cases, mapping the task to get the file ensures that the artifact is generated first and so the file is guaranteed to be up-to-date.

For example:

```
val myTask = taskKey[Unit]("My task.")

myTask := {
  val (art, file) = packagedArtifact.in(Compile, packageBin).value
  println("Artifact definition: " + art)
  println("Packaged file: " + file.getAbsolutePath)
}
```

Defining custom artifacts

In addition to configuring the built-in artifacts, you can declare other artifacts to publish. Multiple artifacts are allowed when using Ivy metadata, but a Maven POM file only supports distinguishing artifacts based on classifiers and these are not recorded in the POM.

Basic `Artifact` construction look like:

```
Artifact("name", "type", "extension")
Artifact("name", "classifier")
Artifact("name", url: URL)
Artifact("name", Map("extra1" -> "value1", "extra2" -> "value2"))
```

For example:

```
Artifact("myproject", "zip", "zip")
Artifact("myproject", "image", "jpg")
Artifact("myproject", "jdk15")
```

See the [Ivy documentation](#) for more details on artifacts. See the [Artifact API](#) for combining the parameters above and specifying [Configurations](#) and extra attributes.

To declare these artifacts for publishing, map them to the task that generates the artifact:

```

val myImageTask = taskKey[File](...)

myImageTask := {
    val artifact: File = makeArtifact(...)
    artifact
}

addArtifact( Artifact("myproject", "image", "jpg"), myImageTask )

```

`addArtifact` returns a sequence of settings (wrapped in a [SettingsDefinition](#)). In a full build configuration, usage looks like:

```

...
lazy val proj = Project(...).
    settings( addArtifact(...).settings :_* )
...

```

Publishing .war files

A common use case for web applications is to publish the `.war` file instead of the `.jar` file.

```

// disable .jar publishing
publishArtifact in (Compile, packageBin) := false

// create an Artifact for publishing the .war file
artifact in (Compile, packageWar) := {
    val previous: Artifact = (artifact in (Compile, packageWar)).value
    previous.copy(`type` = "war", extension = "war")
}

// add the .war file to what gets published
addArtifact(artifact in (Compile, packageWar), packageWar)

```

Using dependencies with artifacts

To specify the artifacts to use from a dependency that has custom or multiple artifacts, use the `artifacts` method on your dependencies. For example:

```

libraryDependencies += "org" % "name" % "rev" artifacts(Artifact("name", "type", "ext"))

```

The `from` and `classifier` methods (described on the [Library Management](#) page) are actually convenience methods that translate to `artifacts`:

```
def from(url: String) = artifacts( Artifact(name, new URL(url)) )
def classifier(c: String) = artifacts( Artifact(name, c) )
```

That is, the following two dependency declarations are equivalent:

```
libraryDependencies += "org.testng" % "testng" % "5.7" classifier "jdk15"
```

```
libraryDependencies += "org.testng" % "testng" % "5.7" artifacts(Artifact("testng", "jdk15"))
```

Dependency Management Flow

sbt 0.12.1 addresses several issues with dependency management. These fixes were made possible by specific, reproducible examples, such as a situation where the resolution cache got out of date (gh-532). A brief summary of the current work flow with dependency management in sbt follows.

Background

`update` resolves dependencies according to the settings in a build file, such as `libraryDependencies` and `resolvers`. Other tasks use the output of `update` (an `UpdateReport`) to form various classpaths. Tasks that in turn use these classpaths, such as `compile` or `run`, thus indirectly depend on `update`. This means that before `compile` can run, the `update` task needs to run. However, resolving dependencies on every `compile` would be unnecessarily slow and so `update` must be particular about when it actually performs a resolution.

Caching and Configuration

1. Normally, if no dependency management configuration has changed since the last successful resolution and the retrieved files are still present, sbt does not ask Ivy to perform resolution.
2. Changing the configuration, such as adding or removing dependencies or changing the version or other attributes of a dependency, will automatically cause resolution to be performed. Updates to locally published dependencies should be detected in sbt 0.12.1 and later and will force an update. Dependent tasks like `compile` and `run` will get updated classpaths.
3. Directly running the `update` task (as opposed to a task that depends on it) will force resolution to run, whether or not configuration changed. This should be done in order to refresh remote SNAPSHOT dependencies.
4. When `offline := true`, remote SNAPSHOTs will not be updated by a resolution, even an explicitly requested update. This should effectively support working without a connection to remote repositories. Reproducible examples demonstrating otherwise are appreciated. Obviously, `update` must have successfully run before going offline.

5. Overriding all of the above, `skip in update := true` will tell sbt to never perform resolution. Note that this can cause dependent tasks to fail. For example, compilation may fail if jars have been deleted from the cache (and so needed classes are missing) or a dependency has been added (but will not be resolved because skip is true). Also, update itself will immediately fail if resolution has not been allowed to run since the last clean.

General troubleshooting steps

- A. Run `update` explicitly. This will typically fix problems with out of date SNAPSHOTS or locally published artifacts.
- B. If a file cannot be found, look at the output of update to see where Ivy is looking for the file. This may help diagnose an incorrectly defined dependency or a dependency that is actually not present in a repository.
- C. `last update` contains more information about the most recent resolution and download. The amount of debugging output from Ivy is high, so you may want to use lastGrep (run `help lastGrep` for usage).
- D. Run `clean` and then `update`. If this works, it could indicate a bug in sbt, but the problem would need to be reproduced in order to diagnose and fix it.
- E. Before deleting all of the Ivy cache, first try deleting files in `~/.ivy2/cache` related to problematic dependencies. For example, if there are problems with dependency `"org.example" % "demo" % "1.0"`, delete `~/.ivy2/cache/org.example/demo/1.0/` and retry update. This avoids needing to redownload all dependencies.
- F. Normal sbt usage should not require deleting files from `~/.ivy2/cache`, especially if the first four steps have been followed. If deleting the cache fixes a dependency management issue, please try to reproduce the issue and submit a test case.

Plugins

These troubleshooting steps can be run for plugins by changing to the build definition project, running the commands, and then returning to the main project. For example:

```
> reload plugins
> update
> reload return
```

Notes

- A. Configure offline behavior for all projects on a machine by putting `offline := true` in `~/.sbt/0.13/global.sbt`. A command that does this for the user would make a nice pull request. Perhaps the setting of offline should go into the output of `about` or should it be a warning in the output of `update` or both?
- B. The cache improvements in 0.12.1 address issues in the change detection for `update` so that it will correctly re-resolve automatically in more situations. A problem with an out of date cache can usually be attributed to a bug in that change detection if explicitly running `update` fixes the problem.
- C. A common solution to dependency management problems in sbt has been to remove `~/.ivy2/cache`. Before doing this with 0.12.1, be sure to follow the steps in the troubleshooting section first. In particular, verify that a clean and an explicit `update` do not solve the issue.
- D. There is no need to mark SNAPSHOT dependencies as `changing()` because sbt configures Ivy to know this already.

Library Management

There's now a [getting started page](#) about library management, which you may want to read first.

Documentation Maintenance Note: it would be nice to remove the overlap between this page and the getting started page, leaving this page with the more advanced topics such as checksums and external Ivy files.

Introduction

There are two ways for you to manage libraries with sbt: manually or automatically. These two ways can be mixed as well. This page discusses the two approaches. All configurations shown here are settings that go either directly in a [.sbt file](#) or are appended to the `settings` of a Project in a [.scala file](#).

Manual Dependency Management

Manually managing dependencies involves copying any jars that you want to use to the `lib` directory. sbt will put these jars on the classpath during compilation, testing, running, and when using the interpreter. You are responsible for adding, removing, updating, and otherwise managing the jars in this directory. No modifications to your project definition are required to use this method unless you would like to change the location of the directory you store the jars in.

To change the directory jars are stored in, change the `unmanagedBase` setting in your project definition. For example, to use `custom_lib/`:

```
unmanagedBase := baseDirectory.value / "custom_lib"
```

If you want more control and flexibility, override the `unmanagedJars` task, which ultimately provides the manual dependencies to sbt. The default implementation is roughly:

```
unmanagedJars in Compile := (baseDirectory.value ** "*.jar").classpath
```

If you want to add jars from multiple directories in addition to the default directory, you can do:

```
unmanagedJars in Compile += {  
  val base = baseDirectory.value  
  val baseDirectories = (base / "libA") +++ (base / "b" / "lib") +++ (base / "libC")  
  val customJars = (baseDirectories ** "*.jar") +++ (base / "d" / "my.jar")  
  customJars.classpath  
}
```

See [Paths](#) for more information on building up paths.

Automatic Dependency Management

This method of dependency management involves specifying the direct dependencies of your project and letting sbt handle retrieving and updating your dependencies. sbt supports three ways of specifying these dependencies:

- Declarations in your project definition
- Maven POM files (dependency definitions only: no repositories)
- Ivy configuration and settings files

sbt uses [Apache Ivy](#) to implement dependency management in all three cases. The default is to use inline declarations, but external configuration can be explicitly selected. The following sections describe how to use each method of automatic dependency management.

Inline Declarations Inline declarations are a basic way of specifying the dependencies to be automatically retrieved. They are intended as a lightweight alternative to a full configuration using Ivy.

Dependencies Declaring a dependency looks like:

```
libraryDependencies += groupId % artifactID % revision
```

or

```
libraryDependencies += groupId % artifactID % revision % configuration
```

See configurations for details on configuration mappings. Also, several dependencies can be declared together:

```
libraryDependencies ++= Seq(  
  groupId %% artifactID % revision,  
  groupId %% otherID % otherRevision  
)
```

If you are using a dependency that was built with sbt, double the first % to be %%:

```
libraryDependencies += groupId %% artifactID % revision
```

This will use the right jar for the dependency built with the version of Scala that you are currently using. If you get an error while resolving this kind of dependency, that dependency probably wasn't published for the version of Scala you are using. See [Cross Build](#) for details.

Ivy can select the latest revision of a module according to constraints you specify. Instead of a fixed revision like "1.6.1", you specify "latest.integration", "2.9.+", or "[1.0,)". See the [Ivy revisions](#) documentation for details.

Resolvers sbt uses the standard Maven2 repository by default.

Declare additional repositories with the form:

```
resolvers += name at location
```

For example:

```
libraryDependencies ++= Seq(  
  "org.apache.derby" % "derby" % "10.4.1.3",  
  "org.specs" % "specs" % "1.6.1"  
)
```

```
resolvers += "Sonatype OSS Snapshots" at "https://oss.sonatype.org/content/repositories/snapshots"
```

sbt can search your local Maven repository if you add it as a repository:

```
resolvers += "Local Maven Repository" at "file://" + Path.userHome.absolutePath + "/.m2/repository"
```

See [Resolvers](#) for details on defining other types of repositories.

Override default resolvers `resolvers` configures additional, inline user resolvers. By default, `sbt` combines these resolvers with default repositories (Maven Central and the local Ivy repository) to form `externalResolvers`. To have more control over repositories, set `externalResolvers` directly. To only specify repositories in addition to the usual defaults, configure `resolvers`.

For example, to use the Sonatype OSS Snapshots repository in addition to the default repositories,

```
resolvers += "Sonatype OSS Snapshots" at "https://oss.sonatype.org/content/repositories/snapshots"
```

To use the local repository, but not the Maven Central repository:

```
externalResolvers := Resolver.withDefaultResolvers(resolvers.value, mavenCentral = false)
```

Override all resolvers for all builds The repositories used to retrieve sbt, Scala, plugins, and application dependencies can be configured globally and declared to override the resolvers configured in a build or plugin definition. There are two parts:

1. Define the repositories used by the launcher.
2. Specify that these repositories should override those in build definitions.

The repositories used by the launcher can be overridden by defining `~/.sbt/repositories`, which must contain a `[repositories]` section with the same format as the `Launcher` configuration file. For example:

```
[repositories]
local
my-maven-repo: http://example.org/repo
my-ivy-repo: http://example.org/ivy-repo/, [organization]/[module]/[revision]/[type]s/[artifact]
```

A different location for the repositories file may be specified by the `sbt.repository.config` system property in the sbt startup script. The final step is to set `sbt.override.build.repos` to true to use these repositories for dependency resolution and retrieval.

Explicit URL If your project requires a dependency that is not present in a repository, a direct URL to its jar can be specified as follows:

```
libraryDependencies += "slinky" % "slinky" % "2.1" from "http://slinky2.googlecode.com/svn/a
```

The URL is only used as a fallback if the dependency cannot be found through the configured repositories. Also, the explicit URL is not included in published metadata (that is, the pom or ivy.xml).

Disable Transitivity By default, these declarations fetch all project dependencies, transitively. In some instances, you may find that the dependencies listed for a project aren't necessary for it to build. Projects using the Felix OSGI framework, for instance, only explicitly require its main jar to compile and run. Avoid fetching artifact dependencies with either `intransitive()` or `notTransitive()`, as in this example:

```
libraryDependencies += "org.apache.felix" % "org.apache.felix.framework" % "1.8.0" intransi
```

Classifiers You can specify the classifier for a dependency using the `classifier` method. For example, to get the jdk15 version of TestNG:

```
libraryDependencies += "org.testng" % "testng" % "5.7" classifier "jdk15"
```

For multiple classifiers, use multiple `classifier` calls:

```
libraryDependencies +=  
  "org.lwjgl.lwjgl" % "lwjgl-platform" % lwjglVersion classifier "natives-windows" classifie
```

To obtain particular classifiers for all dependencies transitively, run the `updateClassifiers` task. By default, this resolves all artifacts with the `sources` or `javadoc` classifier. Select the classifiers to obtain by configuring the `transitiveClassifiers` setting. For example, to only retrieve sources:

```
transitiveClassifiers := Seq("sources")
```

Exclude Transitive Dependencies To exclude certain transitive dependencies of a dependency, use the `excludeAll` or `exclude` methods. The `exclude` method should be used when a pom will be published for the project. It requires the organization and module name to exclude. For example,

```
libraryDependencies +=  
  "log4j" % "log4j" % "1.2.15" exclude("javax.jms", "jms")
```

The `excludeAll` method is more flexible, but because it cannot be represented in a `pom.xml`, it should only be used when a `pom` doesn't need to be generated. For example,

```
libraryDependencies +=  
  "log4j" % "log4j" % "1.2.15" excludeAll(  
    ExclusionRule(organization = "com.sun.jdmk"),  
    ExclusionRule(organization = "com.sun.jmx"),  
    ExclusionRule(organization = "javax.jms")  
  )
```

See [ModuleID](#) for API details.

Download Sources Downloading source and API documentation jars is usually handled by an IDE plugin. These plugins use the `updateClassifiers` and `updateSbtClassifiers` tasks, which produce an `Update-Report` referencing these jars.

To have sbt download the dependency's sources without using an IDE plugin, add `withSources()` to the dependency definition. For API jars, add `withJavadoc()`. For example:

```
libraryDependencies +=  
  "org.apache.felix" % "org.apache.felix.framework" % "1.8.0" withSources() withJavadoc()
```

Note that this is not transitive. Use the `update-*classifiers` tasks for that.

Extra Attributes [Extra attributes](#) can be specified by passing key/value pairs to the `extra` method.

To select dependencies by extra attributes:

```
libraryDependencies += "org" % "name" % "rev" extra("color" -> "blue")
```

To define extra attributes on the current project:

```
projectID := {  
  val previous = projectID.value  
  previous.extra("color" -> "blue", "component" -> "compiler-interface")  
}
```

Inline Ivy XML sbt additionally supports directly specifying the configurations or dependencies sections of an Ivy configuration file inline. You can mix this with inline Scala dependency and repository declarations.

For example:

```
ivyXML :=
  <dependencies>
    <dependency org="javax.mail" name="mail" rev="1.4.2">
      <exclude module="activation"/>
    </dependency>
  </dependencies>
```

Ivy Home Directory By default, sbt uses the standard Ivy home directory location `${user.home}/.ivy2/`. This can be configured machine-wide, for use by both the sbt launcher and by projects, by setting the system property `sbt.ivy.home` in the sbt startup script (described in [Setup](#)).

For example:

```
java -Dsbt.ivy.home=/tmp/.ivy2/ ...
```

Checksums sbt ([through Ivy](#)) verifies the checksums of downloaded files by default. It also publishes checksums of artifacts by default. The checksums to use are specified by the *checksums* setting.

To disable checksum checking during update:

```
checksums in update := Nil
```

To disable checksum creation during artifact publishing:

```
checksums in publishLocal := Nil
```

```
checksums in publish := Nil
```

The default value is:

```
checksums := Seq("sha1", "md5")
```


Conflict Management The conflict manager decides what to do when dependency resolution brings in different versions of the same library. By default, the latest revision is selected. This can be changed by setting `conflictManager`, which has type `ConflictManager`. See the [Ivy documentation](#) for details on the different conflict managers. For example, to specify that no conflicts are allowed,

```
conflictManager := ConflictManager.strict
```

With this set, any conflicts will generate an error. To resolve a conflict,

- configure a dependency override if the conflict is for a transitive dependency
- force the revision if it is a direct dependency

Both are explained in the following sections.

Forcing a revision The following direct dependencies will introduce a conflict on the `log4j` version because `spark` requires `log4j 1.2.16`.

```
libraryDependencies += Seq(  
  "org.spark-project" %% "spark-core" % "0.5.1",  
  "log4j" % "log4j" % "1.2.14"  
)
```

The default conflict manager will select the newer version of `log4j`, `1.2.16`. This can be confirmed in the output of `show update`, which shows the newer version as being selected and the older version as not selected:

```
> show update  
[info] compile:  
[info]    log4j:log4j:1.2.16: ...  
...  
[info]    (EVICTED) log4j:log4j:1.2.14  
...
```

To say that we prefer the version we've specified over the version from indirect dependencies, use `force()`:

```
libraryDependencies += Seq(  
  "org.spark-project" %% "spark-core" % "0.5.1",  
  "log4j" % "log4j" % "1.2.14" force()  
)
```

The output of `show update` is now reversed:

```
> show update
[info] compile:
[info]    log4j:log4j:1.2.14: ...
...
[info]    (EVICTED) log4j:log4j:1.2.16
...
```

Note: this is an Ivy-only feature and cannot be included in a published pom.xml.

Forcing a revision without introducing a dependency Use of the `force()` method described in the previous section requires having a direct dependency. However, it may be desirable to force a revision without introducing that direct dependency. Ivy provides overrides for this and in sbt, overrides are configured in sbt with the `dependencyOverrides` setting, which is a set of `ModuleIDs`. For example, the following dependency definitions conflict because spark uses log4j 1.2.16 and scalaxb uses log4j 1.2.17:

```
libraryDependencies += Seq(
  "org.spark-project" %% "spark-core" % "0.5.1",
  "org.scalaxb" %% "scalaxb" % "1.0.0"
)
```

The default conflict manager chooses the latest revision of log4j, 1.2.17:

```
> show update
[info] compile:
[info]    log4j:log4j:1.2.17: ...
...
[info]    (EVICTED) log4j:log4j:1.2.16
...
```

To change the version selected, add an override:

```
dependencyOverrides += "log4j" % "log4j" % "1.2.16"
```

This will not add a direct dependency on log4j, but will force the revision to be 1.2.16. This is confirmed by the output of `show update`:

```
> show update
[info] compile:
[info]    log4j:log4j:1.2.16
...
```

Note: this is an Ivy-only feature and will not be included in a published pom.xml.

Publishing See [Publishing](#) for how to publish your project.

Configurations Ivy configurations are a useful feature for your build when you need custom groups of dependencies, such as for a plugin. Ivy configurations are essentially named sets of dependencies. You can read the [Ivy documentation](#) for details.

The built-in use of configurations in sbt is similar to scopes in Maven. sbt adds dependencies to different classpaths by the configuration that they are defined in. See the description of [Maven Scopes](#) for details.

You put a dependency in a configuration by selecting one or more of its configurations to map to one or more of your project's configurations. The most common case is to have one of your configurations A use a dependency's configuration B. The mapping for this looks like "A->B". To apply this mapping to a dependency, add it to the end of your dependency definition:

```
libraryDependencies += "org.scalatest" %% "scalatest" % "2.1.3" % "test->compile"
```

This says that your project's "test" configuration uses `ScalaTest`'s "compile" configuration. See the [Ivy documentation](#) for more advanced mappings. Most projects published to Maven repositories will use the "compile" configuration.

A useful application of configurations is to group dependencies that are not used on normal classpaths. For example, your project might use a "js" configuration to automatically download jQuery and then include it in your jar by modifying `resources`. For example:

```
ivyConfigurations += config("js") hide
```

```
libraryDependencies += "jquery" % "jquery" % "1.3.2" % "js->default" from "http://jqueryjs.g
```

```
resources += update.value.select(configurationFilter("js"))
```

The `config` method defines a new configuration with name "js" and makes it private to the project so that it is not used for publishing. See [Update Report](#) for more information on selecting managed artifacts.

A configuration without a mapping (no "->") is mapped to "default" or "compile". The -> is only needed when mapping to a different configuration than those. The `ScalaTest` dependency above can then be shortened to:

```
libraryDependencies += "org.scalatest" %% "scalatest" % "2.1.3" % "test"
```

External Maven or Ivy For this method, create the configuration files as you would for Maven (`pom.xml`) or Ivy (`ivy.xml` and optionally `ivysettings.xml`). External configuration is selected by using one of the following expressions.

Ivy settings (resolver configuration)

```
externalIvySettings()
```

or

```
externalIvySettings(baseDirectory.value / "custom-settings-name.xml")
```

or

```
externalIvySettingsURL(url("your_url_here"))
```

Ivy file (dependency configuration)

```
externalIvyFile()
```

or

```
externalIvyFile(Def.setting(baseDirectory.value / "custom-name.xml"))
```

Because Ivy files specify their own configurations, sbt needs to know which configurations to use for the `compile`, `runtime`, and `test` classpaths. For example, to specify that the `Compile` classpath should use the ‘default’ configuration:

```
classpathConfiguration in Compile := config("default")
```

Maven pom (dependencies only)

```
externalPom()
```

or

```
externalPom(Def.setting(baseDirectory.value / "custom-name.xml"))
```

Full Ivy Example For example, a `build.sbt` using external Ivy files might look like:

```
externalIvySettings()

externalIvyFile(Def.setting(baseDirectory.value / "ivyA.xml"))

classpathConfiguration in Compile := Compile

classpathConfiguration in Test := Test

classpathConfiguration in Runtime := Runtime
```

Known limitations Maven support is dependent on Ivy's support for Maven POMs. Known issues with this support:

- Specifying `relativePath` in the `parent` section of a POM will produce an error.
- Ivy ignores repositories specified in the POM. A workaround is to specify repositories inline or in an Ivy `ivysettings.xml` file.

Proxy Repositories

It's often the case that users wish to set up a maven/ivy proxy repository inside their corporate firewall, and have developer sbt instances resolve artifacts through such a proxy. Let's detail what exact changes must be made for this to work.

Overview

The situation arises when many developers inside an organization are attempting to resolve artifacts. Each developer's machine will hit the internet and download an artifact, regardless of whether or not another on the team has already done so. Proxy repositories provide a single point of remote download for an organization. In addition to control and security concerns, Proxy repositories are primarily important for increased speed across a team.

There are many good proxy repository solutions out there, with the big three being (in alphabetical order):

- [Archiva](#)
- [Artifactory](#)
- [Nexus](#)



Figure 1: image

Once you have a proxy repository installed and configured, then it's time to configure sbt for your needs. Read the note at the bottom about proxy issues with ivy repositories.

sbt Configuration

sbt requires configuration in two places to make use of a proxy repository. The first is the `~/.sbt/repositories` file, and the second is the launcher script.

`~/.sbt/repositories`

The repositories file is an external configuration for the Launcher. The exact syntax for the configuration file is detailed in the [sbt Launcher](#).

Here's an example config:

```
[repositories]
  local
  my-ivy-proxy-releases: http://repo.company.com/ivy-releases/, [organization]/[module]/(scala
  my-maven-proxy-releases: http://repo.company.com/maven-releases/
```

This example configuration has three repositories configured for sbt.

The first resolver is `local`, and is used so that artifacts pushed using `publish-local` will be seen in other sbt projects.

The second resolver is `my-ivy-proxy-releases`. This repository is used to resolve sbt *itself* from the company proxy repository, as well as any sbt plugins that may be required. Note that the ivy resolver pattern is important, make sure that yours matches the one shown or you may not be able to resolve sbt plugins.

The final resolver is `my-maven-proxy-releases`. This repository is a proxy for all standard maven repositories, including maven central.

Launcher Script The sbt launcher supports two configuration options that allow the usage of proxy repositories. The first is the `sbt.override.build.repos` setting and the second is the `sbt.repository.config` setting.

sbt.override.build.repos This setting is used to specify that all sbt project added resolvers should be ignored in favor of those configured in the `repositories` configuration. Using this with a properly configured `~/.sbt/repositories` file leads to only your proxy repository used for builds.

It is specified like so:

```
-Dsbt.override.build.repos=true
```

sbt.repository.config If you are unable to create a `~/.sbt/repositories` file, due to user permission errors or for convenience of developers, you can modify the sbt start script directly with the following:

```
-Dsbt.repository.config=<path-to-your-repo-file>
```

This is only necessary if users do not already have their own default repository file.

Proxying Ivy Repositories

The most common mistake made when setting up a proxy repository for sbt is the attempting to *merge* both *maven* and *ivy* repositories into the *same* proxy repository. While some repository managers will allow this, it's not recommended to do so.

Even if your company does not use ivy, sbt uses a custom layout to handle binary compatibility constraints of its own plugins. To ensure that these are resolved correctly, simple set up two virtual/proxy repositories, one for maven and one for ivy.

Here's an example setup:

Publishing

This page describes how to publish your project. Publishing consists of uploading a descriptor, such as an Ivy file or Maven POM, and artifacts, such as a jar or war, to a repository so that other projects can specify your project as a dependency.

The **publish** action is used to publish your project to a remote repository. To use publishing, you need to specify the repository to publish to and the credentials to use. Once these are set up, you can run **publish**.

The **publishLocal** action is used to publish your project to a local Ivy repository. You can then use this project from other projects on the same machine.

Define the repository

To specify the repository, assign a repository to **publishTo** and optionally set the publishing style. For example, to upload to Nexus:

```
publishTo := Some("Sonatype Snapshots Nexus" at "https://oss.sonatype.org/content/repositories")
```

To publish to a local repository:



Figure 2: image

```
publishTo := Some(Resolver.file("file", new File( "path/to/my/maven-repo/releases" )) )
```

Publishing to the users local maven repository:

```
publishTo := Some(Resolver.file("file", new File(Path.userHome.absolutePath+"/.m2/repository")) )
```

If you're using Maven repositories you will also have to select the right repository depending on your artifacts: SNAPSHOT versions go to the /snapshot repository while other versions go to the /releases repository. Doing this selection can be done by using the value of the `version` `SettingKey`:

```
publishTo := {  
  val nexus = "https://oss.sonatype.org/"  
  if (version.value.trim.endsWith("SNAPSHOT"))  
    Some("snapshots" at nexus + "content/repositories/snapshots")  
  else  
    Some("releases" at nexus + "service/local/staging/deploy/maven2")  
}
```

Credentials

There are two ways to specify credentials for such a repository. The first is to specify them inline:

```
credentials += Credentials("Sonatype Nexus Repository Manager", "nexus.scala-tools.org", "admin:admin123")
```

The second and better way is to load them from a file, for example:

```
credentials += Credentials(Path.userHome / ".ivy2" / ".credentials")
```

The credentials file is a properties file with keys `realm`, `host`, `user`, and `password`. For example:

```
realm=Sonatype Nexus Repository Manager  
host=nexus.scala-tools.org  
user=admin  
password=admin123
```

Cross-publishing

To support multiple incompatible Scala versions, enable cross building and do `++ publish` (see [Cross Build](#)). See [Resolvers](#) for other supported repository types.

Published artifacts

By default, the main binary jar, a sources jar, and a API documentation jar are published. You can declare other types of artifacts to publish and disable or modify the default artifacts. See the [Artifacts](#) page for details.

Modifying the generated POM

When `publishMavenStyle` is `true`, a POM is generated by the `makePom` action and published to the repository instead of an Ivy file. This POM file may be altered by changing a few settings. Set `pomExtra` to provide XML (`scala.xml.NodeSeq`) to insert directly into the generated pom. For example:

```
pomExtra :=
  <licenses>
    <license>
      <name>Apache 2</name>
      <url>http://www.apache.org/licenses/LICENSE-2.0.txt</url>
      <distribution>repo</distribution>
    </license>
  </licenses>
```

`makePom` adds to the POM any Maven-style repositories you have declared. You can filter these by modifying `pomRepositoryFilter`, which by default excludes local repositories. To instead only include local repositories:

```
pomIncludeRepository := { (repo: MavenRepository) =>
  repo.root.startsWith("file:")
}
```

There is also a `pomPostProcess` setting that can be used to manipulate the final XML before it is written. It's type is `Node => Node`.

```
pomPostProcess := { (node: Node) =>
  ...
}
```

Publishing Locally

The `publishLocal` command will publish to the local Ivy repository. By default, this is in `${user.home}/.ivy2/local`. Other projects on the same machine can then list the project as a dependency. For example, if the SBT project you are publishing has configuration parameters like:

```
name := "My Project"

organization := "org.me"

version := "0.1-SNAPSHOT"
```

Then another project can depend on it:

```
libraryDependencies += "org.me" %% "my-project" % "0.1-SNAPSHOT"
```

The version number you select must end with **SNAPSHOT**, or you must change the version number each time you publish. Ivy maintains a cache, and it stores even local projects in that cache. If Ivy already has a version cached, it will not check the local repository for updates, unless the version number matches a [changing pattern](#), and **SNAPSHOT** is one such pattern.

Resolvers

Maven

Resolvers for Maven2 repositories are added as follows:

```
resolvers +=
  "Sonatype OSS Snapshots" at "https://oss.sonatype.org/content/repositories/snapshots"
```

This is the most common kind of user-defined resolvers. The rest of this page describes how to define other types of repositories.

Predefined

A few predefined repositories are available and are listed below

- **DefaultMavenRepository** This is the main Maven repository at <http://repo1.maven.org/maven2/> and is included by default
- **JavaNet1Repository** This is the Maven 1 repository at <http://download.java.net/maven/1/>

For example, to use the **java.net** repository, use the following setting in your build definition:

```
resolvers += JavaNet1Repository
```

Predefined repositories will go under Resolver going forward so they are in one place:

```
Resolver.sonatypeRepo("releases") // Or "snapshots"
```

Custom

sbt provides an interface to the repository types available in Ivy: file, URL, SSH, and SFTP. A key feature of repositories in Ivy is using [patterns](#) to configure repositories.

Construct a repository definition using the factory in `sbt.Resolver` for the desired type. This factory creates a `Repository` object that can be further configured. The following table contains links to the Ivy documentation for the repository type and the API documentation for the factory and repository class. The SSH and SFTP repositories are configured identically except for the name of the factory. Use `Resolver.ssh` for SSH and `Resolver.sftp` for SFTP.

Type

Factory

Ivy Docs

Factory API

Repository Class API

Filesystem

Resolver.file

Ivy filesystem

filesystem factory

FileRepository API

SFTP

Resolver.sftp

Ivy sftp

sftp factory

SftpRepository API

SSH

Resolver.ssh

Ivy ssh

ssh factory

SshRepository API

URL

Resolver.url

Ivy url

url factory

URLRepository API

Basic Examples These are basic examples that use the default Maven-style repository layout.

Filesystem Define a filesystem repository in the `test` directory of the current working directory and declare that publishing to this repository must be atomic.

```
resolvers += Resolver.file("my-test-repo", file("test")) transactional()
```

URL Define a URL repository at `"http://example.org/repo-releases/"`.

```
resolvers += Resolver.url("my-test-repo", url("http://example.org/repo-releases/"))
```

To specify an Ivy repository, use:

```
resolvers += Resolver.url("my-test-repo", url)(Resolver.ivyStylePatterns)
```

or customize the layout pattern described in the Custom Layout section below.

SFTP and SSH Repositories The following defines a repository that is served by SFTP from host `"example.org"`:

```
resolvers += Resolver.sftp("my-sftp-repo", "example.org")
```

To explicitly specify the port:

```
resolvers += Resolver.sftp("my-sftp-repo", "example.org", 22)
```

To specify a base path:

```
resolvers += Resolver.sftp("my-sftp-repo", "example.org", "maven2/repo-releases/")
```

Authentication for the repositories returned by `sftp` and `ssh` can be configured by the `as` methods.

To use password authentication:

```
resolvers += Resolver.ssh("my-ssh-repo", "example.org") as("user", "password")
```

or to be prompted for the password:

```
resolvers += Resolver.ssh("my-ssh-repo", "example.org") as("user")
```

To use key authentication:

```
resolvers += {  
    val keyFile: File = ...  
    Resolver.ssh("my-ssh-repo", "example.org") as("user", keyFile, "keyFilePassword")  
}
```

or if no keyfile password is required or if you want to be prompted for it:

```
resolvers += Resolver.ssh("my-ssh-repo", "example.org") as("user", keyFile)
```

To specify the permissions used when publishing to the server:

```
resolvers += Resolver.ssh("my-ssh-repo", "example.org") withPermissions("0644")
```

This is a chmod-like mode specification.

Custom Layout These examples specify custom repository layouts using patterns. The factory methods accept an `Patterns` instance that defines the patterns to use. The patterns are first resolved against the base file or URL. The default patterns give the default Maven-style layout. Provide a different `Patterns` object to use a different layout. For example:

```
resolvers += Resolver.url("my-test-repo", url)( Patterns("[organisation]/[module]/[revision]"
```

You can specify multiple patterns or patterns for the metadata and artifacts separately. You can also specify whether the repository should be Maven compatible (as defined by Ivy). See the [patterns API](#) for the methods to use.

For filesystem and URL repositories, you can specify absolute patterns by omitting the base URL, passing an empty `Patterns` instance, and using `ivys` and `artifacts`:

```
resolvers += Resolver.url("my-test-repo") artifacts  
    "http://example.org/[organisation]/[module]/[revision]/[artifact].[ext]"
```

Update Report

`update` and related tasks produce a value of type `sbt.UpdateReport`. This data structure provides information about the resolved configurations, modules, and artifacts. At the top level, `UpdateReport` provides reports of type `ConfigurationReport` for each resolved configuration. A `ConfigurationReport` supplies reports (of type `ModuleReport`) for each module resolved for a given configuration. Finally, a `ModuleReport` lists each successfully retrieved `Artifact` and the `File` it was retrieved to as well as the `Artifacts` that couldn't be downloaded. This missing `Artifact` list is always empty for `update`, which will fail if it is non-empty. However, it may be non-empty for `updateClassifiers` and `updateSbtClassifiers`.

Filtering a Report and Getting Artifacts

A typical use of `UpdateReport` is to retrieve a list of files matching a filter. A conversion of type `UpdateReport => RichUpdateReport` implicitly provides these methods for `UpdateReport`. The filters are defined by the `DependencyFilter`, `ConfigurationFilter`, `ModuleFilter`, and `ArtifactFilter` types. Using these filter types, you can filter by the configuration name, the module organization, name, or revision, and the artifact name, type, extension, or classifier.

The relevant methods (implicitly on `UpdateReport`) are:

```
def matching(f: DependencyFilter): Seq[File]

def select(configuration: ConfigurationFilter = ...,
            module: ModuleFilter = ...,
            artifact: ArtifactFilter = ...): Seq[File]
```

Any argument to `select` may be omitted, in which case all values are allowed for the corresponding component. For example, if the `ConfigurationFilter` is not specified, all configurations are accepted. The individual filter types are discussed below.

Filter Basics Configuration, module, and artifact filters are typically built by applying a `NameFilter` to each component of a `Configuration`, `ModuleID`, or `Artifact`. A basic `NameFilter` is implicitly constructed from a `String`, with `*` interpreted as a wildcard.

```
import sbt._
// each argument is of type NameFilter
val mf: ModuleFilter = moduleFilter(organization = "*sbt*",
                                   name = "main" | "actions", revision = "1.*" - "1.0")
```



```

// unspecified arguments match everything by default
val mf: ModuleFilter = moduleFilter(organization = "net.databinder")

// specifying "*" is the same as omitting the argument
val af: ArtifactFilter = artifactFilter(name = "*", `type` = "source",
    extension = "jar", classifier = "sources")

val cf: ConfigurationFilter = configurationFilter(name = "compile" | "test")

```

Alternatively, these filters, including a `NameFilter`, may be directly defined by an appropriate predicate (a single-argument function returning a Boolean).

```

import sbt._

// here the function value of type String => Boolean is implicitly converted to a NameFilter
val nf: NameFilter = (s: String) => s.startsWith("dispatch-")

// a Set[String] is a function String => Boolean
val acceptConfigs: Set[String] = Set("compile", "test")
// implicitly converted to a ConfigurationFilter
val cf: ConfigurationFilter = acceptConfigs

val mf: ModuleFilter = (m: ModuleID) => m.organization contains "sbt"

val af: ArtifactFilter = (a: Artifact) => a.classifier.isEmpty

```

ConfigurationFilter A configuration filter essentially wraps a `NameFilter` and is explicitly constructed by the `configurationFilter` method:

```
def configurationFilter(name: NameFilter = ...): ConfigurationFilter
```

If the argument is omitted, the filter matches all configurations. Functions of type `String => Boolean` are implicitly convertible to a `ConfigurationFilter`. As with `ModuleFilter`, `ArtifactFilter`, and `NameFilter`, the `&`, `|`, and `-` methods may be used to combine `ConfigurationFilters`.

```

import sbt._
val a: ConfigurationFilter = Set("compile", "test")
val b: ConfigurationFilter = (c: String) => c.startsWith("r")
val c: ConfigurationFilter = a | b

```

(The explicit types are optional here.)

ModuleFilter A module filter is defined by three `NameFilters`: one for the organization, one for the module name, and one for the revision. Each component filter must match for the whole module filter to match. A module filter is explicitly constructed by the `moduleFilter` method:

```
def moduleFilter(organization: NameFilter = ..., name: NameFilter = ..., revision: NameFilter = ...): ModuleFilter
```

An omitted argument does not contribute to the match. If all arguments are omitted, the filter matches all `ModuleIDs`. Functions of type `ModuleID => Boolean` are implicitly convertible to a `ModuleFilter`. As with `ConfigurationFilter`, `ArtifactFilter`, and `NameFilter`, the `&`, `|`, and `-` methods may be used to combine `ModuleFilters`:

```
import sbt._
val a: ModuleFilter = moduleFilter(name = "dispatch-twitter", revision = "0.7.8")
val b: ModuleFilter = moduleFilter(name = "dispatch-*")
val c: ModuleFilter = b - a
```

(The explicit types are optional here.)

ArtifactFilter An artifact filter is defined by four `NameFilters`: one for the name, one for the type, one for the extension, and one for the classifier. Each component filter must match for the whole artifact filter to match. An artifact filter is explicitly constructed by the `artifactFilter` method:

```
def artifactFilter(name: NameFilter = ..., `type`: NameFilter = ...,
  extension: NameFilter = ..., classifier: NameFilter = ...): ArtifactFilter
```

Functions of type `Artifact => Boolean` are implicitly convertible to an `ArtifactFilter`. As with `ConfigurationFilter`, `ModuleFilter`, and `NameFilter`, the `&`, `|`, and `-` methods may be used to combine `ArtifactFilters`:

```
import sbt._
val a: ArtifactFilter = artifactFilter(classifier = "javadoc")
val b: ArtifactFilter = artifactFilter(`type` = "jar")
val c: ArtifactFilter = b - a
```

(The explicit types are optional here.)

DependencyFilter A `DependencyFilter` is typically constructed by combining other `DependencyFilters` together using `&&`, `||`, and `--`. Configuration, module, and artifact filters are `DependencyFilters` themselves and can be used directly as a `DependencyFilter` or they can build up a `DependencyFilter`. Note that the symbols for the `DependencyFilter` combining methods are doubled up to distinguish them from the combinators of the more specific filters for configurations, modules, and artifacts. These double-character methods will always return a `DependencyFilter`, whereas the single character methods preserve the more specific filter type. For example:

```
import sbt._

val df: DependencyFilter =
  configurationFilter(name = "compile" | "test") &&
  artifactFilter(`type` = "jar") ||
  moduleFilter(name = "dispatch-*")
```

Here, we used `&&` and `||` to combine individual component filters into a dependency filter, which can then be provided to the `UpdateReport.matches` method. Alternatively, the `UpdateReport.select` method may be used, which is equivalent to calling `matches` with its arguments combined with `&&`.

Tasks and Commands

This part of the documentation has pages documenting particular sbt topics in detail. Before reading anything in here, you will need the information in the [Getting Started Guide](#) as a foundation.

Tasks

Tasks and settings are introduced in the [getting started guide](#), which you may wish to read first. This page has additional details and background and is intended more as a reference.

Introduction

Both settings and tasks produce values, but there are two major differences between them:

1. Settings are evaluated at project load time. Tasks are executed on demand, often in response to a command from the user.
2. At the beginning of project loading, settings and their dependencies are fixed. Tasks can introduce new tasks during execution, however.

Features

There are several features of the task system:

1. By integrating with the settings system, tasks can be added, removed, and modified as easily and flexibly as settings.
2. [Input Tasks](#) use [parser combinators](#) to define the syntax for their arguments. This allows flexible syntax and tab-completions in the same way as [Commands](#).
3. Tasks produce values. Other tasks can access a task's value by calling `value` on it within a task definition.
4. Dynamically changing the structure of the task graph is possible. Tasks can be injected into the execution graph based on the result of another task.
5. There are ways to handle task failure, similar to `try/catch/finally`.
6. Each task has access to its own Logger that by default persists the logging for that task at a more verbose level than is initially printed to the screen.

These features are discussed in detail in the following sections.

Defining a Task

Hello World example (sbt) `build.sbt`:

```
lazy val hello = taskKey[Unit]("Prints 'Hello World'")

hello := println("hello world!")
```

Run “sbt hello” from command line to invoke the task. Run “sbt tasks” to see this task listed.

Define the key To declare a new task, define a lazy val of type `TaskKey`:

```
lazy val sampleTask = taskKey[Int]("A sample task.")
```

The name of the `val` is used when referring to the task in Scala code and at the command line. The string passed to the `taskKey` method is a description of the task. The type parameter passed to `taskKey` (here, `Int`) is the type of value produced by the task.

We'll define a couple of other keys for the examples:

```
lazy val intTask = taskKey[Int]("An int task")
lazy val stringTask = taskKey[String]("A string task")
```

The examples themselves are valid entries in a `build.sbt` or can be provided as part of a sequence to `Project.settings` (see [.scala build definition](#)).

Implement the task There are three main parts to implementing a task once its key is defined:

1. Determine the settings and other tasks needed by the task. They are the task's inputs.
2. Define the code that implements the task in terms of these inputs.
3. Determine the scope the task will go in.

These parts are then combined just like the parts of a setting are combined.

Defining a basic task A task is defined using `:=`

```
intTask := 1 + 2

stringTask := System.getProperty("user.name")

sampleTask := {
  val sum = 1 + 2
  println("sum: " + sum)
  sum
}
```

As mentioned in the introduction, a task is evaluated on demand. Each time `sampleTask` is invoked, for example, it will print the sum. If the username changes between runs, `stringTask` will take different values in those separate runs. (Within a run, each task is evaluated at most once.) In contrast, settings are evaluated once on project load and are fixed until the next reload.

Tasks with inputs Tasks with other tasks or settings as inputs are also defined using `:=`. The values of the inputs are referenced by the `value` method. This method is special syntax and can only be called when defining a task, such as in the argument to `:=`. The following defines a task that adds one to the value produced by `intTask` and returns the result.

```
sampleTask := intTask.value + 1
```

Multiple settings are handled similarly:

```
stringTask := "Sample: " + sampleTask.value + ", int: " + intTask.value
```

Task Scope As with settings, tasks can be defined in a specific scope. For example, there are separate `compile` tasks for the `compile` and `test` scopes. The scope of a task is defined the same as for a setting. In the following example, `test:sampleTask` uses the result of `compile:intTask`.

```
sampleTask in Test := (intTask in Compile).value * 3
```

On precedence As a reminder, infix method precedence is by the name of the method and postfix methods have lower precedence than infix methods.

1. Assignment methods have the lowest precedence. These are methods with names ending in `=`, except for `!=`, `<=`, `>=`, and names that start with `=`.
2. Methods starting with a letter have the next highest precedence.
3. Methods with names that start with a symbol and aren't included in
 1. have the highest precedence. (This category is divided further according to the specific character it starts with. See the Scala specification for details.)

Therefore, the the previous example is equivalent to the following:

```
(sampleTask in Test).:=( (intTask in Compile).value * 3 )
```

Additionally, the braces in the following are necessary:

```
helloTask := { "echo Hello" ! }
```

Without them, Scala interprets the line as `(helloTask.:=("echo Hello")).!` instead of the desired `helloTask.:=("echo Hello".!)`.

Separating implementations The implementation of a task can be separated from the binding. For example, a basic separate definition looks like:

```
// Define a new, standalone task implementation
lazy val intTaskImpl: Initialize[Task[Int]] =
  Def.task { sampleTask.value - 3 }

// Bind the implementation to a specific key
intTask := intTaskImpl.value
```

Note that whenever `.value` is used, it must be within a task definition, such as within `Def.task` above or as an argument to `:=`.

Modifying an Existing Task In the general case, modify a task by declaring the previous task as an input.

```
// initial definition
intTask := 3

// overriding definition that references the previous definition
intTask := intTask.value + 1
```

Completely override a task by not declaring the previous task as an input. Each of the definitions in the following example completely overrides the previous one. That is, when `intTask` is run, it will only print `#3`.

```
intTask := {
  println("#1")
  3
}

intTask := {
  println("#2")
  5
}

intTask := {
  println("#3")
  sampleTask.value - 3
}
```

Getting values from multiple scopes

Introduction The general form of an expression that gets values from multiple scopes is:

```
<setting-or-task>.all(<scope-filter>).value
```

The `all` method is implicitly added to tasks and settings. It accepts a `ScopeFilter` that will select the `Scopes`. The result has type `Seq[T]`, where `T` is the key's underlying type.

Example A common scenario is getting the sources for all subprojects for processing all at once, such as passing them to `scaladoc`. The task that we want to obtain values for is `sources` and we want to get the values in all non-root projects and in the `Compile` configuration. This looks like:

```

lazy val core = project

lazy val util = project

lazy val root = project.settings(
  sources := {
    val filter = ScopeFilter( inProjects(core, util), inConfigurations(Compile) )
    // each sources definition is of type Seq[File],
    // giving us a Seq[Seq[File]] that we then flatten to Seq[File]
    val allSources: Seq[Seq[File]] = sources.all(filter).value
    allSources.flatten
  }
)

```

The next section describes various ways to construct a `ScopeFilter`.

ScopeFilter A basic `ScopeFilter` is constructed by the `ScopeFilter.apply` method. This method makes a `ScopeFilter` from filters on the parts of a `Scope`: a `ProjectFilter`, `ConfigurationFilter`, and `TaskFilter`. The simplest case is explicitly specifying the values for the parts:

```

val filter: ScopeFilter =
  ScopeFilter(
    inProjects( core, util ),
    inConfigurations( Compile, Test )
  )

```

Unspecified filters If the task filter is not specified, as in the example above, the default is to select scopes without a specific task (global). Similarly, an unspecified configuration filter will select scopes in the global configuration. The project filter should usually be explicit, but if left unspecified, the current project context will be used.

More on filter construction The example showed the basic methods `inProjects` and `inConfigurations`. This section describes all methods for constructing a `ProjectFilter`, `ConfigurationFilter`, or `TaskFilter`. These methods can be organized into four groups:

- Explicit member list (`inProjects`, `inConfigurations`, `inTasks`)
- Global value (`inGlobalProject`, `inGlobalConfiguration`, `inGlobalTask`)
- Default filter (`inAnyProject`, `inAnyConfiguration`, `inAnyTask`)
- Project relationships (`inAggregates`, `inDependencies`)

See the [API documentation](#) for details.

Combining ScopeFilters `ScopeFilters` may be combined with the `&&`, `||`, `--`, and `-` methods:

- `a && b` Selects scopes that match both `a` and `b`
- `a || b` Selects scopes that match either `a` or `b`
- `a -- b` Selects scopes that match `a` but not `b`
- `-b` Selects scopes that do not match `b`

For example, the following selects the scope for the `Compile` and `Test` configurations of the `core` project and the global configuration of the `util` project:

```
val filter: ScopeFilter =  
  ScopeFilter( inProjects(core), inConfigurations(Compile, Test)) ||  
  ScopeFilter( inProjects(util), inGlobalConfiguration )
```

More operations The `all` method applies to both settings (values of type `Initialize[T]`) and tasks (values of type `Initialize[Task[T]]`). It returns a setting or task that provides a `Seq[T]`, as shown in this table:

Target

Result

`Initialize[T]`

`Initialize[Seq[T]]`

`Initialize[Task[T]]`

`Initialize[Task[Seq[T]]]`

This means that the `all` method can be combined with methods that construct tasks and settings.

Missing values Some scopes might not define a setting or task. The `?` and `??` methods can help in this case. They are both defined on settings and tasks and indicate what to do when a key is undefined.

`?`

On a setting or task with underlying type `T`, this accepts no arguments and returns a setting or task (respectively) of type `Option[T]`. The result is `None` if the setting/task is undefined and `Some[T]` with the value if it is.

`??`

On a setting or task with underlying type `T`, this accepts an argument of type `T` and uses this argument if the setting/task is undefined.

The following contrived example sets the maximum errors to be the maximum of all aggregates of the current project.

```

maxErrors := {
  // select the transitive aggregates for this project, but not the project itself
  val filter: ScopeFilter =
    ScopeFilter( inAggregates(ThisProject, includeRoot=false) )
  // get the configured maximum errors in each selected scope,
  // using 0 if not defined in a scope
  val allVersions: Seq[Int] =
    (maxErrors ?? 0).all(filter).value
  allVersions.max
}

```

Multiple values from multiple scopes The target of `all` is any task or setting, including anonymous ones. This means it is possible to get multiple values at once without defining a new task or setting in each scope. A common use case is to pair each value obtained with the project, configuration, or full scope it came from.

- **resolvedScoped:** Provides the full enclosing `ScopedKey` (which is a `Scope + AttributeKey[_]`)
- **thisProject:** Provides the `Project` associated with this scope (undefined at the global and build levels)
- **thisProjectRef:** Provides the `ProjectRef` for the context (undefined at the global and build levels)
- **configuration:** Provides the `Configuration` for the context (undefined for the global configuration)

For example, the following defines a task that prints non-Compile configurations that define sbt plugins. This might be used to identify an incorrectly configured build (or not, since this is a fairly contrived example):

```

// Select all configurations in the current project except for Compile
lazy val filter: ScopeFilter = ScopeFilter(
  inProjects(ThisProject),
  inAnyConfiguration -- inConfigurations(Compile)
)

// Define a task that provides the name of the current configuration
// and the set of sbt plugins defined in the configuration
lazy val pluginsWithConfig: Initialize[Task[ (String, Set[String]) ]] =
  Def.task {
    ( configuration.value.name, definedSbtPlugins.value )
  }

checkPluginsTask := {

```

```

val oddPlugins: Seq[(String, Set[String])] =
  pluginsWithConfig.all(filter).value
// Print each configuration that defines sbt plugins
for( (config, plugins) <- oddPlugins if plugins.nonEmpty )
  println(s"$config defines sbt plugins: ${plugins.mkString(", ")}")
}

```

Advanced Task Operations

The examples in this section use the task keys defined in the previous section.

Streams: Per-task logging Per-task loggers are part of a more general system for task-specific data called Streams. This allows controlling the verbosity of stack traces and logging individually for tasks as well as recalling the last logging for a task. Tasks also have access to their own persisted binary or text data.

To use Streams, get the value of the `streams` task. This is a special task that provides an instance of `TaskStreams` for the defining task. This type provides access to named binary and text streams, named loggers, and a default logger. The default `Logger`, which is the most commonly used aspect, is obtained by the `log` method:

```

myTask := {
  val s: TaskStreams = streams.value
  s.log.debug("Saying hi...")
  s.log.info("Hello!")
}

```

You can scope logging settings by the specific task's scope:

```

logLevel in myTask := Level.Debug

traceLevel in myTask := 5

```

To obtain the last logging output from a task, use the `last` command:

```

$ last myTask
[debug] Saying hi...
[info] Hello!

```

The verbosity with which logging is persisted is controlled using the `persistLogLevel` and `persistTraceLevel` settings. The `last` command displays what was logged according to these levels. The levels do not affect already logged information.

Dynamic Computations with `Def.taskDyn`

It can be useful to use the result of a task to determine the next tasks to evaluate. This is done using `Def.taskDyn`. The result of `taskDyn` is called a dynamic task because it introduces dependencies at runtime. The `taskDyn` method supports the same syntax as `Def.task` and `:=` except that you return a task instead of a plain value.

For example,

```
val dynamic = Def.taskDyn {
  // decide what to evaluate based on the value of `stringTask`
  if(stringTask.value == "dev")
    // create the dev-mode task: this is only evaluated if the
    // value of stringTask is "dev"
    Def.task {
      3
    }
  else
    // create the production task: only evaluated if the value
    // of the stringTask is not "dev"
    Def.task {
      intTask.value + 5
    }
}

myTask := {
  val num = dynamic.value
  println(s"Number selected was $num")
}
```

The only static dependency of `myTask` is `stringTask`. The dependency on `intTask` is only introduced in non-dev mode.

Note: A dynamic task cannot refer to itself or a circular dependency will result. In the example above, there would be a circular dependency if the code passed to `taskDyn` referenced `myTask`.

Handling Failure This section discusses the `failure`, `result`, and `andFinally` methods, which are used to handle failure of other tasks.

failure The `failure` method creates a new task that returns the `Incomplete` value when the original task fails to complete normally. If the original task succeeds, the new task fails. `Incomplete` is an exception with information about

any tasks that caused the failure and any underlying exceptions thrown during task execution.

For example:

```
intTask := error("Failed.")

intTask := {
  println("Ignoring failure: " + intTask.failure.value)
  3
}
```

This overrides the `intTask` so that the original exception is printed and the constant 3 is returned.

`failure` does not prevent other tasks that depend on the target from failing. Consider the following example:

```
intTask := if(shouldSucceed) 5 else error("Failed.")

// Return 3 if intTask fails. If intTask succeeds, this task will fail.
aTask := intTask.failure.value - 2

// A new task that increments the result of intTask.
bTask := intTask.value + 1

cTask := aTask.value + bTask.value
```

The following table lists the results of each task depending on the initially invoked task:

invoked task
intTask result
aTask result
bTask result
cTask result
overall result
intTask
failure
not run
not run

not run
failure
aTask
failure
success
not run
not run
success
bTask
failure
not run
failure
not run
failure
cTask
failure
success
failure
failure
failure
intTask
success
not run
not run
not run
success
aTask
success
failure
not run
not run

```

failure
bTask
success
not run
success
not run
success
cTask
success
failure
success
failure
failure

```

The overall result is always the same as the root task (the directly invoked task). A **failure** turns a success into a failure, and a failure into an **Incomplete**. A normal task definition fails when any of its inputs fail and computes its value otherwise.

result The **result** method creates a new task that returns the full `Result[T]` value for the original task. **Result** has the same structure as `Either[Incomplete, T]` for a task result of type `T`. That is, it has two subtypes:

- **Inc**, which wraps **Incomplete** in case of failure
- **Value**, which wraps a task's result in case of success.

Thus, the task created by **result** executes whether or not the original task succeeds or fails.

For example:

```

intTask := error("Failed.")

intTask := intTask.result.value match {
  case Inc(inc: Incomplete) =>
    println("Ignoring failure: " + inc)
    3
  case Value(v) =>

```

```

        println("Using successful result: " + v)
      }
    }
  }
}

```

This overrides the original `intTask` definition so that if the original task fails, the exception is printed and the constant 3 is returned. If it succeeds, the value is printed and returned.

andFinally The `andFinally` method defines a new task that runs the original task and evaluates a side effect regardless of whether the original task succeeded. The result of the task is the result of the original task. For example:

```

intTask := error("I didn't succeed.")

lazy val intTaskImpl = intTask andFinally { println("andFinally") }

intTask := intTaskImpl.value

```

This modifies the original `intTask` to always print “andFinally” even if the task fails.

Note that `andFinally` constructs a new task. This means that the new task has to be invoked in order for the extra block to run. This is important when calling `andFinally` on another task instead of overriding a task like in the previous example. For example, consider this code:

```

intTask := error("I didn't succeed.")

lazy val intTaskImpl = intTask andFinally { println("andFinally") }

otherIntTask := intTaskImpl.value

```

If `intTask` is run directly, `otherIntTask` is never involved in execution. This case is similar to the following plain Scala code:

```

def intTask(): Int =
  error("I didn't succeed.")

def otherIntTask(): Int =
  try { intTask() }
  finally { println("finally") }

intTask()

```

It is obvious here that calling `intTask()` will never result in “finally” being printed.

Input Tasks

Input Tasks parse user input and produce a task to run. [Parsing Input](#) describes how to use the parser combinators that define the input syntax and tab completion. This page describes how to hook those parser combinators into the input task system.

Input Keys

A key for an input task is of type `InputKey` and represents the input task like a `SettingKey` represents a setting or a `TaskKey` represents a task. Define a new input task key using the `inputKey.apply` factory method:

```
// goes in project/Build.scala or in build.sbt  
val demo = inputKey[Unit]("A demo input task.")
```

The definition of an input task is similar to that of a normal task, but it can also use the result of a

[Parser](#) applied to user input. Just as the special `value` method gets the value of a setting or task, the special `parsed` method gets the result of a `Parser`.

Basic Input Task Definition

The simplest input task accepts a space-delimited sequence of arguments. It does not provide useful tab completion and parsing is basic. The built-in parser for space-delimited arguments is constructed via the `spaceDelimited` method, which accepts as its only argument the label to present to the user during tab completion.

For example, the following task prints the current Scala version and then echoes the arguments passed to it on their own line.

```
demo := {  
  // get the result of parsing  
  val args: Seq[String] = spaceDelimited("<arg>").parsed  
  // Here, we also use the value of the `scalaVersion` setting  
  println("The current Scala version is " + scalaVersion.value)  
  println("The arguments to demo were:")  
  args foreach println  
}
```

Input Task using Parsers

The Parser provided by the `spaceDelimited` method does not provide any flexibility in defining the input syntax. Using a custom parser is just a matter of defining your own `Parser` as described on the [Parsing Input](#) page.

Constructing the Parser The first step is to construct the actual `Parser` by defining a value of one of the following types:

- `Parser[I]`: a basic parser that does not use any settings
- `Initialize[Parser[I]]`: a parser whose definition depends on one or more settings
- `Initialize[State => Parser[I]]`: a parser that is defined using both settings and the current [state](#)

We already saw an example of the first case with `spaceDelimited`, which doesn't use any settings in its definition. As an example of the third case, the following defines a contrived `Parser` that uses the project's Scala and sbt version settings as well as the state. To use these settings, we need to wrap the Parser construction in `Def.setting` and get the setting values with the special `value` method:

```
import complete.DefaultParsers._

val parser: Initialize[State => Parser[(String,String)]] =
Def.setting {
  (state: State) =>
    ( token("scala" <~ Space) ~ token(scalaVersion.value) ) |
    ( token("sbt" <~ Space) ~ token(sbtVersion.value) ) |
    ( token("commands" <~ Space) ~
      token(state.remainingCommands.size.toString) )
}
```

This Parser definition will produce a value of type `(String,String)`. The input syntax defined isn't very flexible; it is just a demonstration. It will produce one of the following values for a successful parse (assuming the current Scala version is 2.10.3, the current sbt version is 0.13.5, and there are 3 commands left to run):

Again, we were able to access the current Scala and sbt version for the project because they are settings. Tasks cannot be used to define the parser.

Constructing the Task Next, we construct the actual task to execute from the result of the `Parser`. For this, we define a task as usual, but we can access the result of parsing via the special `parsed` method on `Parser`.

The following contrived example uses the previous example's output (of type `(String,String)`) and the result of the `package` task to print some information to the screen.

```
demo := {  
  val (tpe, value) = parser.parsed  
  println("Type: " + tpe)  
  println("Value: " + value)  
  println("Packaged: " + packageBin.value.getAbsolutePath)  
}
```

The `InputTask` type

It helps to look at the `InputTask` type to understand more advanced usage of input tasks. The core input task type is:

```
class InputTask[T](val parser: State => Parser[Task[T]])
```

Normally, an input task is assigned to a setting and you work with `Initialize[InputTask[T]]`.

Breaking this down,

1. You can use other settings (via `Initialize`) to construct an input task.
2. You can use the current `State` to construct the parser.
3. The parser accepts user input and provides tab completion.
4. The parser produces the task to run.

So, you can use settings or `State` to construct the parser that defines an input task's command line syntax. This was described in the previous section. You can then use settings, `State`, or user input to construct the task to run. This is implicit in the input task syntax.

Using other input tasks

The types involved in an input task are composable, so it is possible to reuse input tasks. The `.parsed` and `.evaluated` methods are defined on `InputTasks` to make this more convenient in common situations:

- Call `.parsed` on an `InputTask[T]` or `Initialize[InputTask[T]]` to get the `Task[T]` created after parsing the command line
- Call `.evaluated` on an `InputTask[T]` or `Initialize[InputTask[T]]` to get the value of type `T` from evaluating that task

In both situations, the underlying `Parser` is sequenced with other parsers in the input task definition. In the case of `.evaluated`, the generated task is evaluated.

The following example applies the `run` input task, a literal separator parser `--`, and `run` again. The parsers are sequenced in order of syntactic appearance, so that the arguments before `--` are passed to the first `run` and the ones after are passed to the second.

```
val run2 = inputKey[Unit](
    "Runs the main class twice with different argument lists separated by --")

val separator: Parser[String] = "--"

run2 := {
    val one = (run in Compile).evaluated
    val sep = separator.parsed
    val two = (run in Compile).evaluated
}
```

For a main class `Demo` that echoes its arguments, this looks like:

```
$ sbt
> run2 a b -- c d
[info] Running Demo c d
[info] Running Demo a b
c
d
a
b
```

Preapplying input

Because `InputTasks` are built from `Parsers`, it is possible to generate a new `InputTask` by applying some input programmatically. (It is also possible to generate a `Task`, which is covered in the next section.) Two convenience methods are provided on `InputTask[T]` and `Initialize[InputTask[T]]` that accept the `String` to apply.

- `partialInput` applies the input and allows further input, such as from the command line
- `fullInput` applies the input and terminates parsing, so that further input is not accepted

In each case, the input is applied to the input task's parser. Because input tasks handle all input after the task name, they usually require initial whitespace to be provided in the input.

Consider the example in the previous section. We can modify it so that we:

- Explicitly specify all of the arguments to the first `run`. We use `name` and `version` to show that settings can be used to define and modify parsers.
- Define the initial arguments passed to the second `run`, but allow further input on the command line.

Note: the current implementation of `:=` doesn't actually support applying input derived from settings yet.

```
lazy val run2 = inputKey[Unit]("Runs the main class twice: " +
  "once with the project name and version as arguments"
  "and once with command line arguments preceded by hard coded values.")

// The argument string for the first run task is ' <name> <version>'
lazy val firstInput: Initialize[String] =
  Def.setting(s" ${name.value} ${version.value}")

// Make the first arguments to the second run task ' red blue'
lazy val secondInput: String = " red blue"

run2 := {
  val one = (run in Compile).fullInput(firstInput.value).evaluated
  val two = (run in Compile).partialInput(secondInput).evaluated
}
```

For a main class `Demo` that echoes its arguments, this looks like:

```
$ sbt
> run2 green
[info] Running Demo demo 1.0
[info] Running Demo red blue green
demo
1.0
red
blue
green
```

Get a Task from an InputTask

The previous section showed how to derive a new `InputTask` by applying input. In this section, applying input produces a `Task`. The `toTask` method on

`Initialize[InputTask[T]]` accepts the `String` input to apply and produces a task that can be used normally. For example, the following defines a plain task `runFixed` that can be used by other tasks or run directly without providing any input, :

```
lazy val runFixed = taskKey[Unit]("A task that hard codes the values to `run`")

runFixed := {
    val _ = (run in Compile).toTask(" blue green").value
    println("Done!")
}
```

For a main class `Demo` that echoes its arguments, running `runFixed` looks like:

```
$ sbt
> runFixed
[info] Running Demo blue green
blue
green
Done!
```

Each call to `toTask` generates a new task, but each task is configured the same as the original `InputTask` (in this case, `run`) but with different input applied. For example, :

```
lazy val runFixed2 = taskKey[Unit]("A task that hard codes the values to `run`")

fork in run := true

runFixed2 := {
    val x = (run in Compile).toTask(" blue green").value
    val y = (run in Compile).toTask(" red orange").value
    println("Done!")
}
```

The different `toTask` calls define different tasks that each run the project's main class in a new jvm. That is, the `fork` setting configures both, each has the same classpath, and each run the same main class. However, each task passes different arguments to the main class. For a main class `Demo` that echoes its arguments, the output of running `runFixed2` might look like:

```
$ sbt
> runFixed2
[info] Running Demo blue green
```

```
[info] Running Demo red orange
blue
green
red
orange
Done!
```

Commands

What is a “command”?

A “command” looks similar to a task: it’s a named operation that can be executed from the sbt console.

However, a command’s implementation takes as its parameter the entire state of the build (represented by [State](#)) and computes a new [State](#). This means that a command can look at or modify other sbt settings, for example. Typically, you would resort to a command when you need to do something that’s impossible in a regular task.

Introduction

There are three main aspects to commands:

1. The syntax used by the user to invoke the command, including:
 - Tab completion for the syntax
 - The parser to turn input into an appropriate data structure
2. The action to perform using the parsed data structure. This action transforms the build [State](#).
3. Help provided to the user

In sbt, the syntax part, including tab completion, is specified with parser combinators. If you are familiar with the parser combinators in Scala’s standard library, these are very similar. The action part is a function `(State, T) => State`, where `T` is the data structure produced by the parser. See the [Parsing Input](#) page for how to use the parser combinators.

[State](#) provides access to the build state, such as all registered [Commands](#), the remaining commands to execute, and all project-related information. See [States and Actions](#) for details on State.

Finally, basic help information may be provided that is used by the `help` command to display command help.

Defining a Command

A command combines a function `State => Parser[T]` with an action `(State, T) => State`. The reason for `State => Parser[T]` and not simply `Parser[T]` is that often the current `State` is used to build the parser. For example, the currently loaded projects (provided by `State`) determine valid completions for the `project` command. Examples for the general and specific cases are shown in the following sections.

See [Command.scala](#) for the source API details for constructing commands.

General commands General command construction looks like:

```
val action: (State, T) => State = ...
val parser: State => Parser[T] = ...
val command: Command = Command("name")(parser)(action)
```

No-argument commands There is a convenience method for constructing commands that do not accept any arguments.

```
val action: State => State = ...
val command: Command = Command.command("name")(action)
```

Single-argument command There is a convenience method for constructing commands that accept a single argument with arbitrary content.

```
// accepts the state and the single argument
val action: (State, String) => State = ...
val command: Command = Command.single("name")(action)
```

Multi-argument command There is a convenience method for constructing commands that accept multiple arguments separated by spaces.

```
val action: (State, Seq[String]) => State = ...

// <arg> is the suggestion printed for tab completion on an argument
val command: Command = Command.args("name", "<arg>")(action)
```


Full Example

The following example is a valid `project/Build.scala` that adds commands to a project. To try it out:

1. Copy the following build definition into `project/Build.scala` for a new project.
2. Run `sbt` on the project.
3. Try out the `hello`, `helloAll`, `failIfTrue`, `color`, and `printState` commands.
4. Use tab-completion and the code below as guidance.

```
import sbt._
import Keys._

// imports standard command parsing functionality
import complete.DefaultParsers._

object CommandExample extends Build {
  // Declare a single project, adding several new commands, which are discussed below.
  lazy override val projects = Seq(root)
  lazy val root = Project("root", file(".")) settings(
    commands ++= Seq(hello, helloAll, failIfTrue, changeColor, printState)
  )

  // A simple, no-argument command that prints "Hi",
  // leaving the current state unchanged.
  def hello = Command.command("hello") { state =>
    println("Hi!")
    state
  }

  // A simple, multiple-argument command that prints "Hi" followed by the arguments.
  // Again, it leaves the current state unchanged.
  def helloAll = Command.args("helloAll", "<name>") { (state, args) =>
    println("Hi " + args.mkString(" "))
    state
  }

  // A command that demonstrates failing or succeeding based on the input
  def failIfTrue = Command.single("failIfTrue") {
    case (state, "true") => state.fail
    case (state, _) => state
  }
```

```

}

// Demonstration of a custom parser.
// The command changes the foreground or background terminal color
// according to the input.
lazy val change = Space ~> (reset | setColor)
lazy val reset = token("reset" ^^^ "\033[0m")
lazy val color = token( Space ~> ("blue" ^^^ "4" | "green" ^^^ "2") )
lazy val select = token( "fg" ^^^ "3" | "bg" ^^^ "4" )
lazy val setColor = (select ~ color) map { case (g, c) => "\033[" + g + c + "m" }

def changeColor = Command("color")(_ => change) { (state, ansicode) =>
  print(ansicode)
  state
}

// A command that demonstrates getting information out of State.
def printState = Command.command("printState") { state =>
  import state._
  println(definedCommands.size + " registered commands")
  println("commands to run: " + show(remainingCommands))
  println()

  println("original arguments: " + show(configuration.arguments))
  println("base directory: " + configuration.baseDirectory)
  println()

  println("sbt version: " + configuration.provider.id.version)
  println("Scala version (for sbt): " + configuration.provider.scalaProvider.version)
  println()

  val extracted = Project.extract(state)
  import extracted._
  println("Current build: " + currentRef.build)
  println("Current project: " + currentRef.project)
  println("Original setting count: " + session.original.size)
  println("Session setting count: " + session.append.size)

  state
}

def show[T](s: Seq[T]) =
  s.map("'" + _ + "'").mkString("[", ", ", "]")
}

```

Parsing and tab completion

This page describes the parser combinators in sbt. These parser combinators are typically used to parse user input and provide tab completion for [Input Tasks](#) and [Commands](#). If you are already familiar with Scala's parser combinators, the methods are mostly the same except that their arguments are strict. There are two additional methods for controlling tab completion that are discussed at the end of the section.

Parser combinators build up a parser from smaller parsers. A `Parser[T]` in its most basic usage is a function `String => Option[T]`. It accepts a `String` to parse and produces a value wrapped in `Some` if parsing succeeds or `None` if it fails. Error handling and tab completion make this picture more complicated, but we'll stick with `Option` for this discussion.

The following examples assume the imports: :

```
import sbt._
import complete.DefaultParsers._
```

Basic parsers

The simplest parser combinators match exact inputs:

```
// A parser that succeeds if the input is 'x', returning the Char 'x'
// and failing otherwise
val singleChar: Parser[Char] = 'x'

// A parser that succeeds if the input is "blue", returning the String "blue"
// and failing otherwise
val litString: Parser[String] = "blue"
```

In these examples, implicit conversions produce a literal `Parser` from a `Char` or `String`. Other basic parser constructors are the `charClass`, `success` and `failure` methods:

```
// A parser that succeeds if the character is a digit, returning the matched Char
// The second argument, "digit", describes the parser and is used in error messages
val digit: Parser[Char] = charClass( (c: Char) => c.isDigit, "digit")

// A parser that produces the value 3 for an empty input string, fails otherwise
val alwaysSucceed: Parser[Int] = success( 3 )

// Represents failure (always returns None for an input String).
// The argument is the error message.
val alwaysFail: Parser[Nothing] = failure("Invalid input.")
```

Built-in parsers

sbt comes with several built-in parsers defined in [sbt.complete.DefaultParsers](#). Some commonly used built-in parsers are:

- `Space`, `NotSpace`, `OptSpace`, and `OptNotSpace` for parsing spaces or non-spaces, required or not.
- `StringBasic` for parsing text that may be quoted.
- `IntBasic` for parsing a signed `Int` value.
- `Digit` and `HexDigit` for parsing a single decimal or hexadecimal digit.
- `Bool` for parsing a `Boolean` value

See the [DefaultParsers API](#) for details.

Combining parsers

We build on these basic parsers to construct more interesting parsers. We can combine parsers in a sequence, choose between parsers, or repeat a parser.

```
// A parser that succeeds if the input is "blue" or "green",
// returning the matched input
val color: Parser[String] = "blue" | "green"

// A parser that matches either "fg" or "bg"
val select: Parser[String] = "fg" | "bg"

// A parser that matches "fg" or "bg", a space, and then the color, returning the matched value
// ~ is an alias for Tuple2.
val setColor: Parser[String ~ Char ~ String] =
  select ~ ' ' ~ color

// Often, we don't care about the value matched by a parser, such as the space above
// For this, we can use ~> or <~, which keep the result of
// the parser on the right or left, respectively
val setColor2: Parser[String ~ String] = select ~ (' ' ~> color)

// Match one or more digits, returning a list of the matched characters
val digits: Parser[Seq[Char]] = charClass(_.isDigit, "digit").+

// Match zero or more digits, returning a list of the matched characters
val digits0: Parser[Seq[Char]] = charClass(_.isDigit, "digit").*

// Optionally match a digit
val optDigit: Parser[Option[Char]] = charClass(_.isDigit, "digit").?
```

Transforming results

A key aspect of parser combinators is transforming results along the way into more useful data structures. The fundamental methods for this are `map` and `flatMap`. Here are examples of `map` and some convenience methods implemented on top of `map`.

```
// Apply the `digits` parser and apply the provided function to the matched
// character sequence
val num: Parser[Int] = digits map { (chars: Seq[Char]) => chars.mkString.toInt }

// Match a digit character, returning the matched character or return '0' if the input is not a digit
val digitWithDefault: Parser[Char] = charClass(_.isDigit, "digit") ?? '0'

// The previous example is equivalent to:
val digitDefault: Parser[Char] =
  charClass(_.isDigit, "digit").? map { (d: Option[Char]) => d.getOrElse '0' }

// Succeed if the input is "blue" and return the value 4
val blue = "blue" ^^ 4

// The above is equivalent to:
val blueM = "blue" map { (s: String) => 4 }
```

Controlling tab completion

Most parsers have reasonable default tab completion behavior. For example, the string and character literal parsers will suggest the underlying literal for an empty input string. However, it is impractical to determine the valid completions for `charClass`, since it accepts an arbitrary predicate. The `examples` method defines explicit completions for such a parser:

```
val digit = charClass(_.isDigit, "digit").examples("0", "1", "2")
```

Tab completion will use the examples as suggestions. The other method controlling tab completion is `token`. The main purpose of `token` is to determine the boundaries for suggestions. For example, if your parser is:

```
("fg" | "bg") ~ ' ' ~ ("green" | "blue")
```

then the potential completions on empty input are: `console fg green fg blue bg green bg blue`

Typically, you want to suggest smaller segments or the number of suggestions becomes unmanageable. A better parser is:

```
token( ("fg" | "bg") ~ ' ' ) ~ token("green" | "blue")
```

Now, the initial suggestions would be (with `_` representing a space): `console fg_ bg_`

Be careful not to overlap or nest tokens, as in `token("green" ~ token("blue"))`. The behavior is unspecified (and should generate an error in the future), but typically the outer most token definition will be used.

State and actions

`State` is the entry point to all available information in sbt. The key methods are:

- `definedCommands: Seq[Command]` returns all registered Command definitions
- `remainingCommands: Seq[String]` returns the remaining commands to be run
- `attributes: AttributeMap` contains generic data.

The action part of a command performs work and transforms `State`. The following sections discuss `State => State` transformations. As mentioned previously, a command will typically handle a parsed value as well: `(State, T) => State`.

Command-related data

A Command can modify the currently registered commands or the commands to be executed. This is done in the action part by transforming the (immutable) State provided to the command. A function that registers additional power commands might look like:

```
val powerCommands: Seq[Command] = ...

val addPower: State => State =
  (state: State) =>
    state.copy(definedCommands =
      (state.definedCommands ++ powerCommands).distinct
    )
```

This takes the current commands, appends new commands, and drops duplicates. Alternatively, State has a convenience method for doing the above:

```
val addPower2 = (state: State) => state ++ powerCommands
```

Some examples of functions that modify the remaining commands to execute:

```
val appendCommand: State => State =
  (state: State) =>
    state.copy(remainingCommands = state.remainingCommands :+ "cleanup")

val insertCommand: State => State =
  (state: State) =>
    state.copy(remainingCommands = "next-command" +: state.remainingCommands)
```

The first adds a command that will run after all currently specified commands run. The second inserts a command that will run next. The remaining commands will run after the inserted command completes.

To indicate that a command has failed and execution should not continue, return `state.fail`.

```
(state: State) => {
  val success: Boolean = ...
  if(success) state else state.fail
}
```

Project-related data

Project-related information is stored in `attributes`. Typically, commands won't access this directly but will instead use a convenience method to extract the most useful information:

```
val state: State
val extracted: Extracted = Project.extract(state)
import extracted._
```

`Extracted` provides:

- Access to the current build and project (`currentRef`)
- Access to initialized project setting data (`structure.data`)
- Access to session `Settings` and the original, permanent settings from `.sbt` and `.scala` files (`session.append` and `session.original`, respectively)
- Access to the current `Eval` instance for evaluating Scala expressions in the build context.

Project data

All project data is stored in `structure.data`, which is of type `sbt.Settings[Scope]`. Typically, one gets information of type `T` in the following way:

```
val key: SettingKey[T]
val scope: Scope
val value: Option[T] = key in scope get structure.data
```

Here, a `SettingKey[T]` is typically obtained from `Keys` and is the same type that is used to define settings in `.sbt` files, for example. `Scope` selects the scope the key is obtained for. There are convenience overloads of `in` that can be used to specify only the required scope axes. See `Structure.scala` for where `in` and other parts of the settings interface are defined. Some examples:

```
import Keys._
val extracted: Extracted
import extracted._

// get name of current project
val nameOpt: Option[String] = name in currentRef get structure.data

// get the package options for the `test:packageSrc` task or Nil if none are defined
val pkgOpts: Seq[PackageOption] = packageOptions in (currentRef, Test, packageSrc) get structure.data
```

`BuildStructure` contains information about build and project relationships. Key members are:

```
units: Map[URI, LoadedBuildUnit]
root: URI
```

A `URI` identifies a build and `root` identifies the initial build loaded. `LoadedBuildUnit` provides information about a single build. The key members of `LoadedBuildUnit` are:

```
// Defines the base directory for the build
localBase: File

// maps the project ID to the Project definition
defined: Map[String, ResolvedProject]
```

`ResolvedProject` has the same information as the `Project` used in a `project/Build.scala` except that `ProjectReferences` are resolved to `ProjectRefs`.

Classpaths

Classpaths in sbt 0.10+ are of type `Seq[Attributed[File]]`. This allows tagging arbitrary information to classpath entries. sbt currently uses this to associate an `Analysis` with an entry. This is how it manages the information needed for multi-project incremental recompilation. It also associates the `ModuleID` and `Artifact` with managed entries (those obtained by dependency management). When you only want the underlying `Seq[File]`, use `files`:

```
val attributedClasspath: Seq[Attributed[File]] = ...
val classpath: Seq[File] = attributedClasspath.files
```

Running tasks

It can be useful to run a specific project task from a `command` (*not from another task*) and get its result. For example, an IDE-related command might want to get the classpath from a project or a task might analyze the results of a compilation. The relevant method is `Project.evaluateTask`, which has the following signature:

```
def evaluateTask[T](taskKey: ScopedKey[Task[T]], state: State,
  checkCycles: Boolean = false, maxWorkers: Int = ...): Option[Result[T]]
```

For example,

```
val eval: State => State = (state: State) => {

  // This selects the main 'compile' task for the current project.
  // The value produced by 'compile' is of type inc.Analysis,
  // which contains information about the compiled code.
  val taskKey = Keys.compile in Compile

  // Evaluate the task
  // None if the key is not defined
  // Some(Inc) if the task does not complete successfully (Inc for incomplete)
  // Some(Value(v)) with the resulting value
  val result: Option[Result[inc.Analysis]] = Project.evaluateTask(taskKey, state)
  // handle the result
  result match
  {
    case None => // Key wasn't defined.
    case Some(Inc(inc)) => // error detail, inc is of type Incomplete, use Incomplete.s
    case Some(Value(v)) => // do something with v: inc.Analysis
  }
}
```

For getting the test classpath of a specific project, use this key:

```
val projectRef: ProjectRef = ...
val taskKey: Task[Seq[Attributed[File]]] =
  Keys.fullClasspath in (projectRef, Test)
```

Using State in a task

To access the current State from a task, use the `state` task as an input. For example,

```
myTask := ... state.value ...
```

Tasks/Settings: Motivation

This page motivates the task and settings system. You should already know how to use tasks and settings, which are described in the [getting started guide](#) and on the [Tasks](#) page.

An important aspect of the task system is to combine two common, related steps in a build:

1. Ensure some other task is performed.
2. Use some result from that task.

Earlier versions of sbt configured these steps separately using

1. Dependency declarations
2. Some form of shared state

To see why it is advantageous to combine them, compare the situation to that of deferring initialization of a variable in Scala. This Scala code is a bad way to expose a value whose initialization is deferred:

```
// Define a variable that will be initialized at some point
// We don't want to do it right away, because it might be expensive
var foo: Foo = _

// Define a function to initialize the variable
def makeFoo(): Unit = ... initialize foo ...
```

Typical usage would be:

```
makeFoo()  
doSomething(foo)
```

This example is rather exaggerated in its badness, but I claim it is nearly the same situation as our two step task definitions. Particular reasons this is bad include:

1. A client needs to know to call `makeFoo()` first.
2. `foo` could be changed by other code. There could be a `def makeFoo2()`, for example.
3. Access to `foo` is not thread safe.

The first point is like declaring a task dependency, the second is like two tasks modifying the same state (either project variables or files), and the third is a consequence of unsynchronized, shared state.

In Scala, we have the built-in functionality to easily fix this: `lazy val`.

```
lazy val foo: Foo = ... initialize foo ...
```

with the example usage:

```
doSomething(foo)
```

Here, `lazy val` gives us thread safety, guaranteed initialization before access, and immutability all in one, DRY construct. The task system in sbt does the same thing for tasks (and more, but we won't go into that here) that `lazy val` did for our bad example.

A task definition must declare its inputs and the type of its output. sbt will ensure that the input tasks have run and will then provide their results to the function that implements the task, which will generate its own result. Other tasks can use this result and be assured that the task has run (once) and be thread-safe and typesafe in the process.

The general form of a task definition looks like:

```
myTask := {  
  val a: A = aTask.value  
  val b: B = bTask.value  
  ... do something with a, b and generate a result ...  
}
```

(This is only intended to be a discussion of the ideas behind tasks, so see the [sbt Tasks](#) page for details on usage.) Here, `aTask` is assumed to produce a result of type `A` and `bTask` is assumed to produce a result of type `B`.

Application

As an example, consider generating a zip file containing the binary jar, source jar, and documentation jar for your project. First, determine what tasks produce the jars. In this case, the input tasks are `packageBin`, `packageSrc`, and `packageDoc` in the main `Compile` scope. The result of each of these tasks is the `File` for the jar that they generated. Our zip file task is defined by mapping these package tasks and including their outputs in a zip file. As good practice, we then return the `File` for this zip so that other tasks can map on the zip task.

```
zip := {  
  val bin: File = (packageBin in Compile).value  
  val src: File = (packageSrc in Compile).value  
  val doc: File = (packageDoc in Compile).value  
  val out: File = zipPath.value  
  val inputs: Seq[(File,String)] = Seq(bin, src, doc) x Path.flat  
  IO.zip(inputs, out)  
  out  
}
```

The `val inputs` line defines how the input files are mapped to paths in the zip. See [Mapping Files](#) for details. The explicit types are not required, but are included for clarity.

The `zipPath` input would be a custom task to define the location of the zip file. For example:

```
zipPath := target.value / "out.zip"
```

Plugins and Best Practices

This part of the documentation has pages documenting particular sbt topics in detail. Before reading anything in here, you will need the information in the [Getting Started Guide](#) as a foundation.

General Best Practices

This page describes best practices for working with sbt.

`project/` vs. `~/.sbt/`

Anything that is necessary for building the project should go in `project/`. This includes things like the web plugin. `~/.sbt/` should contain local customizations and commands for working with a build, but are not necessary. An example is an IDE plugin.

Local settings

There are two options for settings that are specific to a user. An example of such a setting is inserting the local Maven repository at the beginning of the resolvers list:

```
resolvers := {  
  val localMaven = "Local Maven Repository" at "file://" + Path.userHome.absolutePath + "/.m2/r  
  localMaven += resolvers.value  
}
```

1. Put settings specific to a user in a global `.sbt` file, such as `~/.sbt/0.13/global.sbt`. These settings will be applied to all projects.
2. Put settings in a `.sbt` file in a project that isn't checked into version control, such as `<project>/local.sbt`. sbt combines the settings from multiple `.sbt` files, so you can still have the standard `<project>/build.sbt` and check that into version control.

`.sbtrc`

Put commands to be executed when sbt starts up in a `.sbtrc` file, one per line. These commands run before a project is loaded and are useful for defining aliases, for example. sbt executes commands in `$HOME/.sbtrc` (if it exists) and then `<project>/.sbtrc` (if it exists).

Generated files

Write any generated files to a subdirectory of the output directory, which is specified by the `target` setting. This makes it easy to clean up after a build and provides a single location to organize generated files. Any generated files that are specific to a Scala version should go in `crossTarget` for efficient cross-building.

For generating sources and resources, see [Generating Files](#).

Don't hard code

Don't hard code constants, like the output directory `target/`. This is especially important for plugins. A user might change the `target` setting to point to `build/`, for example, and the plugin needs to respect that. Instead, use the setting, like:

```
myDirectory := target.value / "sub-directory"
```

Don't “mutate” files

A build naturally consists of a lot of file manipulation. How can we reconcile this with the task system, which otherwise helps us avoid mutable state? One approach, which is the recommended approach and the approach used by sbt's default tasks, is to only write to any given file once and only from a single task.

A build product (or by-product) should be written exactly once by only one task. The task should then, at a minimum, provide the File created as its result. Another task that wants to use Files should map the task, simultaneously obtaining the File reference and ensuring that the task has run (and thus the file is constructed). Obviously you cannot do much about the user or other processes modifying the files, but you can make the I/O that is under the build's control more predictable by treating file contents as immutable at the level of Tasks.

For example:

```
lazy val makeFile = taskKey[File]("Creates a file with some content.")

// define a task that creates a file,
// writes some content, and returns the File
makeFile := {
  val f: File = file("/tmp/data.txt")
  IO.write(f, "Some content")
  f
}

// The result of makeFile is the constructed File,
// so useFile can map makeFile and simultaneously
// get the File and declare the dependency on makeFile
useFile :=
  doSomething( makeFile.value )
```

This arrangement is not always possible, but it should be the rule and not the exception.

Use absolute paths

Construct only absolute Files. Either specify an absolute path

```
file("/home/user/A.scala")
```

or construct the file from an absolute base:

```
base / "A.scala"
```

This is related to the no hard coding best practice because the proper way involves referencing the `baseDirectory` setting. For example, the following defines the `myPath` setting to be the `<base>/licenses/` directory.

```
myPath := baseDirectory.value / "licenses"
```

In Java (and thus in Scala), a relative `File` is relative to the current working directory. The working directory is not always the same as the build root directory for a number of reasons.

The only exception to this rule is when specifying the base directory for a `Project`. Here, sbt will resolve a relative `File` against the build root directory for you for convenience.

Parser combinators

1. Use `token` everywhere to clearly delimit tab completion boundaries.
2. Don't overlap or nest tokens. The behavior here is unspecified and will likely generate an error in the future.
3. Use `flatMap` for general recursion. sbt's combinators are strict to limit the number of classes generated, so use `flatMap` like:

```
lazy val parser: Parser[Int] =  
  token(IntBasic) flatMap { i =>  
    if(i <= 0)  
      success(i)  
    else  
      token(Space ~> parser)  
  }
```

This example defines a parser a whitespace-delimited list of integers, ending with a negative number, and returning that final, negative number.

Plugins

There's a [getting started page](#) focused on using existing plugins, which you may want to read first.

A plugin is a way to use external code in a build definition. A plugin can be a library used to implement a task (you might use [Knockoff](#) to write a markdown processing task). A plugin can define a sequence of sbt settings that are automatically added to all projects or that are explicitly declared for selected projects. For example, a plugin might add a `proguard` task and associated

(overridable) settings. Finally, a plugin can define new commands (via the `commands` setting).

sbt 0.13.5 introduces auto plugin with improved dependency management among the plugins and explicitly scoped auto importing. Going forward, our recommendation is to migrate to the auto plugins. The [Plugins Best Practices](#) page describes the currently evolving guidelines to writing sbt plugins. See also the general [best practices](#).

Using an auto plugin

A common situation is when using a binary plugin published to a repository. If you're adding sbt-assembly, create `project/assembly.sbt` with the following:

```
addSbtPlugin("com.eed3si9n" % "sbt-assembly" % "0.11.2")
```

Alternatively, you can create `project/plugins.sbt` with all of the desired sbt plugins, any general dependencies, and any necessary repositories:

```
addSbtPlugin("org.example" % "plugin" % "1.0")

addSbtPlugin("org.example" % "another-plugin" % "2.0")

// plain library (not an sbt plugin) for use in the build definition
libraryDependencies += "org.example" % "utilities" % "1.3"

resolvers += "Example Plugin Repository" at "http://example.org/repo/"
```

Many of the auto plugins automatically adds settings into projects, however, some may require explicit enablement. Here's an example:

```
lazy val util = (project in file("util")).
  enablePlugins(FooPlugin, BarPlugin).
  disablePlugins(plugins.IvyPlugin).
  settings(
    name := "hello-util"
  )
```

See [using plugins](#) in the Getting Started guide for more details on using plugins.

By Description

A plugin definition is a project under `project/` folder. This project's classpath is the classpath used for build definitions in `project/` and any `.sbt` files in the project's base directory. It is also used for the `eval` and `set` commands.

Specifically,

1. Managed dependencies declared by the `project/` project are retrieved and are available on the build definition classpath, just like for a normal project.
2. Unmanaged dependencies in `project/lib/` are available to the build definition, just like for a normal project.
3. Sources in the `project/` project are the build definition files and are compiled using the classpath built from the managed and unmanaged dependencies.
4. Project dependencies can be declared in `project/plugins.sbt` (similarly to `build.sbt` file in a normal project) or `project/project/Build.scala` (similarly to `project/Build.scala` in a normal project) and will be available to the build definition sources. Think of `project/project/` as the build definition for the build definition (worth to repeat it here again: “sbt is recursive”, remember?).

The build definition classpath is searched for `sbt/sbt.plugins` descriptor files containing the names of `sbt.AutoPlugin` or `sbt.Plugin` implementations.

The `reload plugins` command changes the current build to the (root) project's `project/` build definition. This allows manipulating the build definition project like a normal project. `reload return` changes back to the original build. Any session settings for the plugin definition project that have not been saved are dropped.

An auto plugin is a module that defines settings to automatically inject to projects. In addition an auto plugin provides the following feature:

- Automatically import selective names to `.sbt` files and the `eval` and `set` commands.
- Specify plugin dependencies to other auto plugins.
- Automatically activate itself when all dependencies are present.
- Specify `projectSettings`, `buildSettings`, and `globalSettings` as appropriate.

Plugin dependencies

When a traditional plugin wanted to reuse some functionality from an existing plugin, it would pull in the plugin as a library dependency, and then it would either:

1. add the setting sequence from the dependency as part of its own setting sequence, or
2. tell the build users to include them in the right order.

This becomes complicated as the number of plugins increase within an application, and becomes more error prone. The main goal of auto plugin is to alleviate this setting dependency problem. An auto plugin can depend on other auto plugins and ensure these dependency settings are loaded first.

Suppose we have the `SbtLessPlugin` and the `SbtCoffeeScriptPlugin`, which in turn depends on the `SbtJsTaskPlugin`, `SbtWebPlugin`, and `JvmPlugin`. Instead of manually activating all of these plugins, a project can just activate the `SbtLessPlugin` and `SbtCoffeeScriptPlugin` like this:

```
lazy val root = (project in file(".")).
  enablePlugins(SbtLessPlugin, SbtCoffeeScriptPlugin)
```

This will pull in the right setting sequence from the plugins in the right order. The key notion here is you declare the plugins you want, and sbt can fill in the gap.

A plugin implementation is not required to produce an auto plugin, however. It is a convenience for plugin consumers and because of the automatic nature, it is not always appropriate.

Global plugins The `~/.sbt/0.13/plugins/` directory is treated as a global plugin definition project. It is a normal sbt project whose classpath is available to all sbt project definitions for that user as described above for per-project plugins.

Creating an auto plugin

A minimal sbt plugin is a Scala library that is built against the version of Scala that sbt runs (currently, 2.10.3) or a Java library. Nothing special needs to be done for this type of library. A more typical plugin will provide sbt tasks, commands, or settings. This kind of plugin may provide these settings automatically or make them available for the user to explicitly integrate.

To make an auto plugin, create a project and configure `sbtPlugin` to `true`.

```
sbtPlugin := true
```

Then, write the plugin code and publish your project to a repository. The plugin can be used as described in the previous section.

First, in an appropriate namespace, define your auto plugin object by extending `sbt.AutoPlugin`.

projectSettings and buildSettings With auto plugins, all provided settings (e.g. `assemblySettings`) are provided by the plugin directly via the `projectSettings` method. Here's an example plugin that adds a command named `hello` to sbt projects:

```
package sbthello

import sbt._
import Keys._

object HelloPlugin extends AutoPlugin {
  override lazy val projectSettings = Seq(commands += helloCommand)
  lazy val helloCommand =
    Command.command("hello") { (state: State) =>
      println("Hi!")
      state
    }
}
```

This example demonstrates how to take a `Command` (here, `helloCommand`) and distribute it in a plugin. Note that multiple commands can be included in one plugin (for example, use `commands ++= Seq(a,b)`). See [Commands](#) for defining more useful commands, including ones that accept arguments and affect the execution state.

If the plugin needs to append settings at the build-level (that is, in `ThisBuild`) there's a `buildSettings` method. The settings returned here are guaranteed to be added to a given build scope only once regardless of how many projects for that build activate this `AutoPlugin`.

```
override def buildSettings: Seq[Setting[_]] = Nil
```

The `globalSettings` is appended once to the global settings (in `Global`). These allow a plugin to automatically provide new functionality or new defaults. One main use of this feature is to globally add commands, such as for IDE plugins.

```
override def globalSettings: Seq[Setting[_]] = Nil
```

Use `globalSettings` to define the default value of a setting.

Implementing plugin dependencies Next step is to define the plugin dependencies.

```

package sbtleast

import sbt._
import Keys._
object SbtLessPlugin extends AutoPlugin {
  override def requires = SbtJsTaskPlugin
  override lazy val projectSettings = ...
}

```

The `requires` method returns a value of type `Plugins`, which is a DSL for constructing the dependency list. The `requires` method typically contains one of the following values:

- `empty` (No plugins, this is the default)
- other auto plugins
- `&&` operator (for defining multiple dependencies)

Root plugins and triggered plugins Some plugins should always be explicitly enabled on projects. we call these root plugins, i.e. plugins that are “root” nodes in the plugin dependency graph. An auto plugin is by default a root plugin.

Auto plugins also provide a way for plugins to automatically attach themselves to projects if their dependencies are met. We call these triggered plugins, and they are created by overriding the `trigger` method.

For example, we might want to create a triggered plugin that can append commands automatically to the build. To do this, set the `requires` method to return `empty` (this is the default), and override the `trigger` method with `allRequirements`.

```

package sbthello

import sbt._
import Keys._

object HelloPlugin2 extends AutoPlugin {
  override def trigger = allRequirements
  override lazy val buildSettings = Seq(commands += helloCommand)
  lazy val helloCommand =
    Command.command("hello") { (state: State) =>
      println("Hi!")
      state
    }
}

```

The build user still needs to include this plugin in `project/plugins.sbt`, but it is no longer needed to be included in `build.sbt`. This becomes more interesting when you do specify a plugin with requirements. Let's modify the `SbtLessPlugin` so that it depends on another plugin:

```
package sbtless
import sbt._
import Keys._
object SbtLessPlugin extends AutoPlugin {
  override def trigger = allRequirements
  override def requires = SbtJsTaskPlugin
  override lazy val projectSettings = ...
}
```

As it turns out, `PlayScala` plugin (in case you didn't know, the `Play` framework is an `sbt` plugin) lists `SbtJsTaskPlugin` as one of its required plugins. So, if we define a `build.sbt` with:

```
lazy val root = (project in file(".")).
  enablePlugins(PlayScala)
```

then the setting sequence from `SbtLessPlugin` will be automatically appended somewhere after the settings from `PlayScala`.

This allows plugins to silently, and correctly, extend existing plugins with more features. It also can help remove the burden of ordering from the user, allowing the plugin authors greater freedom and power when providing features for their users.

Controlling the import with `autoImport` When an auto plugin provides a stable field such as `val` or `object` named `autoImport`, the contents of the field are wildcard imported in `set`, `eval`, and `.sbt` files.

```
package sbthello

import sbt._
import Keys._

object HelloPlugin3 extends AutoPlugin {
  object autoImport {
    val greeting = settingKey[String]("greeting")
  }
  import autoImport._
  override def trigger = allRequirements
```

```

override lazy val buildSettings = Seq(
  greeting := "Hi!",
  commands += helloCommand)
lazy val helloCommand =
  Command.command("hello") { (state: State) =>
    println(greeting.value)
    state
  }
}

```

Typically, `autoImport` is used to provide new keys - `SettingKeys`, `TaskKeys`, or `InputKeys` - or core methods without requiring an import or qualification.

Example Plugin An example of a typical plugin:

build.sbt:

```
sbtPlugin := true
```

```
name := "sbt-obfuscate"
```

```
organization := "org.example"
```

ObfuscatePlugin.scala:

```
package sbtobfuscate
```

```
import sbt._
```

```

object ObfuscatePlugin extends AutoPlugin {
  // by defining autoImport, the settings are automatically imported into user's `*.sbt`
  object autoImport {
    // configuration points, like the built-in `version`, `libraryDependencies`, or `compilerOptions`
    val obfuscate = taskKey[Seq[File]]("Obfuscates files.")
    val obfuscateLiterals = settingKey[Boolean]("Obfuscate literals.")

    // default values for the tasks and settings
    lazy val baseObfuscateSettings: Seq[Def.Setting[_]] = Seq(
      obfuscate := {
        Obfuscate(sources.value, (obfuscateLiterals in obfuscate).value)
      },
      obfuscateLiterals in obfuscate := false
    )
  }
}

```

```

import autoImport._
override def requires = sbt.plugins.JvmPlugin

// This plugin is automatically enabled for projects which are JvmPlugin.
override def trigger = allRequirements

// a group of settings that are automatically added to projects.
override val projectSettings =
  inConfig(Compile)(baseObfuscateSettings) ++
  inConfig(Test)(baseObfuscateSettings)
}

object Obfuscate {
  def apply(sources: Seq[File]): Seq[File] := sources
}

```

Usage example A build definition that uses the plugin might look like.
obfuscate.sbt:

```
obfuscateLiterals in obfuscate := true
```

Global plugins example The simplest global plugin definition is declaring a library or plugin in ~/.sbt/0.13/plugins/build.sbt:

```
libraryDependencies += "org.example" %% "example-plugin" % "0.1"
```

This plugin will be available for every sbt project for the current user.

In addition:

- Jars may be placed directly in ~/.sbt/0.13/plugins/lib/ and will be available to every build definition for the current user.
- Dependencies on plugins built from source may be declared in ~/.sbt/0.13/plugins/project/Build.scala as described at [.scala build definition](#).
- A Plugin may be directly defined in Scala source files in ~/.sbt/0.13/plugins/, such as ~/.sbt/0.13/plugins/MyPlugin.scala. ~/.sbt/0.13/plugins//build.sbt should contain `sbtPlugin := true`. This can be used for quicker turnaround when developing a plugin initially:

1. Edit the global plugin code
2. reload the project you want to use the modified plugin in

3. sbt will rebuild the plugin and use it for the project. Additionally, the plugin will be available in other projects on the machine without recompiling again. This approach skips the overhead of `publishLocal` and `cleaning` the plugins directory of the project using the plugin.

These are all consequences of `~/sbt/0.13/plugins/` being a standard project whose classpath is added to every sbt project's build definition.

Using a library in a build definition example

As an example, we'll add the Grizzled Scala library as a plugin. Although this does not provide sbt-specific functionality, it demonstrates how to declare plugins.

1a) Manually managed

1. Download the jar manually from https://oss.sonatype.org/content/repositories/releases/org/clapper/grizzled-scala_2.8.1/1.0.4/grizzled-scala_2.8.1-1.0.4.jar
2. Put it in `project/lib/`

1b) Automatically managed: direct editing approach Edit `project/plugins.sbt` to contain:

```
libraryDependencies += "org.clapper" %% "grizzled-scala" % "1.0.4"
```

If sbt is running, do `reload`.

1c) Automatically managed: command-line approach We can change to the plugins project in `project/` using `reload plugins`.

```
$ sbt
> reload plugins
[info] Set current project to default (in build file:/Users/sbt/demo2/project/)
>
```

Then, we can add dependencies like usual and save them to `project/plugins.sbt`. It is useful, but not required, to run `update` to verify that the dependencies are correct.


```
> set libraryDependencies += "org.clapper" %% "grizzled-scala" % "1.0.4"
...
> update
...
> session save
...
```

To switch back to the main project use `reload return`:

```
> reload return
[info] Set current project to root (in build file:/Users/sbt/demo2/)
```

1d) Project dependency This variant shows how to use sbt's external project support to declare a source dependency on a plugin. This means that the plugin will be built from source and used on the classpath.

Edit `project/plugins.sbt`

```
lazy val root = (project in file(".")).dependsOn(assemblyPlugin)

lazy val assemblyPlugin = uri("git://github.com/sbt/sbt-assembly")
```

If sbt is running, run `reload`.

Note that this approach can be useful used when developing a plugin. A project that uses the plugin will rebuild the plugin on `reload`. This saves the intermediate steps of `publishLocal` and `update`. It can also be used to work with the development version of a plugin from its repository.

It is however recommended to explicitly specify the commit or tag by appending it to the repository as a fragment:

```
lazy val assemblyPlugin = uri("git://github.com/sbt/sbt-assembly#0.9.1")
```

One caveat to using this method is that the local sbt will try to run the remote plugin's build. It is quite possible that the plugin's own build uses a different sbt version, as many plugins cross-publish for several sbt versions. As such, it is recommended to stick with binary artifacts when possible.

2) Use the library Grizzled Scala is ready to be used in build definitions. This includes the `eval` and `set` commands and `.sbt` and `project/*.scala` files.

```
> eval grizzled.sys.os
```

In a `build.sbt` file:

```
import grizzled.sys._
import OperatingSystem._

libraryDependencies ++=
  if(os == Windows)
    Seq("org.example" % "windows-only" % "1.0")
  else
    Seq.empty
```

Best Practices

If you're a plugin writer, please consult the [Plugins Best Practices](#) page; it contains a set of guidelines to help you ensure that your plugin is consistent and plays well with other plugins.

Plugins Best Practices

This page is intended primarily for sbt plugin authors. This page assumes you've read [using plugins](#) and [Plugins](#).

A plugin developer should strive for consistency and ease of use. Specifically:

- Plugins should play well with other plugins. Avoiding namespace clashes (in both sbt and Scala) is paramount.
- Plugins should follow consistent conventions. The experiences of an sbt *user* should be consistent, no matter what plugins are pulled in.

Here are some current plugin best practices.

Note: Best practices are evolving, so check back frequently.

Get your plugins known

Make sure people can find your plugin. Here are some of the recommended steps:

1. Mention `[@scala_sbt]`(https://twitter.com/scala_sbt) in your announcement, and we will RT it.
2. Announce it on implicit.ly using [n8han/herald](#).
3. Send a pull req to [sbt/website](#) and add your plugin on [the plugins list](#).

Don't use default package

Users who have their build files in some package will not be able to use your plugin if it's defined in default (no-name) package.

Use settings and tasks. Avoid commands.

Your plugin should fit in naturally with the rest of the sbt ecosystem. The first thing you can do is to avoid defining [commands](#), and use settings and [tasks](#) and task-scoping instead (see below for more on task-scoping). Most of the interesting things in sbt like `compile`, `test` and `publish` are provided using tasks. Tasks can take advantage of duplication reduction and parallel execution by the task engine. With features like [ScopeFilter](#), many of the features that previously required commands are now possible using tasks.

Settings can be composed from other settings and tasks. Tasks can be composed from other tasks and input tasks. Commands, on the other hand, cannot be composed from any of the above. In general, use the minimal thing that you need. One legitimate use of commands may be using plugin to access the build definition itself not the code. `sbt-inspectr` was implemented using a [command](#) before it became `inspect tree`.

Use `sbt.AutoPlugin`

sbt is in the process of migrating to `sbt.AutoPlugin` from `sbt.Plugin`. The new mechanism features a set of user-level controls and dependency declarations that cleans up a lot of long-standing issues with plugins.

Reuse existing keys

sbt has a number of [predefined keys](#). Where possible, reuse them in your plugin. For instance, don't define:

```
val sourceFiles = settingKey[Seq[File]]("Some source files")
```

Instead, simply reuse sbt's existing `sources` key.

Avoid namespace clashes

Sometimes, you need a new key, because there is no existing sbt key. In this case, use a plugin-specific prefix.

```

package sbtobfuscate

import sbt._, Keys._

object ObfuscatePlugin extends sbt.AutoPlugin {
  object autoImport {
    lazy val obfuscateStylesheet = settingKey[File]("obfuscate stylesheet")
  }
}

```

In this approach, every `lazy val` starts with `obfuscate`. A user of the plugin would refer to the settings like this:

```
obfuscateStylesheet := file("something.txt")
```

Provide core feature in a plain old Scala object

The core feature of sbt's `package` task, for example, is implemented in `sbt.Package`, which can be called via its `apply` method. This allows greater reuse of the feature from other plugins such as `sbt-assembly`, which in return implements `sbtassembly.Assembly` object to implement its core feature.

Follow their lead, and provide core feature in a plain old Scala object.

Configuration advices

If your plugin introduces either a new set of source code or its own library dependencies, only then you want your own configuration.

You probably won't need your own configuration Configurations should *not* be used to namespace keys for a plugin. If you're merely adding tasks and settings, don't define your own configuration. Instead, reuse an existing one *or* scope by the main task (see below).

```

package sbtwhatever

import sbt._, Keys._

object WhateverPlugin extends sbt.AutoPlugin {
  override def requires = plugins.JvmPlugin
  override def trigger = allRequirements

  object autoImport {

```

```

// BAD sample
lazy val Whatever = config("whatever") extend(Compile)
lazy val dude = settingKey[String]("A plugin specific key")
}
import autoImport._
override lazy val projectSettings = Seq(
  dude in Whatever := "your opinion man" // DON'T DO THIS
)
}

```

When to define your own configuration If your plugin introduces either a new set of source code or its own library dependencies, only then you want your own configuration. For instance, suppose you've built a plugin that performs fuzz testing that requires its own fuzzing library and fuzzing source code. `scalaSource` key can be reused similar to `Compile` and `Test` configuration, but `scalaSource` scoped to `Fuzz` configuration (denoted as `scalaSource in Fuzz`) can point to `src/fuzz/scala` so it is distinct from other Scala source directories. Thus, these three definitions use the same *key*, but they represent distinct *values*. So, in a user's `build.sbt`, we might see:

```
scalaSource in Fuzz := baseDirectory.value / "source" / "fuzz" / "scala"
```

```
scalaSource in Compile := baseDirectory.value / "source" / "main" / "scala"
```

In the fuzzing plugin, this is achieved with an `inConfig` definition:

```

package sbtfuzz

import sbt._, Keys._

object FuzzPlugin extends sbt.AutoPlugin {
  override def requires = plugins.JvmPlugin
  override def trigger = allRequirements

  object autoImport {
    lazy val Fuzz = config("fuzz") extend(Compile)
  }
  import autoImport._

  lazy val baseFuzzSettings: Seq[Def.Setting[_]] = Seq(
    test := {
      println("fuzz test")
    }
  )
}

```

```

    override lazy val projectSettings = inConfig(Fuzz)(baseFuzzSettings)
}

```

When defining a new type of configuration, e.g.

```

lazy val Fuzz = config("fuzz") extend(Compile)

```

should be used to create a configuration. Configurations actually tie into dependency resolution (with Ivy) and can alter generated pom files.

Playing nice with configurations Whether you ship with a configuration or not, a plugin should strive to support multiple configurations, including those created by the build user. Some tasks that are tied to a particular configuration can be re-used in other configurations. While you may not see the need immediately in your plugin, some project may and will ask you for the flexibility.

Provide raw settings and configured settings Split your settings by the configuration axis like so:

```

package sbtobfuscate

import sbt._, Keys._

object ObfuscatePlugin extends sbt.AutoPlugin {
  override def requires = plugins.JvmPlugin
  override def trigger = allRequirements

  object autoImport {
    lazy val obfuscate = taskKey[Seq[File]]("obfuscate the source")
    lazy val obfuscateStylesheet = settingKey[File]("obfuscate stylesheet")
  }
  import autoImport._
  lazy val baseObfuscateSettings: Seq[Def.Setting[_]] = Seq(
    obfuscate := Obfuscate((sources in obfuscate).value),
    sources in obfuscate := sources.value
  )
  override lazy val projectSettings = inConfig(Compile)(baseObfuscateSettings)
}

// core feature implemented here
object Obfuscate {
  def apply(sources: Seq[File]): Seq[File] = {
    sources
  }
}

```

The `baseObfuscateSettings` value provides base configuration for the plugin's tasks. This can be re-used in other configurations if projects require it. The `obfuscateSettings` value provides the default `Compile` scoped settings for projects to use directly. This gives the greatest flexibility in using features provided by a plugin. Here's how the raw settings may be reused:

```
import sbtobfuscate.ObfuscatePlugin

lazy val app = (project in file("app")).
  settings(inConfig(Test)(ObfuscatePlugin.baseObfuscateSettings): _*)
```

Using a “main” task scope for settings Sometimes you want to define some settings for a particular “main” task in your plugin. In this instance, you can scope your settings using the task itself. See the `baseObfuscateSettings`:

```
lazy val baseObfuscateSettings: Seq[Def.Setting[_]] = Seq(
  obfuscate := Obfuscate((sources in obfuscate).value),
  sources in obfuscate := sources.value
)
```

In the above example, `sources in obfuscate` is scoped under the main task, `obfuscate`.

Mucking with `globalSettings`

There may be times when you need to muck with `globalSettings`. The general rule is *be careful what you touch*.

When overriding global settings, care should be taken to ensure previous settings from other plugins are not ignored. e.g. when creating a new `onLoad` handler, ensure that the previous `onLoad` handler is not removed.

```
package sbtsomething

import sbt._, Keys._

object MyPlugin extends AutoPlugin {
  override def requires = plugins.JvmPlugin
  override def trigger = allRequirements

  override val globalSettings: Seq[Def.Setting[_]] = Seq(
    onLoad in Global := (onLoad in Global).value andThen { state =>
      ... return new state ...
    }
  )
}
```

Sbt Launcher

The sbt launcher provides a generic container that can load and run programs resolved using the Ivy dependency manager. Sbt uses this as its own deployment mechanism.

Getting Started with the Sbt Launcher

The sbt launcher component is a self-contained jar that boots a Scala application or server without Scala or the application already existing on the system. The only prerequisites are the launcher jar itself, an optional configuration file, and a java runtime version 1.6 or greater.

Overview

A user downloads the launcher jar and creates a script to run it. In this documentation, the script will be assumed to be called `launch`. For unix, the script would look like: `java -jar sbt-launcher.jar "$@"`

The user can now launch servers and applications which provide sbt launcher configuration.

Applications To launch an application, the user then downloads the configuration file for the application (call it `my.app.configuration`) and creates a script to launch it (call it `myapp`): `launch @my.app.configuration "$@"`

The user can then launch the application using `myapp arg1 arg2 ...`

More on launcher configuration can be found at [Launcher Configuration](#)

Servers The sbt launcher can be used to launch and discover running servers on the system. The launcher can be used to launch servers similarly to applications. However, if desired, the launcher can also be used to ensure that only one instance of a server is running at time. This is done by having clients always use the launcher as a *service locator*.

To discover where a server is running (or launch it if it is not running), the user downloads the configuration file for the server (call it `my.server.configuration`) and creates a script to discover the server (call it `find-myserver`): `launch --locate @my.server.properties`.

This command will print out one string, the URI at which to reach the server, e.g. `sbt://127.0.0.1:65501`. Clients should use the IP/port to connect to to the server and initiate their connection.

When using the `locate` feature, the sbt launcher makes these following restrictions to servers:

- The Server must have a starting class that extends the `xsbti.ServerMain` class
- The Server must have an entry point (URI) that clients can use to detect the server
- The server must have defined a lock file which the launcher can use to ensure that only one instance is running at a time
- The filesystem on which the lock file resides must support locking.
- The server must allow the launcher to open a socket against the port without sending any data. This is used to check if a previous server is still alive.

Resolving Applications/Servers Like the launcher used to distribute `sbt`, the downloaded launcher jar will retrieve Scala and the application according to the provided configuration file. The versions may be fixed or read from a different configuration file (the location of which is also configurable). The location to which the Scala and application jars are downloaded is configurable as well. The repositories searched are configurable. Optional initialization of a properties file on launch is configurable.

Once the launcher has downloaded the necessary jars, it loads the application/server and calls its entry point. The application is passed information about how it was called: command line arguments, current working directory, Scala version, and application ID (organization, name, version). In addition, the application can ask the launcher to perform operations such as obtaining the Scala jars and a `ClassLoader` for any version of Scala retrievable from the repositories specified in the configuration file. It can request that other applications be downloaded and run. When the application completes, it can tell the launcher to exit with a specific exit code or to reload the application with a different version of Scala, a different version of the application, or different arguments.

There are some other options for setup, such as putting the configuration file inside the launcher jar and distributing that as a single download. The rest of this documentation describes the details of configuring, writing, distributing, and running the application.

Creating a Launched Application This section shows how to make an application that is launched by this launcher. First, declare a dependency on the launcher-interface. Do not declare a dependency on the launcher itself. The launcher interface consists strictly of Java interfaces in order to avoid binary incompatibility between the version of Scala used to compile the launcher and the version used to compile your application. The launcher interface class will be provided by the launcher, so it is only a compile-time dependency. If you are building with `sbt`, your dependency definition would be:

Make the entry point to your class implement 'xsbti.AppMain'. An example that uses some of the information:

```
package xsbt.test
class Main extends xsbti.AppMain
{
  def run(configuration: xsbti.AppConfiguration) =
  {
    // get the version of Scala used to launch the application
    val scalaVersion = configuration.provider.scalaProvider.version

    // Print a message and the arguments to the application
    println("Hello world! Running Scala " + scalaVersion)
    configuration.arguments.foreach(println)

    // demonstrate the ability to reboot the application into different versions of Scala
    // and how to return the code to exit with
    scalaVersion match
    {
      case "2.9.3" =>
        new xsbti.Reboot {
          def arguments = configuration.arguments
          def baseDirectory = configuration.baseDirectory
          def scalaVersion = "2.10.2"
          def app = configuration.provider.id
        }
      case "2.10.2" => new Exit(1)
      case _ => new Exit(0)
    }
  }
  class Exit(val code: Int) extends xsbti.Exit
}
```

Next, define a configuration file for the launcher. For the above class, it might look like:

Then, `publishLocal` or `+publishLocal` the application to make it available. For more information, see [Launcher Configuration](#).

Running an Application As mentioned above, there are a few options to actually run the application. The first involves providing a modified jar for download. The second two require providing a configuration file for download.

- Replace the `/sbt/sbt.boot.properties` file in the launcher jar and distribute the modified jar. The user would need a script to run `java -jar your-launcher.jar arg1 arg2`

- The user downloads the launcher jar and you provide the configuration file.
 - The user needs to run `java -Dsbt.boot.properties=your.boot.properties -jar launcher.jar`.
 - The user already has a script to run the launcher (call it ‘launch’). The user needs to run `launch @your.boot.properties your-arg-1 your-arg-2`

Execution Let’s review what’s happening when the launcher starts your application.

On startup, the launcher searches for its configuration and then parses it. Once the final configuration is resolved, the launcher proceeds to obtain the necessary jars to launch the application. The `boot.directory` property is used as a base directory to retrieve jars to. Locking is done on the directory, so it can be shared system-wide. The launcher retrieves the requested version of Scala to

```
${boot.directory}/${scala.version}/lib/
```

If this directory already exists, the launcher takes a shortcut for startup performance and assumes that the jars have already been downloaded. If the directory does not exist, the launcher uses Apache Ivy to resolve and retrieve the jars. A similar process occurs for the application itself. It and its dependencies are retrieved to

```
${boot.directory}/${scala.version}/${app.org}/${app.name}/.
```

Once all required code is downloaded, the class loaders are set up. The launcher creates a class loader for the requested version of Scala. It then creates a child class loader containing the jars for the requested ‘app.components’ and with the paths specified in `app.resources`. An application that does not use components will have all of its jars in this class loader.

The main class for the application is then instantiated. It must be a public class with a public no-argument constructor and must conform to `xsbti.AppMain`. The `run` method is invoked and execution passes to the application. The argument to the ‘run’ method provides configuration information and a callback to obtain a class loader for any version of Scala that can be obtained from a repository in `[repositories]`. The return value of the `run` method determines what is done after the application executes. It can specify that the launcher should restart the application or that it should exit with the provided exit code.

Sbt Launcher Architecture

The sbt launcher is a mechanism whereby modules can be loaded from ivy and executed within a jvm. It abstracts the mechanism of grabbing and caching jars, allowing users to focus on what application they want and control its versions.

The launcher's primary goal is to take configuration for applications, mostly just ivy coordinates and a main class, and start the application. The launcher resolves the ivy module, caches the required runtime jars and starts the application.

The sbt launcher provides the application with the means to load a different application when it completes, exit normally, or load additional applications from inside another.

The sbt launcher provides these core functions:

- Module Resolution
- Classloader Caching and Isolation
- File Locking
- Service Discovery and Isolation

Module Resolution

The primary purpose of the sbt launcher is to resolve applications and run them. This is done through the `[app]` configuration section. See [launcher configuration](#) for more information on how to configure module resolution.

Module resolution is performed using the Ivy dependency management library. This library supports loading artifacts from Maven repositories as well.

Classloader Caching and Isolation

The sbt launcher's classloading structure is different than just starting an application in the standard java mechanism. Every application loaded by the launcher is given its own classloader. This classloader is a child of the Scala classloader used by the application. The Scala classloader can see all of the `xsbti.*` classes from the launcher itself.

Here's an example classloader layout from an sbt launched application.

In this diagram, three different applications were loaded. Two of these use the same version of Scala (2.9.2). In this case, sbt can share the same classloader for these applications. This has the benefit that any JIT optimisations performed on scala classes can be re-used between applications thanks to the shared classloader.



Figure 3: image

Caching

The sbt launcher creates a secondary cache on top of Ivy's own cache. This helps isolate applications from errors resulting from unstable revisions, like `-SNAPSHOT`. For any launched application, the launcher creates a directory to store all its jars. Here's an example layout.

Locking

In addition to providing a secondary cache, the launcher also provides a mechanism of safely doing file-based locks. This is used in two places directly by the launcher:

1. Locking the boot directory.
2. Ensuring located servers have at most one active process.

This feature requires a filesystem which supports locking. It is exposed via the `xsbti.GlobalLock` interface.

Note: This is both a thread and file lock. Not only are we limiting access to a single process, but also a single thread within that process.

Service Discovery and Isolation

The launcher also provides a mechanism to ensure that only one instance of a server is running, while dynamically starting it when a client requests. This is done through the `--locate` flag on the launcher. When the launcher is started with the `--locate` flag it will do the following:

1. Lock on the configured server lock file.
2. Read the server properties to find the URI of the previous server.
3. If the port is still listening to connection requests, print this URI on the command line.
4. If the port is not listening, start a new server and write the URI on the command line.
5. Release all locks and shutdown.

The configured `server.lock` file is thus used to prevent multiple servers from running. Sbt itself uses this to prevent more than one server running on any given project directory by configuring `server.lock` to be `${user.dir}/.sbtserver`.

Sbt Launcher Configuration

The launcher may be configured in one of the following ways in increasing order of precedence:

- Replace the `/sbt/sbt.boot.properties` file in the launcher jar
- Put a configuration file named `sbt.boot.properties` on the classpath. Put it in the classpath root without the `/sbt` prefix.
- Specify the location of an alternate configuration on the command line, either as a path or an absolute URI. This can be done by either specifying the location as the system property `sbt.boot.properties` or as the first argument to the launcher prefixed by `@`. The system property has lower precedence. Resolution of a relative path is first attempted against the current working directory, then against the user's home directory, and then against the directory containing the launcher jar.

An error is generated if none of these attempts succeed.

Example

The default configuration file for sbt as an application looks like:

Let's look at all the launcher configuration sections in detail:

1. Scala Configuration The `[scala]` section is used to configure the version of Scala. It has one property:

- **version** - The version of scala an application uses, or `auto` if the application is not cross-versioned.
- **classifiers** - The (optional) list of additional scala artifacts to resolve, e.g. sources.

2. Application Identification The `[app]` section configures how the launcher will look for your application using the Ivy dependency manager. It consists of the following properties:

- **org** - The organization associated with the Ivy module. (`groupId` in maven vernacular)
- **name** - The name of the Ivy module. (`artifactId` in maven vernacular)
- **version** - The revision of the Ivy module.
- **class** - The name of the “entry point” into the application. An entry point must be a class which meets one of the following criteria

- Extends the xsbti.AppMain interface.
- Extends the xsbti.ServerMain interfaces.
- Contains a method with the signature `static void main(String[])`
- Contains a method with the signature `static int main(String[])`
- Contains a method with the signature `static xsbti.Exit main(String[])`
- **components** - An optional list of additional components that Ivy should resolve.
- **cross-versioned** - An optional string denoting how this application is published. If `app.cross-versioned` is `binary`, the resolved module ID is `{app.name+'_'+CrossVersion.binaryScalaVersion(scala.version)}`. If `app.cross-versioned` is `true` or `full`, the resolved module ID is `{app.name+'_'+scala.version}`. The `scala.version` property must be specified and cannot be `auto` when `cross-versioned`.
- **resources** - An optional list of jar files that should be added to the application's classpath.
- **classifiers** - An optional list of additional classifiers that should be resolved with this application, e.g. `sources`.

3. Repositories Section The `[repositories]` section configures where and how Ivy will look for your application. Each line denotes a repository where Ivy will look.

Note: This section configured the default location where Ivy will look, but this can be overridden via user configuration.

There are several built-in strings that can be used for common repositories:

- `local` - the local ivy repository `~/.ivy2/local`.
- `maven-local` - The local maven repository `~/.ivy2/local`.
- `maven-central` - The maven central repository `repo.maven.org`.

Besides built in repositories, other repositories can be configured using the following syntax:

```
name: url(, pattern)(,descriptorOptional)(,skipConsistencyCheck)
```

The **name** property is an identifier which Ivy uses to cache modules resolved from this location. The **name** should be unique across all repositories.

The **url** property is the base **url** where Ivy should look for modules.

The **pattern** property is an optional specification of *how* Ivy should look for modules. By default, the launcher assumes repositories are in the maven style format.

The **skipConsistencyCheck** string is used to tell ivy not to validate checksums and signatures of files it resolves.

4. The Boot section The `[boot]` section is used to configure where the sbt launcher will store its cache and configuration information. It consists of the following properties:

- **directory** - The directory defined here is used to store all cached JARs resolved launcher.
- **properties** - (optional) A properties file to use for any **read** variables.

5. The Ivy section The `[ivy]` section is used to configure the Ivy dependency manager for resolving applications. It consists of the following properties:

- **ivy-home** - The home directory for Ivy. This determines where the ivy-local repository is located, and also where the ivy cache is stored. Defaults to `~/.ivy2`
- **ivy.cache-directory** - provides an alternative location for the Ivy cache used by the launcher. This does not automatically set the Ivy cache for the application, but the application is provided this location through the `AppConfiguration` instance.
- **checksums** - The comma-separated list of checksums that Ivy should use to verify artifacts have correctly resolved, e.g. `md5` or `sha1`.
- **override-build-repos** - If this is set, then the `isOverrideRepositories` method on `xsbti.Launcher` interface will return its value. The use of this method is application specific, but in the case of sbt denotes that the configuration of repositories in the launcher should override those used by any build. Applications should respect this convention if they can.
- **repository-config** - This specifies a configuration location where ivy repositories can also be configured. If this file exists, then its contents override the `[repositories]` section.

6. The Server Section When using the `--locate` feature of the launcher, this section configures how a server is started. It consists of the following properties:

- **lock** - The file that controls access to the running server. This file will contain the active port used by a server and must be located on a filesystem that supports locking.
- **jvmargs** - A file that contains line-separated JVM arguments that where : use when starting the server.
- **jvmprops** - The location of a properties file that will define override properties in the server. All properties defined in this file will be set as `-D java` properties.

Variable Substitution

Property values may include variable substitutions. A variable substitution has one of these forms:

- `${variable.name}`
- `${variable.name-default}`

where `variable.name` is the name of a system property. If a system property by that name exists, the value is substituted. If it does not exist and a default is specified, the default is substituted after recursively substituting variables in it. If the system property does not exist and no default is specified, the original string is not substituted.

There is also a special variable substitution:

- `read(property.name)[default]`

This will look in the file configured by `boot.properties` for a value. If there is no `boot.properties` file configured, or the property does not exist, then the default value is chosen.

Syntax

The configuration file is line-based, read as UTF-8 encoded, and defined by the following grammar. 'nl' is a newline or end of file and 'text' is plain text without newlines or the surrounding delimiters (such as parentheses or square brackets):

```
configuration: scala app repositories boot log appProperties
scala: "[" "scala" "]" nl version nl classifiers nl
app: "[" "app" "]" nl org nl name nl version nl components nl class nl crossVersioned nl res
repositories: "[" "repositories" "]" nl (repository nl)*
boot: "[" "boot" "]" nl directory nl bootProperties nl search nl promptCreate nl promptFill
log: "[" "log" "]" nl logLevel nl
appProperties: "[" "app-properties" "]" nl (property nl)*
ivy: "[" "ivy" "]" nl homeDirectory nl checksums nl overrideRepos nl repoConfig nl
directory: "directory" ":" path
bootProperties: "boot-properties" ":" path
search: "search" ":" ( "none" | "nearest" | "root-first" | "only" ) ( "," path)*
logLevel: "level" ":" ( "debug" | "info" | "warn" | "error" )
promptCreate: "prompt-create" ":" label
promptFill: "prompt-fill" ":" boolean
quickOption: "quick-option" ":" boolean
```

```

version: "version" ":" versionSpecification
versionSpecification: readProperty | fixedVersion
readProperty: "read" "(" propertyName ")" "[" default "]"
fixedVersion: text
classifiers: "classifiers" ":" text ("," text)*
homeDirectory: "ivy-home" ":" path
checksums: "checksums" ":" checksum ("," checksum)*
overrideRepos: "override-build-repos" ":" boolean
repoConfig: "repository-config" ":" path
org: "org" ":" text
name: "name" ":" text
class: "class" ":" text
components: "components" ":" component ("," component)*
crossVersioned: "cross-versioned" ":" ("true" | "false" | "none" | "binary" | "full")
resources: "resources" ":" path ("," path)*
repository: ( predefinedRepository | customRepository ) nl
predefinedRepository: "local" | "maven-local" | "maven-central"
customRepository: label ":" url [ ("," ivyPattern) ["," artifactPattern] ["," mavenCompatible)
property: label ":" propertyDefinition ("," propertyDefinition)*
propertyDefinition: mode "=" (set | prompt)
mode: "quick" | "new" | "fill"
set: "set" "(" value ")"
prompt: "prompt" "(" label ")" ("[" default "]" )?
boolean: "true" | "false"
nl: "\r\n" | "\n" | "\r"
path: text
propertyName: text
label: text
default: text
checksum: text
ivyPattern: text
artifactPattern: text
url: text
component: text

```

Developer's Guide

This is the set of documentation about the Architecture of sbt. This covers all the core components of sbt as well as the general notion of how they all work together. This documentation is suitable for those who wish to have a deeper understanding of sbt's core, but already understand the fundamentals of `Setting[_]`, `Task[_]` and constructing builds.

Core Principles

This document details the core principles overarching sbt's design and code style. Sbt's core principles can be stated quite simply:

1. Everything should have a **Type**, enforced as much as is practical.
2. Dependencies should be **explicit**.
3. Once learned, a concept should hold throughout **all** parts of sbt.
4. Parallel is the default.

With these principles in mind, let's walk through the core design of sbt.

Introduction to build state

This is the first piece you hit when starting sbt. Sbt's command engine is the means by which it processes user requests using the build state. The command engine is essentially a means of applying **state transformations** on the build state, to execute user requests.

In sbt, commands are functions that take the current build state (`sbt.State`) and produce the next state. In other words, they are essentially functions of `sbt.State => sbt.State`. However, in reality, Commands are actually string processors which take some string input and act on it, returning the next build state.

So, the entirety of sbt is driven off the `sbt.State` class. Since this class needs to be resilient in the face of custom code and plugins, it needs a mechanism to store the state from any potential client. In dynamic languages, this can be done directly on objects.

A naive approach in Scala is to use a `Map<String,Any>`. However, this violates tenant #1: Everything should have a **Type**. So, sbt defines a new type of map called an `AttributeMap`. An `AttributeMap` is a key-value storage mechanism where keys are both strings *and* expected **Types** for their value.

Here is what the typesafe `AttributeKey` key looks like :

```
sealed trait AttributeKey[T] {  
  /** The label is the identifier for the key and is camelCase by convention. */  
  def label: String  
  /** The runtime evidence for ``T`` */  
  def manifest: Manifest[T]  
}
```

These keys store both a `label` (string) and some runtime type information (`manifest`). To put or get something on the `AttributeMap`, we first need to construct one of these keys. Let's look at the basic definition of the `AttributeMap` :

```

trait AttributeMap {
  /** Gets the value of type ``T`` associated with the key ``k`` or ``None`` if no value is associ
  * If a key with the same label but a different type is defined, this method will return ``None``
  def get[T](k: AttributeKey[T]): Option[T]

  /** Adds the mapping ``k -> value`` to this map, replacing any existing mapping for ``k``.
  * Any mappings for keys with the same label but different types are unaffected. */
  def put[T](k: AttributeKey[T], value: T): AttributeMap
}

```

Now that there's a definition of what build state is, there needs to be a way to dynamically construct it. In sbt, this is done through the `Setting[_]` sequence.

Settings Architecture

A `Setting` represents the means of constructing the value of one particular `AttributeKey[_]` in the `AttributeMap` of build state. A setting consists of two pieces:

1. The `AttributeKey[T]` where the value of the setting should be assigned.
2. An `Initialize[T]` object which is able to construct the value for this setting.

Sbt's initialization time is basically just taking a sequence of these `Setting[_]` objects and running their initialization objects and then storing the value into the `AttributeMap`. This means overwriting an existing value at a key is as easy as appending a `Setting[_]` to the end of the sequence which does so.

Where it gets interesting is that `Initialize[T]` can depend on other `AttributeKey[_]`s in the build state. Each `Initialize[_]` can pull values from any `AttributeKey[_]` in the build state's `AttributeMap` to compute its value. Sbt ensures a few things when it comes to `Initialize[_]` dependencies:

1. There can be no circular dependencies
2. If one `Initialize[_]` depends on another `Initialize[_]` key, then *all* associated `Initialize[_]` blocks for that key must have run before we load the value.

Let's look at what gets stored for the setting :

```
normalizedName := normalize(name.value)
```



Figure 4: image

Here, a `Setting[_]` is constructed that understands it depends on the value in the `name` `AttributeKey`. Its `initialize` block first grabs the value of the `name` key, then runs the function `normalize` on it to compute its value.

This represents the core mechanism of how to construct sbt's build state. Conceptually, at some point we have a graph of dependencies and initialization functions which we can use to construct the first build state. Once this is completed, we can then start to process user requests.

Task Architecture

The next layer in sbt is around these user request, or tasks. When a user configures a build, they are defining a set of repeatable tasks that they can run on their project. Things like `compile` or `test`. These tasks *also* have a dependency graph, where e.g. the `test` task requires that `compile` has run before it can successfully execute.

Sbt's defines a class `Task[T]`. The `T` type parameter represents the type of data returned by a task. Remember the tenets of sbt? "All things have types" and "Dependencies are explicit" both hold true for tasks. Sbt promotes a style of task dependencies that is closer to functional programming: Return data for your users rather than using shared mutable state.

Most build tools communicate over the filesystem, and indeed sbt, by necessity, does some of this. However, for stable parallelization it is far better to keep tasks isolated on the filesystem and communicate directly through types.

Similarly to how a `Setting[_]` stores both dependencies and an initialization function, a `Task[_]` stores both its `Task[_]` dependencies and its behavior (a function).

TODO - More on `Task[_]`

TODO - Transition into `InputTask[_]`, rehash `Command`

TODO - Transition into `Scope`.

Settings Core

This page describes the core settings engine a bit. This may be useful for using it outside of sbt. It may also be useful for understanding how sbt works internally.

The documentation is comprised of two parts. The first part shows an example settings system built on top of the settings engine. The second part comments on how sbt's settings system is built on top of the settings engine. This may help illuminate what exactly the core settings engine provides and what is needed to build something like the sbt settings system.

Example

Setting up To run this example, first create a new project with the following build.sbt file:

```
libraryDependencies += "org.scala-sbt" %% "collections" % sbtVersion.value

resolvers += sbtResolver.value
```

Then, put the following examples in source files `SettingsExample.scala` and `SettingsUsage.scala`. Finally, run sbt and enter the REPL using `console`. To see the output described below, enter `SettingsUsage`.

Example Settings System The first part of the example defines the custom settings system. There are three main parts:

1. Define the `Scope` type.
2. Define a function that converts that `Scope` (plus an `AttributeKey`) to a `String`.
3. Define a delegation function that defines the sequence of `Scopes` in which to look up a value.

There is also a fourth, but its usage is likely to be specific to sbt at this time. The example uses a trivial implementation for this part.

`SettingsExample.scala`

```
import sbt._

/** Define our settings system */

// A basic scope indexed by an integer.
final case class Scope(index: Int)

// Extend the Init trait.
// (It is done this way because the Scope type parameter is used everywhere in Init.
// Lots of type constructors would become binary, which as you may know requires lots of type
// when you want a type function with only one parameter.
// That would be a general pain.)
object SettingsExample extends Init[Scope]
{
  // Provides a way of showing a Scope+AttributeKey[_]
  val showFullKey: Show[ScopedKey[_]] = new Show[ScopedKey[_]] {
    def apply(key: ScopedKey[_]) = key.scope.index + "/" + key.key.label
  }
}
```



```

}

// A sample delegation function that delegates to a Scope with a lower index.
val delegates: Scope => Seq[Scope] = { case s @ Scope(index) =>
  s +: (if(index <= 0) Nil else delegates(Scope(index-1))) }

}

// Not using this feature in this example.
val scopeLocal: ScopeLocal = _ => Nil

// These three functions + a scope (here, Scope) are sufficient for defining our settings.
}

```

Example Usage This part shows how to use the system we just defined. The end result is a `Settings[Scope]` value. This type is basically a mapping `Scope -> AttributeKey[T] -> Option[T]`. See the [Settings API documentation](#) for details. `SettingsUsage.scala`:

```

/** Usage Example */

import sbt._
import SettingsExample._
import Types._

object SettingsUsage {

  // Define some keys
  val a = AttributeKey[Int]("a")
  val b = AttributeKey[Int]("b")

  // Scope these keys
  val a3 = ScopedKey(Scope(3), a)
  val a4 = ScopedKey(Scope(4), a)
  val a5 = ScopedKey(Scope(5), a)

  val b4 = ScopedKey(Scope(4), b)

  // Define some settings
  val mySettings: Seq[Setting[_]] = Seq(
    setting( a3, value( 3 ) ),
    setting( b4, map(a4)(_ * 3)),
    update(a5)(_ + 1)
  )

  // "compiles" and applies the settings.

```

```

    // This can be split into multiple steps to access intermediate results if desired.
    // The 'inspect' command operates on the output of 'compile', for example.
    val applied: Settings[Scope] = make(mySettings)(delegates, scopeLocal, showFullKey)

    // Show results.
    for(i <- 0 to 5; k <- Seq(a, b)) {
        println( k.label + i + " = " + applied.get( Scope(i), k) )
    }
}

```

This produces the following output when run:

```

a0 = None
b0 = None
a1 = None
b1 = None
a2 = None
b2 = None
a3 = Some(3)
b3 = None
a4 = Some(3)
b4 = Some(9)
a5 = Some(4)
b5 = Some(9)

```

- For the None results, we never defined the value and there was no value to delegate to.
- For a3, we explicitly defined it to be 3.
- a4 wasn't defined, so it delegates to a3 according to our delegates function.
- b4 gets the value for a4 (which delegates to a3, so it is 3) and multiplies by 3
- a5 is defined as the previous value of a5 + 1 and since no previous value of a5 was defined, it delegates to a4, resulting in 3+1=4.
- b5 isn't defined explicitly, so it delegates to b4 and is therefore equal to 9 as well

sbt Settings Discussion

Scopes sbt defines a more complicated scope than the one shown here for the standard usage of settings in a build. This scope has four components: the project axis, the configuration axis, the task axis, and the extra axis. Each component may be [Global](#) (no specific value), [This](#) (current context), or [Select](#) (containing a specific value). sbt resolves `This_` to either [Global](#) or [Select](#) depending on the context.

For example, in a project, a [This](#) project axis becomes a [Select](#) referring to the defining project. All other axes that are [This](#) are translated to [Global](#). Functions like `inConfig` and `inTask` transform `This` into a [Select](#) for a specific value. For example, `inConfig(Compile)(someSettings)` translates the configuration axis for all settings in `someSettings` to be `Select(Compile)` if the axis value is [This](#).

So, from the example and from sbt's scopes, you can see that the core settings engine does not impose much on the structure of a scope. All it requires is a delegates function `Scope => Seq[Scope]` and a `display` function. You can choose a scope type that makes sense for your situation.

Constructing settings The `app`, `value`, `update`, and related methods are the core methods for constructing settings. This example obviously looks rather different from sbt's interface because these methods are not typically used directly, but are wrapped in a higher-level abstraction.

With the core settings engine, you work with HLists to access other settings. In sbt's higher-level system, there are wrappers around HList for `TupleN` and `FunctionN` for $N = 1-9$ (except `Tuple1` isn't actually used). When working with arbitrary arity, it is useful to make these wrappers at the highest level possible. This is because once wrappers are defined, code must be duplicated for every N . By making the wrappers at the top-level, this requires only one level of duplication.

Additionally, sbt uniformly integrates its task engine into the settings system. The underlying settings engine has no notion of tasks. This is why sbt uses a `SettingKey` type and a `TaskKey` type. Methods on an underlying `TaskKey[T]` are basically translated to operating on an underlying `SettingKey[Task[T]]` (and they both wrap an underlying `AttributeKey`).

For example, `a := 3` for a `SettingKey a` will very roughly translate to `setting(a, value(3))`. For a `TaskKey a`, it will roughly translate to `setting(a, value(task { 3 }))`. See [main/Structure.scala](#) for details.

Settings definitions sbt also provides a way to define these settings in a file (`build.sbt` and `Build.scala`). This is done for `build.sbt` using basic parsing and then passing the resulting chunks of code to `compile/Eval.scala`. For all definitions, sbt manages the classpaths and recompilation process to obtain the settings. It also provides a way for users to define project, task, and configuration delegation, which ends up being used by the delegates function.

Setting Initialization

This page outlines the mechanisms by which sbt loads settings for a particular build, including the hooks where users can control the ordering of everything.

As stated elsewhere, sbt constructs its initialization graph and task graph via `Setting[_]` objects. A setting is something which can take the values stored at other Keys in the build state, and generates a new value for a particular build key. Sbt converts all registered `Setting[_]` objects into a giant linear sequence and *compiles* them into the a task graph. This task graph is then used to execute your build.

All of sbt's loading semantics are contained within the `Load.scala` file. It is approximately the following:




Figure 5: image

The blue circles represent actions happening when sbt loads a project. We can see that sbt performs the following actions in load:

1. Compile the user-level project (`~/.sbt/<version>/`)
 - a. Load any plugins defined by this project (`~/.sbt/<version>/plugins/*.sbt` and `~/.sbt/<version>/plugins/project/*.scala`)
 - b. Load all settings defined (`~/.sbt/<version>/*.sbt` and `~/.sbt/<version>/plugins/*.scala`)
2. Compile the current project (`<working-directory>/project`)
 - a. Load all defined plugins (`project/plugins.sbt` and `project/project/*.scala`)

- b. Load/Compile the project (`project/*.scala`)
3. Load project `*.sbt` files (`build.sbt` and friends).

Each of these loads defines several sequences of settings. The diagram shows the two most important:

- **buildSettings** - These are settings defined to be in `ThisBuild` or directly against the `Build` object. They are initialized *once* for the build. You can add these, e.g. in `project/build.scala` :

```
object MyBuild extends Build {
  override val settings = Seq(foo := "hi")
}
```

or in a `build.sbt` file :

```
foo in ThisBuild := "hi"
```

- **projectSettings** - These are settings specific to a project. They are specific to a *particular sub project* in the build. A plugin may be contributing its settings to more than one project, in which case the values are duplicated for each project. You add project specific settings, eg. in `project/build.scala` :

```
object MyBuild extends Build {
  val test = project.in(file(".")).settings(...)
}
```

After loading/compiling all the build definitions, sbt has a series of `Seq[Setting[_]]` that it must order. As shown in the diagram, the default inclusion order for sbt is:

1. All AutoPlugin settings
2. All settings defined in `project/Build.scala`
3. All settings defined in the user directory (`~/.sbt/<version>/*.sbt`)
4. All local configurations (`build.sbt`)

Controlling Initialization

The order which sbt uses to load settings is configurable at a *project* level. This means that we can't control the order of settings added to `Build/Global` namespace, but we can control how each project loads, e.g. plugins and `.sbt` files. To do so, use the `AddSettings` class :

```

import sbt._
import Keys._

import AddSettings._

object MyOwnOrder extends Build {
  // here we load config from a txt file.
  lazy val root = project.in(file(".")).settingSets( autoPlugins, buildScalaFiles, sbtFiles
}

```

In the above project, we’ve modified the order of settings to be:

1. All AutoPlugin settings.
2. All settings defined in the `project/Build.scala` file (shown above).
3. All settings found in the `silly.txt` file.

What we’ve excluded:

- All settings from the user directory (`~/.sbt/<version>`)
- All `*.sbt` settings.

The `AddSettings` object provides the following “groups” of settings you can use for ordering:

- `autoPlugins` All the ordered settings of plugins after they’ve gone through dependency resolution
- `buildScalaFiles` The full sequence of settings defined directly in `project/*.scala` builds.
- `sbtFiles(*)` Specifies the exact setting DSL files to include (files must use the `.sbt` file format)
- `userSettings` All the settings defined in the user directory `~/.sbt/<version>/`.
- `defaultSbtFiles` Include all local `*.sbt` file settings.

Note: Be very careful when reordering settings. It’s easy to accidentally remove core functionality.

For example, let’s see what happens if we move the `build.sbt` files *before* the `buildScalaFile`.

Let’s create an example project the following definition. `project/build.scala`:

```

object MyTestBuild extends Build {
  val testProject = project.in(file(".")).settingSets(autoPlugins, defaultSbtFiles, buildSc
    version := scalaBinaryVersion.value match {
      case "2.10" => "1.0-SNAPSHOT"
      case v => "1.0-for- $\{v\}$ -SNAPSHOT"
    }
  )
}

```

This build defines a version string which appends the scala version if the current scala version is not the in the 2.10.x series. Now, when issuing a release we want to lock down the version. Most tools assume this can happen by writing a `version.sbt` file. `version.sbt` :

```

version := "1.0.0"

```

However, when we load this new build, we find that the `version` in `version.sbt` has been **overridden** by the one defined in `project/Build.scala` because of the order we defined for settings, so the new `version.sbt` file has no effect.

Build Loaders

Build loaders are the means by which sbt resolves, builds, and transforms build definitions. Each aspect of loading may be customized for special applications. Customizations are specified by overriding the *buildLoaders* methods of your build definition's Build object. These customizations apply to external projects loaded by the build, but not the (already loaded) Build in which they are defined. Also documented on this page is how to manipulate inter-project dependencies from a setting.

Custom Resolver

The first type of customization introduces a new resolver. A resolver provides support for taking a build URI and retrieving it to a local directory on the filesystem. For example, the built-in resolver can checkout a build using git based on a git URI, use a build in an existing local directory, or download and extract a build packaged in a jar file. A resolver has type:

```

ResolveInfo => Option[() => File]

```

The resolver should return `None` if it cannot handle the URI or `Some` containing a function that will retrieve the build. The `ResolveInfo` provides a staging directory that can be used or the resolver can determine its own target directory.

Whichever is used, it should be returned by the loading function. A resolver is registered by passing it to *BuildLoader.resolve* and overriding *Build.buildLoaders* with the result:

```
...
object Demo extends Build {
  ...
  override def buildLoaders =
    BuildLoader.resolve(demoResolver) ::
    Nil

  def demoResolver: BuildLoader.ResolveInfo => Option[() => File] = ...
}
```

API Documentation Relevant API documentation for custom resolvers:

- [ResolveInfo](#)
- [BuildLoader](#)

Full Example

```
import sbt._
import Keys._

object Demo extends Build
{
  // Define a project that depends on an external project with a custom URI
  lazy val root = Project("root", file(".")).dependsOn(
    uri("demo:a")
  )

  // Register the custom resolver
  override def buildLoaders =
    BuildLoader.resolve(demoResolver) ::
    Nil

  // Define the custom resolver, which handles the 'demo' scheme.
  // The resolver's job is to produce a directory containing the project to load.
  // A subdirectory of info.staging can be used to create new local
  // directories, such as when doing 'git clone ...'
  def demoResolver(info: BuildLoader.ResolveInfo): Option[() => File] =
    if(info.uri.getScheme != "demo")
      None
}
```



```

else
{
  // Use a subdirectory of the staging directory for the new local build.
  // The subdirectory name is derived from a hash of the URI,
  // and so identical URIs will resolve to the same directory (as desired).
  val base = RetrieveUnit.temporary(info.staging, info.uri)

  // Return a closure that will do the actual resolution when requested.
  Some(() => resolveDemo(base, info.uri.getSchemeSpecificPart))
}

// Construct a sample project on the fly with the name specified in the URI.
def resolveDemo(base: File, ssp: String): File =
{
  // Only create the project if it hasn't already been created.
  if(!base.exists)
    IO.write(base / "build.sbt", template.format(ssp))
  base
}

def template = """
name := "%s"

version := "1.0"
"""
}

```

Custom Builder

Once a project is resolved, it needs to be built and then presented to sbt as an instance of `sbt.BuildUnit`. A custom builder has type:

```
BuildInfo => Option[() => BuildUnit]
```

A builder returns `None` if it does not want to handle the build identified by the `BuildInfo`. Otherwise, it provides a function that will load the build when evaluated. Register a builder by passing it to `BuildLoader.build` and overriding `Build.buildLoaders` with the result:

```

...
object Demo extends Build {
  ...
  override def buildLoaders =
    BuildLoader.build(demoBuilder) ::

```

```

Nil

def demoBuilder: BuildLoader.BuildInfo => Option[() => BuildUnit] = ...
}

```

API Documentation Relevant API documentation for custom builders:

- [BuildInfo](#)
- [BuildLoader](#)
- [BuildUnit](#)

Example This example demonstrates the structure of how a custom builder could read configuration from a pom.xml instead of the standard .sbt files and project/ directory.

```

... imports ...

object Demo extends Build
{
  lazy val root = Project("root", file(".")) dependsOn( file("basic-pom-project") )

  override def buildLoaders =
    BuildLoader.build(demoBuilder) ::
    Nil

  def demoBuilder: BuildInfo => Option[() => BuildUnit] = info =>
    if(pomFile(info).exists)
      Some(() => pomBuild(info))
    else
      None

  def pomBuild(info: BuildInfo): BuildUnit =
  {
    val pom = pomFile(info)
    val model = readPom(pom)

    val n = Project.normalizeProjectID(model.getName)
    val base = Option(model.getProjectDirectory) getOrElse info.base
    val root = Project(n, base) settings( pomSettings(model) : _*)
    val build = new Build { override def projects = Seq(root) }
    val loader = this.getClass.getClassLoader
    val definitions = new LoadedDefinitions(info.base, Nil, loader, build :: Nil, Nil)
    val plugins = new LoadedPlugins(info.base / "project", Nil, loader, Nil, Nil)
  }
}

```

```

    new BuildUnit(info.uri, info.base, definitions, plugins)
  }

  def readPom(file: File): Model = ...
  def pomSettings(m: Model): Seq[Setting[_]] = ...
  def pomFile(info: BuildInfo): File = info.base / "pom.xml"

```

Custom Transformer

Once a project has been loaded into an `sbt.BuildUnit`, it is transformed by all registered transformers. A custom transformer has type:

```
TransformInfo => BuildUnit
```

A transformer is registered by passing it to `BuildLoader.transform` and overriding `Build.buildLoaders` with the result:

```

...
object Demo extends Build {
  ...
  override def buildLoaders =
    BuildLoader.transform(demoTransformer) ::
    Nil

  def demoBuilder: BuildLoader.TransformInfo => BuildUnit = ...
}

```

API Documentation Relevant API documentation for custom transformers:

- [TransformInfo](#)
- [BuildLoader](#)
- [BuildUnit](#)

Manipulating Project Dependencies in Settings The `buildDependencies` setting, in the Global scope, defines the aggregation and classpath dependencies between projects. By default, this information comes from the dependencies defined by `Project` instances by the `aggregate` and `dependsOn` methods. Because `buildDependencies` is a setting and is used everywhere dependencies need to be known (once all projects are loaded), plugins and build definitions can transform it to manipulate inter-project dependencies at setting evaluation time. The only requirement is that no new projects are introduced because all projects are loaded before settings get evaluated. That is, all `Projects` must have been declared directly in a `Build` or referenced as the argument to `Project.aggregate` or `Project.dependsOn`.

The BuildDependencies type

The type of the `buildDependencies` setting is [BuildDependencies](#). `BuildDependencies` provides mappings from a project to its aggregate or classpath dependencies. For classpath dependencies, a dependency has type `ClasspathDep[ProjectRef]`, which combines a `ProjectRef` with a configuration (see [ClasspathDep](#) and [ProjectRef](#)). For aggregate dependencies, the type of a dependency is just `ProjectRef`.

The API for `BuildDependencies` is not extensive, covering only a little more than the minimum required, and related APIs have more of an internal, unpolished feel. Most manipulations consist of modifying the relevant map (classpath or aggregate) manually and creating a new `BuildDependencies` instance.

Example As an example, the following replaces a reference to a specific build URI with a new URI. This could be used to translate all references to a certain git repository to a different one or to a different mechanism, like a local directory.

```
buildDependencies in Global := {
  val deps = (buildDependencies in Global).value
  val oldURI = uri("...") // the URI to replace
  val newURI = uri("...") // the URI replacing oldURI
  def substitute(dep: ClasspathDep[ProjectRef]): ClasspathDep[ProjectRef] =
    if(dep.project.build == oldURI)
      ResolvedClasspathDependency(ProjectRef(newURI, dep.project.project), dep.configuration)
    else
      dep
  val newcp =
    for( (proj, deps) <- deps.cp ) yield
      (proj, deps map substitute)
  new BuildDependencies(newcp, deps.aggregate)
}
```

It is not limited to such basic translations, however. The configuration a dependency is defined in may be modified and dependencies may be added or removed. Modifying `buildDependencies` can be combined with modifying `libraryDependencies` to convert binary dependencies to and from source dependencies, for example.

Creating Command Line Applications Using sbt

There are several components of sbt that may be used to create a command line application. The [launcher](#) and the [command system](#) are the two main ones illustrated here.

As described on the [launcher page](#), a launched application implements the `xsbti.AppMain` interface and defines a brief configuration file that users pass to the launcher to run the application. To use the command system, an application sets up a [State](#) instance that provides [command implementations](#) and the initial commands to run. A minimal hello world example is given below.

Hello World Example

There are three files in this example:

1. `build.sbt`
2. `Main.scala`
3. `hello.build.properties`

To try out this example:

1. Put the first two files in a new directory
2. Run `sbt publishLocal` in that directory
3. Run `sbt @path/to/hello.build.properties` to run the application.

Like for sbt itself, you can specify commands from the command line (batch mode) or run them at an prompt (interactive mode).

Build Definition: `build.sbt` The `build.sbt` file should define the standard settings: name, version, and organization. To use the sbt command system, a dependency on the `command` module is needed. To use the task system, add a dependency on the `task-system` module as well.

```
organization := "org.example"

name := "hello"

version := "0.1-SNAPSHOT"

libraryDependencies += "org.scala-sbt" % "command" % "0.12.0"
```

Application: `Main.scala` The application itself is defined by implementing [xsbti.AppMain](#). The basic steps are

1. Provide command definitions. These are the commands that are available for users to run.

2. Define initial commands. These are the commands that are initially scheduled to run. For example, an application will typically add anything specified on the command line (what sbt calls batch mode) and if no commands are defined, enter interactive mode by running the 'shell' command.
3. Set up logging. The default setup in the example rotates the log file after each user interaction and sends brief logging to the console and verbose logging to the log file.

```
package org.example

import sbt._
import java.io.{File, PrintWriter}

final class Main extends xsbti.AppMain
{
  /** Defines the entry point for the application.
   * The call to `initialState` sets up the application.
   * The call to runLogged starts command processing. */
  def run(configuration: xsbti.AppConfiguration): xsbti.MainResult =
    MainLoop.runLogged( initialState(configuration) )

  /** Sets up the application by constructing an initial State instance with the supported
   * and initial commands to run. See the State API documentation for details. */
  def initialState(configuration: xsbti.AppConfiguration): State =
  {
    val commandDefinitions = hello +: BasicCommands.allBasicCommands
    val commandsToRun = Hello +: "iflast shell" +: configuration.arguments.map(_.trim)
    State( configuration, commandDefinitions, Set.empty, None, commandsToRun, State.newHis
      AttributeMap.empty, initialGlobalLogging, State.Continue )
  }

  // defines an example command. see the Commands page for details.
  val Hello = "hello"
  val hello = Command.command(Hello) { s =>
    s.log.info("Hello!")
    s
  }

  /** Configures logging to log to a temporary backing file as well as to the console.
   * An application would need to do more here to customize the logging level and
   * provide access to the backing file (like sbt's last command and logLevel setting).*/
  def initialGlobalLogging: GlobalLogging =
    GlobalLogging.initial(MainLogging.globalDefault _, File.createTempFile("hello", "log")
  }
```

Launcher configuration file: `hello.build.properties` The launcher needs a configuration file in order to retrieve and run an application.
`hello.build.properties`:

```
[scala]
  version: 2.9.1
```

```
[app]
  org: org.example
  name: hello
  version: 0.1-SNAPSHOT
  class: org.example.Main
  components: xsbti
  cross-versioned: true
```

```
[repositories]
  local
  maven-central
  typesafe-ivy-releases: http://repo.typesafe.com/typesafe/ivy-releases/, [organization]/[mod
```

Nightly Builds

The latest development versions of 0.13.5 are available as nightly builds on [Typesafe Snapshots](#).

To use a nightly build, the instructions are the same for [normal manual setup](#) except:

1. Download the launcher jar from one of the subdirectories of `nightly-launcher`. They should be listed in chronological order, so the most recent one will be last.
2. The version number is the name of the subdirectory and is of the form `0.13.5.x-yyyyMMdd-HH:mm:ss`. Use this in a `build.properties` file.
3. Call your script something like `sbt-nightly` to retain access to a stable sbt launcher. The documentation will refer to the script as `sbt`, however.

Related to the third point, remember that an `sbt.version` setting in `<build-base>/project/build.properties` determines the version of sbt to use in a project. If it is not present, the default version associated with the launcher is used. This means that you must set `sbt.version=yyyyMMdd-HH:mm:ss` in an existing `<build-base>/project/build.properties`. You can verify the right version of sbt is being used to build a project by running `about`.

To reduce problems, it is recommended to not use a launcher jar for one nightly version to launch a different nightly version of sbt.

How to...

See [Detailed Table of Contents](#) for the list of all the how-tos.

Classpaths

Include a new type of managed artifact on the classpath, such as `mar`

The `classpathTypes` setting controls the types of managed artifacts that are included on the classpath by default. To add a new type, such as `mar`,

```
classpathTypes += "mar"
```

Get the classpath used for compilation

See the default types included by running `show classpathTypes` at the sbt prompt.

The `dependencyClasspath` task scoped to `Compile` provides the classpath to use for compilation. Its type is `Seq[Attributed[File]]`, which means that each entry carries additional metadata. The `files` method provides just the raw `Seq[File]` for the classpath. For example, to use the files for the compilation classpath in another task, :

```
example := {  
  val cp: Seq[File] = (dependencyClasspath in Compile).value.files  
  ...  
}
```

Note: This classpath does not include the class directory, which may be necessary for compilation in some situations.

Get the runtime classpath, including the project's compiled classes

The `fullClasspath` task provides a classpath including both the dependencies and the products of project. For the runtime classpath, this means the main resources and compiled classes for the project as well as all runtime dependencies.

The type of a classpath is `Seq[Attributed[File]]`, which means that each entry carries additional metadata. The `files` method provides just the raw `Seq[File]` for the classpath. For example, to use the files for the runtime classpath in another task, :


```
example := {
  val cp: Seq[File] = (fullClasspath in Runtime).value.files
  ...
}
```

Get the test classpath, including the project's compiled test classes

The `fullClasspath` task provides a classpath including both the dependencies and the products of a project. For the test classpath, this includes the main and test resources and compiled classes for the project as well as all dependencies for testing.

The type of a classpath is `Seq[Attributed[File]]`, which means that each entry carries additional metadata. The `files` method provides just the raw `Seq[File]` for the classpath. For example, to use the files for the test classpath in another task, :

```
example := {
  val cp: Seq[File] = (fullClasspath in Test).value.files
  ...
}
```

Use packaged jars on classpaths instead of class directories

By default, `fullClasspath` includes a directory containing class files and resources for a project. This in turn means that tasks like `compile`, `test`, and `run` have these class directories on their classpath. To use the packaged artifact (such as a jar) instead, configure `exportJars` :

```
exportJars := true
```

This will use the result of `packageBin` on the classpath instead of the class directory.

Note: Specifically, `fullClasspath` is the concatenation of `dependencyClasspath` and `exportedProducts`. When `exportJars` is true, `exportedProducts` is the output of `packageBin`. When `exportJars` is false, `exportedProducts` is just `products`, which is by default the directory containing class files and resources.

Get all managed jars for a configuration

The result of the `update` task has type [UpdateReport](#), which contains the results of dependency resolution. This can be used to extract the files for specific types of artifacts in a specific configuration. For example, to get the jars and zips of dependencies in the `Compile` configuration, :

```
example := {  
  val artifactTypes = Set("jar", "zip")  
  val files: Seq[File] =  
    Classpaths.managedJars(Compile, artifactTypes, update.value)  
  ...  
}
```

Get the files included in a classpath

A classpath has type `Seq[Attributed[File]]`, which means that each entry carries additional metadata. The `files` method provides just the raw `Seq[File]` for the classpath. For example, :

```
val cp: Seq[Attributed[File]] = ...  
val files: Seq[File] = cp.files
```

Get the module and artifact that produced a classpath entry

A classpath has type `Seq[Attributed[File]]`, which means that each entry carries additional metadata. This metadata is in the form of an [AttributeMap](#). Useful keys for entries in the map are `artifact.key`, `module.key`, and `analysis`. For example,

```
val classpath: Seq[Attributed[File]] = ???  
for(entry <- classpath) yield {  
  val art: Option[Artifact] = entry.get(artifact.key)  
  val mod: Option[ModuleID] = entry.get(module.key)  
  val an: Option[inc.Analysis] = entry.get(analysis)  
  ...  
}
```

Note: Entries may not have some or all metadata. Only entries from source dependencies, such as internal projects, have an incremental compilation [Analysis](#). Only entries for managed dependencies have an [Artifact](#) and [ModuleID](#).

Customizing paths

This page describes how to modify the default source, resource, and library directories and what files get included from them.

Change the default Scala source directory

The directory that contains the main Scala sources is by default `src/main/scala`. For test Scala sources, it is `src/test/scala`. To change this, modify `scalaSource` in the `Compile` (for main sources) or `Test` (for test sources). For example,

```
scalaSource in Compile := baseDirectory.value / "src"
```

```
scalaSource in Test := baseDirectory.value / "test-src"
```

Note: The Scala source directory can be the same as the Java source directory.

Change the default Java source directory

The directory that contains the main Java sources is by default `src/main/java`. For test Java sources, it is `src/test/java`. To change this, modify `javaSource` in the `Compile` (for main sources) or `Test` (for test sources).

For example,

```
javaSource in Compile := baseDirectory.value / "src"
```

```
javaSource in Test := baseDirectory.value / "test-src"
```

Note: The Scala source directory can be the same as the Java source directory.

Change the default resource directory

The directory that contains the main resources is by default `src/main/resources`. For test resources, it is `src/test/resources`. To change this, modify `resourceDirectory` in either the `Compile` or `Test` configuration.

For example,

```
resourceDirectory in Compile := baseDirectory.value / "resources"
```

```
resourceDirectory in Test := baseDirectory.value / "test-resources"
```

Change the default (unmanaged) library directory

The directory that contains the unmanaged libraries is by default `lib/`. To change this, modify `unmanagedBase`. This setting can be changed at the project level or in the `Compile`, `Runtime`, or `Test` configurations.

When defined without a configuration, the directory is the default directory for all configurations. For example, the following declares `jars/` as containing libraries:

```
unmanagedBase := baseDirectory.value / "jars"
```

When set for `Compile`, `Runtime`, or `Test`, `unmanagedBase` is the directory containing libraries for that configuration, overriding the default. For example, the following declares `lib/main/` to contain jars only for `Compile` and not for running or testing: :

```
unmanagedBase in Compile := baseDirectory.value / "lib" / "main"
```

Disable using the project's base directory as a source directory

By default, sbt includes `.scala` files from the project's base directory as main source files. To disable this, configure `sourcesInBase`:

```
sourcesInBase := false
```

Add an additional source directory

sbt collects `sources` from `unmanagedSourceDirectories`, which by default consists of `scalaSource` and `javaSource`. Add a directory to `unmanagedSourceDirectories` in the appropriate configuration to add a source directory. For example, to add `extra-src` to be an additional directory containing main sources,

```
unmanagedSourceDirectories in Compile += baseDirectory.value / "extra-src"
```

Note: This directory should only contain unmanaged sources, which are sources that are manually created and managed. See [Generating Files](#) for working with automatically generated sources.

Add an additional resource directory

sbt collects `resources` from `unmanagedResourceDirectories`, which by default consists of `resourceDirectory`. Add a directory to `unmanagedResourceDirectories` in the appropriate configuration to add another resource directory. For example, to add `extra-resources` to be an additional directory containing main resources,

```
unmanagedResourceDirectories in Compile += baseDirectory.value / "extra-resources"
```

Note: This directory should only contain unmanaged resources, which are resources that are manually created and managed. See [Generating Files](#) for working with automatically generated resources.

Include/exclude files in the source directory

When sbt traverses `unmanagedSourceDirectories` for sources, it only includes directories and files that match `includeFilter` and do not match `excludeFilter`. `includeFilter` and `excludeFilter` have type `java.io.FileFilter` and sbt [provides some useful combinators](#) for constructing a `FileFilter`. For example, in addition to the default hidden files exclusion, the following also ignores files containing `impl` in their name,

```
excludeFilter in unmanagedSources := HiddenFileFilter || "*impl*"
```

To have different filters for main and test libraries, configure `Compile` and `Test` separately:

```
includeFilter in (Compile, unmanagedSources) := "*.scala" || "*.java"
```

```
includeFilter in (Test, unmanagedSources) := HiddenFileFilter || "*impl*"
```

Note: By default, sbt includes `.scala` and `.java` sources, excluding hidden files.

Include/exclude files in the resource directory

When sbt traverses `unmanagedResourceDirectories` for resources, it only includes directories and files that match `includeFilter` and do not match `excludeFilter`. `includeFilter` and `excludeFilter` have type `java.io.FileFilter` and sbt [provides some useful combinators](#) for constructing a `FileFilter`. For example, in addition to the default hidden files exclusion, the following also ignores files containing `impl` in their name,

```
excludeFilter in unmanagedSources := HiddenFileFilter || "*impl*
```

To have different filters for main and test libraries, configure `Compile` and `Test` separately:

```
includeFilter in (Compile, unmanagedSources) := "*.txt"
```

```
includeFilter in (Test, unmanagedSources) := "*.html"
```

Note: By default, sbt includes all files that are not hidden.

Include only certain (unmanaged) libraries

When sbt traverses `unmanagedBase` for resources, it only includes directories and files that match `includeFilter` and do not match `excludeFilter`. `includeFilter` and `excludeFilter` have type `java.io.FileFilter` and sbt [provides some useful combinators](#) for constructing a `FileFilter`. For example, in addition to the default hidden files exclusion, the following also ignores zips,

```
excludeFilter in unmanagedJars := HiddenFileFilter || "*.zip"
```

To have different filters for main and test libraries, configure `Compile` and `Test` separately:

```
includeFilter in (Compile, unmanagedJars) := "*.jar"
```

```
includeFilter in (Test, unmanagedJars) := "*.jar" || "*.zip"
```

Note: By default, sbt includes jars, zips, and native dynamic libraries, excluding hidden files.

Generating files

sbt provides standard hooks for adding source and resource generation tasks.

Generate sources

A source generation task should generate sources in a subdirectory of `sourceManaged` and return a sequence of files generated. The signature of a source generation function (that becomes a basis for a task) is usually as follows:

```
def makeSomeSources(base: File): Seq[File]
```

The key to add the task to is called `sourceGenerators`. Because we want to add the task, and not the value after its execution, we use `taskValue` instead of the usual `value`. `sourceGenerators` should be scoped according to whether the generated files are main (`Compile`) or test (`Test`) sources. This basic structure looks like:

```
sourceGenerators in Compile += <task of type Seq[File]>.taskValue
```

For example, assuming a method `def makeSomeSources(base: File): Seq[File]`,

```
sourceGenerators in Compile += Def.task {  
  makeSomeSources((sourceManaged in Compile).value / "demo")  
}.taskValue
```

As a specific example, the following source generator generates `Test.scala` application object that once executed, prints "Hi" to the console:

```
sourceGenerators in Compile += Def.task {  
  val file = (sourceManaged in Compile).value / "demo" / "Test.scala"  
  IO.write(file, """"object Test extends App { println("Hi") }""")  
  Seq(file)  
}.taskValue
```

Executing `run` will print "Hi".

```
> run  
[info] Running Test  
Hi
```

Change `Compile` to `Test` to make it a test source. For efficiency, you would only want to generate sources when necessary and not every run.

By default, generated sources are not included in the packaged source artifact. To do so, add them as you would other mappings. See [Adding files to a package](#). A source generator can return both Java and Scala sources mixed together in the same sequence. They will be distinguished by their extension later.

Generate resources

A resource generation task should generate resources in a subdirectory of `resourceManaged` and return a sequence of files generated. Like a source generation function, the signature of a resource generation function (that becomes a basis for a task) is usually as follows:

```
def makeSomeResources(base: File): Seq[File]
```

The key to add the task to is called `resourceGenerators`. Because we want to add the task, and not the value after its execution, we use `taskValue` instead of the usual `value`. It should be scoped according to whether the generated files are main (`Compile`) or test (`Test`) resources. This basic structure looks like:

```
resourceGenerators in Compile += <task of type Seq[File]>.taskValue
```

For example, assuming a method `def makeSomeResources(base: File): Seq[File]`,

```
resourceGenerators in Compile += Def.task {  
  makeSomeResources((resourceManaged in Compile).value / "demo")  
}.taskValue
```

As a specific example, the following generates a properties file `myapp.properties` containing the application name and version:

```
resourceGenerators in Compile += Def.task {  
  val file = (resourceManaged in Compile).value / "demo" / "myapp.properties"  
  val contents = "name=%s\nversion=%s".format(name.value, version.value)  
  IO.write(file, contents)  
  Seq(file)  
}.taskValue
```

Change `Compile` to `Test` to make it a test resource. Normally, you would only want to generate resources when necessary and not every run.

By default, generated resources are not included in the packaged source artifact. To do so, add them as you would other mappings. See [Adding files to a package](#).

Inspect the build

Show or search help for a command, task, or setting

The `help` command is used to show available commands and search the help for commands, tasks, or settings. If run without arguments, `help` lists the available commands.


```

> help

help          Displays this help message or prints detailed help on
               requested commands (run 'help <command>').
about         Displays basic information about sbt and the build.
reload        (Re)loads the project in the current directory
...

> help compile

```

If the argument passed to **help** is the name of an existing command, setting or task, the help for that entity is displayed. Otherwise, the argument is interpreted as a regular expression that is used to search the help of all commands, settings and tasks.

The **tasks** command is like **help**, but operates only on tasks. Similarly, the **settings** command only operates on settings.

See also **help help**, **help tasks**, and **help settings**.

List available tasks

The **tasks** command, without arguments, lists the most commonly used tasks. It can take a regular expression to search task names and descriptions. The verbosity can be increased to show or search less commonly used tasks. See **help tasks** for details.

The **settings** command, without arguments, lists the most commonly used settings. It can take a regular expression to search setting names and descriptions. The verbosity can be increased to show or search less commonly used settings. See **help settings** for details.

List available settings

The **inspect** command displays several pieces of information about a given setting or task, including the dependencies of a task/setting as well as the tasks/settings that depend on the it. For example,

```

> inspect test:compile
...
[info] Dependencies:
[info]   test:compile::compileInputs
[info]   test:compile::streams
[info] Reverse dependencies:
[info]   test:definedTestNames

```

```

[info] test:definedSbtPlugins
[info] test:printWarnings
[info] test:discoveredMainClasses
[info] test:definedTests
[info] test:exportedProducts
[info] test:products
...

```

See the [Inspecting Settings](#) page for details.

Display tree of setting/task dependencies

In addition to displaying immediate forward and reverse dependencies as described in the previous section, the `inspect` command can display the full dependency tree for a task or setting. For example,

```

> inspect tree clean
[info] *:clean = Task[Unit]
[info] +-*:cleanFiles = List(<project>/lib_managed, <project>/target)
[info] | +-{.}/*:managedDirectory = lib_managed
[info] | +-*:target = target
[info] |   +-*:baseDirectory = <project>
[info] |   +-*:thisProject = Project(id: demo, base: <project>, ...
[info] |
[info] +-*:cleanKeepFiles = List(<project>/target/.history)
[info]   +-*:history = Some(<project>/target/.history)
...

```

For each task, `inspect tree` show the type of the value generated by the task. For a setting, the `toString` of the setting is displayed. See the [Inspecting Settings](#) page for details on the `inspect` command.

Display the description and type of a setting or task

While the `help`, `settings`, and `tasks` commands display a description of a task, the `inspect` command also shows the type of a setting or task and the value of a setting. For example:

```

> inspect update
[info] Task: sbt.UpdateReport
[info] Description:
[info] Resolves and optionally retrieves dependencies, producing a report.
...

```

```
> inspect scalaVersion
[info] Setting: java.lang.String = 2.9.2
[info] Description:
[info]   The version of Scala used for building.
...
```

See the [Inspecting Settings](#) page for details.

Display the delegation chain of a setting or task

See the [Inspecting Settings](#) page for details.

Display related settings or tasks

The `inspect` command can help find scopes where a setting or task is defined. The following example shows that different options may be specified to the Scala for testing and API documentation generation.

```
> inspect scalacOptions
...
[info] Related:
[info]   compile:doc::scalacOptions
[info]   test:scalacOptions
[info]   */*:scalacOptions
[info]   test:doc::scalacOptions
```

See the [Inspecting Settings](#) page for details.

Show the list of projects and builds

The `projects` command displays the currently loaded projects. The projects are grouped by their enclosing build and the current project is indicated by an asterisk. For example,

```
> projects
[info] In file:/home/user/demo/
[info]   * parent
[info]     sub
[info] In file:/home/user/dep/
[info]   sample
```

Show the current session (temporary) settings

`session list` displays the settings that have been added at the command line for the current project. For example,

```
> session list
  1. maxErrors := 5
  2. scalacOptions += "-explaintypes"
```

`session list-all` displays the settings added for all projects. For details, see `help session`.

Show basic information about sbt and the current build

```
> about
[info] This is sbt 0.12.0
[info] The current project is {file:~/code/sbt.github.com/}default
[info] The current project is built against Scala 2.9.2
[info] Available Plugins: com.jsuereth.ghpages.GhPages, com.jsuereth.git.GitPlugin, com.jsuereth.ghpages.GhPages, com.jsuereth.git.GitPlugin, com.jsuereth.ghpages.GhPages, com.jsuereth.git.GitPlugin
[info] sbt, sbt plugins, and build definitions are using Scala 2.9.2
```

Show the value of a setting

The `inspect` command shows the value of a setting as part of its output, but the `show` command is dedicated to this job. It shows the output of the setting provided as an argument. For example,

```
> show organization
[info] com.github.sbt
```

The `show` command also works for tasks, described next.

Show the result of executing a task

```
> show update
... <output of update> ...
[info] Update report:
[info] Resolve time: 122 ms, Download time: 5 ms, Download size: 0 bytes
[info] compile:
[info]     org.scala-lang:scala-library:2.9.2: ...
```

The **show** command will execute the task provided as an argument and then print the result. Note that this is different from the behavior of the **inspect** command (described in other sections), which does not execute a task and thus can only display its type and not its generated value.

```
> show compile:dependencyClasspath
...
[info] ArrayBuffer(Attributed(~/.sbt/0.12.0/boot/scala-2.9.2/lib/scala-library.jar))
```

Show the classpath used for compilation or testing

For the test classpath,

```
> show test:dependencyClasspath
...
[info] ArrayBuffer(Attributed(~/.code/sbt.github.com/target/scala-2.9.2/classes), Attributed(
```

Show the main classes detected in a project

sbt detects the classes with public, static main methods for use by the **run** method and to tab-complete the **runMain** method. The **discoveredMainClasses** task does this discovery and provides as its result the list of class names. For example, the following shows the main classes discovered in the main sources:

```
> show compile:discoveredMainClasses
... <runs compile if out of date> ...
[info] List(org.example.Main)
```

Show the test classes detected in a project

sbt detects tests according to fingerprints provided by test frameworks. The **definedTestNames** task provides as its result the list of test names detected in this way. For example,

```
> show test:definedTestNames
... < runs test:compile if out of date > ...
[info] List(org.example.TestA, org.example.TestB)
```

Interactive mode

Use tab completion

By default, sbt's interactive mode is started when no commands are provided on the command line or when the `shell` command is invoked.

As the name suggests, tab completion is invoked by hitting the tab key. Suggestions are provided that can complete the text entered to the left of the current cursor position. Any part of the suggestion that is unambiguous is automatically appended to the current text. Commands typically support tab completion for most of their syntax.

As an example, entering `tes` and hitting tab:

```
> tes<TAB>
```

results in sbt appending a `t`:

```
> test
```

To get further completions, hit tab again:

```
> test<TAB>
testFrameworks  testListeners  testLoader      testOnly      testOptions  test:
```

Now, there is more than one possibility for the next character, so sbt prints the available options. We will select `testOnly` and get more suggestions by entering the rest of the command and hitting tab twice:

```
> testOnly<TAB><TAB>
--          sbt.DagSpecification  sbt.EmptyRelationTest  sbt.KeyTest          sbt.Relati
```

The first tab inserts an unambiguous space and the second suggests names of tests to run. The suggestion of `--` is for the separator between test names and options provided to the test framework. The other suggestions are names of test classes for one of sbt's modules. Test name suggestions require tests to be compiled first. If tests have been added, renamed, or removed since the last test compilation, the completions will be out of date until another successful compile.

Show more tab completion suggestions

Some commands have different levels of completion. Hitting tab multiple times increases the verbosity of completions. (Presently, this feature is only used by the `set` command.)

Modify the default JLine keybindings

JLine, used by both Scala and sbt, uses a configuration file for many of its keybindings. The location of this file can be changed with the system property `jline.keybindings`. The default keybindings file is included in the sbt launcher and may be used as a starting point for customization.

Configure the prompt string

By default, sbt only displays `>` to prompt for a command. This can be changed through the `shellPrompt` setting, which has type `State => String`. `State` contains all state for sbt and thus provides access to all build information for use in the prompt string.

Examples:

```
// set the prompt (for this build) to include the project id.
shellPrompt in ThisBuild := { state => Project.extract(state).currentRef.project + "> " }

// set the prompt (for the current project) to include the username
shellPrompt := { state => System.getProperty("user.name") + "> " }
```

Use history

Interactive mode remembers history even if you exit sbt and restart it. The simplest way to access history is to press the up arrow key to cycle through previously entered commands. Use `Ctrl+r` to incrementally search history backwards. The following commands are supported:

- `!` Show history command help.
- `!!` Execute the previous command again.
- `!:` Show all previous commands.
- `!:n` Show the last `n` commands.
- `!n` Execute the command with index `n`, as shown by the `!:` command.
- `!-n` Execute the `n`th command before this one.
- `!string` Execute the most recent command starting with `'string'`
- `!?string` Execute the most recent command containing `'string'`

Change the location of the interactive history file

By default, interactive history is stored in the `target/` directory for the current project (but is not removed by a `clean`). History is thus separate for each subproject. The location can be changed with the `historyPath` setting, which

has type `Option[File]`. For example, history can be stored in the root directory for the project instead of the output directory:

```
historyPath := Some(baseDirectory.value / ".history")
```

The history path needs to be set for each project, since sbt will use the value of `historyPath` for the current project (as selected by the `project` command).

Use the same history for all projects

The previous section describes how to configure the location of the history file. This setting can be used to share the interactive history among all projects in a build instead of using a different history for each project. The way this is done is to set `historyPath` to be the same file, such as a file in the root project's `target/` directory:

```
historyPath :=  
  Some( (target in LocalRootProject).value / ".history")
```

The `in LocalRootProject` part means to get the output directory for the root project for the build.

Disable interactive history

If, for whatever reason, you want to disable history, set `historyPath` to `None` in each project it should be disabled in:

```
> historyPath := None
```

Run commands before entering interactive mode

Interactive mode is implemented by the `shell` command. By default, the `shell` command is run if no commands are provided to sbt on the command line. To run commands before entering interactive mode, specify them on the command line followed by `shell`. For example,

```
$ sbt clean compile shell
```

This runs `clean` and then `compile` before entering the interactive prompt. If either `clean` or `compile` fails, sbt will exit without going to the prompt. To enter the prompt whether or not these initial commands succeed, prepend `-shell`, which means to run `shell` if any command fails. For example,

```
$ sbt -shell clean compile shell
```


Configure and use logging

View the logging output of the previously executed command

When a command is run, more detailed logging output is sent to a file than to the screen (by default). This output can be recalled for the command just executed by running `last`.

For example, the output of `run` when the sources are uptodate is:

```
> run
[info] Running A
Hi!
[success] Total time: 0 s, completed Feb 25, 2012 1:00:00 PM
```

The details of this execution can be recalled by running `last`:

```
> last
[debug] Running task... Cancelable: false, max worker threads: 4, check cycles: false
[debug]
[debug] Initial source changes:
[debug]   removed: Set()
[debug]   added: Set()
[debug]   modified: Set()
[debug] Removed products: Set()
[debug] Modified external sources: Set()
[debug] Modified binary dependencies: Set()
[debug] Initial directly invalidated sources: Set()
[debug]
[debug] Sources indirectly invalidated by:
[debug]   product: Set()
[debug]   binary dep: Set()
[debug]   external source: Set()
[debug] Initially invalidated: Set()
[debug] Copy resource mappings:
[debug]
[info] Running A
[debug] Starting sandboxed run...
[debug] Waiting for threads to exit or System.exit to be called.
[debug]   Classpath:
[debug]     /tmp/e/target/scala-2.9.2/classes
[debug]     /tmp/e/.sbt/0.12.0/boot/scala-2.9.2/lib/scala-library.jar
[debug] Waiting for thread runMain to exit
[debug]   Thread runMain exited.
[debug] Interrupting remaining threads (should be all daemons).
```

```
[debug] Sandboxed run complete..  
[debug] Exited with code 0  
[success] Total time: 0 s, completed Jan 1, 2012 1:00:00 PM
```

Configuration of the logging level for the console and for the backing file are described in following sections.

View the previous logging output of a specific task

When a task is run, more detailed logging output is sent to a file than to the screen (by default). This output can be recalled for a specific task by running `last <task>`. For example, the first time `compile` is run, output might look like:

```
> compile  
[info] Updating {file:/.../demo/}example...  
[info] Resolving org.scala-lang#scala-library;2.9.2 ...  
[info] Done updating.  
[info] Compiling 1 Scala source to .../demo/target/scala-2.9.2/classes...  
[success] Total time: 0 s, completed Jun 1, 2012 1:11:11 PM
```

The output indicates that both dependency resolution and compilation were performed. The detailed output of each of these may be recalled individually. For example,

```
> last compile  
[debug]  
[debug] Initial source changes:  
[debug]   removed:Set()  
[debug]   added: Set(/home/mark/tmp/a/b/A.scala)  
[debug]   modified: Set()  
...
```

and:

```
> last update  
[info] Updating {file:/.../demo/}example...  
[debug] post 1.3 ivy file: using exact as default matcher  
[debug] :: resolving dependencies :: example#example_2.9.2;0.1-SNAPSHOT  
[debug]   confs: [compile, runtime, test, provided, optional, compile-internal, runtime-internal]  
[debug]   validate = true  
[debug]   refresh = false  
[debug] resolving dependencies for configuration 'compile'  
...
```

Show warnings from the previous compilation

The Scala compiler does not print the full details of warnings by default. Compiling code that uses the deprecated `error` method from `Predef` might generate the following output:

```
> compile
[info] Compiling 1 Scala source to <...>/classes...
[warn] there were 1 deprecation warnings; re-run with -deprecation for details
[warn] one warning found
```

The details aren't provided, so it is necessary to add `-deprecation` to the options passed to the compiler (`scalacOptions`) and recompile. An alternative when using Scala 2.10 and later is to run `printWarnings`. This task will display all warnings from the previous compilation. For example,

```
> printWarnings
[warn] A.scala:2: method error in object Predef is deprecated: Use sys.error(message) instead
[warn]   def x = error("Failed.")
[warn]             ^
```

Change the logging level globally

The quickest way to change logging levels is by using the `error`, `warn`, `info`, or `debug` commands. These set the default logging level for commands and tasks. For example,

```
> warn
```

will by default show only warnings and errors. To set the logging level before any commands are executed on startup, use `--` before the logging level. For example,

```
$ sbt --warn
> compile
[warn] there were 2 feature warning(s); re-run with -feature for details
[warn] one warning found
[success] Total time: 4 s, completed ...
>
```

The logging level can be overridden at a finer granularity, which is described next.

Change the logging level for a specific task, configuration, or project

The amount of logging is controlled by the `logLevel` setting, which takes values from the `Level` enumeration. Valid values are `Error`, `Warn`, `Info`, and `Debug` in order of increasing verbosity. The logging level may be configured globally, as described in the previous section, or it may be applied to a specific project, configuration, or task. For example, to change the logging level for compilation to only show warnings and errors:

```
> set logLevel in compile := Level.Warn
```

To enable debug logging for all tasks in the current project,

```
> set logLevel := Level.Warn
```

A common scenario is that after running a task, you notice that you need more information than was shown by default. A `logLevel` based solution typically requires changing the logging level and running a task again. However, there are two cases where this is unnecessary. First, warnings from a previous compilation may be displayed using `printWarnings` for the main sources or `test:printWarnings` for test sources. Second, output from the previous execution is available either for a single task or for in its entirety. See the section on `printWarnings` and the sections on previous output.

Configure printing of stack traces

By default, sbt hides the stack trace of most exceptions thrown during execution. It prints a message that indicates how to display the exception. However, you may want to show more of stack traces by default.

The setting to configure is `traceLevel`, which is a setting with an `Int` value. When `traceLevel` is set to a negative value, no stack traces are shown. When it is zero, the stack trace is displayed up to the first sbt stack frame. When positive, the stack trace is shown up to that many stack frames.

For example, the following configures sbt to show stack traces up to the first sbt frame:

```
> set every traceLevel := 0
```

The `every` part means to override the setting in all scopes. To change the trace printing behavior for a single project, configuration, or task, scope `traceLevel` appropriately:

```
> set traceLevel in Test := 5
> set traceLevel in update := 0
> set traceLevel in ThisProject := -1
```

Print the output of tests immediately instead of buffering

By default, sbt buffers the logging output of a test until the whole class finishes. This is so that output does not get mixed up when executing in parallel. To disable buffering, set the `logBuffered` setting to false:

```
logBuffered := false
```

Add a custom logger

The setting `extraLoggers` can be used to add custom loggers. A custom logger should implement `[AbstractLogger]`. `extraLoggers` is a function `ScopedKey[_] => Seq[AbstractLogger]`. This means that it can provide different logging based on the task that requests the logger.

```
extraLoggers := {  
  val currentFunction = extraLoggers.value  
  (key: ScopedKey[_]) => {  
    myCustomLogger(key) +: currentFunction(key)  
  }  
}
```

Here, we take the current function `currentFunction` for the setting and provide a new function. The new function prepends our custom logger to the ones provided by the old function.

Log messages in a task

The special task `streams` provides per-task logging and I/O via a `Streams` instance. To log, a task uses the `log` member from the `streams` task:

```
myTask := {  
  val log = streams.value.log  
  log.warn("A warning.")  
}
```

Project metadata

Set the project name

A project should define `name` and `version`. These will be used in various parts of the build, such as the names of generated artifacts. Projects that are published to a repository should also override `organization`.

```
name := "Your project name"
```

For published projects, this name is normalized to be suitable for use as an artifact name and dependency ID. This normalized name is stored in `normalizedName`.

Set the project version

```
version := "1.0"
```

Set the project organization

```
organization := "org.example"
```

By convention, this is a reverse domain name that you own, typically one specific to your project. It is used as a namespace for projects.

A full/formal name can be defined in the `organizationName` setting. This is used in the generated `pom.xml`. If the organization has a web site, it may be set in the `organizationHomepage` setting. For example:

```
organizationName := "Example, Inc."
```

```
organizationHomepage := Some(url("http://example.org"))
```

Set the project's homepage and other metadata

```
homepage := Some(url("http://scala-sbt.org"))
```

```
startYear := Some(2008)
```

```
description := "A build tool for Scala."
```

```
licenses += "GPLv2" -> url("http://www.gnu.org/licenses/gpl-2.0.html")
```

Configure packaging

Use the packaged jar on classpaths instead of class directory

By default, a project exports a directory containing its resources and compiled class files. Set `exportJars` to true to export the packaged jar instead. For example,

```
exportJars := true
```

The jar will be used by `run`, `test`, `console`, and other tasks that use the full classpath.

Add manifest attributes

By default, sbt constructs a manifest for the binary package from settings such as `organization` and `mainClass`. Additional attributes may be added to the `packageOptions` setting scoped by the configuration and package task.

Main attributes may be added with `Package.ManifestAttributes`. There are two variants of this method, once that accepts repeated arguments that map an attribute of type `java.util.jar.Attributes.Name` to a String value and other that maps attribute names (type String) to the String value.

For example,

```
packageOptions in (Compile, packageBin) +=  
  Package.ManifestAttributes( java.util.jar.Attributes.Name.SEALED -> "true" )
```

Other attributes may be added with `Package.JarManifest`.

```
packageOptions in (Compile, packageBin) += {  
  import java.util.jar.{Attributes, Manifest}  
  val manifest = new Manifest  
  manifest.getAttributes("foo/bar/").put(Attributes.Name.SEALED, "false")  
  Package.JarManifest( manifest )  
}
```

Or, to read the manifest from a file:

```
packageOptions in (Compile, packageBin) += {  
  val manifest = Using.fileInputStream( in => new java.util.jar.Manifest(in) )  
  Package.JarManifest( manifest )  
}
```

Change the file name of a package

The `artifactName` setting controls the name of generated packages. See the [Artifacts](#) page for details.

Modify the contents of the package

The contents of a package are defined by the `mappings` task, of type `Seq[(File,String)]`. The `mappings` task is a sequence of mappings from a file to include in the package to the path in the package. See [Mapping Files](#) for convenience functions for generating these mappings. For example, to add the file `in/example.txt` to the main binary jar with the path “`out/example.txt`”,

```
mappings in (Compile, packageBin) += {  
  (baseDirectory.value / "in" / "example.txt") -> "out/example.txt"  
}
```

Note that `mappings` is scoped by the configuration and the specific package task. For example, the mappings for the test source package are defined by the `mappings in (Test, packageSrc)` task.

Running commands

Pass arguments to a command or task in batch mode

`sbt` interprets each command line argument provided to it as a command together with the command's arguments. Therefore, to run a command that takes arguments in batch mode, quote the command and its arguments. For example,

```
$ sbt 'project X' clean '~ compile'
```

Provide multiple commands to run consecutively

Multiple commands can be scheduled at once by prefixing each command with a semicolon. This is useful for specifying multiple commands where a single command string is accepted. For example, the syntax for triggered execution is `~ <command>`. To have more than one command run for each triggering, use semicolons. For example, the following runs `clean` and then `compile` each time a source file changes:

```
> ~ ;clean;compile
```

Read commands from a file

The `<` command reads commands from the files provided to it as arguments. Run `help <` at the `sbt` prompt for details.

Define an alias for a command or task

The `alias` command defines, removes, and displays aliases for commands. Run `help alias` at the sbt prompt for details.

Example usage:

```
> alias a=about
> alias
    a = about
> a
[info] This is sbt ...
> alias a=
> alias
> a
[error] Not a valid command: a ...
```

Quickly evaluate a Scala expression

The `eval` command compiles and runs the Scala expression passed to it as an argument. The result is printed along with its type. For example,

```
> eval 2+2
4: Int
```

Variables defined by an `eval` are not visible to subsequent `evals`, although changes to system properties persist and affect the JVM that is running sbt. Use the Scala REPL (`console` and related commands) for full support for evaluating Scala code interactively.

Configure and use Scala

Set the Scala version used for building the project

The `scalaVersion` configures the version of Scala used for compilation. By default, sbt also adds a dependency on the Scala library with this version. See the next section for how to disable this automatic dependency. If the Scala version is not specified, the version sbt was built against is used. It is recommended to explicitly specify the version of Scala.

For example, to set the Scala version to “2.11.1”,

```
scalaVersion := "2.11.1"
```

Disable the automatic dependency on the Scala library

sbt adds a dependency on the Scala standard library by default. To disable this behavior, set the `autoScalaLibrary` setting to false.

```
autoScalaLibrary := false
```

Temporarily switch to a different Scala version

To set the Scala version in all scopes to a specific value, use the `++` command. For example, to temporarily use Scala 2.10.4, run:

```
> ++ 2.10.4
```

Use a local Scala installation for building a project

Defining the `scalaHome` setting with the path to the Scala home directory will use that Scala installation. sbt still requires `scalaVersion` to be set when a local Scala version is used. For example,

```
scalaVersion := "2.10.0-local"
```

```
scalaHome := Some(file("/path/to/scala/home/"))
```

Build a project against multiple Scala versions

See [cross building](#).

Enter the Scala REPL with a project's dependencies on the classpath, but not the compiled project classes

The `consoleQuick` action retrieves dependencies and puts them on the classpath of the Scala REPL. The project's sources are not compiled, but sources of any source dependencies are compiled. To enter the REPL with test dependencies on the classpath but without compiling test sources, run `test:consoleQuick`. This will force compilation of main sources.

Enter the Scala REPL with a project's dependencies and compiled code on the classpath

The `console` action retrieves dependencies and compiles sources and puts them on the classpath of the Scala REPL. To enter the REPL with test dependencies and compiled test sources on the classpath, run `test:console`.

Enter the Scala REPL with plugins and the build definition on the classpath

```
> consoleProject
```

For details, see the [consoleProject](#) page.

Define the initial commands evaluated when entering the Scala REPL

Set `initialCommands` in `console` to set the initial statements to evaluate when `console` and `consoleQuick` are run. To configure `consoleQuick` separately, use `initialCommands` in `consoleQuick`. For example,

```
initialCommands in console := """println("Hello from console")"""

initialCommands in consoleQuick := """println("Hello from consoleQuick")"""
```

The `consoleProject` command is configured separately by `initialCommands` in `consoleProject`. It does not use the value from `initialCommands` in `console` by default. For example,

```
initialCommands in consoleProject := """println("Hello from consoleProject")"""
```

Define the commands evaluated when exiting the Scala REPL

Set `cleanupCommands` in `console` to set the statements to evaluate after exiting the Scala REPL started by `console` and `consoleQuick`. To configure `consoleQuick` separately, use `cleanupCommands` in `consoleQuick`. For example,

```
cleanupCommands in console := """println("Bye from console")"""

cleanupCommands in consoleQuick := """println("Bye from consoleQuick")"""
```

The `consoleProject` command is configured separately by `cleanupCommands` in `consoleProject`. It does not use the value from `cleanupCommands` in `console` by default. For example,

```
cleanupCommands in consoleProject := """println("Bye from consoleProject")"""
```

Use the Scala REPL from project code

sbt runs tests in the same JVM as sbt itself and Scala classes are not in the same class loader as the application classes. This is also the case in `console` and when `run` is not forked. Therefore, when using the Scala interpreter, it is important to set it up properly to avoid an error message like:

```
Failed to initialize compiler: class scala.runtime.VolatileBooleanRef not found.
** Note that as of 2.8 scala does not assume use of the java classpath.
** For the old behavior pass -usejavacp to scala, or if using a Settings
** object programmatically, settings.usejavacp.value = true.
```

The key is to initialize the Settings for the interpreter using *embeddedDefaults*. For example:

```
val settings = new Settings
settings.embeddedDefaults[MyType]
val interpreter = new Interpreter(settings, ...)
```

Here, `MyType` is a representative class that should be included on the interpreter's classpath and in its application class loader. For more background, see the [original proposal](#) that resulted in *embeddedDefaults* being added.

Similarly, use a representative class as the type argument when using the *break* and *breakIf* methods of *ILoop*, as in the following example:

```
def x(a: Int, b: Int) = {
  import scala.tools.nsc.interpreter.ILoop
  ILoop.breakIf[MyType](a != b, "a" -> a, "b" -> b )
}
```

Generate API documentation

Select javadoc or scaladoc

sbt will run `javadoc` if there are only Java sources in the project. If there are any Scala sources, sbt will run `scaladoc`. (This situation results from `scaladoc` not processing Javadoc comments in Java sources nor linking to Javadoc.)

Set the options used for generating scaladoc independently of compilation

Scope `scalacOptions` to the `doc` task to configure `scaladoc`. Use `:=` to definitively set the options without appending to the options for `compile`. Scope to `Compile` for main sources or to `Test` for test sources. For example,

```
scalacOptions in (Compile,doc) := Seq("-groups", "-implicit")
```

Add options for scaladoc to the compilation options

Scope `scalacOptions` to the `doc` task to configure `scaladoc`. Use `+=` or `++=` to append options to the base options. To append a single option, use `+=`. To append a `Seq[String]`, use `++=`. Scope to `Compile` for main sources or to `Test` for test sources. For example,

```
scalacOptions in (Compile,doc) ++= Seq("-groups", "-implicit")
```

Set the options used for generating javadoc independently of compilation

Scope `javacOptions` to the `doc` task to configure `javadoc`. Use `:=` to definitively set the options without appending to the options for `compile`. Scope to `Compile` for main sources or to `Test` for test sources.

Add options for javadoc to the compilation options

Scope `javacOptions` to the `doc` task to configure `javadoc`. Use `+=` or `++=` to append options to the base options. To append a single option, use `+=`. To append a `Seq[String]`, use `++=`. Scope to `Compile` for main sources or to `Test` for test sources. For example,

```
javacOptions in (Compile,doc) ++= Seq("-notimestamp", "-linksource")
```

Enable automatic linking to the external Scaladoc of managed dependencies

Set `autoAPIMappings := true` for sbt to tell `scaladoc` where it can find the API documentation for managed dependencies. This requires that dependencies have this information in its metadata and you are using `scaladoc` for Scala 2.10.2 or later.

Enable manual linking to the external Scaladoc of managed dependencies

Add mappings of type `(File, URL)` to `apiMappings` to manually tell `scaladoc` where it can find the API documentation for dependencies. (This requires `scaladoc` for Scala 2.10.2 or later.) These mappings are used in addition to

`autoAPIMappings`, so this manual configuration is typically done for unmanaged dependencies. The `File` key is the location of the dependency as passed to the classpath. The `URL` value is the base URL of the API documentation for the dependency. For example,

```
apiMappings += (  
  (unmanagedBase.value / "a-library.jar") ->  
    url("http://example.org/api/")  
)
```

Define the location of API documentation for a library

Set `apiURL` to define the base URL for the Scaladocs for your library. This will enable clients of your library to automatically link against the API documentation using `autoAPIMappings`. (This only works for Scala 2.10.2 and later.) For example,

```
apiURL := Some(url("http://example.org/api/"))
```

This information will get included in a property of the published `pom.xml`, where it can be automatically consumed by sbt.

Triggered execution

Run a command when sources change

You can make a command run when certain files change by prefixing the command with `~`. Monitoring is terminated when **enter** is pressed. This triggered execution is configured by the `watch` setting, but typically the basic settings `watchSources` and `pollInterval` are modified as described in later sections.

The original use-case for triggered execution was continuous compilation:

```
> ~ test:compile
```

```
> ~ compile
```

You can use the triggered execution feature to run any command or task, however. The following will poll for changes to your source code (main or test) and run `testOnly` for the specified test.

```
> ~ testOnly example.TestA
```

Run multiple commands when sources change

The command passed to `~` may be any command string, so multiple commands may be run by separating them with a semicolon. For example,

```
> ~ ;a ;b
```

This runs `a` and then `b` when sources change.

Configure the sources that are checked for changes

- `watchSources` defines the files for a single project that are monitored for changes. By default, a project watches resources and Scala and Java sources.
- `watchTransitiveSources` then combines the `watchSources` for the current project and all execution and classpath dependencies (see [.scala build definition](#) for details on inter-project dependencies).

To add the file `demo/example.txt` to the files to watch,

```
watchSources += baseDirectory.value / "demo" / "examples.txt"
```

Set the time interval between checks for changes to sources

`pollInterval` selects the interval between polling for changes in milliseconds. The default value is 500 ms. To change it to 1 s,

```
pollInterval := 1000 // in ms
```

Examples

This section of the documentation has example sbt build definitions and code. Contributions are welcome!

You may want to read the [Getting Started Guide](#) as a foundation for understanding the examples.

.sbt build examples

Listed here are some examples of settings (each setting is independent). See [.sbt build definition](#) for details.

Note that blank lines are used to separate individual settings. Avoid using blank lines within a single multiline expression. As explained in [.sbt build definition](#), each setting is otherwise a normal Scala expression with expected type `sbt.SettingDefinition`.

```
// set the name of the project
name := "My Project"

version := "1.0"

organization := "org.myproject"

// set the Scala version used for the project
scalaVersion := "2.9.0-SNAPSHOT"

// set the main Scala source directory to be <base>/src
scalaSource in Compile := baseDirectory.value / "src"

// set the Scala test source directory to be <base>/test
scalaSource in Test := baseDirectory.value / "test"

// add a test dependency on ScalaCheck
libraryDependencies += "org.scala-tools.testing" %% "scalacheck" % "1.8" % "test"

// add compile dependencies on some dispatch modules
libraryDependencies ++= Seq(
  "net.databinder" %% "dispatch-meetup" % "0.7.8",
  "net.databinder" %% "dispatch-twitter" % "0.7.8"
)

// Set a dependency based partially on a val.
{
  val libosmVersion = "2.5.2-RC1"
  libraryDependencies += ("net.sf.travelingsales" % "osmlib" % libosmVersion from
    "http://downloads.sourceforge.net/project/travelingsales/libosm/"+libosmVersion+"/libosm")
}

// reduce the maximum number of errors shown by the Scala compiler
maxErrors := 20

// increase the time between polling for file changes when using continuous execution
```



```

pollInterval := 1000

// append several options to the list of options passed to the Java compiler
javacOptions ++= Seq("-source", "1.5", "-target", "1.5")

// append -deprecation to the options passed to the Scala compiler
scalacOptions += "-deprecation"

// define the statements initially evaluated when entering 'console', 'consoleQuick', or 'c
initialCommands := """
  import System.{currentTimeMillis => now}
  def time[T](f: => T): T = {
    val start = now
    try { f } finally { println("Elapsed: " + (now - start)/1000.0 + " s") }
  }
"""

// set the initial commands when entering 'console' or 'consoleQuick', but not 'consoleProj
initialCommands in console := "import myproject._"

// set the main class for packaging the main jar
// 'run' will still auto-detect and prompt
// change Compile to Test to set it for the test jar
mainClass in (Compile, packageBin) := Some("myproject.MyMain")

// set the main class for the main 'run' task
// change Compile to Test to set it for 'test:run'
mainClass in (Compile, run) := Some("myproject.MyMain")

// add <base>/input to the files that '~' triggers on
watchSources += baseDirectory.value / "input"

// add a maven-style repository
resolvers += "name" at "url"

// add a sequence of maven-style repositories
resolvers ++= Seq("name" at "url")

// define the repository to publish to
publishTo := Some("name" at "url")

// set Ivy logging to be at the highest level
ivyLoggingLevel := UpdateLogging.Full

// disable updating dynamic revisions (including -SNAPSHOT versions)
offline := true

```

```

// set the prompt (for this build) to include the project id.
shellPrompt in ThisBuild := { state => Project.extract(state).currentRef.project + "> " }

// set the prompt (for the current project) to include the username
shellPrompt := { state => System.getProperty("user.name") + "> " }

// disable printing timing information, but still print [success]
showTiming := false

// disable printing a message indicating the success or failure of running a task
showSuccess := false

// change the format used for printing task completion time
timingFormat := {
  import java.text.DateFormat
  DateFormat.getDateInstance(DateFormat.SHORT, DateFormat.SHORT)
}

// disable using the Scala version in output paths and artifacts
crossPaths := false

// fork a new JVM for 'run' and 'test:run'
fork := true

// fork a new JVM for 'test:run', but not 'run'
fork in Test := true

// add a JVM option to use when forking a JVM for 'run'
javaOptions += "-Xmx2G"

// only use a single thread for building
parallelExecution := false

// Execute tests in the current project serially
// Tests from other projects may still run concurrently.
parallelExecution in Test := false

// set the location of the JDK to use for compiling Java code.
// if 'fork' is true, this is used for 'run' as well
javaHome := Some(file("/usr/lib/jvm/sun-jdk-1.6"))

// Use Scala from a directory on the filesystem instead of retrieving from a repository
scalaHome := Some(file("/home/user/scala/trunk/"))

// don't aggregate clean (See FullConfiguration for aggregation details)

```

```

aggregate in clean := false

// only show warnings and errors on the screen for compilations.
// this applies to both test:compile and compile and is Info by default
logLevel in compile := Level.Warn

// only show warnings and errors on the screen for all tasks (the default is Info)
// individual tasks can then be more verbose using the previous setting
logLevel := Level.Warn

// only store messages at info and above (the default is Debug)
// this is the logging level for replaying logging with 'last'
persistLogLevel := Level.Debug

// only show 10 lines of stack traces
traceLevel := 10

// only show stack traces up to the first sbt stack frame
traceLevel := 0

// add SWT to the unmanaged classpath
unmanagedJars in Compile += Attributed.blank(file("/usr/share/java/swt.jar"))

// publish test jar, sources, and docs
publishArtifact in Test := true

// disable publishing of main docs
publishArtifact in (Compile, packageDoc) := false

// change the classifier for the docs artifact
artifactClassifier in packageDoc := Some("doc")

// Copy all managed dependencies to <build-root>/lib_managed/
// This is essentially a project-local cache and is different
// from the lib_managed/ in sbt 0.7.x. There is only one
// lib_managed/ in the build root (not per-project).
retrieveManaged := true

/* Specify a file containing credentials for publishing. The format is:
realm=Sonatype Nexus Repository Manager
host=nexus.scala-tools.org
user=admin
password=admin123
*/
credentials += Credentials(Path.userHome / ".ivy2" / ".credentials")

```

```

// Directly specify credentials for publishing.
credentials += Credentials("Sonatype Nexus Repository Manager", "nexus.scala-tools.org", "ac

// Exclude transitive dependencies, e.g., include log4j without including logging via jdmk,
libraryDependencies +=
  "log4j" % "log4j" % "1.2.15" excludeAll(
    ExclusionRule(organization = "com.sun.jdmk"),
    ExclusionRule(organization = "com.sun.jmx"),
    ExclusionRule(organization = "javax.jms")
  )

```

.scala build example

.scala builds are written in Scala, so this example would be placed as project/Build.scala, not build.sbt. The build can be split into multiple files.

```

import sbt._
import Keys._

object BuildSettings {
  val buildOrganization = "odp"
  val buildVersion      = "2.0.29"
  val buildScalaVersion = "2.9.0-1"

  val buildSettings = Seq (
    organization := buildOrganization,
    version      := buildVersion,
    scalaVersion := buildScalaVersion,
    shellPrompt  := ShellPrompt.buildShellPrompt
  )
}

// Shell prompt which show the current project,
// git branch and build version
object ShellPrompt {
  object devnull extends ProcessLogger {
    def info (s: => String) {}
    def error (s: => String) { }
    def buffer[T] (f: => T): T = f
  }

  def currBranch = (
    ("git status -sb" lines_! devnull headOption)
      .getOrElse "-" stripPrefix "## "
  )
}

```

```

val buildShellPrompt = {
  (state: State) => {
    val currProject = Project.extract (state).currentProject.id
    "%s:%s:%s> ".format (
      currProject, currBranch, BuildSettings.buildVersion
    )
  }
}

object Resolvers {
  val sunrepo = "Sun Maven2 Repo" at "http://download.java.net/maven/2"
  val sunrepoGF = "Sun GF Maven2 Repo" at "http://download.java.net/maven/glassfish"
  val oraclerepo = "Oracle Maven2 Repo" at "http://download.oracle.com/maven"

  val oracleResolvers = Seq (sunrepo, sunrepoGF, oraclerepo)
}

object Dependencies {
  val logbackVer = "0.9.16"
  val grizzlyVer = "1.9.19"

  val logbackcore = "ch.qos.logback" % "logback-core" % logbackVer
  val logbackclassic = "ch.qos.logback" % "logback-classic" % logbackVer

  val jacksonjson = "org.codehaus.jackson" % "jackson-core-lgpl" % "1.7.2"

  val grizzlyframework = "com.sun.grizzly" % "grizzly-framework" % grizzlyVer
  val grizzlyhttp = "com.sun.grizzly" % "grizzly-http" % grizzlyVer
  val grizzlyrcm = "com.sun.grizzly" % "grizzly-rcm" % grizzlyVer
  val grizzlyutils = "com.sun.grizzly" % "grizzly-utils" % grizzlyVer
  val grizzlyportunif = "com.sun.grizzly" % "grizzly-portunif" % grizzlyVer

  val sleepycat = "com.sleepycat" % "je" % "4.0.92"

  val apachenet = "commons-net" % "commons-net" % "2.0"
  val apachecodec = "commons-codec" % "commons-codec" % "1.4"

  val scalatest = "org.scalatest" % "scalatest_2.9.0" % "1.4.1" % "test"
}

object CDAP2Build extends Build {
  import Resolvers._
  import Dependencies._
  import BuildSettings._

```

```

// Sub-project specific dependencies
val commonDeps = Seq (
  logbackcore,
  logbackclassic,
  jacksonjson,
  scalatest
)

val serverDeps = Seq (
  grizzlyframework,
  grizzlyhttp,
  grizzlyrcm,
  grizzlyutils,
  grizzlyportunif,
  sleepycat,
  scalatest
)

val pricingDeps = Seq (apachenet, apachecodec, scalatest)

lazy val cdap2 = Project (
  "cdap2",
  file ("."),
  settings = buildSettings
) aggregate (common, server, compact, pricing, pricing_service)

lazy val common = Project (
  "common",
  file ("cdap2-common"),
  settings = buildSettings ++ Seq (libraryDependencies += commonDeps)
)

lazy val server = Project (
  "server",
  file ("cdap2-server"),
  settings = buildSettings ++ Seq (resolvers := oracleResolvers,
    libraryDependencies += serverDeps)
) dependsOn (common)

lazy val pricing = Project (
  "pricing",
  file ("cdap2-pricing"),
  settings = buildSettings ++ Seq (libraryDependencies += pricingDeps)
) dependsOn (common, compact, server)

```

```

lazy val pricing_service = Project (
  "pricing-service",
  file ("cdap2-pricing-service"),
  settings = buildSettings
) dependsOn (pricing, server)

lazy val compact = Project (
  "compact",
  file ("compact-hashmap"),
  settings = buildSettings
)
}

```

External Builds

- [Mojolly Backchat Build](#)
- [Scalaz Build](#)
- Source Code Generation
- Generates Scaladoc and Scala X-Ray HTML Sources, with a unified view of source from all sub-projects
- Builds an archive with the artifacts from all modules
- “Roll your own” approach to appending the Scala version to the module id of dependencies to allow using snapshot releases of Scala.

Advanced configurations example

This is an example [.scala build definition](#) that demonstrates using Ivy configurations to group dependencies.

The `utils` module provides utilities for other modules. It uses Ivy configurations to group dependencies so that a dependent project doesn’t have to pull in all dependencies if it only uses a subset of functionality. This can be an alternative to having multiple utilities modules (and consequently, multiple utilities jars).

In this example, consider a `utils` project that provides utilities related to both Scalate and Saxon. It therefore needs both Scalate and Saxon on the compilation classpath and a project that uses all of the functionality of ‘utils’ will need these dependencies as well. However, project `a` only needs the utilities related to Scalate, so it doesn’t need Saxon. By depending only on the `scalate` configuration of `utils`, it only gets the Scalate-related dependencies.

```

import sbt._
import Keys._

object B extends Build {

```

```

***** Projects *****

// An example project that only uses the Scalate utilities.
lazy val a = Project("a", file("a")) dependsOn(utls % "compile->scalate")

// An example project that uses the Scalate and Saxon utilities.
// For the configurations defined here, this is equivalent to doing dependsOn(utls),
// but if there were more configurations, it would select only the Scalate and Saxon
// dependencies.
lazy val b = Project("b", file("b")) dependsOn(utls % "compile->scalate,saxon")

// Defines the utilities project
lazy val utls = Project("utls", file("utls")) settings(utlsSettings : _*)

def utlsSettings: Seq[Setting[_]] =
    // Add the src/common/scala/ compilation configuration.
    inConfig(Common)(Defaults.configSettings) ++
    // Publish the common artifact
    addArtifact(artifact in (Common, packageBin), packageBin in Common) ++ Seq(
        // We want our Common sources to have access to all of the dependencies on the classpath
        // for compile and test, but when depended on, it should only require dependencies
        classpathConfiguration in Common := CustomCompile,
        // Modify the default Ivy configurations.
        // 'overrideConfigs' ensures that Compile is replaced by CustomCompile
        ivyConfigurations := overrideConfigs(Scalate, Saxon, Common, CustomCompile)(ivyConfigurations),
        // Put all dependencies without an explicit configuration into Common (optional)
        defaultConfiguration := Some(Common),
        // Declare dependencies in the appropriate configurations
        libraryDependencies ++= Seq(
            "org.fusesource.scalate" % "scalate-core" % "1.5.0" % "scalate",
            "org.squeryl" %% "squeryl" % "0.9.4" % "scalate",
            "net.sf.saxon" % "saxon" % "8.7" % "saxon"
        )
    )

***** Configurations *****

lazy val Scalate = config("scalate") extend(Common) describedAs("Dependencies for using Scalate")
lazy val Common = config("common") describedAs("Dependencies required in all configurations")
lazy val Saxon = config("saxon") extend(Common) describedAs("Dependencies for using Saxon")

// Define a customized compile configuration that includes
// dependencies defined in our other custom configurations
lazy val CustomCompile = config("compile") extend(Saxon, Common, Scalate)
}

```


Advanced command example

This is an advanced example showing some of the power of the new settings system. It shows how to temporarily modify all declared dependencies in the build, regardless of where they are defined. It directly operates on the final `Seq[Setting[_]]` produced from every setting involved in the build.

The modifications are applied by running *canonicalize*. A *reload* or using *set* reverts the modifications, requiring *canonicalize* to be run again.

This particular example shows how to transform all declared dependencies on ScalaCheck to use version 1.8. As an exercise, you might try transforming other dependencies, the repositories used, or the scalac options used. It is possible to add or remove settings as well.

This kind of transformation is possible directly on the settings of Project, but it would not include settings automatically added from plugins or build.sbt files. What this example shows is doing it unconditionally on all settings in all projects in all builds, including external builds.

```
import sbt._
import Keys._

object Canon extends Plugin {
  // Registers the canonicalize command in every project
  override def settings = Seq(commands += canonicalize)

  // Define the command. This takes the existing settings (including any session settings)
  // and applies 'f' to each Setting[_]
  def canonicalize = Command.command("canonicalize") { (state: State) =>
    val extracted = Project.extract(state)
    import extracted._
    val transformed = session.mergeSettings map ( s => f(s) )
    val newStructure = Load.reapply(transformed, structure)
    Project.setProject(session, newStructure, state)
  }

  // Transforms a Setting[_].
  def f(s: Setting[_]): Setting[_] = s.key.key match {
    // transform all settings that modify libraryDependencies
    case Keys.libraryDependencies.key =>
      // hey scalac. T == Seq[ModuleID]
      s.asInstanceOf[Setting[Seq[ModuleID]]].mapInit(mapLibraryDependencies)
      // preserve other settings
    case _ => s
  }

  // This must be idempotent because it gets applied after every transformation.
}
```

```

// That is, if the user does:
// libraryDependencies += a
// libraryDependencies += b
// then this method will be called for Seq(a) and Seq(a,b)
def mapLibraryDependencies(key: ScopedKey[Seq[ModuleID]], value: Seq[ModuleID]): Seq[ModuleID] =
  value map mapSingle

// This is the fundamental transformation.
// Here we map all declared ScalaCheck dependencies to be version 1.8
def mapSingle(module: ModuleID): ModuleID =
  if(module.name == "scalacheck") module.copy(revision = "1.8")
  else module
}

```

Frequently Asked Questions

Project Information

How do I get help? Please use [Stack Overflow](#) for questions. Use the [sbt-dev mailing list](#) for comments and discussions about sbt development.

- Please state the problem or question clearly and provide enough context. Code examples and build transcripts are often useful when appropriately edited.
- Providing small, reproducible examples is a good way to get help quickly.
- Include relevant information such as the version of sbt and Scala being used.

How do I report a bug? Please use the [issue tracker](#) to report confirmed bugs. Do not use it to ask questions or to determine if something is a bug.

How can I help?

- Answer questions on [Stack Overflow](#).
- Fix mistakes that you notice on the wiki.
- Make [bug reports](#) that are clear and reproducible.
- Discuss development on the [sbt-dev mailing list](#).
- Fix issues that affect you. [Fork, fix, and submit a pull request](#).
- Implement features that are important to you. See [Contributing to sbt](#) for more details.

Usage

My last command didn't work but I can't see an explanation. Why? sbt 0.13.5 by default suppresses most stack traces and debugging information. It has the nice side effect of giving you less noise on screen, but as a newcomer it can leave you lost for explanation. To see the previous output of a command at a higher verbosity, type `last <task>` where `<task>` is the task that failed or that you want to view detailed output for. For example, if you find that your `update` fails to load all the dependencies as you expect you can enter:

```
> last update
```

and it will display the full output from the last run of the `update` command.

How do I disable ansi codes in the output? Sometimes sbt doesn't detect that ansi codes aren't supported and you get output that looks like:

```
[0m[ [0minfo [0m] [0mSet current project to root
```

or ansi codes are supported but you want to disable colored output. To completely disable ansi codes, set the `sbt.log.format` system property to `false`. For example,

How can I start a Scala interpreter (REPL) with sbt project configuration (dependencies, etc.)? You may run `sbt console`.

Build definitions

What are the `:=`, `+=`, and `++=` methods? These are methods on keys used to construct a `Setting` or a `Task`. The Getting Started Guide covers all these methods, see [.sbt build definition](#) and [more kinds of setting](#) for example.

What is the `%` method? It's used to create a `ModuleID` from strings, when specifying managed dependencies. Read the Getting Started Guide about [library dependencies](#).

What is `ModuleID`, `Project`, ...? To figure out an unknown type or method, have a look at the [Getting Started Guide](#) if you have not. Also try the [index](#) of commonly used methods, values, and types, the [API Documentation](#) and the [hyperlinked sources](#).

How do I add files to a jar package? The files included in an artifact are configured by default by a task `mappings` that is scoped by the relevant package task. The `mappings` task returns a sequence `Seq[(File,String)]` of mappings from the file to include to the path within the jar. See [mapping files](#) for details on creating these mappings.

For example, to add generated sources to the packaged source artifact:

```
mappings in (Compile, packageSrc) += {
  import Path.{flat, relativeTo}
  val base = (sourceManaged in Compile).value
  val srcs = (managedSources in Compile).value
  srcs x (relativeTo(base) | flat)
}
```

This takes sources from the `managedSources` task and relativizes them against the `managedSource` base directory, falling back to a flattened mapping. If a source generation task doesn't write the sources to the `managedSource` directory, the mapping function would have to be adjusted to try relativizing against additional directories or something more appropriate for the generator.

How can I generate source code or resources? See [Generating Files](#).

How can a task avoid redoing work if the input files are unchanged?

There is basic support for only doing work when input files have changed or when the outputs haven't been generated yet. This support is primitive and subject to change.

The relevant methods are two overloaded methods called [FileFunction.cached](#). Each requires a directory in which to store cached data. Sample usage is:

```
// define a task that takes some inputs
// and generates files in an output directory
myTask := {
  // wraps a function taskImpl in an uptodate check
  // taskImpl takes the input files, the output directory,
  // generates the output files and returns the set of generated files
  val cachedFun = FileFunction.cached(cacheDirectory.value / "my-task") { (in: Set[File]) =>
    taskImpl(in, target.value) : Set[File]
  }
  // Applies the cached function to the inputs files
  cachedFun(inputs.value)
}
```

There are two additional arguments for the first parameter list that allow the file tracking style to be explicitly specified. By default, the input tracking style is `FileInfo.lastModified`, based on a file's last modified time, and the output tracking style is `FileInfo.exists`, based only on whether the file exists. The other available style is `FileInfo.hash`, which tracks a file based on a hash of its contents. See the [FileInfo API](#) for details.

A more advanced version of `FileFunction.cached` passes a data structure of type [ChangeReport](#) describing the changes to input and output files since the last evaluation. This version of `cached` also expects the set of files generated as output to be the result of the evaluated function.

Extending sbt

How can I add a new configuration? The following example demonstrates adding a new set of compilation settings and tasks to a new configuration called `samples`. The sources for this configuration go in `src/samples/scala/`. Unspecified settings delegate to those defined for the `compile` configuration. For example, if `scalacOptions` are not overridden for `samples`, the options for the main sources are used.

Options specific to `samples` may be declared like:

```
scalacOptions in Samples += "-deprecation"
```

This uses the main options as base options because of `+=`. Use `:=` to ignore the main options:

```
scalacOptions in Samples := "-deprecation" :: Nil
```

The example adds all of the usual compilation related settings and tasks to `samples`:

```
samples:run
samples:runMain
samples:compile
samples:console
samples:consoleQuick
samples:scalacOptions
samples:fullClasspath
samples:package
samples:packageSrc
...
```

How do I add a test configuration? See the [Additional test configurations](#) section of [Testing](#).

How can I create a custom run task, in addition to run? This answer is extracted from a [mailing list discussion](#).

Read the Getting Started Guide up to [custom settings](#) for background.

A basic run task is created by:

```
lazy val myRunTask = taskKey[Unit]("A custom run task.")

// this can go either in a `build.sbt` or the settings member
// of a Project in a full configuration
fullRunTask(myRunTask, Test, "foo.Foo", "arg1", "arg2")
```

If you want to be able to supply arguments on the command line, replace `TaskKey` with `InputKey` and `fullRunTask` with `fullRunInputTask`. The `Test` part can be replaced with another configuration, such as `Compile`, to use that configuration's classpath.

This run task can be configured individually by specifying the task key in the scope. For example:

```
fork in myRunTask := true

javaOptions in myRunTask += "-Xmx6144m"
```

How should I express a dependency on an outside tool such as `proguard`? Tool dependencies are used to implement a task and are not needed by project source code. These dependencies can be declared in their own configuration and classpaths. These are the steps:

1. Define a new [configuration](#).
2. Declare the tool [dependencies](#) in that configuration.
3. Define a classpath that pulls the dependencies from the [Update Report](#) produced by `update`.
4. Use the classpath to implement the task.

As an example, consider a `proguard` task. This task needs the ProGuard jars in order to run the tool. First, define and add the new configuration:

```
val ProguardConfig = config("proguard") hide

ivyConfigurations += ProguardConfig
```

Then,

```
// Add proguard as a dependency in the custom configuration.
// This keeps it separate from project dependencies.
libraryDependencies +=
  "net.sf.proguard" % "proguard" % "4.4" % ProguardConfig.name

// Extract the dependencies from the UpdateReport.
managedClasspath in proguard := {
  // these are the types of artifacts to include
  val artifactTypes: Set[String] = (classpathTypes in proguard).value
  Classpaths.managedJars(proguardConfig, artifactTypes, update.value)
}

// Use the dependencies in a task, typically by putting them
// in a ClassLoader and reflectively calling an appropriate
// method.
proguard := {
  val cp: Seq[File] = (managedClasspath in proguard).value
  // ... do something with , which includes proguard ...
}
```

Defining the intermediate classpath is optional, but it can be useful for debugging or if it needs to be used by multiple tasks. It is also possible to specify artifact types inline. This alternative `proguard` task would look like:

```
proguard := {
  val artifactTypes = Set("jar")
  val cp: Seq[File] =
    Classpaths.managedJars(proguardConfig, artifactTypes, update.value)
  // ... do something with , which includes proguard ...
}
```

How would I change sbt's classpath dynamically? It is possible to register additional jars that will be placed on sbt's classpath (since version 0.10.1). Through [State](#), it is possible to obtain a `xsbti.ComponentProvider`, which manages application components. Components are groups of files in the `~/.sbt/boot/` directory and, in this case, the application is sbt. In addition to the base classpath, components in the “extra” component are included on sbt's classpath.

(Note: the additional components on an application's classpath are declared by the `components` property in the `[main]` section of the launcher configuration file `boot.properties`.)

Because these components are added to the `~/.sbt/boot/` directory and `~/.sbt/boot/` may be read-only, this can fail. In this case, the user has generally intentionally set sbt up this way, so error recovery is not typically necessary (just a short error message explaining the situation.)

Example of dynamic classpath augmentation The following code can be used where a `State => State` is required, such as in the `onLoad` setting (described below) or in a `command`. It adds some files to the “extra” component and reloads sbt if they were not already added. Note that reloading will drop the user’s session state.

```
def augment(extra: Seq[File])(s: State): State = {
  // Get the component provider
  val cs: xsbti.ComponentProvider = s.configuration.provider.components()

  // Adds the files in 'extra' to the "extra" component
  // under an exclusive machine-wide lock.
  // The returned value is 'true' if files were actually copied and 'false'
  // if the target files already exists (based on name only).
  val copied: Boolean = s.locked(cs.lockFile, cs.addToComponent("extra", extra.toArray))

  // If files were copied, reload so that we use the new classpath.
  if(copied) s.reload else s
}
```

How can I take action when the project is loaded or unloaded? The single, global setting `onLoad` is of type `State => State` (see [State and Actions](#)) and is executed once, after all projects are built and loaded. There is a similar hook `onUnload` for when a project is unloaded. Project unloading typically occurs as a result of a `reload` command or a `set` command. Because the `onLoad` and `onUnload` hooks are global, modifying this setting typically involves composing a new function with the previous value. The following example shows the basic structure of defining `onLoad`:

```
// Compose our new function 'f' with the existing transformation.
{
  val f: State => State = ...
  onLoad in Global := {
    val previous = (onLoad in Global).value
    f compose previous
  }
}
```


Example of project load/unload hooks The following example maintains a count of the number of times a project has been loaded and prints that number:

```
{
  // the key for the current count
  val key = AttributeKey[Int]("loadCount")
  // the State transformer
  val f = (s: State) => {
    val previous = s.get key.getOrElse 0
    println("Project load count: " + previous)
    s.put(key, previous + 1)
  }
  onLoad in Global := {
    val previous = (onLoad in Global).value
    f compose previous
  }
}
```

Errors

On project load, “Reference to uninitialized setting” Setting initializers are executed in order. If the initialization of a setting depends on other settings that has not been initialized, sbt will stop loading.

In this example, we try to append a library to `libraryDependencies` before it is initialized with an empty sequence.

```
object MyBuild extends Build {
  val root = Project(id = "root", base = file("."),
    settings = Seq(
      libraryDependencies += "commons-io" % "commons-io" % "1.4" % "test"
    )
  ).disablePlugins(plugins.IvyModule)
}
```

To correct this, include the `IvyModule` plugin settings, which includes `libraryDependencies := Seq()`. So, we just drop the explicit disabling.

```
object MyBuild extends Build {
  val root = Project(id = "root", base = file("."),
    settings = Seq(
      libraryDependencies += "commons-io" % "commons-io" % "1.4" % "test"
    )
  )
}
```

A more subtle variation of this error occurs when using [scoped settings](#).

```
// error: Reference to uninitialized setting
settings = Seq(
  libraryDependencies += "commons-io" % "commons-io" % "1.2" % "test",
  fullClasspath := fullClasspath.value.filterNot(_.data.name.contains("commons-io"))
)
```

This setting varies between the test and compile scopes. The solution is use the scoped setting, both as the input to the initializer, and the setting that we update.

```
fullClasspath in Compile := (fullClasspath in Compile).value.filterNot(_.data.name.contains("commons-io"))
```

Dependency Management

How do I resolve a checksum error? This error occurs when the published checksum, such as a sha1 or md5 hash, differs from the checksum computed for a downloaded artifact, such as a jar or pom.xml. An example of such an error is:

```
[warn] problem while downloading module descriptor:
http://repo1.maven.org/maven2/commons-fileupload/commons-fileupload/1.2.2/commons-fileupload-1.2.2.jar
invalid sha1: expected=ad3fda4adc95eb0d061341228cc94845ddb9a6fe computed=0ce5d4a03b07c8b00ab0
```

The invalid checksum should generally be reported to the repository owner (as [was done](#) for the above error). In the meantime, you can temporarily disable checking with the following setting:

```
checksums in update := Nil
```

See [library management](#) for details.

I've added a plugin, and now my cross-compilations fail! This problem crops up frequently. Plugins are only published for the Scala version that sbt uses (currently, 2.9.1). You can still *use* plugins during cross-compilation, because sbt only looks for a 2.9.1 version of the plugin.

... unless you specify the plugin in the wrong place!

A typical mistake is to put global plugin definitions in `~/.sbt/plugins.sbt`. **THIS IS WRONG.** `.sbt` files in `~/.sbt` are loaded for *each* build—that is, for *each* cross-compilation. So, if you build for Scala 2.9.0, sbt will try to find

a version of the plugin that's compiled for 2.9.0—and it usually won't. That's because it doesn't *know* the dependency is a plugin.

To tell sbt that the dependency is an sbt plugin, make sure you define your global plugins in a `.sbt` file in `~/.sbt/plugins/`. sbt knows that files in `~/.sbt/plugins` are only to be used by sbt itself, not as part of the general build definition. If you define your plugins in a file under *that* directory, they won't foul up your cross-compilations. Any file name ending in `.sbt` will do, but most people use `~/.sbt/plugins/build.sbt` or `~/.sbt/plugins/plugins.sbt`.

Miscellaneous

How do I use the Scala interpreter in my code? sbt runs tests in the same JVM as sbt itself and Scala classes are not in the same class loader as the application classes. Therefore, when using the Scala interpreter, it is important to set it up properly to avoid an error message like:

```
Failed to initialize compiler: class scala.runtime.VolatileBooleanRef not found.
** Note that as of 2.8 scala does not assume use of the java classpath.
** For the old behavior pass -usejavacp to scala, or if using a Settings
** object programmatically, settings.usejavacp.value = true.
```

The key is to initialize the Settings for the interpreter using *embeddedDefaults*. For example:

```
val settings = new Settings
settings.embeddedDefaults[MyType]
val interpreter = new Interpreter(settings, ...)
```

Here, `MyType` is a representative class that should be included on the interpreter's classpath and in its application class loader. For more background, see the [original proposal](#) that resulted in *embeddedDefaults* being added.

Similarly, use a representative class as the type argument when using the *break* and *breakIf* methods of *ILoop*, as in the following example:

```
def x(a: Int, b: Int) = {
  import scala.tools.nsc.interpreter.ILoop
  ILoop.breakIf[MyType](a != b, "a" -> a, "b" -> b )
}
```

0.7 to 0.10+ Migration

How do I migrate from 0.7 to 0.10+? See the [migration page](#) first and then the following questions.

Where has 0.7's `lib_managed` gone? By default, sbt 0.13.5 loads managed libraries from your ivy cache without copying them to a `lib_managed` directory. This fixes some bugs with the previous solution and keeps your project directory small. If you want to insulate your builds from the ivy cache being cleared, set `retrieveManaged := true` and the dependencies will be copied to `lib_managed` as a build-local cache (while avoiding the issues of `lib_managed` in 0.7.x).

This does mean that existing solutions for sharing libraries with your favoured IDE may not work. Refer to [Community Plugins page](#) for a list of currently available plugins for your IDE.

What are the commands I can use in 0.13.5 vs. 0.7? For a list of commands, run `help`. For details on a specific command, run `help <command>`. To view a list of tasks defined on the current project, run `tasks`. Alternatively, see the [Running](#) page in the Getting Started Guide for descriptions of common commands and tasks.

If in doubt start by just trying the old command as it may just work. The built in TAB completion will also assist you, so you can just press TAB at the beginning of a line and see what you get.

The following commands work pretty much as in 0.7 out of the box:

```
reload
update
compile
test
testOnly
publishLocal
exit
```

Why have the resolved dependencies in a multi-module project changed since 0.7? sbt 0.10 fixes a flaw in how dependencies get resolved in multi-module projects. This change ensures that only one version of a library appears on a classpath.

Use `last update` to view the debugging output for the last `update` run. Use `show update` to view a summary of files comprising managed classpaths.

My tests all run really fast but some are broken that weren't in 0.7!

Be aware that compilation and tests run in parallel by default in sbt 0.13.5. If your test code isn't thread-safe then you may want to change this behaviour by adding one of the following to your `build.sbt`:

```
// Execute tests in the current project serially.
// Tests from other projects may still run concurrently.
parallelExecution in Test := false

// Execute everything serially (including compilation and tests)
parallelExecution := false
```

What happened to the web development and Web Start support since 0.7? Web application support was split out into a plugin. See the [xsbt-web-plugin](#) project.

For an early version of an xsbt Web Start plugin, visit the [xsbt-webstart](#) project.

How are inter-project dependencies different in 0.13.5 vs. 0.7? In 0.13.5, there are three types of project dependencies (classpath, execution, and configuration) and they are independently defined. These were combined in a single dependency type in 0.7.x. A declaration like:

```
lazy val a = project("a", "A")
lazy val b = project("b", "B", a)
```

meant that the B project had a classpath and execution dependency on A and A had a configuration dependency on B. Specifically, in 0.7.x:

1. Classpath: Classpaths for A were available on the appropriate classpath for B.
2. Execution: A task executed on B would be executed on A first.
3. Configuration: For some settings, if they were not overridden in A, they would default to the value provided in B.

In 0.13.5, declare the specific type of dependency you want. Read about [multi-project builds](#) in the Getting Started Guide for details.

Where did class/object X go since 0.7? 0.7

0.13.5

FileUtilities

IO

Path class and object

Path object, File, RichFile

</td>

PathFinder class

Seq[File], PathFinder class, PathFinder object

Where can I find plugins for 0.13.5? See [Community Plugins](#) for a list of currently available plugins.

Index

This is an index of common methods, types, and values you might find in an sbt build definition. For command names, see [Running](#). For available plugins, see [the plugins list](#).

Values and Types

Dependency Management

- [ModuleID](#) is the type of a dependency definition. See [Library Management](#).
- [Artifact](#) represents a single artifact (such as a jar or a pom) to be built and published. See [Library Management](#) and [Artifacts](#).
- A [Resolver](#) can resolve and retrieve dependencies. Many types of Resolvers can publish dependencies as well. A repository is a closely linked idea that typically refers to the actual location of the dependencies. However, sbt is not very consistent with this terminology and repository and resolver are occasionally used interchangeably.
- A [\[ModuleConfiguration\]](#) defines a specific resolver to use for a group of dependencies.
- A [Configuration](#) is a useful Ivy construct for grouping dependencies. See [ivy-configurations](#). It is also used for [scoping settings](#).
- [Compile](#), [Test](#), [Runtime](#), [Provided](#), and [Optional](#) are predefined [configurations](#).

Settings and Tasks

- A [Setting](#) describes how to initialize a specific setting in the build. It can use the values of other settings or the previous value of the setting being initialized.
- A [SettingsDefinition](#) is the actual type of an expression in a build.sbt. This allows either a single [Setting](#) or a sequence of settings ([SettingList](#)) to be defined at once. The types in a [.scala build definition](#) always use just a plain [Setting](#).

- [Initialize](#) describes how to initialize a setting using other settings, but isn't bound to a particular setting yet. Combined with an initialization method and a setting to initialize, it produces a full [Setting](#).
- [TaskKey](#), [SettingKey](#), and [InputKey](#) are keys that represent a task or setting. These are not the actual tasks, but keys that are used to refer to them. They can be scoped to produce [ScopedTask](#), [ScopedSetting](#), and [ScopedInput](#). These form the base types that provide the Settings methods.
- [InputTask](#) parses and tab completes user input, producing a task to run.
- [Task](#) is the type of a task. A task is an action that runs on demand. This is in contrast to a setting, which is run once at project initialization.

Process

- A [ProcessBuilder](#) is the type used to define a process. It provides combinators for building up processes from smaller processes.
- A [Process](#) represents the actual forked process.
- The [Process companion object](#) provides methods for constructing primitive processes.

Build Structure

- [Build](#) is the trait implemented for a [.scala build definition](#), which defines project relationships and settings.
- [Plugin](#) is the trait implemented for sbt [plugins](#).
- [Project](#) is both a trait and a companion object that declares a single module in a build. See [.scala build definition](#).
- [Keys](#) is an object that provides all of the built-in keys for settings and tasks.
- [State](#) contains the full state for a build. It is mainly used by [Commands](#) and sometimes [Input Tasks](#). See also [State and Actions](#).

Methods

Settings and Tasks See the [Getting Started Guide](#) for details.

- `:=`, `+=`, `++=` These construct a [Setting](#), which is the fundamental type in the [settings](#) system.
- `value` This uses the value of another setting or task in the definition of a new setting or task. This method is special (it is a macro) and cannot be used except in the argument of one of the setting definition methods above (`:=`, `...`) or in the standalone construction methods `Def.setting` and `Def.task`. See [more about settings](#) for details.

- `in` specifies the [Scope](#) or part of the [Scope](#) of a setting being referenced. See [scopes](#).

File and IO See [RichFile](#), [PathFinder](#), and [Paths](#) for the full documentation.

- `/` When called on a single `File`, this is `new File(x,y)`. For `Seq[File]`, this is applied for each member of the sequence..
- `*` and `**` are methods for selecting children (`*`) or descendants (`**`) of a `File` or `Seq[File]` that match a filter.
- `|`, `||`, `&&`, `&`, `-`, and `--` are methods for combining filters, which are often used for selecting `Files`. See [NameFilter](#) and [FileFilter](#). Note that methods with these names also exist for other types, such as collections (like `Seq`) and [Parser](#) (see [Parsing Input](#)).
- `pair` Used to construct mappings from a `File` to another `File` or to a `String`. See [Mapping Files](#).
- `get` forces a [PathFinder](#) (a call-by-name data structure) to a strict `Seq[File]` representation. This is a common name in Scala, used by types like `Option`.

Dependency Management See [Library Management](#) for full documentation.

- `%` This is used to build up a [ModuleID](#).
- `%%` This is similar to `%` except that it identifies a dependency that has been [cross built](#).
- `from` Used to specify the fallback URL for a dependency
- `classifier` Used to specify the classifier for a dependency.
- `at` Used to define a Maven-style resolver.
- `intransitive` Marks a [dependency](#) or [Configuration](#) as being intransitive.
- `hide` Marks a [Configuration](#) as internal and not to be included in the published metadata.

Parsing These methods are used to build up [Parsers](#) from smaller [Parsers](#). They closely follow the names of the standard library's parser combinators. See [Parsing Input](#) for the full documentation. These are used for [Input Tasks](#) and [Commands](#).

- `~`, `~>`, `<~` Sequencing methods.
- `??`, `?` Methods for making a `Parser` optional. `?` is postfix.
- `id` Used for turning a `Char` or `String` literal into a `Parser`. It is generally used to trigger an implicit conversion to a `Parser`.
- `|`, `||` Choice methods. These are common method names in Scala.

- `^^^` Produces a constant value when a Parser matches.
- `+`, `*` Postfix repetition methods. These are common method names in Scala.
- `map`, `flatMap` Transforms the result of a Parser. These are common method names in Scala.
- `filter` Restricts the inputs that a Parser matches on. This is a common method name in Scala.
- `-` Prefix negation. Only matches the input when the original parser doesn't match the input.
- `examples`, `token` Tab completion
- `!!!` Provides an error message to use when the original parser doesn't match the input.

Processes These methods are used to [fork external processes](#). Note that this API has been included in the Scala standard library for version 2.9. [Process-Builder](#) is the builder type and [Process](#) is the type representing the actual forked process. The methods to combine processes start with `#` so that they share the same precedence.

- `run`, `!`, `!!`, `!<`, `lines`, `lines_!` are different ways to start a process once it has been defined. The lines variants produce a `Stream[String]` to obtain the output lines.
- `#<`, `#<<`, `#>` are used to get input for a process from a source or send the output of a process to a sink.
- `#|` is used to pipe output from one process into the input of another.
- `#||`, `#&&`, `###` sequence processes in different ways.